

LEARNING BETTER PHYSICS:
A MACHINE LEARNING APPROACH TO LATTICE GAUGE THEORY

by

Samuel Alfred Foreman

A thesis submitted in partial fulfillment
of the requirements for the
Doctor of Philosophy
degree in Physics
in the Graduate College of
The University of Iowa

August 2019

Thesis Supervisor: Professor Yannick Meurice

Acknowledgements

First and foremost I would like to thank my supervisor Yannick Meurice for his support and encouragement throughout my entire graduate research experience. I would also like to thank both James Osborn, and Xiao-Yong Jin for their incredible patience, guidance, and expertise. Finally, I would like to thank Argonne National Laboratory and the US Department of Energy for their financial support and hospitality.

This research was supported in part by the Office of Science Graduate Student Research (SCGSR) program. The SCGSR program is administered by the Oak Ridge Institute for Science and Education (ORISE) for the DOE. ORISE is managed by ORAU under contract number DESC0014664.

Additional support was provided by DOE grant DE-SC0010113.

Abstract

In this work we explore how lattice gauge theory stands to benefit from new developments in machine learning, and look at two specific examples that illustrate this point. We begin with a brief overview of selected topics in machine learning for those who may be unfamiliar, and provide a simple example that helps to show how these ideas are carried out in practice.

After providing the relevant background information, we then introduce an example of renormalization group (RG) transformations, inspired by the tensor RG, that can be used for arbitrary image sets, and look at applying this idea to equilibrium configurations of the two-dimensional Ising model.

The second main idea presented in this thesis involves using machine learning to improve the efficiency of Markov Chain Monte Carlo (MCMC) methods. Explicitly, we describe a new technique for performing Hamiltonian Monte Carlo (HMC) simulations using an alternative leapfrog integrator that is parameterized by weights in a neural network. This work is based on the L2HMC ('Learning to Hamiltonian Monte Carlo') algorithm introduced in [1].

Public Abstract

In recent years there has been a growing interest in expanding the scientific applications of machine learning. In this work we will explore two specific examples of how lattice gauge theory stands to benefit from these new developments. Lattice gauge theory is a sub-discipline of high energy physics that makes heavy use of computer simulations to test existing models and generate new results. In order to perform these simulations, spacetime is discretized as a lattice whose behavior can then be controlled using the laws of physics.

We begin by introducing a new technique that is capable of extracting information about a physical system by ‘looking’ at pictures of the system at different temperatures, and provide an analytical framework that explains how this is done behind the scenes. This result demonstrates that **our approach is capable of learning non-trivial information about the system without being told explicitly how to do so**, and without having to provide any information about the underlying physics.

Next, we look at how machine learning can be used to improve the efficiency of the *Hamiltonian Monte Carlo* (HMC) algorithm, a widely used technique in lattice gauge theory for generating *gauge configurations*. These gauge configurations are essentially ‘snapshots’ of the spacetime lattice, and are used to make predictions about quantum theory.

Currently, these configurations are generated via *Hamiltonian Monte Carlo simulations*, an algorithm that can be summarized as follows:

1. Start with a random initial configuration.
2. Propose a new configuration by (approximately) evolving the current state through time¹.
3. Check if this new configuration is better than the previous one. If so, we accept it, otherwise, we retain the current configuration.²

The difference between the current and proposed configurations can be controlled through the *step size*, a parameter that is (typically) fixed for the duration of the simulation. We can improve the likelihood of

¹Using *Hamilton’s equations*, which describe how a system changes in time.

²Technically, configurations which are ‘worse’ will still be accepted occasionally, but this becomes increasingly unlikely as the drop in ‘quality’ increases.

accepting a new configuration by using a smaller step size, but this leads to configurations that are highly correlated with each other, which is undesirable. Alternatively, we can take larger steps to try and reduce correlations, but this causes more of the proposed configurations to be rejected. Immediately we see that computational resources are being wasted each time we propose a new configuration that gets rejected, and is a major source of inefficiency in the algorithm.

The approach presented in this work attempts to combat this issue by using machine learning to reduce the number of wasted calculations. Explicitly, this is done by modifying the equations that govern how our system evolves in time, and then training the algorithm to identify those modifications that ultimately produce ‘better’ configurations. In doing so, we are able to reduce the number of unnecessary calculations (which don’t produce new configurations), thereby improving the efficiency of the algorithm as a whole.

Contents

LIST OF FIGURES	x
1 INTRODUCTION	1
2 MACHINE LEARNING: AN OVERVIEW	3
2.1 Introduction	3
2.2 Supervised Learning and Gradient Descent	4
2.2.1 Example: Linear Regression	5
2.2.2 Probabilistic Interpretation	6
2.3 Feed-forward Neural Networks	8
2.3.1 Parameter optimization & Backpropagation	9
2.3.1.1 Derivatives of the Error Function	11
2.3.1.2 Output Layer	12
2.3.1.3 Hidden Layers	12
2.4 Convolutional Neural Networks	13
2.4.1 Local Receptive Fields	15
2.4.2 Shared Weights	16
2.4.3 Pooling Layers	17
2.4.4 Hyperparameters	17
2.5 Conclusion	18
3 UNSUPERVISED LEARNING: ISING WORMS	19
3.1 Introduction	19
3.2 From Loops to Images	21
3.3 PCA and Criticality	24

3.4	TRG Coarse-graining	25
3.5	Image Coarse-graining	28
3.6	Partial Data Collapse for Blocked Images	29
3.7	TRG Calculation of $\langle N_b \rangle$	32
3.8	Technical Results	34
3.8.1	Loop Representation	34
3.8.2	Heat Capacity	36
3.8.3	Monte Carlo Implementation	36
3.8.4	Tests	37
3.8.5	Conjecture About λ_{\max}	37
3.8.6	Illustration of Alternate Blockings	43
3.9	Possible Applications: From Images to Loops	43
3.10	Conclusions	45
4	MARKOV CHAIN MONTE CARLO (MCMC)	47
4.1	Markov Chains	47
4.1.1	Metropolis-Hastings Algorithm	48
4.2	Aside: Bayesian Analysis	50
4.3	Hamiltonian Monte Carlo	50
4.3.1	Hamiltonian Dynamics	51
4.3.1.1	Properties of Hamiltonian Dynamics	52
5	SEMI-SUPERVISED LEARNING: L2HMC	54
5.1	Introduction	54
5.2	Generalizing the Leapfrog Integrator	54
5.2.1	Metropolis-Hastings Accept/Reject	56
5.3	Forward Direction ($d = 1$):	57
5.4	Backward Direction ($d = -1$):	58
5.5	Determinant of the Jacobian	58
5.6	Network Architecture	58
5.7	Training Procedure	60

5.8 Gaussian Mixture Model	61
5.8.1 Example	62
5.9 2D $U(1)$ Lattice Gauge Theory	65
5.9.1 Modified Network Architecture	66
5.9.2 Annealing Schedule	69
5.9.3 Modified Loss Function for $U(1)$ Gauge Model	69
5.9.4 Issues with the Average Plaquette	70
5.10 Conclusion	72
A APPENDIX: L2HMC SOURCE CODE	74
A.1 README.md	74
A.1.1 l2hmc-qcd	74
A.1.1.1 Overview	74
A.1.1.2 Modified Implementation	74
A.1.1.3 Features	74
A.1.1.4 Organization	75
A.1.1.5 Contact	75
A.1.1.6 Citation	75
A.2 l2hmc-qcd/args/args.txt	76
A.3 l2hmc-qcd/main.py	78
A.4 l2hmc-qcd/inference.py	85
A.5 l2hmc-qcd/models/model.py	93
A.6 l2hmc-qcd/dynamics/gauge_dynamics.py	103
A.7 l2hmc-qcd/lattice/lattice.py	113
A.8 l2hmc-qcd/network/conv_net.py	117
A.9 l2hmc-qcd/network/generic_net.py	122
A.10 l2hmc-qcd/network/network.py	124
A.11 l2hmc-qcd/network/network_utils.py	126
A.12 l2hmc-qcd/trainers/trainer.py	130
A.13 l2hmc-qcd/runners/runner.py	133
A.14 l2hmc-qcd/loggers/run_logger.py	136

A.15 l2hmc-qcd/loggers/train_logger.py	145
A.16 l2hmc-qcd/plotters/gauge_model_plotter.py	148
A.17 l2hmc-qcd/plotters/leapfrog_plotters.py	156
A.18 l2hmc-qcd/utils/parse_args.py	161
A.19 l2hmc-qcd/utils/parse_inference_args.py	168
A.20 l2hmc-qcd/utils/file_io.py	170
A.21 l2hmc-qcd/loggers/summary_utils.py	175
A.22 l2hmc-qcd/utils/data_loader.py	179
B APPENDIX: SYSTEMATIC DEBUGGING RESULTS OF THE AVERAGE PLAQUETTE	182
REFERENCES	189

List of Figures

Figure 2.1: Network diagram for a two-layer fully-connected neural network. The input, hidden, and output variables are shown as nodes while the weight parameters are the links between them.	9
Figure 2.2: Illustration of the local receptive fields in the input image and their corresponding weights in the first hidden layer. Additionally, we can see how these local receptive fields are slid across the input image to construct additional units in the hidden layer.	16
Figure 2.3: Illustration of a 2×2 max pooling layer.	17
Figure 3.1: (a) Picture of an eye with 4096 pixels; (b) black and white version with a graycut at 0.72; (c) boundaries of the black domains.	22
Figure 3.2: (a) Legal worm configuration on an $L \times L$ lattice with periodic boundary conditions and; (b) its equivalent representation as a $2L \times 2L$ black and white pixel image.	24
Figure 3.3: λ_{max} and $\frac{3}{2} \langle \Delta_{Nb}^2 \rangle / V$ (per unit volume) vs. T , illustrating the relation between the eigenvalue corresponding to the first principal component and the logarithmic divergence of the specific heat. The inset shows a qualitative agreement near the critical temperature.	26
Figure 3.4: Illustration of the tensor blocking discussed in the text. Each dot is a tensor at a lattice site with four lines coming out, each representing a tensor index. Lines connecting dots represent tensor contractions.	27
Figure 3.5: (a) T_{1100} vs. $T - T_c^{(2s)}$ for six successive iterations of the blocking transformation, beginning with an initial lattice $L = 64$; (b) T_{1100} vs. $(T - T_c^{(2s)})/L_{\text{eff}}$ illustrating the data collapse, where $T_c^{(2s)}$ is the critical temperature of the two state projection, beginning at iteration 0 on an $L = 64$ lattice.	28
Figure 3.6: Illustration of the coarse-graining procedure in which the original lattice sites (i) are replaced by blocked sites ($2i$) (grey circles) with twice the original lattice spacing. In the coarse-grained lattice, an elementary block (blue) consists of four sites on the original lattice, eight external bonds (red), and four blocked external bonds (green).	29
Figure 3.7: (a) Illustration of the $1 + 1 \rightarrow 0$ blocking procedure discussed in the text: original configuration; (b) introduction of the blocks and replacement of single or double bounds according to the $1 + 1 \rightarrow 0$ rule; (c) construction of the corresponding blocked image.	30

Figure 3.18: $\langle N_b \rangle$ and $\langle \Delta_{N_b}^2 \rangle$ vs. grayscale cutoff value for 500 randomly chosen images from the CIFAR-10 dataset.	45
Figure 4.1: Visualizing HMC for a 1D Gaussian (example from [49], figure adapted with permission from [50]). Each Hamiltonian Markov transition lifts the initial state onto a random level set of the Hamiltonian, $\mathcal{H}^{(-1)}(E)$, which can then be explored with a Hamiltonian trajectory before projecting back down to the target parameter space	53
Figure 5.1: Example of how the determinant of the Jacobian can deform the energy landscape.	57
Figure 5.2: Illustration showing the generic (fully-connected) network architecture for training S_v , Q_v , and T_v . Figure adapted with permission from [50].	59
Figure 5.3: Flowchart illustrating the generic fully-connected network architecture including the intermediate variables computed at each hidden layer of the network.	60
Figure 5.4: Comparison of trajectories generated using L2HMC (top), and traditional HMC with $\epsilon = 0.25$ (middle) and $\epsilon = 0.5$ (bottom). Note that L2HMC is able to successfully mix between modes, whereas HMC is not.	63
Figure 5.5: Autocorrelation vs. gradient evaluations (i.e. MD steps). Note that L2HMC (blue) has a significantly reduced autocorrelation after the same number of gradient evaluations when compared to either of the two HMC trajectories	64
Figure 5.6: Illustration of an elementary plaquette on the lattice.	66
Figure 5.7: (left) Example of topological freezing in the 2D $U(1)$ lattice gauge theory, generated from generic HMC sampling for a 16×16 lattice. Note that for the majority of the simulation $Q = -2$, making it virtually impossible to get a reasonable estimate of χ . (right) Topological charge vs. step generated using the trained L2HMC sampler.	67
Figure 5.8: Convolutional structure used for learning localized features of rectangular lattice.	67
Figure 5.9: Illustration taken from TensorBoard showing an overview of the network architecture for VNet. Note that the architecture is identical for XNet.	68
Figure 5.10: Detailed view of additional convolutional structure included to better account for rectangular geometry of lattice inputs.	68
Figure 5.11: (left): Average plaquette $\langle \phi_p \rangle$ vs. step for $L = 8$, $N_{LF} = 7$, and $\alpha_Q = 0.5$. Here the solid red line indicates the true value of the average plaquette (in the infinite volume limit, and is equal to 0.89338 . . . (right): Difference between the observed and expected value of the average plaquette δ_{ϕ_p} vs. step. Note that $\delta_{\phi_p} \neq 0$	71
Figure 5.12: Same quantities as in Fig 5.11, with $\alpha_Q = 0$. Note that the discrepancy δ_{ϕ_p} is no longer present.	71
Figure B.1: δ_{ϕ_p} vs MD step with $N_{LF} = 5$ for $\beta = 5.0$ (top row) and $\beta = 6.0$ (bottom row). The results from the trained L2HMC (generic HMC) sampler are shown in the left (right) column. As can be seen, the difference δ_{ϕ_p} remains roughly consistent for all values of α_Q	182

1 | Introduction

In recent years there has been an explosive growth in the fields of machine learning (ML) and data science. This trend has had a significant impact on a wide variety of industries and has been used to produce incredible new technologies ranging from self-driving cars to personalized recommendation engines and facial recognition software, to name just a few. Just as importantly, (but not as glamorously) many scientific researchers have also begun looking for new ways to apply these ideas to their fields, producing new interdisciplinary efforts and mutually-beneficial collaborations.

One field in particular that stands to benefit from this new direction is high energy physics, which relies heavily on computational science through the use of data analysis and computer simulations. Historically, much of this work been done either through the use of ‘brute-force’ calculations, requiring tremendous computational resources, or ‘by-hand’ which tends to be pain-stakingly slow and is often subject to a whole variety of potential issues (e.g. incomplete information, poor statistics, incorrect models, unjustified approximations, human-error, etc.). Machine learning, on the other hand, aims to sidestep many of these problems entirely by automatically ‘learning’ information from data, and is capable of discovering meaningful patterns that are oftentimes imperceivable to humans. With projects such as the Large Hadron Collider at CERN producing roughly 15 petabytes of data per year, the need for new and better methods for dealing with this data continues to grow.

This thesis is primarily composed of the work completed during my graduate career and includes some relevant background information that may be useful for those looking to learn more information about how ideas from machine learning can be applied to problems in lattice gauge theory and lattice quantum chromodynamics (QCD). In particular, Chapter 2 provides a thorough background on many of the machine learning tools used throughout the remainder of the thesis and provides concrete examples on their use. For example, topics such as such as supervised learning, gradient descent / backpropagation, feed-forward and convolutional neural networks are covered.

Chapter 3 covers an example of how unsupervised learning (specifically, principal component analysis) can be applied to extract information about the phase transition of the two-dimensional Ising model. By rep-

resenting equilibrium configurations of the system as two-dimensional greyscale images, principal component analysis allows us to obtain a direct relationship between the specific heat capacity and the eigenvalue of the dominant principal component. In Sec. 3.4 and Sec. 3.5, a renormalization group transformation is proposed that can be applied to generic sets of images, which when applied to the images under consideration, leads to a finite-size scaling analysis of the critical point.

In Chapter 4, a brief overview of Markov Chain Monte Carlo (MCMC) methods in general is discussed, and relevant notation introduced. In sections 4.3.1, 5.1, and 5.9, some of the current problems faced by HMC are discussed, particularly within the context of simulations in lattice gauge theory and lattice QCD.

Chapter 5 describes a new technique for applying supervised learning to help improve the efficiency of Hamiltonian / Hybrid Monte Carlo (HMC) simulations. This new approach is called ‘Learning to Hamiltonian Monte Carlo’ (L2HMC), and is based off of the work described in [1]. The details of the L2HMC algorithm are presented in Sec 5.2, and an example of this algorithm applied to a two-dimensional Gaussian Mixture Model is included in Sec. 5.8. Building on these results, we proceed to look at applying this approach to a two-dimensional $U(1)$ lattice gauge theory, the details of which are laid-out in Sec. 5.9. Finally, in Sec. 5.9.1- 5.9.3 we discuss some of the modifications that were introduced when applying this approach to the gauge model under consideration, and provide insight into why these modifications were both necessary and advantageous.

2 | Machine Learning: An Overview

2.1 Introduction

Broadly speaking, the subject of *machine learning* (ML), refers to the automated detection of meaningful patterns in data [2], and encompasses two major classes of problems [3]: *estimation* and *prediction*. To better understand the difference between the two, we begin with an example. Suppose we observe some measurable quantity x (e.g. measurements of the acceleration of an object in a gravitational field) of the system under consideration that is related to some parameters w (e.g. the gravitational constant) of a model $p(x|w)$ that describes the probability of observing x given w . We can perform an experiment to obtain a dataset x , and use this data to “fit” our model. This procedure of fitting the model typically corresponds to finding the \hat{w} that maximizes the probability of observing x , i.e. $\hat{w} \equiv \text{argmax}_w\{p(x|w)\}$. In this context, *estimation* problems are those concerned with the accuracy of \hat{w} , whereas *prediction* problems are concerned with the ability of the model to predict new observations, i.e. the accuracy of $p(x|\hat{w})$.

For our purposes, we mainly restrict our attention to *prediction*-type problems, which includes the types of problems most frequently associated with machine learning (i.e. regression, classification, etc.). There are three generic approaches commonly used to tackle these types of problems, namely:

1. Supervised learning
2. Unsupervised learning
3. Reinforcement learning

While it is true that most machine learning problems fit nicely into one of these three categories, this is not always the case. Some of the most successful approaches to historically difficult problems have employed ideas from combinations of the three in new and unexpected ways (e.g. *variational auto-encoders* (VAEs) [4], and *generative adversarial networks* (GANs) [5]).

2.2 Supervised Learning and Gradient Descent

Supervised learning is the task of inferring a function from labeled training data. Our training data consists of (input, output) pairs: $\{(x^{(i)}, y^{(i)})\}, i = 1, \dots, n$, with $x^{(i)} \in \mathbb{R}^p$ being the inputs (or features) and $y^{(i)} \in \mathbb{R}^d$, being the outputs (or target) (often $d = 1$) that we wish to predict.¹ For concreteness, suppose we have a collection of n two-dimensional greyscale images, $x^{(1)}, x^{(2)}, \dots, x^{(n)}$, each of which contains p pixels flattened into a vector, and their associated labels $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ matching the type of animal present in the image (e.g. dog, cat, frog, etc). Our goal then is to find a function f that is able to accurately approximate the output $y^{(i)}$ for a given input $x^{(i)}$. If our output is discrete, as in the present case, the problem is said to be a *classification* problem, whereas continuous outputs are referred to as *regression* problems. In order to measure how well our function performs, we often split the available data into two disjoint sets called the training set and test set respectively. The idea is to train our function (using a suitably chosen procedure) on the training data and then use the (previously unseen) data from the test set to evaluate the performance. The ability to accurately predict the desired output for a new, unused input is known as *generalizability* and is an important metric for measuring the quality of a given model.

In order to find this desired function f_w , we must introduce a way to measure the similarity between the expected ($y^{(i)}$) and predicted output $f_w(x^{(i)})$. One way to do this is to introduce a *loss* (or *cost*) function, $\mathcal{L} = \mathcal{L}[f_w(x^{(i)}), y^{(i)}] \equiv \mathcal{L}(w)$, with the idea being that the loss is small when $f_w(x^{(i)}) \simeq y^{(i)}$. In doing so, we are able to improve the quality of our desired function f_w by adjusting the values of w in such a way that the loss function is minimized. One of the most popular techniques for carrying out this optimization is an algorithm known as **gradient descent**. Explicitly, given some initial (often random) values of the parameters w , gradient descent repeatedly updates their values by “stepping” in the direction of steepest decrease of \mathcal{L} :

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} \mathcal{L}(w), \quad (2.1)$$

simultaneously for all values of $j = 0, \dots, n$.

Here, α is a *hyperparameter* called the **learning rate**, and it is responsible for determining the “step size” the algorithm takes with each update. It is important to note that the value of the learning rate must be chosen appropriately since large steps can potentially cause stability issues (resulting from ‘over-shooting’ the minima), whereas taking steps which are too small can dramatically increase the amount of updates necessary to obtain the minimum.

¹Here superscript (i) indexes samples in the training set.

2.2.1 Example: Linear Regression

To make all of these ideas explicit we consider the example of *linear regression*. We begin by assuming that, to good approximation, y can be expressed as a linear function of x :

$$f_w(x) = \sum_{i=0}^n w_i x_i = w^T x \quad (2.2)$$

where we've defined $x_0 \equiv 1$ as the intercept (or *bias*) term.

Next, we must choose a loss function \mathcal{L} . For this example, we can choose the least-squares cost function defined to be

$$\mathcal{L}(w) = \frac{1}{2} \sum_{i=1}^m (f_w(x_i) - y_i)^2. \quad (2.3)$$

In order to implement gradient descent using this cost function, we first calculate the partial derivatives for the case of a single training example (x, y) . This gives

$$\frac{\partial}{\partial w_j} \mathcal{L}(w) = \frac{\partial}{\partial w_j} \frac{1}{2} (f_w(x) - y)^2 \quad (2.4)$$

$$= (f_w(x) - y) \cdot \frac{\partial}{\partial w_j} \left(\sum_{i=0}^n w_i x_i - y \right) \quad (2.5)$$

$$= (f_w(x) - y) x_j, \quad (2.6)$$

and so

$$w_j := w_j + \alpha \left(y^{(i)} - f_w(x^{(i)}) \right) x_j^{(i)}. \quad (2.7)$$

To expand this result to the case of multiple training samples, we have two options:

1. Repeat the above update for every j until convergence, i.e. look at every example in the training set for every update step. This approach is known as **batch gradient descent**.
2. Repeatedly run through the training set, and for each individual training example, update the parameters according to the gradient of the error with respect to *that single training example only*. This approach is known as **stochastic gradient descent**.

Note that we can generalize this approach by constructing the *design matrix* X to be the $m \times n^2$ matrix $X_{ij} = x_j^{(i)}$, i.e. each row containing an individual training example. Similarly, we can write the target as an m -dimensional vector containing all the target values from the training set: $y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]^T$. In terms

²technically $m \times n + 1$, accounting for the intercept term

of these then, our cost function becomes

$$\mathcal{L}(w) = \frac{1}{2}(Xw - \mathbf{y})^T(Xw - \mathbf{y}) \quad (2.8)$$

Again, computing the required gradient gives

$$\nabla_w \mathcal{L}(w) = \nabla_w \frac{1}{2}(Xw - \mathbf{y})^T(Xw - \mathbf{y}) \quad (2.9)$$

$$= \frac{1}{2} \nabla_w (w^T X^T X w - w^T X^T \mathbf{y} - \mathbf{y}^T X w + \mathbf{y}^T \mathbf{y}) \quad (2.10)$$

$$= \frac{1}{2} \nabla_w \text{tr}(w^T X^T X w - w^T X^T \mathbf{y} - \mathbf{y}^T X w + \mathbf{y}^T \mathbf{y}) \quad (2.11)$$

$$= \frac{1}{2} \nabla_w (\text{tr} w^T X^T X w - 2 \text{tr} \mathbf{y}^T X w) \quad (2.12)$$

$$= \frac{1}{2} (X^T X w + X^T X w - 2 X^T \mathbf{y}) \quad (2.13)$$

$$= X^T X w - X^T \mathbf{y} \quad (2.14)$$

Setting this equal to zero and solving gives the **normal equations**

$$X^T X w = X^T \mathbf{y}, \quad (2.15)$$

from which we get

$$\hat{w} = (X^T X)^{-1} X^T \mathbf{y}. \quad (2.16)$$

2.2.2 Probabilistic Interpretation

To motivate the reasoning behind our choice of the cost function for linear regression, we adopt a probability approach. Assume

$$y^{(i)} = w^T x^{(i)} + \varepsilon^{(i)} \quad (2.17)$$

with $\varepsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$ is an error term that captures either unmodeled effects (e.g. caused by features that are not accounted for) or random noise. Assume $\varepsilon^{(i)}$ are i. i. d. The probability density is given by

$$p(\varepsilon^{(i)}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(\varepsilon^{(i)})^2}{2\sigma^2}\right) \quad (2.18)$$

This implies that

$$p(y^{(i)} | x^{(i)}; w) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right) \quad (2.19)$$

and we say that $p(y^{(i)}|x^{(i)}; w)$ is the distribution of $y^{(i)}$ given $x^{(i)}$ and parameterized by w , or equivalently, $y^{(i)}|x^{(i)}; w \sim \mathcal{N}(w^T x^{(i)}; \sigma^2)$. The question we have now is: given the design matrix X and the weights w , what is the distribution of the $y^{(i)}$'s? The probability of the data is given by $p(\mathbf{y}|X; w)$, which is typically viewed as a function of \mathbf{y} (and perhaps X) for a fixed value of w . When we want to explicitly view this as a function of w , we will instead call it the **likelihood function**,

$$L(w) = L(w; X, \mathbf{y}) \equiv p(\mathbf{y}|X; w) \quad (2.20)$$

By the independence assumption on the $\varepsilon^{(i)}$'s (and consequently, also the $y^{(i)}$'s and the $x^{(i)}$'s) this can also be written as

$$L(w) = \prod_{i=1}^m p(y^{(i)}|x^{(i)}; w) \quad (2.21)$$

$$= \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right) \quad (2.22)$$

By the principle of **maximum likelihood** [6], we should choose w so as to make the data as high probability as possible, i.e. we should choose w to maximize $L(w)$. Instead of maximizing $L(w)$ directly, we can also maximize any strictly increasing function of $L(w)$. In particular, we choose to maximize the **log likelihood** $\ell(w)$:

$$\ell(w) \equiv \log L(w) \quad (2.23)$$

$$= \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right) \quad (2.24)$$

$$= \sum_{i=1}^m \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right) \quad (2.25)$$

$$= m \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2. \quad (2.26)$$

Hence, maximizing $\ell(w)$ gives the same answer as minimizing

$$\frac{1}{2} \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 \quad (2.27)$$

which is equivalent to $\mathcal{L}(w)$, the original loss function from our least squares example.

2.3 Feed-forward Neural Networks

We can construct a Neural Network model as a series of functional transformations [7]. First, we construct M linear combinations of the input variables x_1, \dots, x_D in the form

$$a_j = \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \quad (2.28)$$

where $j = 1, \dots, M$ and now here the superscript (1) indicates the ‘layer’ of the network. The quantities a_j are known as ‘activations’, each of which is transformed via a differentiable, nonlinear activation function $h(\cdot)$ to give

$$z_j = h(a_j) \quad (2.29)$$

which are often referred to as ‘hidden units’ or ‘nodes’ in the context of neural networks. These nonlinear functions are often chosen to be sigmoidal functions such as the tanh function or the rectified linear unit (ReLU), defined as

$$\text{ReLU}(x) \equiv \max(0, x). \quad (2.30)$$

Following Eq. 2.29, these values are again linearly combined to give *output unit activations*

$$a_k = \sum_{j=1}^M w_{kj}^{(2)} z_j + w_{k0}^{(2)} \quad (2.31)$$

where $k = 1, \dots, K$, and K is the total number of outputs. Depending on the context of the problem, the output unit activations may be transformed using an appropriate activation function $\sigma(\cdot)$ to give a set of network outputs $y_k = \sigma(a_k)$. For example, multi-class classification problems often use the standard (unit) softmax activation function $\sigma : \mathbb{R}^K \rightarrow \mathbb{R}^K$ given by

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad \text{for } i = 1, \dots, K \text{ and } \mathbf{z} = (z_1, \dots, z_K) \in \mathbb{R}^K \quad (2.32)$$

which takes as input a vector of K real numbers and normalizes it into a probability distribution consisting of K probabilities. A visualization of this network structure is shown in Fig 2.1. Often the bias parameters are absorbed into the set of weight parameters by defining an additional input variable x_0 whose value is fixed at $x_0 = 1$ so that Eq 2.28 becomes

$$a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i, \quad (2.33)$$

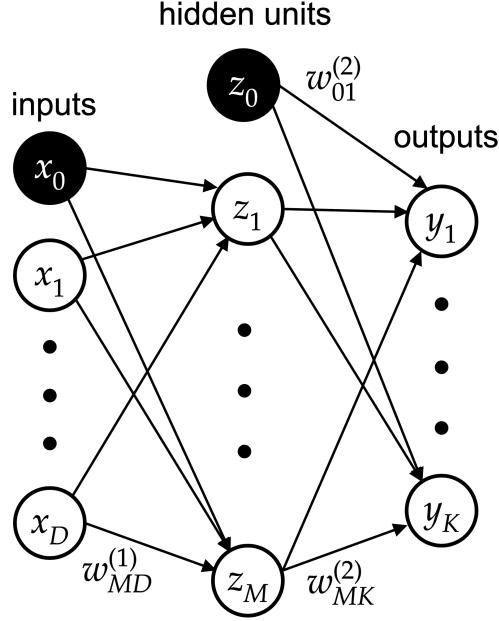


Figure 2.1: Network diagram for a two-layer fully-connected neural network. The input, hidden, and output variables are shown as nodes while the weight parameters are the links between them.

and similarly for the second layer. Written like this then, we can write the overall network function as

$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right) \quad (2.34)$$

Note that this structure can be easily generalized an arbitrary depth by successively ‘stacking’ these hidden layers.

2.3.1 Parameter optimization & Backpropagation

Given a set of input vectors $\{\mathbf{x}_n\}$ for $n = 1, \dots, N$, together with a corresponding set of target vectors $\{\mathbf{t}_n\}$, our goal is to find the values \mathbf{w}^* that minimize some appropriately chosen loss function $\mathcal{L}(\mathbf{w})$ that measures the ‘closeness’ between the desired outputs \mathbf{t}_n and the calculated outputs produced by our network $\mathbf{y}_n = y(\mathbf{x}_n, \mathbf{w})$. Similar to the gradient descent algorithm discussed in Sec. 2.2, training a neural network with gradient descent requires the calculation of the gradient of the loss function \mathcal{L} with respect to the weights w_{ij}^k and biases, collectively denoted θ . Then, according to the learning rate α , each iteration of gradient descent updates the weights according to

$$\theta^{t+1} = \theta^t - \alpha \frac{\partial \mathcal{L}}{\partial \theta} \quad (2.35)$$

where θ^t denotes the parameters of the neural network at iteration t of gradient descent. Since the nodes in the hidden layers have no target output, we can't simply define an error function that is specific to that node. Instead, any error function for that node will be dependent on the values of the parameters in the previous layers and following layers. This coupling of parameters between layers often leads to complicated expressions for the gradients and is often slow in practice. In order to address these issues, an alternative approach is frequently used which relies upon the **backpropagation** algorithm [8–11], outlined below. For consistency, we adopt the following notations:

- $X = \{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$: the dataset of N input-output pairs where \mathbf{x}_i is the input and \mathbf{t}_i is the desired output of the network on input \mathbf{x}_i
- w_{ij}^k : weight for node j in layer k for incoming node i
- $w_{0i}^k = b_i^k$: bias for node i in layer k with a fixed output of $z_0^{k-1} = 1$ for node 0 in layer $k - 1$.
- a_i^k activation (i.e. the product sum plus bias) for node i in layer k
- z_i^k : output for node i in layer k
- r_k : the number of nodes in hidden layer k
- h : activation function for the hidden layer nodes
- h_z : activation function for the output layer nodes

For simplicity, we suppose that the our neural network produces a single, scalar valued output y for a given vector-valued input \mathbf{x} , and choose a loss function given by the mean-squared error, i.e.

$$\mathcal{L}(w) = \frac{1}{2N} \sum_{i=1}^N (y_i - t_i)^2 \quad (2.36)$$

Backpropagation then attempts to minimize this loss function with respect to the neural network's weights by calculating, for each weight w_{ij}^k , the value of $\frac{\partial \mathcal{L}}{\partial w_{ij}^k}$. Since the error function can be decomposed into a sum over individual loss terms for each individual input-output pair, the derivative can be calculated with respect to each input-output pair individually and then combined at the end (due to the fact that the derivative of a

sum of functions is the sum of the derivatives of each function)

$$\frac{\partial \mathcal{L}(X, \theta)}{\partial w_{ij}^k} = \frac{1}{N} \sum_{d=1}^N \frac{\partial}{\partial w_{ij}^k} \left(\frac{1}{2} (y_d - t_d)^2 \right) \quad (2.37)$$

$$= \frac{1}{N} \sum_{d=1}^N \frac{\partial \mathcal{L}_d}{\partial w_{ij}^k}. \quad (2.38)$$

Because of this, we can restrict our attention to only one input-output pair and from this the general form for all input-output pairs can be generated by combining individual gradients. Our loss function then simplifies to

$$\mathcal{L} = \frac{1}{2} (y - t)^2 \quad (2.39)$$

2.3.1.1 Derivatives of the Error Function

We begin by applying the chain rule to the loss function partial derivative

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^k} = \frac{\partial \mathcal{L}}{\partial a_j^k} \frac{\partial a_j^k}{\partial w_{ij}^k} \quad (2.40)$$

The first term in the product is often referred to as the *error* and denoted by $\delta_j^k \equiv \partial \mathcal{L} / \partial a_j^k$. The second term can be calculated from the fact that

$$a_j^k = \sum_{j=0}^{r_{k-1}} w_{ij}^k z_j^{k-1} \quad (2.41)$$

to give

$$\frac{\partial a_j^k}{\partial w_{ij}^k} = \frac{\partial}{\partial w_{ij}^k} \left(\sum_{l=0}^{r_{k-1}} w_{ij}^k z_l^{k-1} \right) \quad (2.42)$$

$$= z_i^{k-1}. \quad (2.43)$$

Thus, the partial derivative of the loss function \mathcal{L} with respect to a weight w_{ij}^k is given by

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^k} = \delta_j^k z_i^{k-1}, \quad (2.44)$$

i.e. the partial derivative of a weight is a product of the error term δ_j^k at node j in layer k , and the output z_i^{k-1} of node i in layer $k-1$. Since the error term δ_j^k depends on both the loss function, and as we will show, the values of the error terms in the next layer, the computation of these terms proceeds backwards from the

output layer down to the input layer which is where the term *backpropagation* (of errors) gets its name.

2.3.1.2 Output Layer

Starting from the final layer, backpropagation calculates the value³ δ_1^m . Expressing the loss function in terms of the value a_1^m (since δ_1^m is a partial derivative with respect to a_1^m) gives

$$\mathcal{L} = \frac{1}{2}(t - y)^2 = \frac{1}{2}(h_z(a_1^m) - t)^2 \quad (2.45)$$

where again $h_z(x)$ is the activation function for the output layer. Applying the partial derivative and using the chain rule gives

$$\delta_1^m = (h_z(a_1^m) - t)h'_z(a_1^m) = (y - t)h'_z(a_1^m) \quad (2.46)$$

and putting it all together, the partial derivative of the loss function \mathcal{L} with respect to a weight in the final layer w_{i1}^m is

$$\frac{\partial E}{\partial w_{i1}^m} = \delta_1^m z_i^{m-1} = (y - t)h'_z(a_1^m)z_i^{m-1}. \quad (2.47)$$

2.3.1.3 Hidden Layers

Now, we calculate the partial derivatives for the hidden layers, again using the multivariate chain rule. Note that for $1 \leq k < m$, the error term δ_j^k is:

$$\delta_j^k = \frac{\partial \mathcal{L}}{\partial a_j^k} = \sum_{l=1}^{r_{k+1}} \frac{\partial \mathcal{L}}{\partial a_l^{k+1}} \frac{\partial a_l^{k+1}}{\partial a_j^k}, \quad (2.48)$$

where l ranges from 1 to r_{k+1} (the number of nodes in the next layer). Plugging in the error term δ_l^{k+1} gives

$$\delta_j^k = \sum_{l=1}^{r_{k+1}} \delta_l^{k+1} \frac{\partial a_l^{k+1}}{\partial a_j^k}. \quad (2.49)$$

Recalling that

$$a_l^{k+1} = \sum_{j=1}^{r_k} w_{jl}^{k+1} h(a_j^k) \quad (2.50)$$

³recall we are considering a one-output neural network, hence the subscript 1 and not j

where $h(x)$ is the activation function in the hidden layers, we are left with

$$\frac{\partial a_l^{k+1}}{\partial a_j^k} = w_{jl}^{k+1} h'(a_j^k). \quad (2.51)$$

Plugging this into the above equation gives a final equation for the error term δ_j^k in the hidden layers, a result known as the *backpropagation formula*:

$$\delta_j^k = \sum_{l=1}^{r_{k+1}} \delta_l^{k+1} w_{jl}^{k+1} h'(a_j^k) = h'(a_j^k) \sum_{l=1}^{r_{k+1}} w_{jl}^{k+1} \delta_l^{k+1} \quad (2.52)$$

Putting this all together we get the partial derivative of the loss function \mathcal{L} with respect to a weight in the hidden layers w_{ij}^k for $1 \leq k < m$:

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^k} = \delta_j^k z_i^{k-1} = h'(a_j^k) z_i^{k-1} \sum_{l=1}^{r_{k+1}} w_{jl}^{k+1} \delta_l^{k+1}. \quad (2.53)$$

For completeness, we include the complete algorithm in Alg. 1.

2.4 Convolutional Neural Networks

Convolutional neural networks (CNNs or ConvNets) is a generic term used to describe any neural network whose architecture includes convolutional layers. Consequently, ConvNets are similar in many ways to the fully-connected feed-forward layers previously discussed, and are almost always used in conjunction with some combination of fully-connected layers. An example architecture can be found in Fig. 5.8. The difference now is that we expect our input data to have some rectangular structure (e.g. two-dimensional images), which are often represented as a three-dimensional volume of shape height \times width \times depth where height and width are, as we would expect, the height and width of the input image, and depth usually refers to the number of *color channels* (i.e. RGB) of the image. For example, if we are dealing with a typical two-dimensional RGB image, then depth = 3, which we could represent as a three-dimensional volume of height \times width \times [red, green, blue] pixels.

Additionally, in contrast to fully-connected networks, ConvNets combine three architectural ideas to ensure some degree of shift and distortion invariance [12]:

- *Local receptive fields*
- *Shared weights* (or weight replication)

Algorithm 1: Backpropagation Algorithm

input :

- A data set $X = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$ consisting of N input-output pairs.
 - The learning rate α .
 - the loss function \mathcal{L} , and the activation function h
-

1. **Forward pass:** For each input-output pair (\mathbf{x}_d, t_d) , store the results y_d , a_j^k , and z_j^k for each node j in layer k by proceeding from the input layer to the output layer m .
2. **Backward pass:** For each input-output pair (\mathbf{x}_d, t_d) , store the results $\partial \mathcal{L}_d / \partial w_{ij}^k$ for each weight w_{ij}^k connecting node i in layer $k - 1$ to node j in layer k by proceeding from layer m , the output layer, to the input layer.
 - Evaluate the error term for the final layer δ_1^m by using Eq. 2.47.
 - Backpropagate the error terms for the hidden layers δ_j^k , working backwards from the final hidden layer $k = m - 1$ by repeatedly using Eq. 2.52.
 - Evaluate the partial derivatives of the individual error \mathcal{L}_d with respect to w_{ij}^k by using Eq. 2.53.
3. **Combine individual gradients:** for each input-output pair $\frac{\partial \mathcal{L}_n}{\partial w_{ij}^k}$ to get the total gradient $\frac{\partial \mathcal{L}(X, \theta)}{\partial w_{ij}^k}$ for the entire set of input-output pairs using:

$$\frac{\partial \mathcal{L}(X, \theta)}{\partial w_{ij}^k} = \frac{1}{N} \sum_{d=1}^N \frac{\partial}{\partial w_{ij}^k} \left(\frac{1}{2} (y_n - t_n)^2 \right) = \frac{1}{N} \sum_{d=1}^N \frac{\partial \mathcal{L}_d}{\partial w_{ij}^k}$$

- Spatial or temporal subsampling (or *pooling*)

2.4.1 Local Receptive Fields

When dealing with high-dimensional inputs such as images, it is impractical to fully-connect each neuron to all neurons in the previous volume because such a network architecture does not take into account the spatial structure of the data. Convolutional networks exploit spatially local correlation by enforcing a sparse local connectivity pattern between neurons of adjacent layers: each neuron is connected to only a small region of the input volume. The spatial extent of this connectivity is a hyperparameter, known as the *local receptive field*. The connections are local in space (along the height and width), but always extend along the entire depth of the input volume.

We begin with an example. Suppose we take as input a 16×16 square of pixels, corresponding to an image. Proceeding as before, we connect the input pixels to a layer of hidden units, with the exception that now we only make connections in small, localized regions of the input image. Explicitly, each neuron in the first hidden layer will be connected to a small region of the input units, say, for example, a 4×4 region, corresponding to 16 input pixels. The region in the input image is known as the *local receptive field* for the hidden unit. This local receptive field can be understood as a little window on the input pixels, with each connection learning a weight. During the forward pass through our network, each filter is convolved across the height and width of the input volume (as illustrated in Fig. 2.2) computing the dot product between the entries of the filter and the input, producing a two-dimensional activation map of that filter. As a result, the network learns filters that activate when it detects some specific type of feature at some spatial position in the input. Stacking these activation maps for all filters along the depth dimension forms the full output volume of the convolution layer. All together, there are three hyperparameters that control the size of the output volume of the convolutional layer: the *depth*, *stride length*, and *zero-padding*. The *depth* of the output volume controls the number of neurons in a layer that connect to the same local receptive field of the input volume. These neurons learn to activate for different features in the input, e.g. various oriented edges or blobs of color. The *stride length*, S (chosen as $S = 1$ in Fig. 2.2), controls how depth columns around the spatial dimensions (height and width) are allocated. This can be understood as the distance by which the local receptive field is moved to construct a new hidden unit. Note that in our example we chose $S = 1$, but it is not uncommon to use different values (typically $S = 1, 2$), depending on the particulars of the problem at hand. Finally, it is sometimes convenient to pad the input with zeros on the border of the input volume. This hyperparameter is known as the *zero-padding*, and denoted by P . The spatial size of the output volume can be computed as a function of the input volume size W , the spatial size of the local receptive field (also known as the *kernel* or

filter size), K , the stride length with which they are applied S , and the amount of zero-padding P used on the border. The equation for calculating how many neurons ‘fit’ inside a given volume is then given by

$$\frac{W - K + 2P}{S} + 1 \quad (2.54)$$

If this number is not an integer, then the strides are incorrect and the neurons cannot be tiled to fit across the input volume in a symmetric manner. In general, setting the zero padding to be $P = (K - 1)/2$ when the stride is $S = 1$ ensures that the input volume and output volume will have the same size spatially.

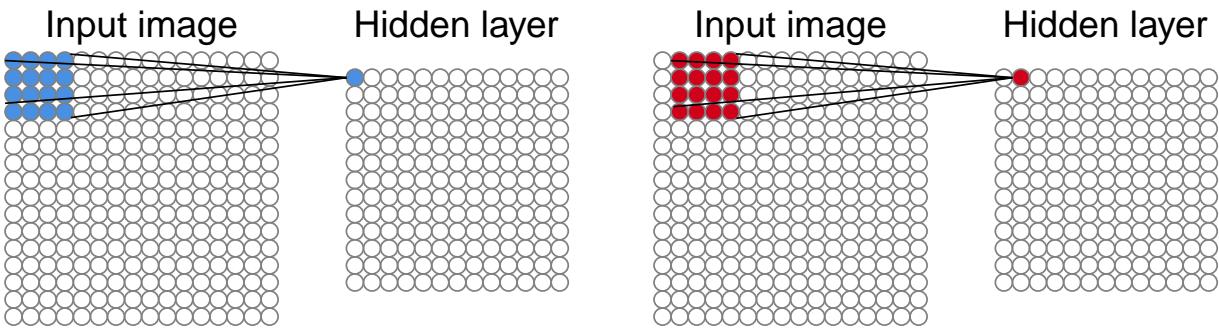


Figure 2.2: Illustration of the local receptive fields in the input image and their corresponding weights in the first hidden layer. Additionally, we can see how these local receptive fields are slid across the input image to construct additional units in the hidden layer.

2.4.2 Shared Weights

A parameter sharing scheme is used in convolutional layers to control the number of free parameters. This scheme relies on the reasonable assumption that if a patch feature is useful to compute at some spatial position, then it should also be useful to compute at other positions. Equivalently, if we denote a single two-dimensional slice of depth as a *depth slice*, we constrain the neurons in each depth slice to use the same weights and biases. Since all neurons in a single depth slice share the same parameters, the forward pass in each depth slice of the convolutional layer can be computed as a convolution of the neuron’s weights with the input volume. Because of this, the sets of weights which are convolved with the input are often referred to as a *filter* or *kernel*. The result of this convolution is then called the *activation map*, and the set of activation maps for each different filter are stacked together along the depth dimension to produce the output volume. This idea of parameter sharing helps contribute to the translation invariance of the CNN architecture [13].

2.4.3 Pooling Layers

Another important concept of CNNs is pooling, which is a form of non-linear down-sampling. This concept is implemented in what we call *pooling layers*, which are usually used immediately following convolutional layers [14]. The pooling layer serves to progressively reduce the spatial size of the representation, to reduce the number of parameters, memory footprint and amount of computation in the network, and hence to also control overfitting [15, 16]. In doing so, the pooling operation provides another form of translation invariance. What these layers do is simplify the information in the output from the convolutional layer by summarizing some region of neurons in the previous layer. The most common form is a pooling layer with filters of size 2×2 applied with a stride of 2 downsamples at every depth slice in the input by 2 along both the height and width, discarding 75% of the activations. One common type of pooling layer is known as *max pooling*, in which a pooling unit simply outputs the maximum activation in the 2×2 input region. This behavior can be seen in Fig. 2.3.

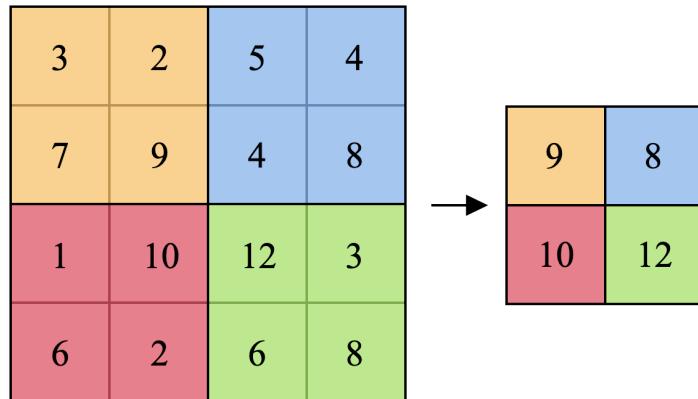


Figure 2.3: Illustration of a 2×2 max pooling layer.

2.4.4 Hyperparameters

Since CNNs use more hyperparameters than typical feed-forward fully-connected neural networks, it is often difficult to determine what values are appropriate for the problem at hand. One idea to keep in mind is that since the size of the feature maps decreases with increasing network depth, layers closer to the input layer will tend to have fewer filters while deeper layers will usually have more. The number of feature maps directly controls the layers' capacity and depends on the number of available examples and the complexity of the problem being studied. It is common to use small filters (e.g. 3×3 or at most 5×5) using a stride of $S = 1$, and padding the input volume with zeros in such a way that the convolutional layer does not alter the spatial

dimensions of the input. Since the pooling layers are in charge of downsampling the spatial dimensions of the input, a common setting is to use max-pooling with 2×2 receptive fields (i.e. $K = 2$), and with a stride of $S = 2$. With this choice of hyperparameters, exactly 75% of the activations in an input volume are discarded (since we downsample by 2 in both height and width).

2.5 Conclusion

In this chapter we have introduced the concept of machine learning and provided a broad overview of some of the topics that will be used in the following chapters. The information provided here is by no means exhaustive, and for further reference I highly suggest the excellent book by Bishop [7]. Moreover, for those interested, there is no shortage of online materials available and a quick google search will provide an endless collection of useful resources. In many ways, we have just barely scratched the surface of machine learning and it would be nearly impossible to include everything in this work. The idea of applying ideas from machine learning to problems in physics has only just recently begun to be explored and it is quickly becoming a very active area of research.

3 | Unsupervised Learning: Ising Worms

Portions of this chapter have appeared in:

“Examples of renormalization group transformations for image sets”, S. Foreman, J. Giedt, Y. Meurice, and J. Unmuth-Yockey, *Phys. Rev. E*, **98**, 052129 – Published 26 November 2018

3.1 Introduction

Machine learning (ML) is a general framework for recognizing patterns in data without detailed human elaboration of the rules for doing so. As an example, a very general function, with many parameters (for example, thousands or millions) can be optimized on a training set, where the desired output is known. The problem is typically nonconvex and plagued by over-fitting problems, and so advanced methods are necessary in order to get reliable answers. One tool that has been exploited is principal component analysis (PCA), which reduces the dimensionality of the data to the most important “directions.” Immediately the practitioner of renormalization group (RG) methods recognizes an analogy, since the RG techniques are also supposed to identify the most important directions in an enlarged space of Hamiltonians. One of the motivations of the present research is to make this analogy more concrete.

A number of papers [17–19] attempt to draw a connection between deep learning and the RG as it appears in physics. However, the analogies between RG flow and depth in a neural network would be strengthened if one could determine conditions under which fixed points can be identified. It would be helpful to show more explicitly how passing from one level to another in a neural network genuinely translates to a renormalization group transformation. There have been steps in the direction of making a full connection. For instance in [18], the principle of *causal influence* is emphasized. That is, when descending in depth, only neighboring nodes should influence the outcome of a lower level node. We have also implemented this in a simple training scheme in earlier work [20]. It can be called “cheap learning” because far fewer variational parameters are involved, due to the constraints of locality. In [19] it is emphasized that deep neural networks outperform shallower networks for reasons which may ultimately be understood in terms of the power of the

renormalization group. Other topics related to machine learning, such as principal component analysis [21] have been previously interpreted in terms of the renormalization group (in this case momentum shells). Machine learning has also been used to identify phase transitions in numerical simulations [22–25]. RG transformations are usually defined in a space of couplings/Hamiltonians, but typically, it is not possible to write down Hamiltonians directly associated with image sets. In this article we propose RG transformations that can be applied to a specific set of images but which could be generalized for other image sets, and can also be understood analytically without any graphical representation. We use the well-studied example of the two-dimensional Ising model on a square lattice. The spin configurations generated with importance sampling provide images with black and white pixels. They have features that can be used to attempt to recognize the temperature used to generate them. However, constructing blocked Hamiltonians in configuration space is a difficult task which involves approximations that are difficult to improve. In other words, it is very difficult to explicitly construct the exact RG transformation mapping the original couplings among the Ising spins into coarse grained ones.

A better control on the RG transformation can be achieved by using the tensor renormalization group (TRG) method [26–32]. The starting point for this reformulation is the character expansion of the Boltzmann weights which is also used in the duality transformation [33]. This leads to an exact expression of the partition function as a sum over closed paths which can be generated with importance sampling using the worm algorithm [34] and then pixelated. These samples will be our sets of images indexed by the temperature used to generate them. The procedure is reviewed in Sec. 3.2.

The goal of a RG analysis is to study systems with large correlation lengths in lattice spacing units and iteratively replace them by coarser ones with a larger effective lattice spacing. This process is useful if we can tune a parameter such as the temperature towards its critical value. Typical image sets such as the MNIST data can be thought as “far from criticality” and the use of RG methods for such a data set may be of limited interest [20]. Criticality may sometimes refer to the choice of parameters used in data analysis [35].

The PCA is a standard method to analyze sets of images. In configuration space, the PCA analysis is identical to the study of the spin-spin correlation matrix. In particular, the largest eigenvalue λ_{\max} is directly connected to the magnetic susceptibility which diverges at criticality [23]. In the loop representation (worms), we will show that λ_{\max} diverges logarithmically at criticality with a constant of proportionality which can be estimated quite precisely ($3/\pi$). This is explained in Sec. 3.3. More generally, it seems reasonable to identify the criticality with the divergence of λ_{\max} .

The advantage of rewriting the high-temperature expansion in terms of tensors is that it allows a very simple blocking (coarse-graining) procedure where a group of sites is replaced by a single site. In the TRG

approach the blocking procedure is local. This leads to simple and exact coarse-graining formulas because we can separate the links into two categories: those links that are inside the blocks and integrated over, and those outside the blocks which are kept fixed and communicate between the blocks [29]. The main goal of this chapter is to relate blocking procedures that can be applied to sets of pixelated images, to approximate TRG transformations. A short summary of the TRG procedure is given in Sec. 3.4.

Having defined criticality, the next step is to define a RG transformation for sets of “legal” loop configurations, also called “worm configurations” later, sampled at various temperatures. In Sec. 3.5, we propose a family of transformations which replaces two parallel links in a block by a single link carrying a specific value x . We call this procedure $1 + 1 \rightarrow x$ hereafter. In the case $1 + 1 \rightarrow 0$, the blocked images follow the same rules (for legal configurations) as the original ones. There is a clear analogy with the 2-state approximation of the TRG method. In the 2-state TRG approximation, the average fraction of occupied links shows a characteristic crossing at a critical point and a collapse when the distance to the critical point is appropriately rescaled at each iteration. The average fraction of occupied links in the blocked worm configurations (with $1 + 1 \rightarrow 0$) shows a somewhat similar behavior in the low temperature phase. However, on the high-temperature side, we observe a “merging” rather than a crossing. In Sec. 3.6, we provide explanations for the similarities and differences between the two procedures.

In Sec. 3.7 we discuss an approximate 2-state TRG method to calculate the average number of bonds through several iterations. The worm configurations can be directly connected to spin configurations using duality [33]: they are the boundary of the positive spin islands. This suggests that the methods discussed here could be applied to generic images. Boundaries of generic grayscale pictures can be defined by converting the picture to black and white pixels. A grayscale picture with gray values between 0 and 1 can be converted into an Ising spin configuration, by introducing a “graycut” below which the value is converted to 0 (spin down) and above which the value is converted to 1 (spin up). It is then possible to construct the boundaries of the spin up domains. This is illustrated in Fig. 3.1. Possible applications are briefly discussed in the Conclusions and illustrated with the CIFAR database in Section 3.9.

3.2 From Loops to Images

In the following we consider the two-dimensional Ising model with spins $\sigma_i = \pm 1$ on a square lattice. The partition function reads

$$Z = \sum_{\{\sigma_i\}} e^{\beta \sum_{\langle i,j \rangle} \sigma_i \sigma_j}, \quad (3.1)$$

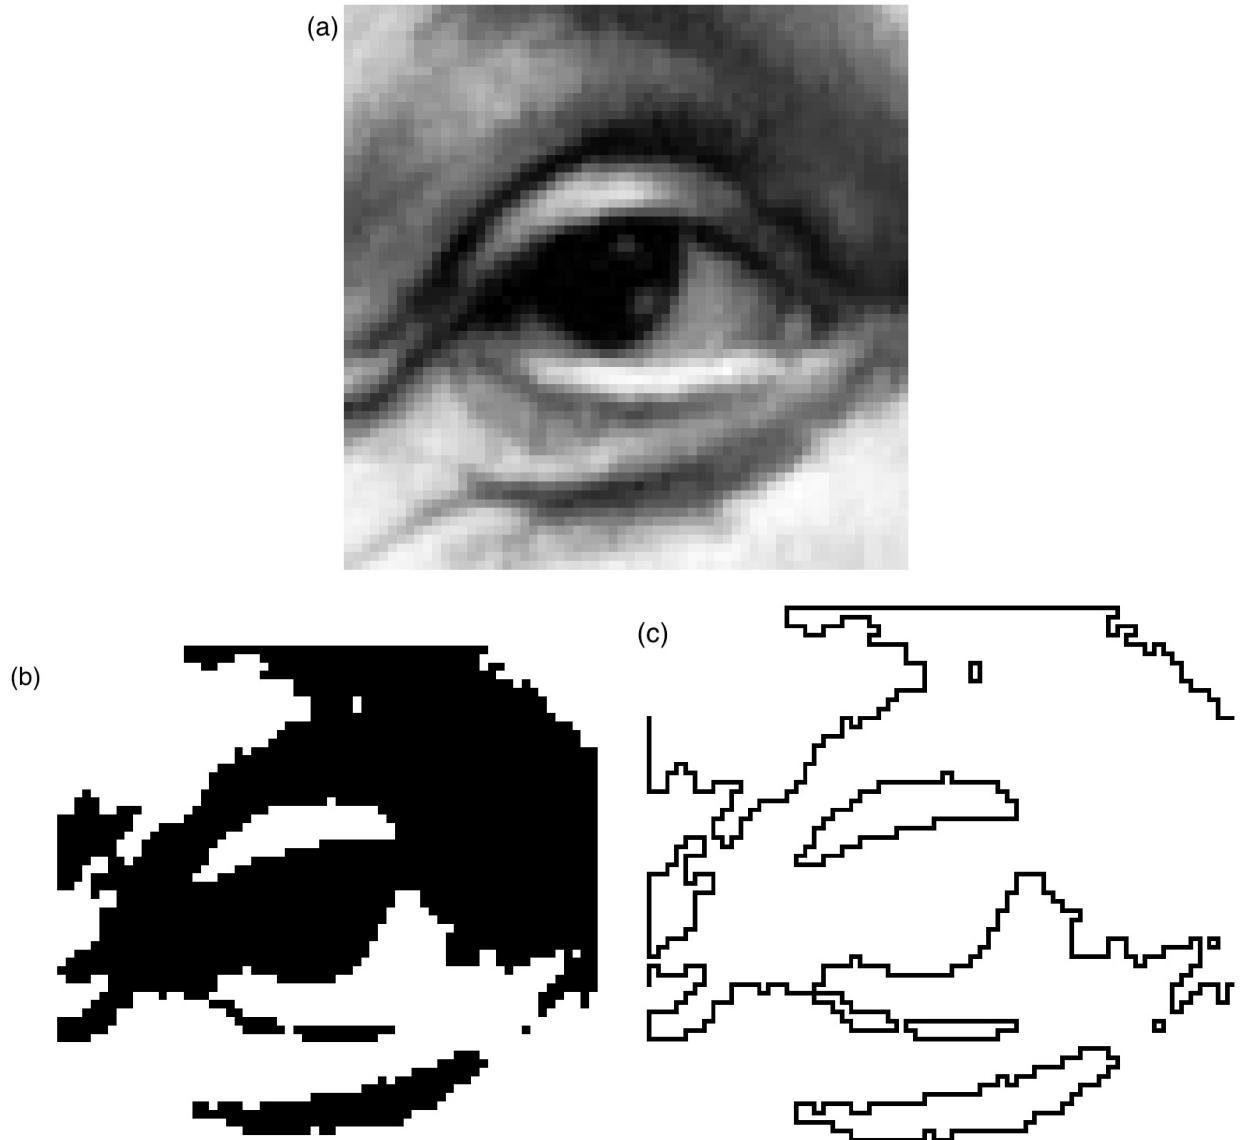


Figure 3.1: (a) Picture of an eye with 4096 pixels; (b) black and white version with a graycut at 0.72; (c) boundaries of the black domains.

where $\langle i, j \rangle$ denotes nearest neighbor sites on the square lattice. In some occasions we will use the notation $T = 1/\beta$ for the temperature. The partition function can be rewritten by using the character expansion [33]

$$\exp(\beta\sigma) = \cosh(\beta) + \sigma \sinh(\beta), \quad (3.2)$$

and integrating over the spins. Factoring out the $\cosh(\beta)$, each link can carry a weight 1 when unoccupied or $t \equiv \tanh(\beta)$ when occupied. The integration over the spins guarantees that an even number of occupied links is coming out of each site [33]. The set of occupied links then form a “legal graph” with N_b occupied links. The partition function can then be written as sum over such legal graphs. If $\mathcal{N}(N_b)$ denotes the number of legal graphs with N_b links we can write:

$$Z = 2^V (\cosh(\beta))^{2V} \sum_{N_b} t^{N_b} \mathcal{N}(N_b). \quad (3.3)$$

Using the fact that $\tanh(\beta) = \exp(-2\tilde{\beta})$, with $\tilde{\beta}$ the inverse dual temperature, Eq. (3.3) has the same form as a spectral decomposition using a density of states and a Boltzmann weight (with $2N_b$ playing the role of the energy). Details of this reformulation can be found in Section 3.8.1.

As shown in Section 3.8.2, we can use derivatives of the logarithm of the partition function to relate $\langle N_b \rangle$ to the average energy, and the bond number fluctuations,

$$\langle \Delta_{N_b}^2 \rangle \equiv \langle (N_b - \langle N_b \rangle)^2 \rangle, \quad (3.4)$$

to the specific heat per site. From the logarithmic singularity of the specific heat we find that

$$\langle \Delta_{N_b}^2 \rangle / V = -\frac{2}{\pi} \ln(|T - T_c|) + \text{regular}. \quad (3.5)$$

In the following we use interchangeably the “bond” terminology, for instance in N_b as in [34] and the link terminology more common in the lattice gauge theory context. In all our numerical simulations we use periodic boundary conditions which guarantees translation invariance.

We will show in Sec. 3.4 that the new form of the partition function in Eq. (3.3) can also be written in an equivalent way as a sum of products of tensors with four indices contracted along the links of the lattice.

The contributions to Eq. (3.3) can be sampled using a worm algorithm [34] outlined in Section 3.8.3. Using this algorithm, we generated multiple configurations at each temperature ($N_{\text{configs}} \approx 10,000$) which are then used for averaging. For example, we can calculate the average number of occupied bonds at a particular

temperature by averaging over all configurations.

Using a legal graph (worm configuration), we can construct an image by introducing a lattice of $2L \times 2L$ pixels with a size of one half lattice spacing. One quarter of these pixels are attached to the sites, one quarter to the horizontal links and one quarter to the vertical links. The remaining quarter are in the middle of the plaquettes and always white. In this representation, each site, link, and plaquette are designated an individual pixel, where occupied links and their respective endpoints are colored black. An example of this representation is shown in Fig. 3.2. We can then flatten each of these images into a vector $\mathbf{v} \in \mathbb{R}^{4V}$, with

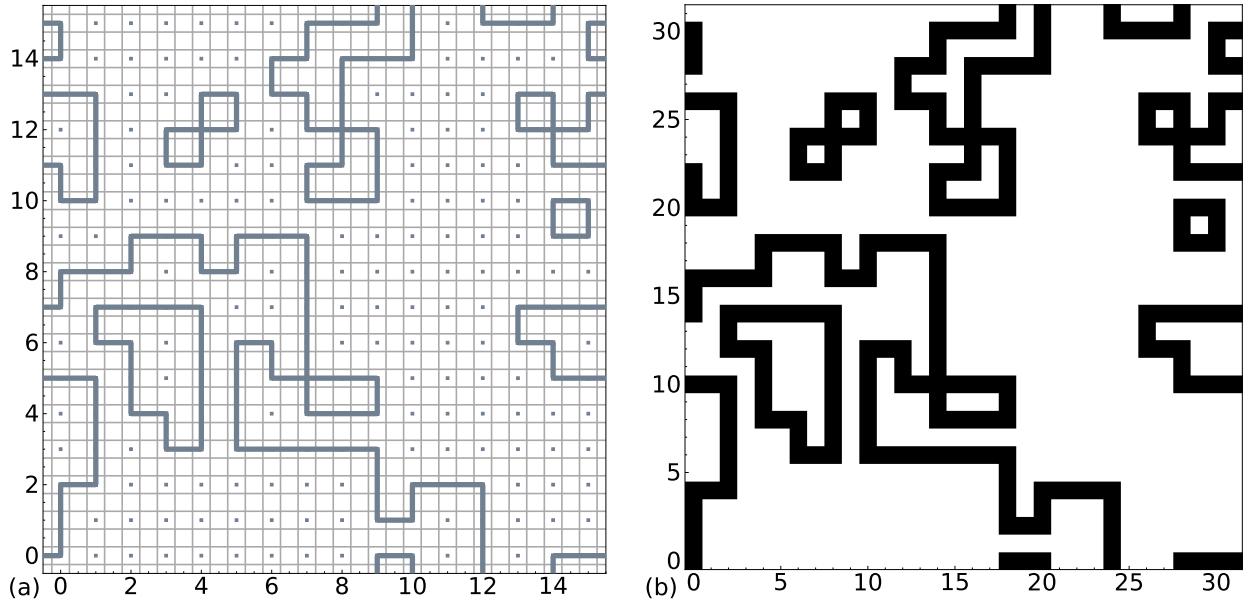


Figure 3.2: (a) Legal worm configuration on an $L \times L$ lattice with periodic boundary conditions and; (b) its equivalent representation as a $2L \times 2L$ black and white pixel image.

$v_i \in \{0, 1\}$. This allows us to write the number of occupied bonds, in a single configuration, N_b as

$$\sum_{j=bonds} v_j = N_b \quad (3.6)$$

3.3 PCA and Criticality

Having now sets of images for a range of temperatures, we can apply PCA [7]. PCA isolates the “most relevant” directions in the dataset. PCA is simply the computation of the eigenvalues λ_α and eigenvectors u_α of the

covariance matrix for a dataset with N configurations corresponding to a given temperature $\{\mathbf{v}^n\}_{n=1}^N$:

$$S_{ij} = \frac{1}{N} \sum_{n=1}^N (\mathbf{v}_i^n - \bar{\mathbf{v}}_i)(\mathbf{v}_j^n - \bar{\mathbf{v}}_j). \quad (3.7)$$

In this equation, each sample \mathbf{v}_j is a vector in \mathbb{R}^{4V} , labeled by the indices $i, j = 1, \dots, 4V$. The PCA extracts solutions to

$$Su_\alpha = \lambda_\alpha u_\alpha \quad (3.8)$$

and orders them, in descending magnitude of λ_α , which are all non-negative. The usefulness of PCA is that one can approximate the data (see for instance the discussion in [7]) by the first M principal components.

Illustrations of the PCA for the MNIST data can be found in Sec. 4 of Ref. [20], where we show the eigenvectors corresponding to the largest eigenvalues and the approximation of the data by subspaces of the largest eigenvalues of dimensions 10, 20 etc.

It should be noted that the PCA is an analysis that can be performed for each temperature separately and not obviously connected to the closeness to criticality. However, we were able to find a relation between the largest PCA eigenvalue denoted λ_{\max} and the logarithmic divergence of the specific heat, namely

$$\lambda_{\max} \simeq \frac{3}{2} \langle \Delta_{N_b}^2 \rangle / V \simeq -\frac{3}{\pi} \ln(|T - T_c|). \quad (3.9)$$

This property was found by an approximate reasoning shown in Section 3.8.2 and relies on two assumptions. The first one is that the eigenvector associated with λ_{\max} is proportional to $\langle \mathbf{v} \rangle$ which is invariant under translation by two pixels in either direction. The second assumption is that in good approximation we can neglect the contributions from sites that are visited twice (four occupied links coming out of one site). Numerically, only 4% of sites are visited twice near the critical temperature which justifies the second assumption. Figure 3.3 provides an independent confirmation of the approximate validity of Eq. (3.9).

3.4 TRG Coarse-graining

So far we have sampled the legal graphs of the high temperature expansion of the Ising model using the worm algorithm. An alternative approach is to use a tractable real-space renormalization group method known as the TRG [28–32].

In order to understand what we want to accomplish by blocking the loop configuration, it is useful to first understand the evolution of a tensor element using the TRG method.

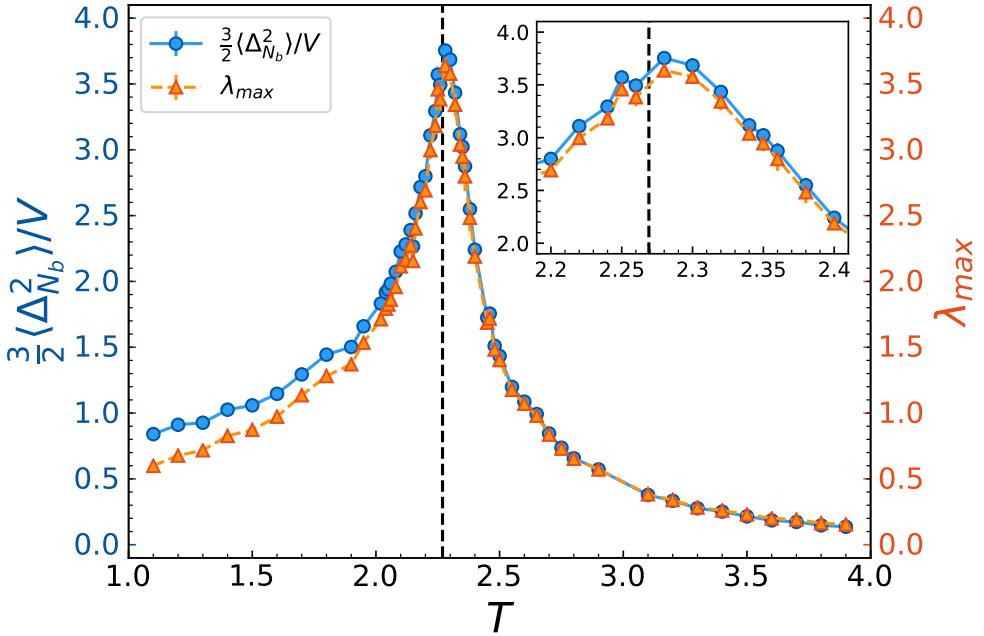


Figure 3.3: λ_{\max} and $\frac{3}{2} \langle \Delta_{N_b}^2 \rangle / V$ (per unit volume) vs. T , illustrating the relation between the eigenvalue corresponding to the first principal component and the logarithmic divergence of the specific heat. The inset shows a qualitative agreement near the critical temperature.

The tensor formulation used here connects easily with the worm formulation used in this paper. After the character expansion has been carried out, one is left with new integer variables on the links of the lattice with constraints on the sites which guarantee the sum of the link variables associated with that site is even. Therefore we build a tensor using this constraint and the surrounding link weights. The tensor has the form

$$T_{xx'yy'}^{(i)}(\beta) = [\tanh(\beta)]^{(n_x+n_{x'}+n_y+n_{y'})/2} \times \delta_{n_x+n_{x'}+n_y+n_{y'} \text{ mod } 2, 0}. \quad (3.10)$$

Here the notation being used is that this tensor is located at the i^{th} site of the lattice, $n_{\hat{i}}$ is the integer variable, taking value 0 or 1, on an adjacent link, and the Kronecker delta, $\delta_{i,j}$ is understood to be satisfied if the sum is even. By contracting these tensors together in the pattern of the lattice one recreates the closed-loop paths generated by the high-temperature expansion and exactly match those paths which are sampled by the worm algorithm.

Using these tensors one can write a partition function for the Ising model that is exactly equal to the original partition function,

$$Z = 2^V (\cosh(\beta))^{2V} \text{Tr} \prod_i T_{xx'yy'}^{(i)} \quad (3.11)$$

where Tr means contractions (sums over 0 and 1) over the links.

The most important aspect of this reformulation is that it can be coarse-grained efficiently. The process is illustrated in Fig. 3.4 where four fundamental tensors have been contracted to form a new “blocked” tensor. This new tensor has a squared number of degrees of freedom for each new effective index. The partition function can be written exactly as

$$Z = 2^V (\cosh(\beta))^{\frac{1}{2}V} \text{Tr} \prod_{2i} T'_{XX'YY'}^{(2i)},$$

where $2i$ denotes the sites of the coarser lattice with twice the lattice spacing of the original lattice. In

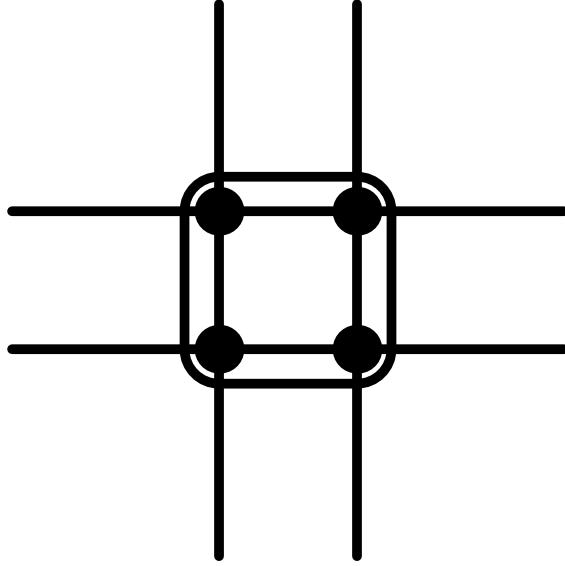


Figure 3.4: Illustration of the tensor blocking discussed in the text. Each dot is a tensor at a lattice site with four lines coming out, each representing a tensor index. Lines connecting dots represent tensor contractions.

practice, this exact procedure cannot be repeated indefinitely and truncations are necessary. This can be accomplished by projecting the product states into a smaller number of states that optimizes the closeness to the exact answer. A two-state projection is discussed in [29] and will be followed hereafter. Note that in this procedure, T_{0000} is factored out and the final expression for the other blocked tensors are given in these units. For definiteness we consider T_{1100} which in the microscopic formulation is the weight associated with a horizontal line in a loop configuration. By looking at the fixed point equation [29], one can see that there is a high temperature fixed point where all the tensor elements except for T_{0000} are zero and a low temperature fixed point where all the tensor elements are one. In between these two limits, there is a non-trivial fixed point illustrated by the crossing of iterated values of T_{1100} in Fig. 3.5. Note that because of the two-state approximation, the critical temperature T_c is slightly higher than the exact one [29]. To be completely specific, the exact T_c for the original model is $2/\ln(1 + \sqrt{2}) = 2.269\dots$ while for the two state

projection with the second projection procedure of Ref. [29], it is $1/0.3948678 = 2.53249\dots$ It is easy to

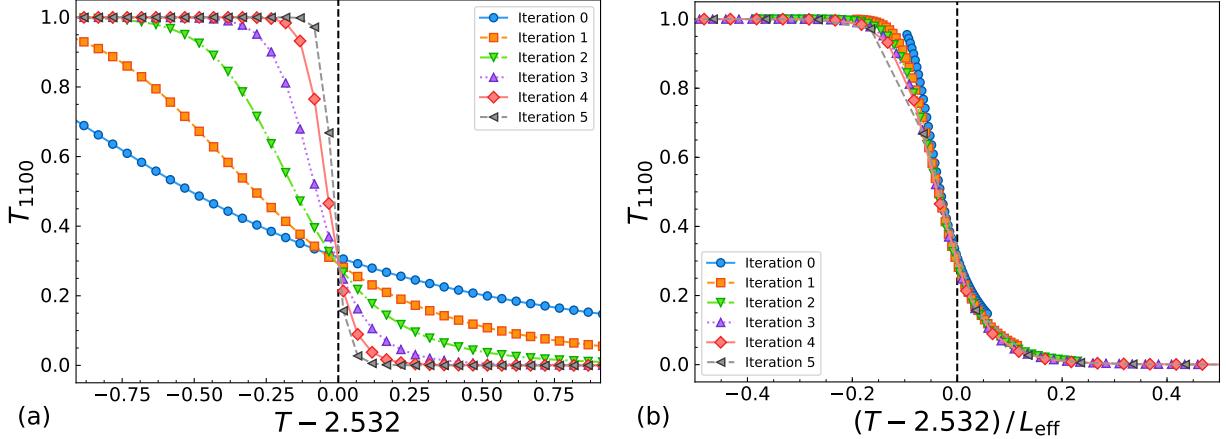


Figure 3.5: (a) T_{1100} vs. $T - T_c^{(2s)}$ for six successive iterations of the blocking transformation, beginning with an initial lattice $L = 64$; (b) T_{1100} vs. $(T - T_c^{(2s)})/L_{\text{eff}}$ illustrating the data collapse, where $T_c^{(2s)}$ is the critical temperature of the two state projection, beginning at iteration 0 on an $L = 64$ lattice.

relate the properties of the iterated curves near the non-trivial fixed point using the linear RG approximation. Below we just state the results, for details and references see [29]. With the blocking procedure used, the scale factor is $b = 2$. The eigenvalue in the relevant direction is $\lambda = b^{1/\nu} = 2$ since $\nu = 1$. In Fig. 3.5 (a), one can see that the height

$$\delta T_{1100} \equiv T_{1100} - T_{1100}^* \quad (3.12)$$

(where T_{1100}^* is the value at the fixed point), nearly doubles each time the blocking procedure is performed, making the slope twice as large each time. A nice data collapse can be reached by offsetting this effect by rescaling the horizontal axis each iteration by $\lambda=2$ as shown in Fig. 3.5 (b). In numerical calculations, we start with a finite L (64 in Fig. 3.5) and then after ℓ iterations, we are left with an effective size $L_{\text{eff}} = L/b^\ell$.

The remainder of this chapter will be dedicated towards obtaining data collapse for $\langle N_b \rangle$ calculated with successive blockings.

3.5 Image Coarse-graining

In an attempt to explicitly connect the ideas from RG theory to similar concepts in machine learning, we will implement a coarse-graining procedure directly on the images but in a way inspired by the TRG construction of Sec. 3.4. The construction relies on visual intuition and will be reanalyzed in the TRG context in Sec. 3.7.

As in the TRG coarse-graining procedure, the image is first divided up into blocks of 2×2 squares, as shown in Fig. 3.6. Each of these 2×2 squares are then replaced, or “blocked”, by a single site with bonds

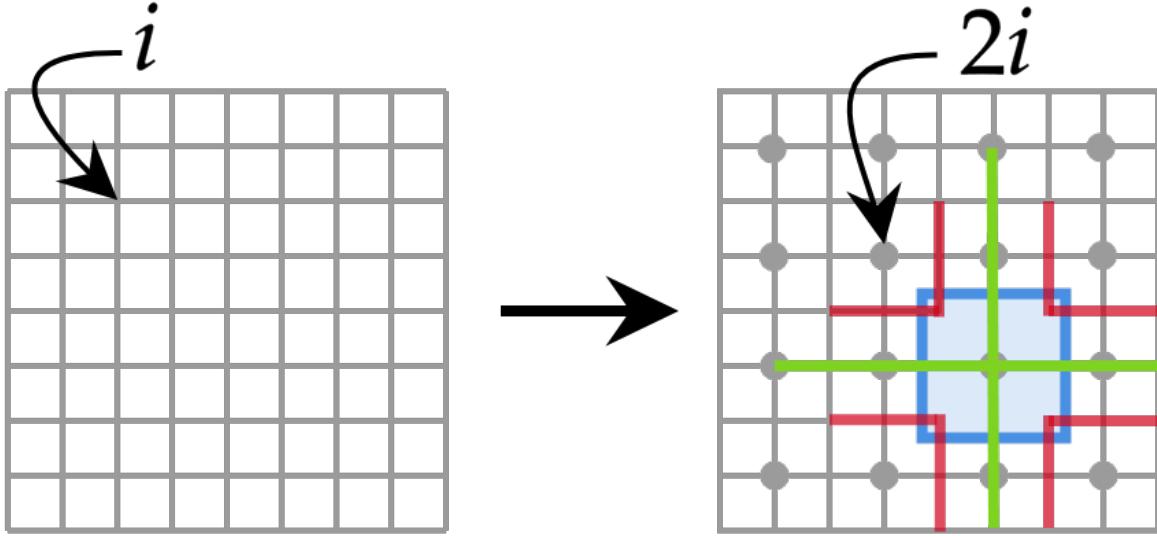


Figure 3.6: Illustration of the coarse-graining procedure in which the original lattice sites (i) are replaced by blocked sites ($2i$) (grey circles) with twice the original lattice spacing. In the coarse-grained lattice, an elementary block (blue) consists of four sites on the original lattice, eight external bonds (red), and four blocked external bonds (green).

determined by the number of occupied external bonds in the original square. In doing so, we reduce the size of each linear dimension by a factor of two, resulting in a new blocked configuration whose volume is one-quarter the original. In particular, if a given block has exactly one external bond in a given direction, the blocked site retains this bond in the blocked configuration. This seems to be a natural choice. However, if a given block has exactly two external bonds in a given direction, we can consider several options. The simplest option is to neglect the double bond entirely, and we denote this blocking scheme by “ $1 + 1 \rightarrow 0$ ”. This approach respects the selection rule (conservation modulo 2) and has the advantage of maintaining the closed-path restriction. In other words with the $1 + 1 \rightarrow 0$ option, the blocked image corresponds to a legal graph and the procedure can be iterated without introducing new parameters. This procedure is illustrated for a specific configuration on a 16×16 lattice in Fig. 3.7. Alternatively, we can include this double bond in the blocked configuration, and give it some weight m between 0 and 2. The examples of $m = 1$ and 2 are denoted “ $1 + 1 \rightarrow 1$ ”, and “ $1 + 1 \rightarrow 2$ ” respectively and are shown in Fig. 3.16. This blocking procedure introduces new elements and iterations require more involved procedures. This is not discussed hereafter.

3.6 Partial Data Collapse for Blocked Images

In this section, we study the properties of $\langle N_b \rangle$ obtained for successive blockings with the $1 + 1 \rightarrow 0$ rule starting with configurations on a 64×64 lattice. A first observation is that the $1 + 1 \rightarrow 0$ blocking preserves

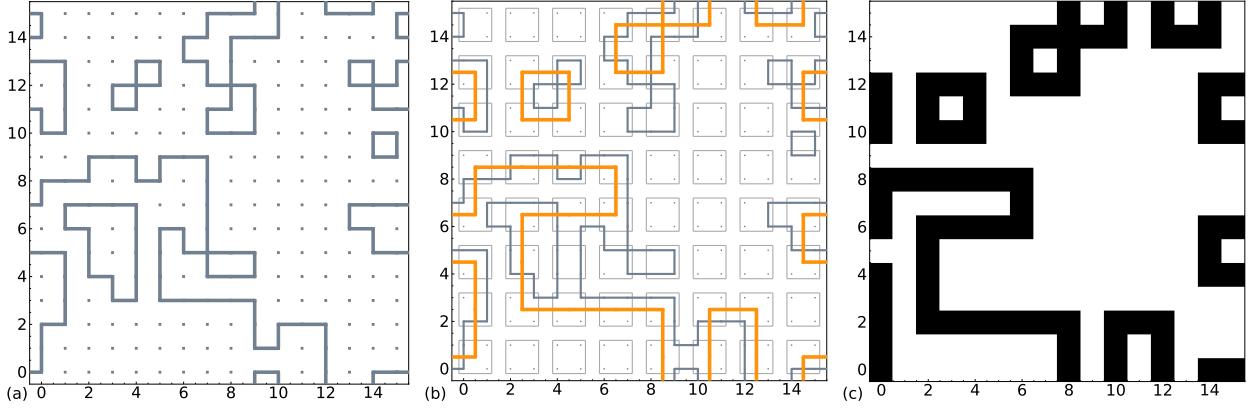


Figure 3.7: (a) Illustration of the $1 + 1 \rightarrow 0$ blocking procedure discussed in the text: original configuration; (b) introduction of the blocks and replacement of single or double bounds according to the $1 + 1 \rightarrow 0$ rule; (c) construction of the corresponding blocked image.

the location of the peak of the fluctuations $\langle \Delta_{N_b}^2 \rangle$. In addition it is possible to stabilize this quantity for a few iterations by dividing by $V_{\text{eff}} \ln(L_{\text{eff}})$. This is illustrated in Fig. 3.8. However, a very different scaling appears for the last two iterations which may be due to the very short effective sizes (4 and 2). This indicates the last two iterations are very different from the previous ones. We now consider $\langle N_b \rangle$ for successive iterations. The results are shown in Fig. 3.9. We see that in the low temperature side, the curves sharpen in a way similar to T_{1100} in Fig. 3.5. However on the high temperature side, we observe a merging rather than a crossing. This can be explained as follows. In the high T regime occasionally a single loop, the size of a plaquette, forms. This is due to other configurations being highly suppressed. With the $1 + 1 \rightarrow 0$ rule, one out of four possible plaquettes becomes a larger plaquette which exactly compensates the change in V_{eff} which is also reduced by a factor of four. There are four kinds of plaquettes (see Fig. 3.7): those inside the blocks (they disappear after blocking), those between two neighboring blocks in the vertical or horizontal direction (these are double links between the blocks and so they disappear with the $1 + 1 \rightarrow 0$ rule), and finally those which share a corner with four blocks (they generate a larger plaquette). This last type can be seen at coordinate (4, 12) in Fig. 3.7. We now attempt to obtain data collapse for $\langle N_b \rangle / V_{\text{eff}}$ by performing a rescaling of the temperature axis with respect to the critical value as in Fig. 3.5. After this rescaling by a factor 2 at each iteration, we observe a reasonable collapse on the low-temperature side. On the high temperature side, since the unrescaled curves merge, the rescaling splits them and there is no collapse on that side. This is illustrated in Fig. 3.9.

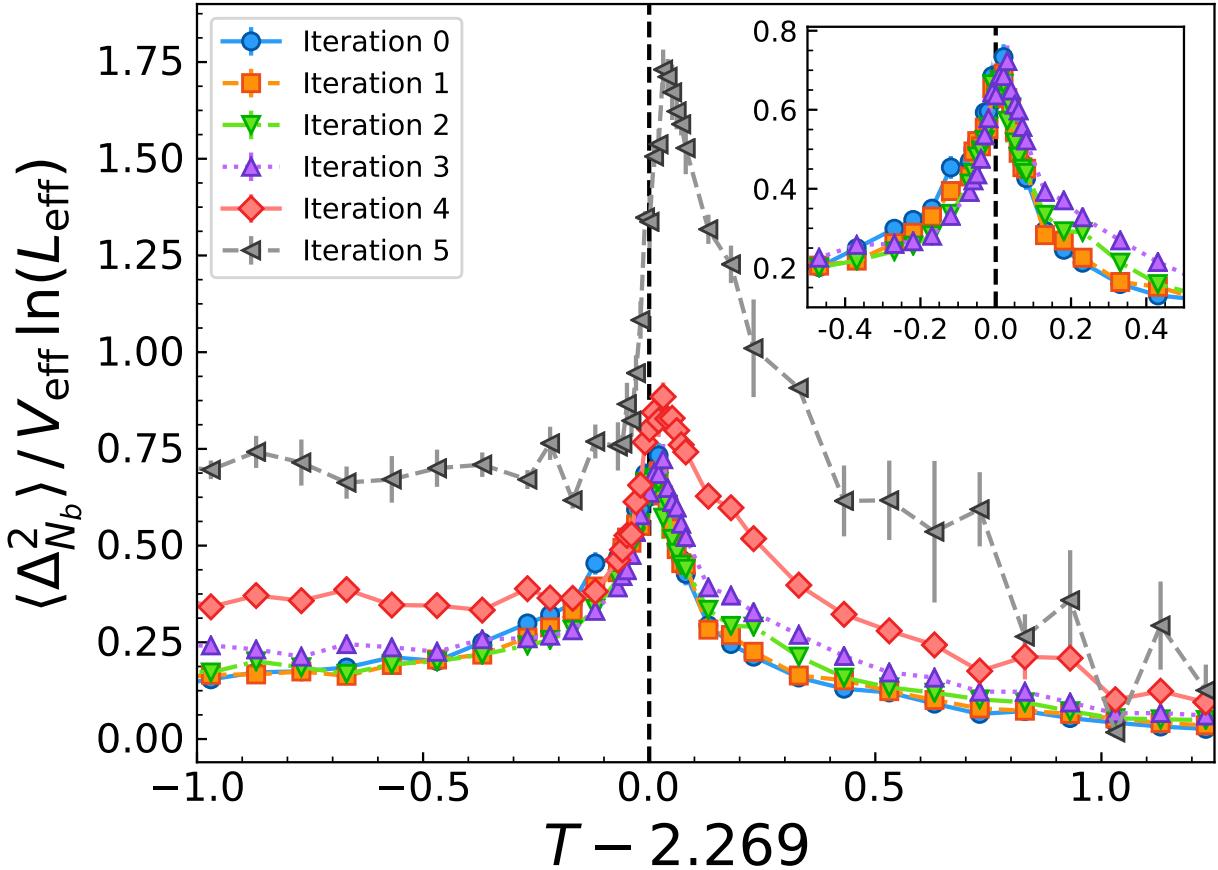


Figure 3.8: Fluctuations in the average number of bonds $\langle \Delta_{N_b}^2 \rangle$ vs. temperature T under iterated blocking steps beginning with an initial lattice size of $L = 64$. The results are scaled by $1/V_{\text{eff}} \log(L_{\text{eff}})$ in order to demonstrate the data collapse near the critical temperature. This collapse is especially apparent in the inset, which shows the results under the first three blocking steps, with $L_{\text{eff}} = 64, 32, 16$, and 8.

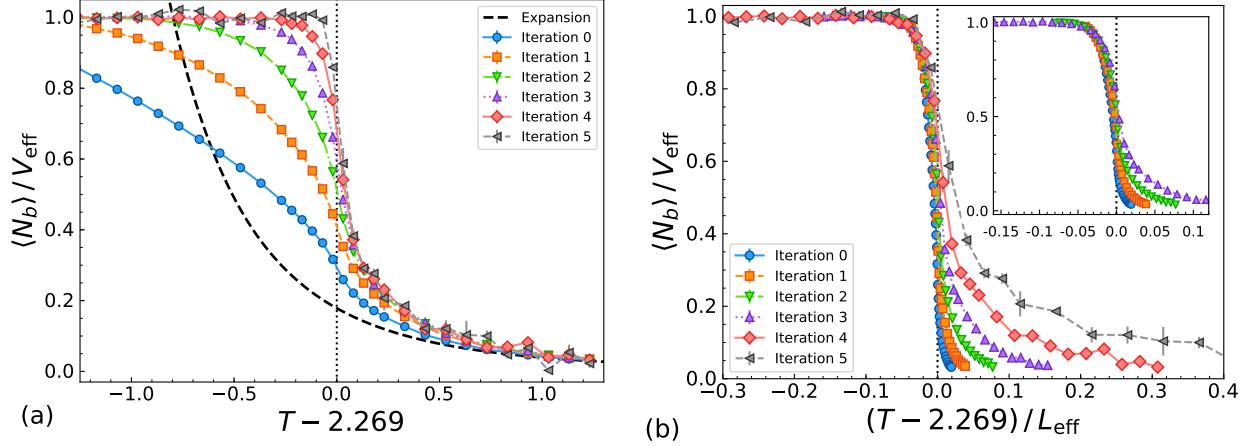


Figure 3.9: (a) Average number of bonds $\langle N_b \rangle$ vs. temperature T under iterated blocking steps beginning with an initial lattice size of $L = 64$. The dashed black line illustrates the high temperature expansion, showing that the dominant configurations are those consisting of small, isolated plaquettes. (b) Average number of bonds $\langle N_b \rangle$ vs. the rescaled temperature $(T - 2.269)/L_{\text{eff}}$ under successive blocking steps. Iteration 0 represents the original lattice before blocking, with $L_{\text{eff}} = 64$.

3.7 TRG Calculation of $\langle N_b \rangle$

Using the tensor method we were able to calculate $\langle N_b \rangle$ to compare with the worm algorithm. Consider the equation for $\langle N_b \rangle$ with $N_b = \sum_l n_l$ the sum over bond numbers at every link:

$$\langle N_b \rangle = \frac{1}{Z} \sum_{\{n\}} \left(\sum_l n_l \right) \left(\prod_l \tanh^{n_l}(\beta) \right) \left(\prod_i \delta_{n_x + n_{x'} + n_y + n_{y'} \bmod 2, 0}^{(i)} \right). \quad (3.13)$$

This expression can be seen as $\langle N_b \rangle = \sum_l \langle n_l \rangle$, and because of translation and 90° rotational invariance, all $\langle n_l \rangle$ are equal. Thus, it is enough to calculate $\langle n_l \rangle$ for one particular link (just call it $\langle n \rangle$) and multiply it by $2V$: $\langle N_b \rangle = 2V\langle n \rangle$.

To calculate $\langle n \rangle$, it amounts to associating an n with one particular link on the lattice. This alters *two* tensors on the lattice such that the two tensors which contain that link as indices are now defined as

$$\tilde{T}_{n_x n_{x'} n_y n_{y'}}^{(1)} = \sqrt{n_x} T_{n_x n_{x'} n_y n_{y'}}(\beta), \quad (3.14)$$

$$\tilde{T}_{n_x n_{x'} n_y n_{y'}}^{(2)} = \sqrt{n_{x'}} T_{n_x n_{x'} n_y n_{y'}}(\beta), \quad (3.15)$$

where x and x' were chosen without loss of generality. It could just as well have been chosen as y and y' . One can see that when these two tensors are contracted along their shared link, the product picks up a factor of n for that link, which when combined with the other tensors in the lattice, and divided by Z , yields $\langle n \rangle$.

Knowing the above, one is free to block and construct the partition function, Z , and $\langle n \rangle$. This can be done by blocking symmetrically in both directions, or by constructing a transfer matrix by contracting only along a time-slice (i.e. a snapshot of the system at fixed t). This is shown in Fig. 3.10. In practice contracting to build a transfer matrix is optimum since one direction of the lattice is never renormalized and allows the easy calculation of $\langle n \rangle$. What was just described is a method to calculate $\langle n \rangle$ for the original, fine

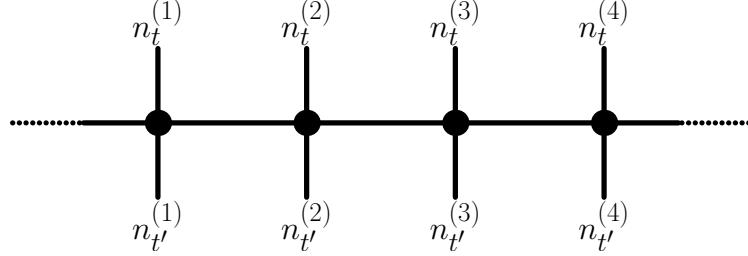


Figure 3.10: A pictorial representation of the transfer matrix made by contracting a fundamental tensor along a single time-slice.

lattice. However, one can also calculate the same quantity for a coarse lattice. The actual blocking method is essentially identical, with a small difference. Instead of contracting the fundamental tensor to the desired lattice size, one contracts a blocked tensor to the desired lattice size.

For example, if one wanted to calculate $\langle N_b \rangle$ for a 32×32 lattice, one could contract the fundamental tensor along a time slice with itself five times. This would give a $2^{32} \times 2^{32}$ transfer matrix which could be used to build the whole partition function. Now, under a single coarse-graining step the 32×32 lattice becomes a 16×16 lattice of blocks. Therefore, to build this, one could contract four fundamental tensors in a block and consider this a new, effective fundamental tensor. This is shown in Fig. 3.4. Then one repeats the same steps to construct the transfer matrix, however only contracting four times with itself to create a matrix representing 16 lattice sites of the blocked tensor.

To actually calculate $\langle n \rangle$ by building the transfer matrix, one can take the final tensor, prior to contracting the dangling spatial indices, and multiply by \sqrt{n} against the indices n_x and $n_{x'}$. This is shown for the unblocked case in Fig. 3.11, however the procedure is identical for the blocked case once the blocked tensor has been constructed. This is also the point where one can choose the level of approximation one will use in the blocking. For instance one could choose that the state $|1\ 1\rangle \rightarrow |0\rangle$ and assign $n = 0$ to that state. Alternatively one could preserve N_b and let $|1\ 1\rangle \rightarrow |2\rangle$ and assign $n = 2$ to that state. This procedure was found to agree with the results obtained by changing the pixels of the worm configurations.

Once the original transfer matrix has been constructed, as well as the matrix with the insertion of n

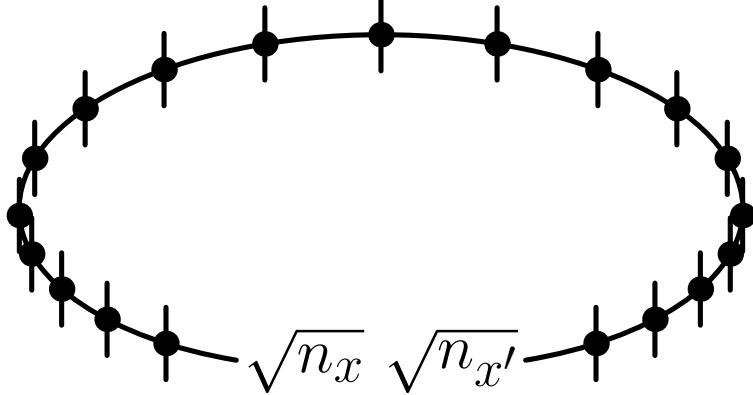


Figure 3.11: By multiplying the remaining free spatial indices by \sqrt{n} and contracting for periodic boundary conditions in space we form an “impure” transfer matrix. Combining the resultant matrix with the original transfer matrix allows one to calculate $\langle n \rangle$.

along a single link, one can combine these to find $\langle n \rangle$. This is done by simple matrix multiplication:

$$\langle n \rangle = \frac{1}{Z} \text{Tr}[\underbrace{\mathbb{T} \cdots \mathbb{T}'}_{N_\tau} \cdots \mathbb{T}], \quad (3.16)$$

with

$$Z = \text{Tr}[\mathbb{T}^{N_\tau}]. \quad (3.17)$$

Here \mathbb{T} represents the transfer matrix built by contracting tensors along a time slice, and \mathbb{T}' represents the single (“impure”) transfer matrix at a time-slice with a single bond multiplied by n . Since the lattice has Euclidean temporal extent, $L = N_\tau$, there are that many matrices multiplied in each case. The values of $\langle N_b \rangle / V_{\text{eff}}$ obtained with this procedure are shown in Fig. 3.12. The rescaling by 2 at each iteration provides a good data collapse on both sides of the transition.

3.8 Technical Results

3.8.1 Loop Representation

We can rewrite the Ising partition function in terms of bonds between neighboring sites $\langle i, j \rangle$. The allowed bond configurations are concisely described by concepts in graph theory, because they form edges (bonds) between neighboring vertices (sites). Making use of well-known identities allows for the partition function to

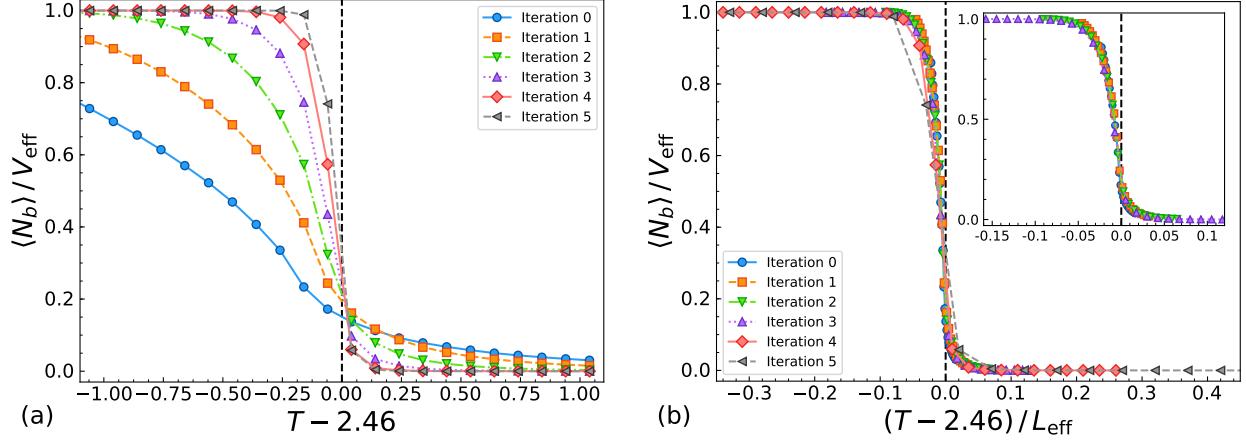


Figure 3.12: (a) $\langle N_b \rangle$ vs $(T - 2.46)$ under successive blocking steps calculated using 2-state HOTRG. (b) $\langle N_b \rangle$ vs $(T - 2.46)/L_{\text{eff}}$ under successive blocking steps calculated using 2-state HOTRG. Note that the value of 2.46 was determined qualitatively by choosing the value which gave the best resulting data collapse.

be written in the following high-temperature expansion:

$$Z = 2^{|V|} \cosh^{|E|} \beta \sum_{\Gamma \in \mathcal{C}(G)} \tanh^{|\Gamma|}(\beta) \quad (3.18)$$

$$= 2^{|V|} \cosh^{|E|} \beta \sum_{|\Gamma|} n(|\Gamma|) \tanh^{|\Gamma|}(\beta) \quad (3.19)$$

The notation is as follows. We have a graph $G = (V, E)$ that describes our lattice, where V are the vertices and E are the edges, which are the bonds between neighboring sites. If we restrict ourselves to subgraphs with only occupied bonds allowed by the Ising model, then the degree of each vertex is even. This is the number of bonds emanating from a particular vertex. The set of edges of such a subgraph is described as being “Eulerian.” The space of all such sets of edges is known as the cycle space $\mathcal{C}(G)$. The notation $|V|$, $|E|$, $|\Gamma|$ denotes the number of elements in each set (cardinality). In the second line, $n(|\Gamma|)$ counts the multiplicity of subgraphs of cardinality $|\Gamma|$, and is zero when $|\Gamma|$ does not correspond to a “legal” subgraph.

We now specialize the presentation to the case of the two-dimensional Ising model on a square lattice with periodic boundary conditions. In this case $|V|$ is $V = L^2$, the volume that we express in lattice units, and $|E| = 2V$. We introduce the notation $t \equiv \tanh(\beta)$ and we call N_b the number of bonds in a graph (values taken by $|\Gamma|$). With these notations we recover Eq. 3.3.

It is this bond formulation that is the basis of both random cluster algorithms [36] and worm algorithms [34]. In this paper we utilize the latter. Both of these classes of algorithms have the benefit of significantly avoiding critical slowing down. This is essential near the critical temperature T_c .

3.8.2 Heat Capacity

One striking feature of the second order transition for the two-dimensional Ising model is the logarithmic divergence of the specific heat density at the critical temperature T_c . In this section, we review the way the specific heat can be calculated with the worm algorithm and we check our answer with the exact finite volume expressions [37].

Using the standard thermodynamical formula for the average energy

$$\langle E \rangle = -\frac{\partial}{\partial \beta} \ln Z, \quad (3.20)$$

with the expression Eq. (3.3) of Z , we get

$$\langle E \rangle = -\tanh(\beta) \left(2V + \frac{\langle N_b \rangle}{\sinh^2(\beta)} \right), \quad (3.21)$$

where we define

$$\langle f(N_b) \rangle \equiv \sum_{N_b} f(N_b) t^{N_b} n(N_b) / \sum_{N_b} t^{N_b} n(N_b). \quad (3.22)$$

We can then use

$$C_V = \frac{\partial \langle E \rangle}{\partial T}, \quad (3.23)$$

to write

$$\frac{C_V}{V} = \beta^2 \left[\frac{2}{\cosh^2(\beta)} - \frac{4 \cosh(2\beta)}{\sinh(2\beta)} \frac{\langle N_b \rangle}{V} + \left(\frac{2}{\sinh(2\beta)} \right)^2 \frac{\langle (N_b - \langle N_b \rangle)^2 \rangle}{V} \right]. \quad (3.24)$$

Since $\frac{\langle N_b \rangle}{V} \leq 2$ (in two dimensions), the only possibly divergent part is the variance of N_b per unit volume $\langle \Delta_{N_b}^2 \rangle$ defined in Eq. (3.4). The singularity near T_c is known from Onsager's solution:

$$\frac{C_V}{V} = -\frac{2}{\pi} \left(\ln(1 + \sqrt{2}) \right)^2 \ln(|T - T_c|) + \text{regular}. \quad (3.25)$$

This implies Eq. (3.5).

3.8.3 Monte Carlo Implementation

We can proceed to sample the closed path configuration space using the worm algorithm [34]. A single Monte Carlo step is outlined below.

1. Randomly select a starting point on the lattice (i_0, j_0) .

2. Propose a move to a neighboring site (i', j') , selected at random.
3. If no link is present between these two sites, a bond is created with acceptance probability $P = \min\{1, \tanh \beta\}$. If the bond is accepted, we update the bond number for the present worm, $n_b = n_b + 1$.
4. If a link already exists between the two sites, it is removed with probability $P = 1$.
5. If $(i', j') = (i_0, j_0)$, i.e. we have a closed path, go to (1.). Otherwise, $(i', j') \neq (i_0, j_0)$, go to (2.)

The number of necessary Monte Carlo steps required to achieve sufficient statistics varies with the lattice size, thermalization time, and temperature. After each step, we calculate the energy in terms of the average number of active bonds N_b , and consider the system to be thermalized when fluctuations between subsequent values of the energy are less than 1×10^{-3} . We then save the resulting configuration, along with the final values for all physical quantities of interest. This process is then repeated many times over a range of different temperatures to generate sufficient statistics for physical observables. All errorbars are calculated using the block jackknife resampling technique.

3.8.4 Tests

The above formulas have been used to calculate C_V . Precise checks were performed by comparing with the exact results obtained from Ref. [37]. The agreement can be seen in Fig. 3.13. Results for other lattice sizes that we have simulated are similar.

3.8.5 Conjecture About λ_{\max}

Using the Monte Carlo algorithm outlined above, we can calculate the average number of occupied bonds at a particular temperature by averaging over all configurations

$$\langle N_b \rangle \equiv \frac{1}{N_{\text{configs}}} \sum_{n=1}^{N_{\text{configs}}} N_b^{(n)} \quad (3.26)$$

$$= \left\langle \sum_{j=\text{bonds}} v_j \right\rangle \quad (3.27)$$

$$= 2V\langle v_b \rangle. \quad (3.28)$$

From this, we have that

$$\langle v_b \rangle = \frac{\langle N_b \rangle}{2V}, \quad (3.29)$$

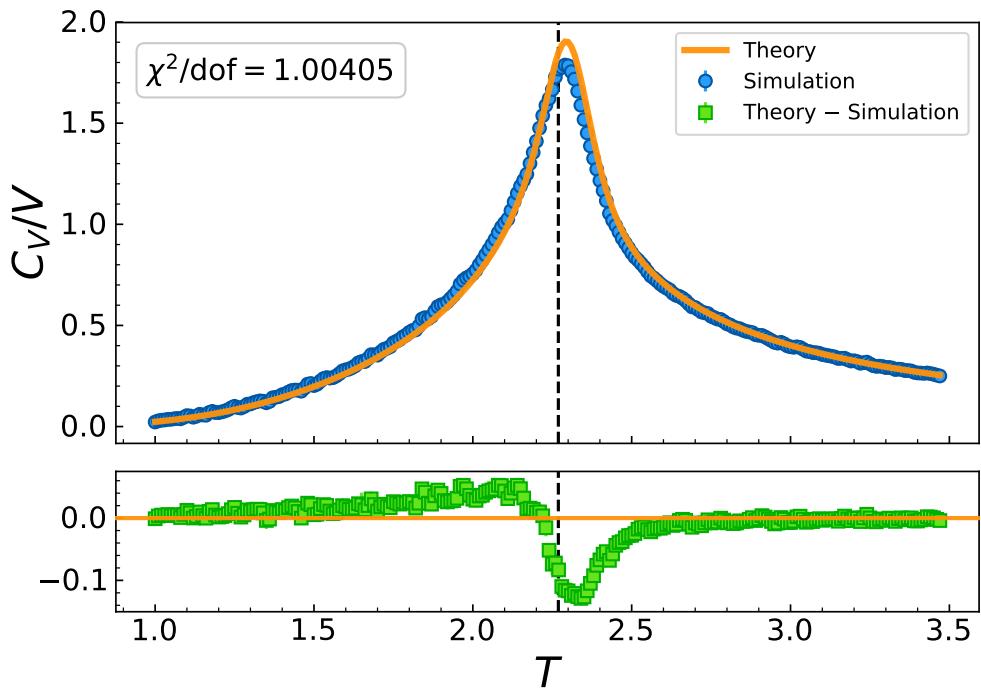


Figure 3.13: Comparison of the worm Monte Carlo computation of the specific heat C_v versus temperature and the exact results using the formula in [37], for an $L = 32$ lattice. Note that χ^2/dof represents the reduced chi-squared statistic. It can be seen that the agreement is excellent, except for a slight deviation at the critical temperature, where Monte Carlo algorithms tend to face difficulties with critical slowing down. This is mostly addressed with the worm algorithm, in terms of having a dynamical scaling exponent that is zero rather than two, but there is (as can be seen), still a residual suppression of fluctuations in the immediate vicinity of T_c .

where we have defined $\langle \mathbf{v}_b \rangle$ to be the average occupation of bonds, $N_b^{(n)}$ to be the number of occupied bonds in the n -th configuration, and we have used Eq. (3.6) in the second line. If we consider graphs with no self-intersections,

$$\sum_{j=bonds} v_j = \sum_{j=sites} v_j. \quad (3.30)$$

For small β (high T) this can be a good approximation,

$$\left\langle \sum_{j=bonds} v_j \right\rangle \approx \left\langle \sum_{j=sites} v_j \right\rangle \implies \quad (3.31)$$

$$\langle \mathbf{v}_b \rangle \approx \frac{1}{2} \langle \mathbf{v}_s \rangle \quad (3.32)$$

This agrees with our intuition, that the average image $\langle \mathbf{v} \rangle$ should resemble a “tablecloth”, where the site pixels are twice as dark as the link pixels. This can be seen clearly in Fig. 3.14. For a general graph, a link is shared by two sites (its endpoints), whereas a site can be shared by either 0, 2, or 4 bonds. If the site is shared by two bonds, it is only visited once, denoted $sites^{(1)}$, and if it is shared by four bonds, it is visited twice, denoted $sites^{(2)}$. This allows us to break up the sum over bonds into two terms

$$\sum_{j=bonds} v_j = \frac{2}{2} \sum_{j=sites^{(1)}} v_j + \frac{4}{2} \sum_{j=sites^{(2)}} v_j \quad (3.33)$$

Rearranging and taking averages gives

$$\left\langle \sum_{j=sites^{(1)}} v_j + \sum_{j=sites^{(2)}} v_j \right\rangle = \left\langle \sum_{j=bonds} v_j - \sum_{j=sites^{(2)}} v_j \right\rangle \quad (3.34)$$

$$= \left\langle \sum_{j=sites} v_j \right\rangle \quad (3.35)$$

$$= V \langle \mathbf{v}_s \rangle \quad (3.36)$$

$$= \left\langle \sum_{j=bonds} v_j \right\rangle - \left\langle \sum_{j=sites^{(2)}} v_j \right\rangle \quad (3.37)$$

$$= 2V \langle \mathbf{v}_b \rangle - \langle N_{sites^{(2)}} \rangle \implies \quad (3.38)$$

$$\frac{\langle N_{sites^{(2)}} \rangle}{V} = 2 \langle \mathbf{v}_b \rangle - \langle \mathbf{v}_s \rangle. \quad (3.39)$$

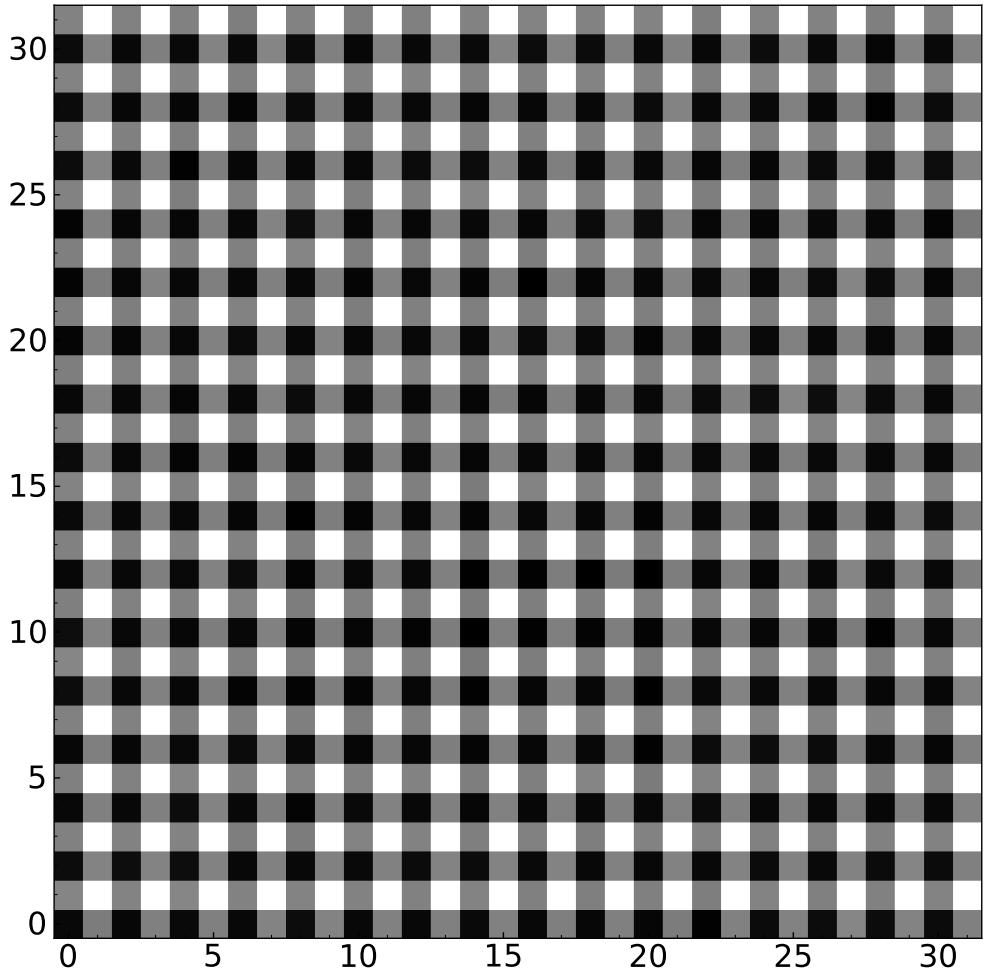


Figure 3.14: Average image $\langle v \rangle$ calculated for the $L = 16$ lattice at $T = 2.0$, illustrating the tablecloth-like appearance.

We can rewrite the last equation using (3.28)

$$\langle \mathbf{v}_s \rangle = \frac{\langle N_b \rangle}{V} - \frac{\langle N_{sites^{(2)}} \rangle}{V}. \quad (3.40)$$

This suggests that a departure from a perfect tablecloth ($\langle \mathbf{v}_s \rangle = 2\langle \mathbf{v}_b \rangle$) contains information. Another useful construct is the *covariance matrix*,

$$C_{ij} = \left\langle (v_i - \langle \mathbf{v} \rangle_i) (v_j - \langle \mathbf{v} \rangle_j)^T \right\rangle \quad (3.41)$$

$$= \frac{1}{N_{\text{configs}}} \sum_{n=1}^{N_{\text{configs}}} \left(v_i^{(n)} - \langle \mathbf{v} \rangle_i \right) \left(v_j^{(n)} - \langle \mathbf{v} \rangle_j \right)^T, \quad (3.42)$$

where we have defined $v_k^{(n)}$ to be the grayscale value of the k -th pixel in the n -th sample configuration, and $\langle \mathbf{v} \rangle_k$ is the average grayscale value of the k -th pixel over the set of configurations.

At some fixed temperature, the covariance matrix, $C \in \mathbb{R}^{N_{\text{configs}} \times 4L^2}$, where N_{configs} is the number of sample configurations (images), with each configuration flattened into a vector of length $4L^2$. We can then perform a singular value decomposition (SVD) on the covariance matrix,

$$C = W \Lambda W^T \quad (3.43)$$

where W is a $4L^2 \times 4L^2$ matrix whose columns (\mathbf{w}_k) are the eigenvectors of C , and Λ is the diagonal matrix of the absolute value of the eigenvalues $\lambda^{(k)}$ of C , arranged in decreasing order. Without loss of generality, we can assume that the eigenvectors \mathbf{w}_k are real and normalized such that $\mathbf{w}_k^T \mathbf{w}_k = 1$. Thus, we can write

$$C \mathbf{w}_k = \lambda^{(k)} \mathbf{w}_k \quad (3.44)$$

$$\mathbf{w}_k^T C \mathbf{w}_k = \lambda^{(k)} \quad (3.45)$$

For our purposes, we are interested in the first principal component, with eigenvalue $\lambda^{(1)} \equiv \lambda_1$ and corresponding eigenvector \mathbf{w}_1 .

We conjectured that the first principal component, \mathbf{w}_1 of the covariance matrix C is directly proportional to the average worm configuration (image) $\langle \mathbf{v} \rangle$, i.e.

$$\mathbf{w}_1 \propto \langle \mathbf{v} \rangle. \quad (3.46)$$

From our results in 3.2, we can write

$$\langle \mathbf{v} \rangle^2 = \langle \mathbf{v} \rangle^T \langle \mathbf{v} \rangle \quad (3.47)$$

$$= 2V\langle \mathbf{v}_b \rangle^2 + V\langle \mathbf{v}_s \rangle^2. \quad (3.48)$$

This suggests that

$$\mathbf{w}_1 = \frac{\langle \mathbf{v} \rangle}{\sqrt{(2\langle \mathbf{v}_b \rangle^2 + \langle \mathbf{v}_s \rangle^2)V}}. \quad (3.49)$$

Moreover, we can write

$$\sum_i \left(v_i^{(n)} - \langle \mathbf{v} \rangle_i \right) \langle \mathbf{v} \rangle_i = \langle \mathbf{v}_b \rangle \sum_{j=bonds} v_j^{(n)} + \langle \mathbf{v}_s \rangle \sum_{j=sites} v_j^{(n)} - \left(2V\langle \mathbf{v}_b \rangle^2 + V\langle \mathbf{v}_s \rangle^2 \right) \quad (3.50)$$

$$= \langle \mathbf{v}_b \rangle N_b^{(n)} + \langle \mathbf{v}_s \rangle N_s^{(n)} - V \left(2\langle \mathbf{v}_b \rangle^2 + \langle \mathbf{v}_s \rangle^2 \right) \quad (3.51)$$

$$= \langle \mathbf{v}_b \rangle \left(N_b^{(n)} - \langle N_b \rangle \right) + \langle \mathbf{v}_s \rangle \left(N_s^{(n)} - \langle N_s \rangle \right) \quad (3.52)$$

$$\equiv \langle \mathbf{v}_b \rangle \Delta_{N_b}^{(n)} + \langle \mathbf{v}_s \rangle \Delta_{N_s}^{(n)}. \quad (3.53)$$

From this, we can extract a relationship between the eigenvalue corresponding to the first principal component, $\lambda^{(1)}$ and the fluctuations Δ_{N_b} and Δ_{N_s} ,

$$\begin{aligned} \mathbf{w}_1^T C \mathbf{w}_1 &= \lambda^{(1)} \\ &= \frac{1}{N_{\text{configs}}} \sum_{n=1}^{N_{\text{configs}}} \frac{\left[\langle \mathbf{v}_b \rangle \Delta_{N_b}^{(n)} \right]^2 + 2\langle \mathbf{v}_b \rangle \langle \mathbf{v}_s \rangle \Delta_{N_b}^{(n)} \Delta_{N_s}^{(n)} + \left[\langle \mathbf{v}_s \rangle \Delta_{N_s}^{(n)} \right]^2}{2\langle \mathbf{v}_b \rangle^2 + \langle \mathbf{v}_s \rangle^2} \end{aligned}$$

Now, if we consider the high temperature approximation where sites only have single visits (no self-intersections), $\langle \mathbf{v}_s \rangle \simeq 2\langle \mathbf{v}_b \rangle$, $\langle N_s \rangle \simeq \langle N_b \rangle$, and $\Delta_{N_b} \simeq \Delta_{N_s}$, we have that $2\langle \mathbf{v}_b \rangle^2 + \langle \mathbf{v}_s \rangle^2 \simeq 6\langle \mathbf{v}_b \rangle^2$ and

$$\lambda_1 \simeq \frac{\langle \mathbf{v}_b \rangle^2}{N_{\text{configs}}} \frac{9}{6\langle \mathbf{v}_b \rangle^2} \sum_{n=1}^{N_{\text{configs}}} \left(\Delta_{N_b}^{(n)} \right)^2 \quad (3.54)$$

$$= \frac{3}{2} \frac{1}{N_{\text{configs}}} \sum_{n=1}^{N_{\text{configs}}} \left(\Delta_{N_b}^{(n)} \right)^2 \quad (3.55)$$

$$= \frac{3}{2} \langle \Delta_{N_b}^2 \rangle. \quad (3.56)$$

A justification for making this approximation can be seen in Fig. 3.15.

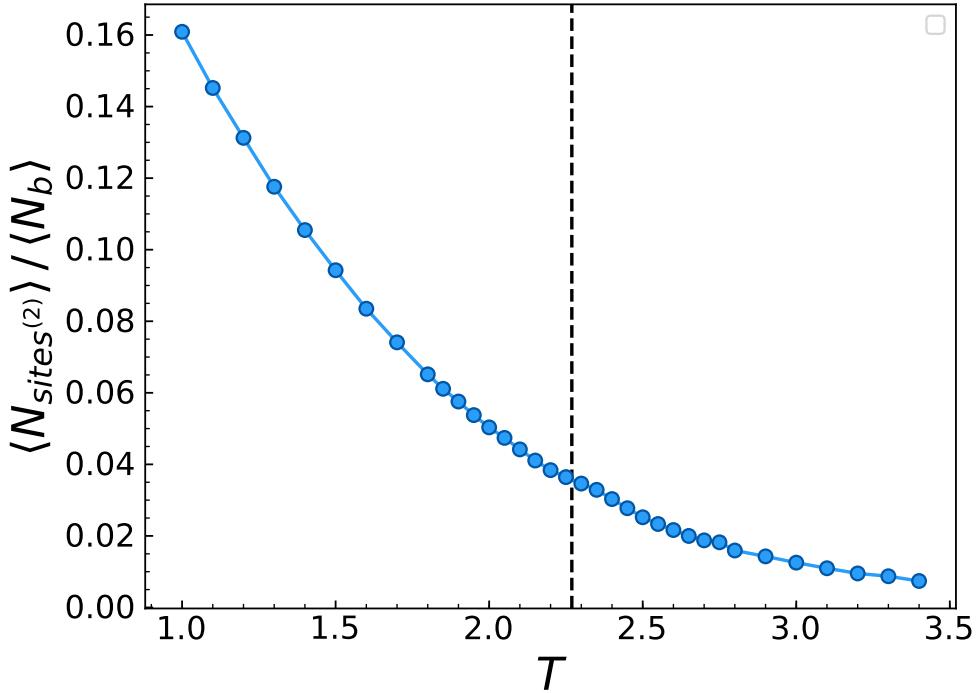


Figure 3.15: Ratio of the number of twice visited sites $\langle N_{sites^{(2)}} \rangle$ to the average number of bonds $\langle N_b \rangle$ versus temperature, for the $L = 32$ lattice. This clearly justifies our approximation $\langle v_s \rangle \simeq 2\langle v_b \rangle$, where we ignore the contribution from twice visited sites.

3.8.6 Illustration of Alternate Blockings

In Fig 3.16 we consider the alternative blocking procedures: $1 + 1 \rightarrow 1$ and $1 + 1 \rightarrow 2$.

3.9 Possible Applications: From Images to Loops

Having better understood how these RG transformations can be used to describe the 2D Ising model near criticality, we began to look for possible applications to real-world datasets. For our analysis, we used the CIFAR-10 [38] image set consisting of 60,000 32×32 color images in 10 classes. First, each of the images were converted to a grayscale with pixel values in the range $[0, 1]$. Next, a grayscale cutoff value was chosen so that all pixels with values below the cutoff would become black, and pixels above the cutoff would become white, resulting in images consisting entirely of black and white pixels. Finally, each of these images were converted to ‘worm-like’ images by drawing the boundaries separating black and white collections of pixels. An example of these preprocessing steps are illustrated in Fig. 3.17. This process was carried out on a mini-batch consisting of 500 randomly selected images from the CIFAR-10 image set. For each image in our mini-batch, we calculated $\langle N_b \rangle$ and $\langle \Delta_{N_b}^2 \rangle$ over a range of grayscale cutoff values in $[0, 1]$ in steps of 0.02. Each of

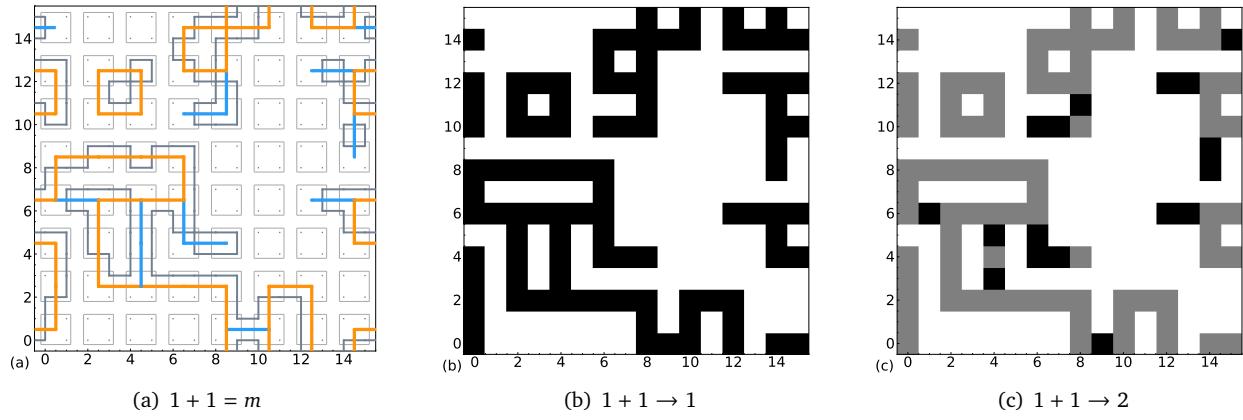


Figure 3.16: Example of the different coarse-graining (“blocking”) procedures applied to a sample worm configuration generated at $T = 2.0$. Note that in (3.16(a)) $m \in \{1, 2\}$, and double bonds are represented by blue lines. (3.16(b)), (3.16(c)), illustrate the results of applying different weights to the so-called “double bonds” in the images representing a blocked configuration. Note that in (3.16(b)) $1 + 1 \rightarrow 1$, double bonds are given the same weight as single bonds, and in (3.16(c)) $1 + 1 \rightarrow 2$, double bonds are given twice the weight as single bonds, appearing twice as dark.

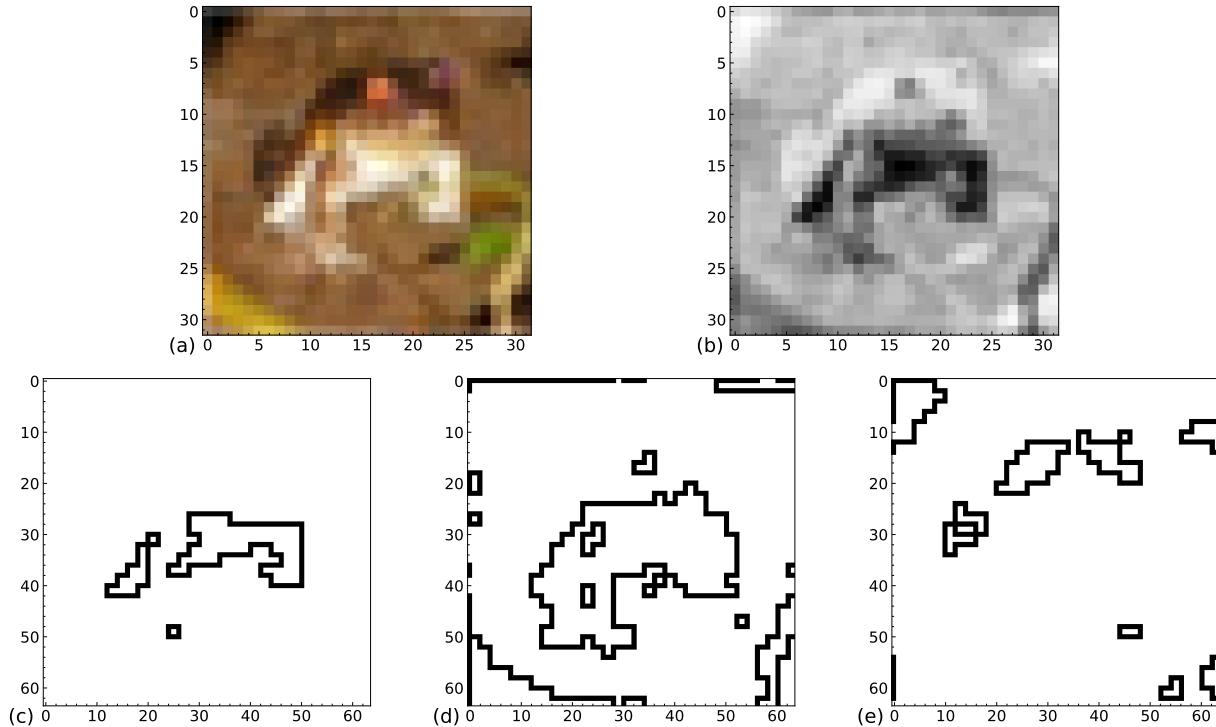


Figure 3.17: Example of preprocessing steps for converting CIFAR-10 images to ‘worm-like’ images, illustrating the resulting image for different values of the grayscale cutoff. (a) Original image from CIFAR-10 dataset. (b) Image converted to grayscale. (c) Resulting image from cutoff values of 0.25, (d) 0.5, and (e) 0.75.

these images were then iteratively blocked using the $(1 + 1 \rightarrow 0)$ blocking procedure described in Sec. 3.4, calculating $\langle N_b \rangle$ and $\langle \Delta_{N_b}^2 \rangle$ for each successive blocking step, as shown in Fig. 3.18. Immediately we see that

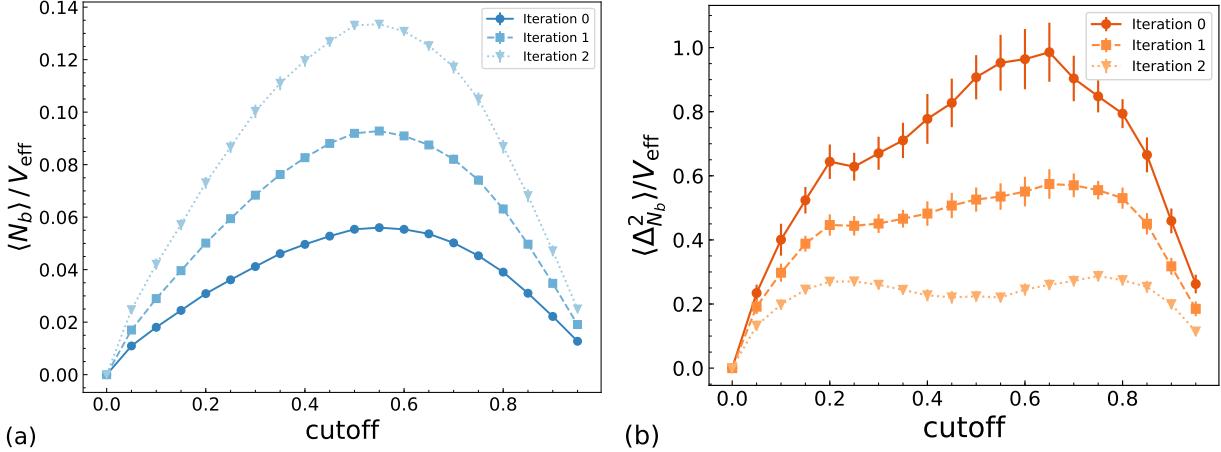


Figure 3.18: $\langle N_b \rangle$ and $\langle \Delta_{N_b}^2 \rangle$ vs. grayscale cutoff value for 500 randomly chosen images from the CIFAR-10 dataset.

there is no identifiable low temperature phase, and that for cutoff values near both 0 and 1, we obtain images which are mostly empty, similar to the high temperature configurations obtained from the worm algorithm. This suggests that there is no such notion of criticality (as characterized by the abrupt transition from a low to high temperature phase) like we found for the two-dimensional Ising model.

3.10 Conclusions

In summary, we have motivated, constructed and applied a RG transformation to sets of worm configurations at various temperature. This transformation is approximate and the coarse-grained configurations are themselves worm configurations. This allows multiple iterations. We monitored the bond density at successive iterations and compared them with a two-state TRG approximation. We found clear similarities in the low-temperature side, where data collapse is observed for both procedures when the distance to the critical point is rescaled at each iteration. In the high-temperature phase, only the TRG approximation shows good data collapse.

Can the procedure developed here be applied to the boundary of arbitrary sets of images as illustrated in the introduction? The gray cutoff could be used as a tunable parameter. However, in the limiting cases of a zero (one) gray cutoff we have uniform black (white) images which are both similar to the high-temperature phase, and we do not expect a phase transition. Applications to the CIFAR dataset are discussed in Section 3.9 and confirm this point of view.

RG methods have been considered for assisting in image recognition [39–41]. By mapping from fine to coarse in several ways, such as the $1 + 1 \rightarrow 0$ and $1 + 1 \rightarrow 1$ in our approach, one begins to see how the inverse process might go in replacing a degraded image with a higher resolution reconstruction. The physics of defining RG transformations and quantifying scheme dependence then guides such reconstructions using physical principles, which are expected to be embedded into real world image characteristics due to principles of universality.

It should be noted that the TRG procedure is often considered as a “local update” of the tensor. A more sophisticated approach consists of using the standard recursion to provide an environment for subsequent updates [28, 42]. An environment tensor E is propagated backward from the coarse to the fine scale. An improved version of the initial iteration can then be performed in an environment. This forward-backward procedure can be repeated and is very reminiscent of the procedure proposed by Hinton and Salakhutdinov [43] in the context of image recognition.

A better understanding of RG concepts in machine learning could enhance physics discovery, especially in the context of simulation and modeling of physical systems at a fundamental level [44]. The general idea is to render computational tools “smart,” i.e., that they would learn features and patterns without the intervention of a human “assistant,” and would, in the best possible scenario, guide the direction of further simulations. This could accelerate and deepen the process of understanding and characterizing the complex systems that are deemed important in pure and applied physics.

4 | Markov Chain Monte Carlo (MCMC)

For consistency with the original paper, we adopt much of the same notation. In order to better understand the theory behind the L2HMC algorithm, we begin with a general description of Markov Chain Monte Carlo (MCMC) methods.

4.1 Markov Chains

Generally speaking, MCMC methods are a class of algorithms that use Markov Chains to sample from a particular probability distribution that is often too complicated to sample from directly. These algorithms are of tremendous importance in a wide variety of fields and are of particular importance for simulations in lattice quantum chromodynamics (QCD) and lattice gauge theory.

A Markov chain can be understood as a sequence of random variables $\{x_1, x_2, \dots, x_N\}$ sampled from a conditional probability distribution $p(x_{t+1}|x_t)$, with the condition that the next sample x_{t+1} depends only on the current state x_t and does not depend on the history of previous states $\{x_1, x_2, \dots, x_{t-1}\}$ [45]. The set in which the x_i take values is often referred to as the *state space* \mathcal{X} of the Markov chain. For finite state spaces, the initial distribution can be associated with a vector $\lambda = (\lambda_1, \dots, \lambda_N)$ defined by

$$p(x_1 = x_i) = \lambda_i, \quad i = 1, \dots, N \tag{4.1}$$

and the transition probabilities can be associated with a transition matrix (kernel) P with elements p_{ij} given by

$$p(x_{t+1} = x_j | x_t = x_i) = p_{ij}, \quad i = 1, \dots, N \quad j = 1, \dots, n \tag{4.2}$$

This says that the $(i, j)^{\text{th}}$ entry of the n^{th} power of P gives the probability of transitioning from state i to state j in n steps. For our purposes, we are usually interested in Markov chains that have stationary transition probabilities, i.e. the conditional distribution of x_{t+1} given x_t does not depend on t . Note that for $i = 1, 2, \dots, N$,

$$\sum_{j=1}^N p_{ij} = 1. \quad (4.3)$$

There are two main (defining) properties of Markov chains that are relevant for our discussion, namely *stationarity* and *reversibility*.

- **Stationarity:** A sequence $\{x_1, x_2, \dots, x_N\}$ of random elements is said to be *stationary* if for every positive integer k the distribution of the k -tuple

$$(x_{t+1}, \dots, x_{t+k})$$

does not depend on t . An initial distribution is said to be stationary if the Markov chain specified by the initial distribution and transition probability distribution is stationary.

- **Reversibility:** A Markov chain is said to be reversible if its transition probability is reversible with respect to its initial distribution. Note that this is a stronger condition than *stationarity* since if a Markov chain is reversible it is also stationary. This condition can be better understood by noticing that for any $i, k > 0 \in \mathbb{Z}$, the distributions of $(x_{i+1}, \dots, x_{i+k})$ and $(x_{i+k}, \dots, x_{i+1})$ are the same.

Additionally, we often desire that our Markov chain is

- **Irreducible:** For any state of the Markov chain, there is a positive probability of visiting all other states, i.e. any two configurations are connected by a finite number of steps..
- **Aperiodic:** The chain will not tend towards a periodic limit-cycle where we can only reach a certain subset of states at any step, i.e. we must be able to reach all states at any given time step [46].

More formally, if we define the *period* of a state $x \in \mathcal{X}$ to be:

$$d(x) = \gcd\{n \in \mathbb{N}^+ : P^n(x, x) > 0\}, \quad (4.4)$$

then, starting at x , the chain can return to x only at multiples of the period d , where d is the largest such integer. We say that x is *aperiodic* if $d(x) = 1$ and *periodic* if $d(x) > 1$.

If a (finite) chain satisfies both of these properties, it will always have a unique stationary distribution.

4.1.1 Metropolis-Hastings Algorithm

The purpose of the Metropolis-Hastings algorithm is to generate a collection of states according to a desired distribution $p(x)$. More formally, let p be the target distribution defined up to a constant over a space \mathcal{X} .

We wish to construct a Markov Chain with stationary distribution equal to the target distribution p . We can obtain samples by simulating a Markov process.

Given an initial distribution π_0 and a transition matrix (kernel) K , we construct the following sequence of random variables

$$X_0 \sim \pi_0, \quad X_{t+1} \sim K(\cdot|X_t). \quad (4.5)$$

As discussed previously, for p to be the stationary distribution of the chain, K must be irreducible and aperiodic, and p has to be a *fixed point* of K . Note we can express the fixed point condition as

$$p(x') = \int K(x'|x)p(x)dx \quad (4.6)$$

We can ensure that stationarity is satisfied by requiring the (stronger) *detailed balance condition*¹

$$p(x')K(x|x') = p(x)K(x'|x). \quad (4.7)$$

We can rewrite the detailed balance condition as

$$\frac{K(x'|x)}{K(x|x')} = \frac{p(x')}{p(x)} \quad (4.8)$$

Given a proposal distribution $q(x'|x)$, we can construct a transition kernel satisfying detailed balance using Metropolis-Hastings accept/reject rules, as outlined in Alg. 2.

Algorithm 2: Metropolis-Hastings Algorithm

input : Transition kernel (proposal distribution), $q(x'|x)$

Initialize $x_0 \sim \pi_0$.

for $t = 0$ **to** N :

 Sample $x' \sim q(\cdot|x_t)$

 Compute the acceptance probability: $A(x'|x_t) = \min\left(1, \frac{p(x')q(x_t|x')}{p(x_t)q(x'|x_t)}\right)$

 With probability A accept the proposed value and set $x_{t+1} = x'$. Otherwise set $x_{t+1} = x_t$.

In practice, a Gaussian distribution is often used as the proposal distribution, i.e. $q(x'|x) = \mathcal{N}(x'|x, \Sigma)$.

In this case the acceptance probability reduces to

$$A(x'|x) = \min\left(1, \frac{p(x')}{p(x)}\right). \quad (4.9)$$

¹Note that this is equivalent to the reversibility condition defined above.

This ‘‘random walk’’ approach is conceptually simple and can be easily implemented but performs poorly for the high-dimensional target distributions commonly encountered in lattice gauge theory and lattice QCD.

4.2 Aside: Bayesian Analysis

While not directly relevant to our discussion, it is interesting to note the connection between MCMC and Bayesian analysis, which has applications in a wide range of disciplines outside of physics [47]. Recall Bayes’ theorem:

$$P(B_i|A) = \frac{P(A|B_i)p(B_i)}{\sum_j P(A|B_j)p(B_j)} \quad (4.10)$$

Here, $p(A|B)$ is the conditional probability of A given B , i.e. the probability that A is true when the condition B is satisfied. For example, suppose we are interested in determining if an email is as spam. The test for spam is that the message contains some flagged words B_i , i.e. B_1 = ‘winner’, B_2 = ‘Nigerian prince’, etc. In this case, $p(A|B_i)$ would then be the probability that the email is spam, given that it contains the word B_i . Suppose $p(A|B_i)$ and $p(B_i)$ are given (e.g. $P(A|B_1) = 10^{-4}$, $P(A|B_2) = 0.05, \dots$) and we want to derive $p(B_i|A)$, the probability of an email containing the word B_i , given that it was classified as spam. To use an analogy from physics, we can identify B_i with some field ϕ , and $p(A|B_i)p(B_i)$ with the path integral weight $e^{-S[\phi]}[d\phi]$. The denominator $\sum_j p(A|B_j)p(B_j) = p(A)$ would then be regarded as the partition function Z . Using MCMC sampling, we can obtain $p(B_i|A) \sim \frac{e^{-S[\phi]}[d\phi]}{Z}$ via Bayes’ theorem, i.e. we can collect many samples and see the resulting distribution.

4.3 Hamiltonian Monte Carlo

We can improve upon this random-walk guess and check strategy by ‘‘guiding’’ the simulation according to the systems natural dynamics using a method known as Hamiltonian (Hybrid) Monte Carlo (HMC).

In HMC, model samples can be obtained by simulating a physical system governed by a Hamiltonian comprised of kinetic and potential energy functions that govern a particles dynamics. By transforming the density function to a potential energy function and introducing the auxiliary momentum variable v , HMC lifts the target distribution onto a joint probability distribution in phase space (x, v) , where x is the original variable of interest (e.g. position in Euclidean space). A new state is then obtained by solving the equations of motion for a fixed period of time using a volume-preserving integrator (most commonly the *leapfrog integrator*). The addition of random (typically normally distributed) momenta encourages long-distance jumps in state space with a single Metropolis-Hastings (MH) step.

Let the ‘position’ of the physical state be denoted by a vector $x \in \mathbb{R}^n$ and the conjugate momenta of the physical state be denoted by a vector $v \in \mathbb{R}^n$. Then the Hamiltonian reads

$$\mathcal{H}(x, v) = U(x) + K(v) \quad (4.11)$$

$$= U(x) + \frac{1}{2}v^T v, \quad (4.12)$$

where $U(x)$ is the potential energy, and $K(v) = \frac{1}{2}v^T v$ the kinetic energy. We assume without loss of generality that the position and momentum variables are independently distributed. That is, we assume the target distribution of the system can be written as $p(x, v) = p(x)p(v)$. Further, instead of sampling $p(x)$ directly, HMC operates by sampling from the canonical distribution $p(x, v) = \frac{1}{Z} \exp(-\mathcal{H}(x, v)) = p(x)p(v)$, for some partition function Z that provides a normalization factor. Additionally, we assume the momentum is distributed according to an identity-covariance Gaussian given by $p(v) \propto \exp(-\frac{1}{2}v^T v)$. For convenience, we will denote the combined state of the system by $\xi \equiv (x, v)$. From this augmented state ξ , HMC produces a proposed state $\xi' = (x', v')$ by approximately integrating Hamiltonian dynamics jointly on x and v . This integration is performed along approximate iso-probability contours of $p(x, v) = p(x)p(v)$ due to the Hamiltonians energy conservation.

4.3.1 Hamiltonian Dynamics

One of the characteristic properties of Hamilton’s equations is that they conserve the value of the Hamiltonian. Because of this, every Hamiltonian trajectory is confined to an energy *level set*,

$$\mathcal{H}^{(-1)}(E) = \{x, v | \mathcal{H}(x, v) = E\}. \quad (4.13)$$

Our state $\xi = (x, v)$ then proceeds to explore this level set by integrating Hamilton’s equations, which are shown as a system of differential equations in Eq. 4.15.

$$\dot{x}_i = \frac{\partial \mathcal{H}}{\partial v_i} = v_i \quad (4.14)$$

$$\dot{v}_i = -\frac{\partial \mathcal{H}}{\partial x_i} = -\frac{\partial U}{\partial x_i} \quad (4.15)$$

It can be shown [48] that the above transformation is volume-preserving and reversible, two necessary factors to guarantee asymptotic convergence of the simulation to the target distribution. The dynamics are simulated

using the leapfrog integrator, which for a single time step consists of:

$$v^{\frac{1}{2}} = v - \frac{\epsilon}{2} \partial_x U(x) \quad (4.16)$$

$$x' = x + \epsilon v^{\frac{1}{2}} \quad (4.17)$$

$$v' = v - \frac{\epsilon}{2} \partial_x U(x'). \quad (4.18)$$

We write the action of the leapfrog integrator in terms of an operator $\mathbf{L} : \mathbf{L}\xi \equiv \mathbf{L}(x, v) \equiv (x', v')$, and introduce a momentum flip operator $\mathbf{F} : \mathbf{F}(x, v) \equiv (x, -v)$. The Metropolis-Hastings acceptance probability for the HMC proposal is given by:

$$A(\mathbf{FL}\xi | \xi) = \min \left(1, \frac{p(\mathbf{FL}\xi)}{p(\xi)} \left| \frac{\partial [\mathbf{FL}\xi]}{\partial \xi^T} \right| \right), \quad (4.19)$$

Where $\left| \frac{\partial [\mathbf{FL}\xi]}{\partial \xi^T} \right|$ denotes the determinant of the Jacobian describing the transformation, and is equal to 1 for traditional HMC. In order to utilize these Hamiltonian trajectories to construct an efficient Markov transition, we need a mechanism for introducing momentum to a given point in the target parameter space.

Fortunately, this can be done by exploiting the probabilistic structure of the system [49]. To lift an initial point in parameter space into one on phase space, we simply sample from the conditional distribution over the momentum,

$$v \sim \pi(x|v). \quad (4.20)$$

Sampling the momentum directly from the conditional distribution ensures that this lift will fall into the typical set in phase space. We can then proceed to explore the joint typical set by integrating Hamilton's equations as demonstrated above to obtain a new configuration $\xi \rightarrow \xi'$. We can then return to the target parameter space by simply projecting away the momentum,

$$(x, v) \rightarrow x \quad (4.21)$$

These three steps when performed in series gives a complete Hamiltonian Markov transition composed of random trajectories that rapidly explore the target distribution, as desired. An example of this process can be seen in Fig 4.1.

4.3.1.1 Properties of Hamiltonian Dynamics

There are three fundamental properties of Hamiltonian dynamics which are crucial to its use in constructing Markov Chain Monte Carlo updates.

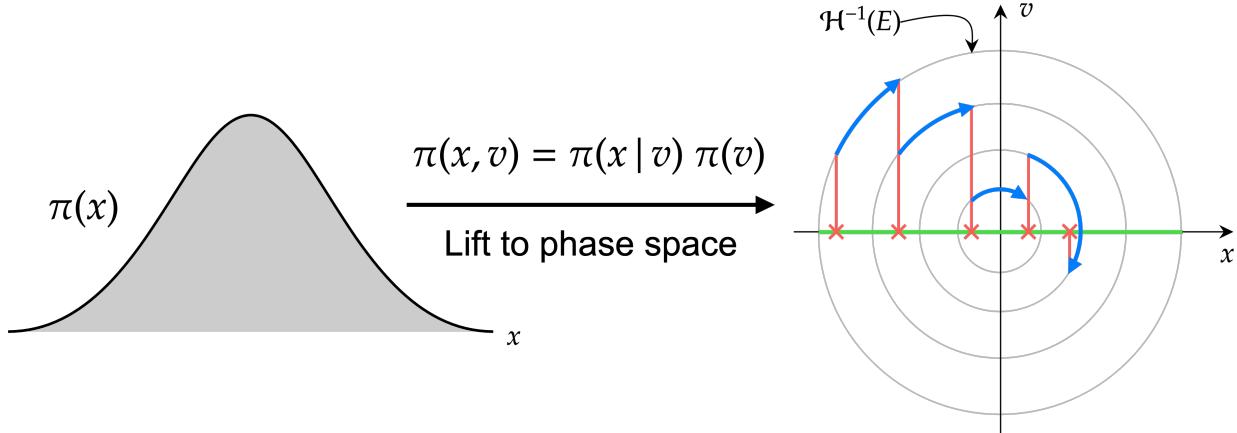


Figure 4.1: Visualizing HMC for a 1D Gaussian (example from [49], figure adapted with permission from [50]). Each Hamiltonian Markov transition lifts the initial state onto a random level set of the Hamiltonian, $\mathcal{H}^{(-1)}(E)$, which can then be explored with a [Hamiltonian trajectory](#) before [projecting back down](#) to the [target parameter space](#).

1. **Reversibility:** Hamiltonian dynamics are *reversible* — the mapping from $\mathbf{L} : \xi(t) \rightarrow \xi' = \xi(t+s)$ is one-to-one, and consequently has an inverse \mathbf{L}^{-1} , obtained by negating the time derivatives in Eq. 4.15.
2. **Conservation of the Hamiltonian:** Moreover, the dynamics *keeps the Hamiltonian invariant*.
3. **Volume preservation:** The final property of Hamiltonian dynamics is that it *preserves volume* in (x, v) phase space (i.e. Liouville's Theorem).

All in all, HMC offers noticeable improvements compared to the ‘random-walk’ approach of generic MCMC, but tends to perform poorly on high-dimensional distributions. This becomes immediately apparent when it is used for simulations in lattice gauge theory and lattice QCD, where large autocorrelations and slow ‘burn-in’ can become prohibitively expensive.

5 | Semi-Supervised Learning: L2HMC

5.1 Introduction

We describe a new technique for performing Hamiltonian Monte-Carlo (HMC) simulations called: ‘Learning to Hamiltonian Monte Carlo’ (L2HMC) [1] which expands upon the traditional HMC by using a generalized version of the leapfrog integrator that is parameterized by weights in a neural network. In order to demonstrate the usefulness of this new approach, we use various metrics for measuring the performance of the trained (L2HMC) sampler vs. the generic HMC sampler.

First, we will look at applying this algorithm to a two-dimensional Gaussian Mixture Model (GMM). The GMM is a notoriously difficult distribution for HMC due to the vanishingly small likelihood of the leapfrog integrator traversing the space between the two modes. Conversely, we see that through the use of a carefully chosen training procedure, the trained L2HMC sampler is able to successfully discover the existence of both modes, and mixes (‘tunnels’) between the two with ease. Additionally, we will observe that the trained L2HMC sampler mixes much faster than the generic HMC sampler, as evidenced through their respective autocorrelation spectra.

This ability to reduce autocorrelations is an important metric for measuring the efficiency of a general MCMC algorithm, and is of great importance for simulations in lattice gauge theory and lattice QCD. Following this, we introduce the two-dimensional $U(1)$ lattice gauge theory and describe important modifications to the algorithm that are of particular relevance for lattice models. Ongoing issues and potential areas for improvement are also discussed, particularly within the context of high-performance computing and long-term goals of the lattice QCD community.

5.2 Generalizing the Leapfrog Integrator

As in the previously described HMC algorithm, we start by augmenting the current state $x \in \mathbb{R}^n$ with a continuous momentum variable $v \in \mathbb{R}^n$ drawn from a standard normal distribution. Additionally, we introduce a binary direction variable $d \in \{-1, 1\}$, drawn from a uniform distribution. The complete augmented state

is then denoted by $\xi \equiv (x, v, d)$, with probability density $p(\xi) = p(x)p(v)p(d)$. To improve the overall performance of our model, for each step t of the leapfrog operator \mathbf{L}_θ , we assign a fixed random binary mask $m^t \in \{0, 1\}^n$ that will determine which variables are affected by each sub-update. The mask m^t is drawn uniformly from the set of binary vectors satisfying $\sum_{i=1}^n m_i^t = \lfloor \frac{n}{2} \rfloor$, i.e. half the entries of m^t are 0 and half are 1. Additionally, we write $\bar{m}^t = \mathbb{1} - m^t$ and $x_{m^t} = x \odot m^t$, where \odot denotes element-wise multiplication, and $\mathbb{1}$ the vector of 1's in each entry.

We begin with a subset of the augmented space, $\zeta_1 \equiv (x, \partial_x U(x), t)$, independent of the momentum v . We introduce three new functions of ζ_1 : T_v , Q_v , and S_v . We can then perform a single time-step of our modified leapfrog integrator \mathbf{L}_θ .

First, we update the momentum v , which depends only on the subset ζ_1 . This update is written

$$v' = v \odot \underbrace{\exp\left(\frac{\epsilon}{2} S_v(\zeta_1)\right)}_{\text{Momentum scaling}} - \frac{\epsilon}{2} \left[\underbrace{\partial_x U(x) \odot \exp(\epsilon Q_v(\zeta_1))}_{\text{Gradient scaling}} + \underbrace{T_v(\zeta_1)}_{\text{Translation}} \right] \quad (5.1)$$

and the corresponding Jacobian is given by: $\exp\left(\frac{\epsilon}{2} \mathbb{1} \cdot S_v(\zeta_1)\right)$. Next, we update x by first updating a subset of the coordinates of x (determined according to the mask m^t), followed by the complementary subset (determined from \bar{m}^t). The first update affects only x_{m^t} and produces x' . This update depends only on the subset $\zeta_2 \equiv (x_{\bar{m}^t}, v, t)$. Following this, we perform the second update which only affects $x'_{\bar{m}^t}$ and depends only on $\zeta_3 \equiv (x'_{m^t}, v, t)$, to produce x'' :

$$x' = x_{\bar{m}^t} + m^t \odot [x \odot \exp(\epsilon S_x(\zeta_2)) + \epsilon (v' \odot \exp(\epsilon Q_x(\zeta_2)) + T_x(\zeta_2))] \quad (5.2)$$

$$x'' = x'_{m^t} + \bar{m}^t \odot [x' \odot \exp(\epsilon S_x(\zeta_3)) + \epsilon (v' \odot \exp(\epsilon Q_x(\zeta_3)) + T_x(\zeta_3))] \quad (5.3)$$

with Jacobians: $\exp(\epsilon m^t \cdot S_x(\zeta_2))$, and $\exp(\epsilon \bar{m}^t \cdot S_x(\zeta_3))$, respectively. Finally, we proceed to update v again, using the subset $\zeta_4 \equiv (x'', \partial_x U'', t)$:

$$v'' = v' \odot \exp\left(\frac{\epsilon}{2} S_v(\zeta_4)\right) - \frac{\epsilon}{2} [\partial_x U \odot \exp(\epsilon Q_v(\zeta_4)) + T_v(\zeta_4)] \quad (5.4)$$

In order to build some intuition about each of these terms, we discuss below some of the subtleties contained in this approach and how they are (carefully) dealt with.

The first thing to notice about these equations is that if $S_i = Q_i = T_i = 0$ ($i = x, v$), we recover the previous equations for the generic leapfrog integrator (as we would expect since we are attempting to *generalize* HMC). We can also see a similarity between the equations for updating v and those for updating

x : each update is generalized by *scaling* the previous value (v or x), and *scaling and translating* the updating value (either $\partial_x U(x)$ or x). It can be shown [1], that the scaling applied to the momentum in Eq 5.1 can enable, among other things, acceleration in low-density zones to facilitate mixing between modes, and that the scaling term applied to the gradient may allow better conditioning of the energy landscape (e.g., by learning a diagonal inertia tensor), or partial ignoring of the energy gradient for rapidly oscillating energies.

Second, note that because the determinant of the Jacobian appears in the Metropolis-Hastings (MH) acceptance probability, we require the Jacobian of each update to be efficiently computable (i.e. independent of the variable actually being updated). For each of the momentum updates, the input is a subset $\zeta = (x, \partial_x U(x), t)$ of the augmented space and the associated Jacobian is $\exp\left(\frac{\varepsilon}{2} \mathbb{1} \cdot S_v(\zeta)\right)$ which is independent of v as desired. For the position updates however, things are complicated by the fact that the input ζ is x -dependent. In order to ensure that the Jacobian of the x update is efficiently computable, it is necessary to break the update into two parts following the approach outlined in *Real-valued Non-Volume Preserving transformations (RealNVP)* [51].

5.2.1 Metropolis-Hastings Accept/Reject

Written in terms of these transformations, the augmented leapfrog operator \mathbf{L}_θ consists of M sequential applications of the single-step leapfrog operator $\mathbf{L}_\theta \xi = \mathbf{L}_\theta(x, v, d) = (x''^{\times M}, v''^{\times M}, d)$, followed by the previously-defined momentum flip operator \mathbf{F} which flips the direction variable d , i.e. $\mathbf{F}\xi = (x, v, -d)$. Using these, we can express a complete molecular dynamics update step as $\mathbf{FL}_\theta \xi = \xi'$, where now the Metropolis-Hastings acceptance probability for this proposal is given by

$$A(\mathbf{FL}\xi | \xi) = \min\left(1, \frac{p(\mathbf{FL}\xi)}{p(\xi)} \left| \frac{\partial [\mathbf{FL}\xi]}{\partial \xi^T} \right| \right), \quad (5.5)$$

Where $\left| \frac{\partial [\mathbf{FL}\xi]}{\partial \xi^T} \right|$ denotes the determinant of the Jacobian describing the transformation.

In contrast to generic HMC where $\left| \frac{\partial [\mathbf{FL}\xi]}{\partial \xi^T} \right| = 1$, we now have non-symplectic transformations (i.e. non-volume preserving) and so we must explicitly account for the determinant of the Jacobian. These non-volume preserving transformations have the effect of deforming the energy landscape, which, depending on the nature of the transformation, may allow for the exploration of regions of space which were previously inaccessible.

To simplify our notation, introduce an additional operator \mathbf{R} that re-samples the momentum and direction, e.g. given $\xi = (x, v, d)$, $\mathbf{R}\xi = (x, v', d')$ where $v' \sim \mathcal{N}(0, I)$, $d' \sim \mathcal{U}(\{-1, 1\})$. A complete sampling step of our algorithm then consists of the following two steps:

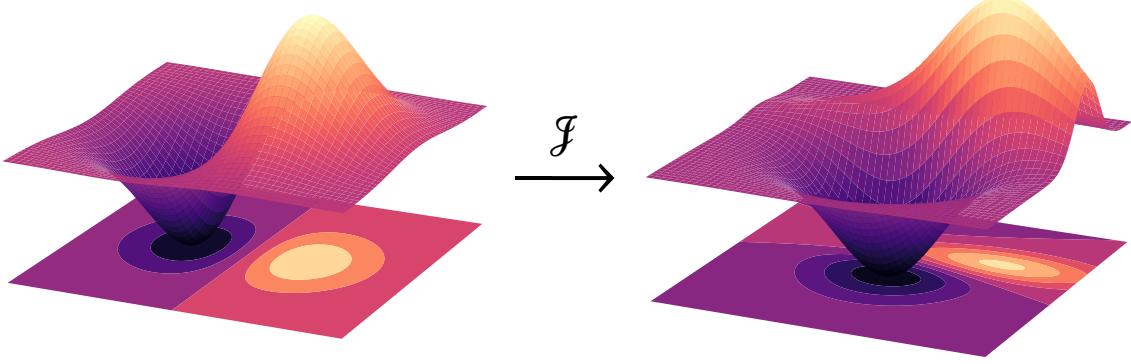


Figure 5.1: Example of how the determinant of the Jacobian can deform the energy landscape.

1. $\xi' = \text{FL}_\theta \xi$ with probability $A(\text{FL}_\theta \xi | \xi)$ (Eq. 4.19), otherwise $\xi' = \xi$.
2. $\xi' = \mathbf{R} \xi$.

Note however, that for MH to be well-defined, this deterministic operator must be *invertible* and *have a tractable Jacobian* (i.e. we can compute its determinant). In order to make this operator invertible, we augment the state space (x, v) into (x, v, d) , where $d \in \{-1, 1\}$ is drawn with equal probability and represent the direction of the update. All of the previous expressions for the augmented leapfrog updates represent the forward ($d = 1$) direction. We can derive the expressions for the backward direction ($d = -1$) by reversing the order of the updates (i.e. $v'' \rightarrow v'$, then $x'' \rightarrow x'$, followed by $x' \rightarrow x$ and finally $v' \rightarrow v$). For completeness, we include in Sec. 5.3 and Sec 5.4 all of the equations (both forward and backward directions) relevant for updating the variables of interest in our augmented leapfrog sampler.

5.3 Forward Direction ($d = 1$):

$$v' = v \odot \exp\left(\frac{\epsilon}{2} S_v(\zeta_1)\right) - \frac{\epsilon}{2} [\partial_x U(x) \odot \exp(\epsilon Q_v(\zeta_1)) + T_v(\zeta_1)] \quad (5.6)$$

$$x' = x_{\bar{m}^t} + m^t \odot [x \odot \exp(\epsilon S_x(\zeta_2)) + \epsilon (v' \odot \exp(\epsilon Q_x(\zeta_2)) + T_x(\zeta_2))] \quad (5.7)$$

$$x'' = x'_{m^t} + \bar{m}^t \odot [x' \odot \exp(\epsilon S_x(\zeta_3)) + \epsilon (v' \odot \exp(\epsilon Q_x(\zeta_3)) + T_x(\zeta_3))] \quad (5.8)$$

$$v'' = v' \odot \exp\left(\frac{\epsilon}{2} S_v(\zeta_4)\right) - \frac{\epsilon}{2} [\partial_x U(x'') \odot \exp(\epsilon Q_v(\zeta_4)) + T_v(\zeta_4)] \quad (5.9)$$

With $\zeta_1 = (x, \partial_x U(x), t)$, $\zeta_2 = (x_{\bar{m}^t}, v, t)$, $\zeta_3 = (x'_{m^t}, v, t)$, $\zeta_4 = (x'', \partial_x U(x''), t)$.

5.4 Backward Direction ($d = -1$):

$$v' = \left\{ v + \frac{\epsilon}{2} [\partial_x U(x) \odot \exp(\epsilon Q_v(\zeta_1)) + T_v(\zeta_1)] \right\} \odot \exp\left(-\frac{\epsilon}{2} S_v(\zeta_1)\right) \quad (5.10)$$

$$x' = x_{m^t} + \bar{m}^t \odot [x - \epsilon(\exp(\epsilon Q_x(\zeta_2)) \odot v' + T_x(\zeta_2))] \odot \exp(-\epsilon S_x(\zeta_2)) \quad (5.11)$$

$$x'' = x_{\bar{m}^t} + m^t \odot [x' - \epsilon(\exp(\epsilon Q_x(\zeta_3)) \odot v' + T_x(\zeta_3))] \odot \exp(-\epsilon S_x(\zeta_3)) \quad (5.12)$$

$$v'' = \left\{ v' + \frac{\epsilon}{2} [\partial_x U(x'') \odot \exp(\epsilon Q_v(\zeta_1)) + T_v(\zeta_1)] \right\} \odot \exp\left(-\frac{\epsilon}{2} S_v(\zeta_4)\right) \quad (5.13)$$

With $\zeta_1 = (x, \partial_x U(x), t)$, $\zeta_2 = (x_{m^t}, v, t)$, $\zeta_3 = (x'_{\bar{m}^t}, v, t)$, $\zeta_4 = (x'', \partial_x U(x''), t)$.

5.5 Determinant of the Jacobian

In terms of the auxiliary functions S_i, Q_i, T_i , we can compute the Jacobian:

$$\log |\mathcal{J}| = \log \left| \frac{\partial [\mathbf{FL}_\theta \xi]}{\partial \xi^T} \right| \quad (5.14)$$

$$= d \sum_{t \leq N_{LF}} \left[\frac{\epsilon}{2} \mathbb{1} \cdot S_v(\zeta_1^t) + \epsilon m^t \cdot S_x(\zeta_2^t) + \epsilon \bar{m}^t \cdot S_x(\zeta_3^t) + \frac{\epsilon}{2} \mathbb{1} \cdot S_v(\zeta_4^t) \right]. \quad (5.15)$$

where N_{LF} is the number of leapfrog steps, and ζ_i^t denotes the intermediary variable ζ_i at time step t and d is the direction of ξ , i.e. $d = 1$ (-1) for the forward (backward) update.

5.6 Network Architecture

As previously mentioned, each of the functions Q , S , and T , are implemented using multi-layer perceptrons with shared weights. It's important to note that we keep separate the network responsible for parameterizing the functions used in the position updates (' X_{net} ', i.e. Q_x, S_x , and T_x), and the network responsible for parameterizing the momentum updates (' V_{net} ', i.e. Q_v, S_v , and T_v). Since both networks are identical, we describe the architecture of V_{net} below, and include a flowchart for X_{net} illustrative purposes in Fig 5.3.

The network takes as input $\zeta_1 = (x, \partial_x U(x), t)$, where $x, v \in \mathbb{R}^n$, and t is encoded as $\tau(t) = \left(\cos\left(\frac{2\pi t}{N_{LF}}\right), \sin\left(\frac{2\pi t}{N_{LF}}\right) \right)$.

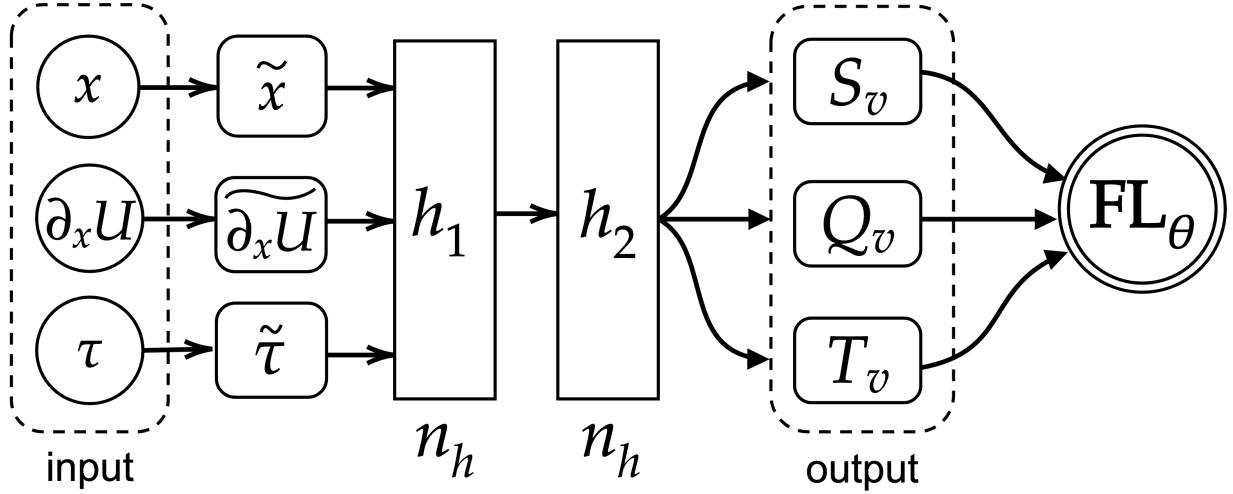


Figure 5.2: Illustration showing the generic (fully-connected) network architecture for training S_v , Q_v , and T_v . Figure adapted with permission from [50].

Each of the inputs is then passed through a fully-connected ('dense' layer), consisting of n_h hidden units

$$\tilde{x} = W^{(x)}x + b^{(x)} \quad (\in \mathbb{R}^{n_h}) \quad (5.16)$$

$$\tilde{v} = W^{(v)}v + b^{(v)} \quad (\in \mathbb{R}^{n_h}) \quad (5.17)$$

$$\tilde{\tau} = W^{(\tau)}\tau + b^{(\tau)} \quad (\in \mathbb{R}^{n_h}). \quad (5.18)$$

Where $W^{(x)}, W^{(v)} \in \mathbb{R}^{n \times n_h}$, $W^{(\tau)} \in \mathbb{R}^{2 \times n_h}$, and $b^{(x)}, b^{(v)}, b^{(\tau)} \in \mathbb{R}^{n_h}$. From these, the network computes

$$h_1 = \sigma(\tilde{x} + \tilde{v} + \tilde{\tau}) \quad (\in \mathbb{R}^{n_h}). \quad (5.19)$$

Where $\sigma(x) = \max(0, x)$ denotes the rectified linear unit (ReLU) activation function. Next, the network computes

$$h_2 = \sigma(W^{(h_1)}h_1 + b^{(h_1)}) \quad (\in \mathbb{R}^{n_h}). \quad (5.20)$$

These weights (h_2) are then used to compute the network's output:

$$S_x = \lambda_S \tanh(W^{(S)}h_2 + b^{(S)}) \quad (\in \mathbb{R}^n) \quad (5.21)$$

$$Q_x = \lambda_Q \tanh(W^{(Q)}h_2 + b^{(Q)}) \quad (\in \mathbb{R}^n) \quad (5.22)$$

$$T_x = W^{(T)}h_2 + b^{(T)} \quad (\in \mathbb{R}^n), \quad (5.23)$$

Where $W^{(s)}, W^{(q)}$, and $W^{(T)} \in \mathbb{R}^{n_h \times n}$ and $b^{(s)}, b^{(q)}$, and $b^{(T)} \in \mathbb{R}^n$. The parameters λ_s and λ_q are additional trainable variables initialized to zero. The network used for parameterizing the functions T_v , Q_v and S_v takes as input $(x, \partial_x U(x), t)$ where again t is encoded as above. The architecture of this network is the same, and produces outputs T_v , Q_v , and S_v .

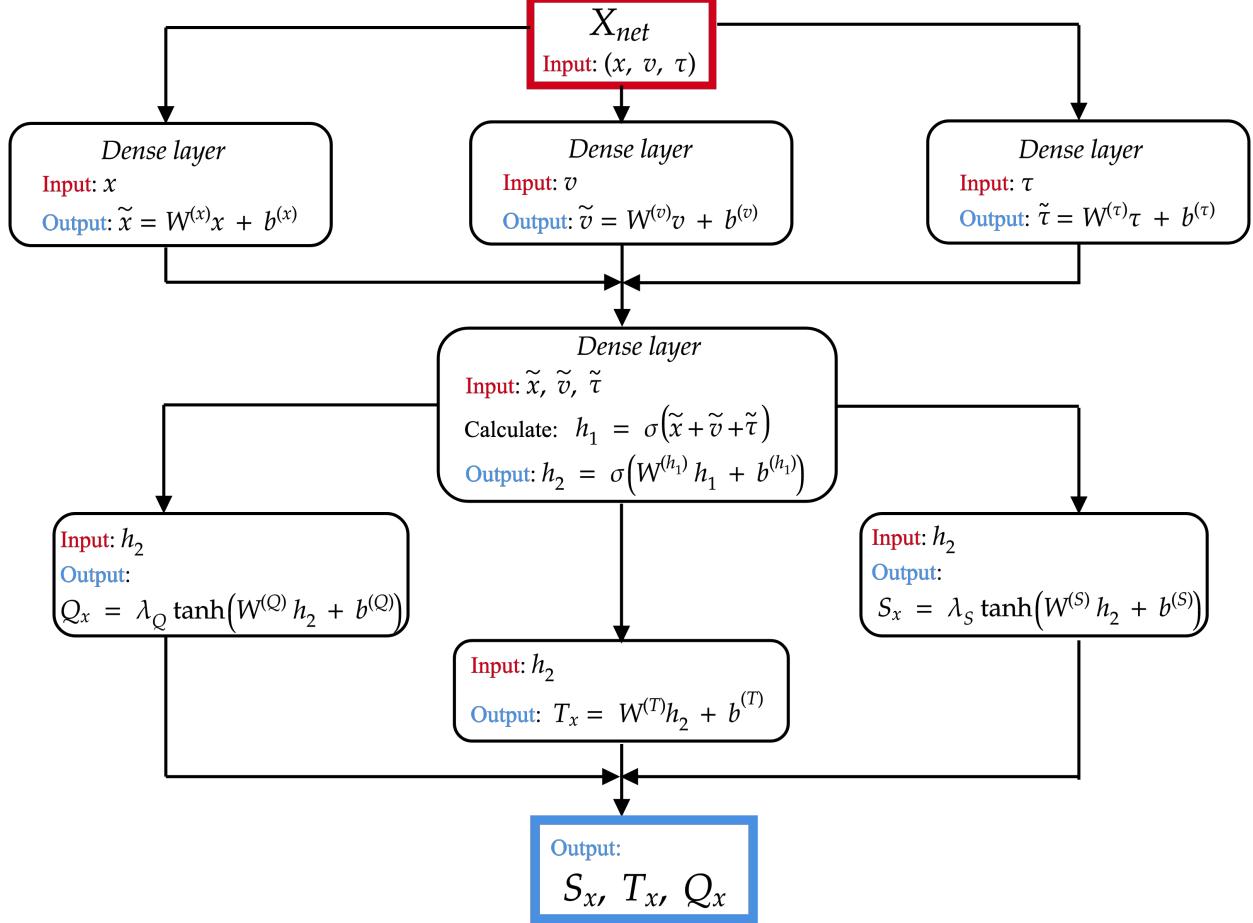


Figure 5.3: Flowchart illustrating the generic fully-connected network architecture including the intermediate variables computed at each hidden layer of the network.

5.7 Training Procedure

By augmenting traditional HMC methods with these trainable functions, we hope to obtain a sampler that has the following key properties:

1. Fast mixing (i.e. able to quickly produce uncorrelated samples).
2. Fast burn-in (i.e. rapid convergence to the target distribution).

3. Ability to mix across energy levels.
4. Ability to mix between modes.

Following the results in [52], we design a loss function with the goal of maximizing the expected squared jumped distance (or analogously, minimizing the lag-one autocorrelation). To do this, we first introduce

$$\delta(\xi, \xi') = \delta((x', v', d'), (x, v, d)) \equiv \|x - x'\|_2^2. \quad (5.24)$$

Then, the expected squared jumped distance is given by $\mathbb{E}_{\xi \sim p(\xi)} [\delta(\mathbf{FL}_\theta \xi, \xi) A(\mathbf{FL}_\theta \xi | \xi)]$. By maximizing this objective function, we are encouraging transitions that efficiently explore a local region of state-space, but may fail to explore regions where very little mixing occurs. To help combat this effect, we define a loss function

$$\ell_\lambda(\xi, \xi', A(\xi' | \xi)) = \frac{\lambda^2}{\delta(\xi, \xi') A(\xi' | \xi)} - \frac{\delta(\xi, \xi') A(\xi' | \xi)}{\lambda^2} \quad (5.25)$$

where λ is a scale parameter describing the characteristic length scale of the problem. Note that the first term helps to prevent the sampler from becoming stuck in a state where it cannot move effectively, and the second term helps to maximize the distance between subsequent moves in the Markov chain.

The sampler is then trained by minimizing ℓ_λ over both the target and initialization distributions. Explicitly, for an initial distribution π_0 over \mathcal{X} , we define the initialization distribution as $q(\xi) = \pi_0(x)\mathcal{N}(v; 0, I)p(d)$, and minimize

$$\mathcal{L}(\theta) \equiv \mathbb{E}_{p(\xi)} [\ell_\lambda(\xi, \mathbf{FL}_\theta \xi, A(\mathbf{FL}_\theta \xi | \xi))] + \lambda_b \mathbb{E}_{q(\xi)} [\ell_\lambda(\xi, \mathbf{FL}_\theta \xi, A(\mathbf{FL}_\theta \xi | \xi))]. \quad (5.26)$$

For completeness, we include the full algorithm [1] used to train L2HMC in Alg. 3.

5.8 Gaussian Mixture Model

The Gaussian Mixture Model (GMM) is a notoriously difficult example for traditional HMC to sample accurately due to the existence of multiple modes. In particular, HMC cannot mix between modes that are reasonably separated without recourse to additional tricks. This is due, in part, to the fact that HMC cannot easily traverse the low-density zones which exist between modes.

In the most general case, we consider a target distribution described by a mixture of $M > 1$ components

Algorithm 3: Training procedure for the L2HMC algorithm.

input :

1. A (potential) energy function, $U : \mathcal{X} \rightarrow \mathbb{R}$ and its gradient $\nabla_x U : \mathcal{X} \rightarrow \mathcal{X}$
2. Initial distribution over the augmented state space, q
3. Number of iterations, N_{train}
4. Number of leapfrog steps, N_{LF}
5. Learning rate schedule, $(\alpha_t)_{t \leq N_{\text{train}}}$
6. Batch size, N_{samples}
7. Scale parameter, λ
8. Regularization strength, λ_b

Initialize the parameters of the sampler, θ

Initialize $\{\xi_{p^{(i)}}\}_{i \leq N_{\text{samples}}}$ from $q(\xi)$

for $t = 0$ to N_{train} :

Sample a minibatch, $\{\xi_q^{(i)}\}_{i \leq N_{\text{samples}}}$	from $q(\xi)$. $\mathcal{L} \leftarrow 0$ for $i = 1$ to N_{LF} : $\xi_p^{(i)} \leftarrow \mathbf{R} \xi_p^{(i)}$ $\mathcal{L} \leftarrow \mathcal{L} + \ell_\lambda \left(\xi_p^{(i)}, \mathbf{FL}_\theta \xi_p^{(i)}, A(\mathbf{FL}_\theta \xi_p^{(i)} \xi_p^{(i)}) \right) + \lambda_b \ell_\lambda \left(\xi_q^{(i)}, \mathbf{FL}_\theta \xi_q^{(i)}, A(\mathbf{FL}_\theta \xi_q^{(i)} \xi_q^{(i)}) \right)$ $\xi_p^{(i)} \leftarrow \mathbf{FL}_\theta \xi_p^{(i)}$ with probability $A(\mathbf{FL}_\theta \xi_p^{(i)} \xi_p^{(i)})$ $\theta \leftarrow \theta - \alpha_t \nabla_\theta \mathcal{L}$
---	--

in \mathbb{R}^D for $D \geq 1$:

$$p(\mathbf{x}) \equiv \sum_{m=1}^M p(m)p(\mathbf{x}|m) \equiv \sum_{m=1}^M \pi_m p(\mathbf{x}|m) \quad \forall \mathbf{x} \in \mathbb{R}^D \quad (5.27)$$

where $\sum_{m=1}^M \pi_m = 1$, $\pi_M \in (0, 1) \forall m = 1, \dots, M$ and each component distribution is a normal probability distribution in \mathbb{R}^D . So $\mathbf{x}|m \sim \mathcal{N}(\boldsymbol{\mu}_m, \boldsymbol{\Sigma}_m)$, where $\boldsymbol{\mu}_m \equiv \mathbb{E}_{p(\mathbf{x}|m)} \{\mathbf{x}\}$ and $\boldsymbol{\Sigma}_m \equiv \mathbb{E}_{p(\mathbf{x}|m)} \{(\mathbf{x} - \boldsymbol{\mu}_m)(\mathbf{x} - \boldsymbol{\mu}_m)^T\} > 0$ are the mean vector and covariance matrix, respectively, of component m .

5.8.1 Example

Consider a simple 2D case consisting of two Gaussians

$$\mathbf{x} \sim \pi_1 \mathcal{N}(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1) + \pi_2 \mathcal{N}(\boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2) \quad (5.28)$$

with $\pi_1 = \pi_2 = 0.5$, $\boldsymbol{\mu}_1 = (-2, 0)$, $\boldsymbol{\mu}_2 = (2, 0)$ and

$$\boldsymbol{\Sigma}_1 = \boldsymbol{\Sigma}_2 = \begin{bmatrix} 0.1 & 0 \\ 0 & 0.1 \end{bmatrix} \quad (5.29)$$

The results of trajectories generated using both traditional HMC and the L2HMC algorithm can be seen in Fig. 5.4. Note that traditional HMC performs poorly and is unable to mix between the two modes, whereas L2HMC is able to correctly sample from the target distribution without getting stuck in either of the individual modes.

The L2HMC sampler was trained using simulated annealing using the schedule shown in Eq 5.30 with a starting temperature of $T = 10$, for 5,000 training steps. By starting with a high temperature, the chain is able to move between both modes ('tunnel') successfully. Once it has learned this, we can lower the temperature back to $T = 1$ and recover the initial distribution while preserving information about tunneling in the networks "memory".

$$T(n) = (T_i - T_f) \cdot \left(1 - \frac{n}{N_{\text{train}}}\right) + T_f \quad (5.30)$$

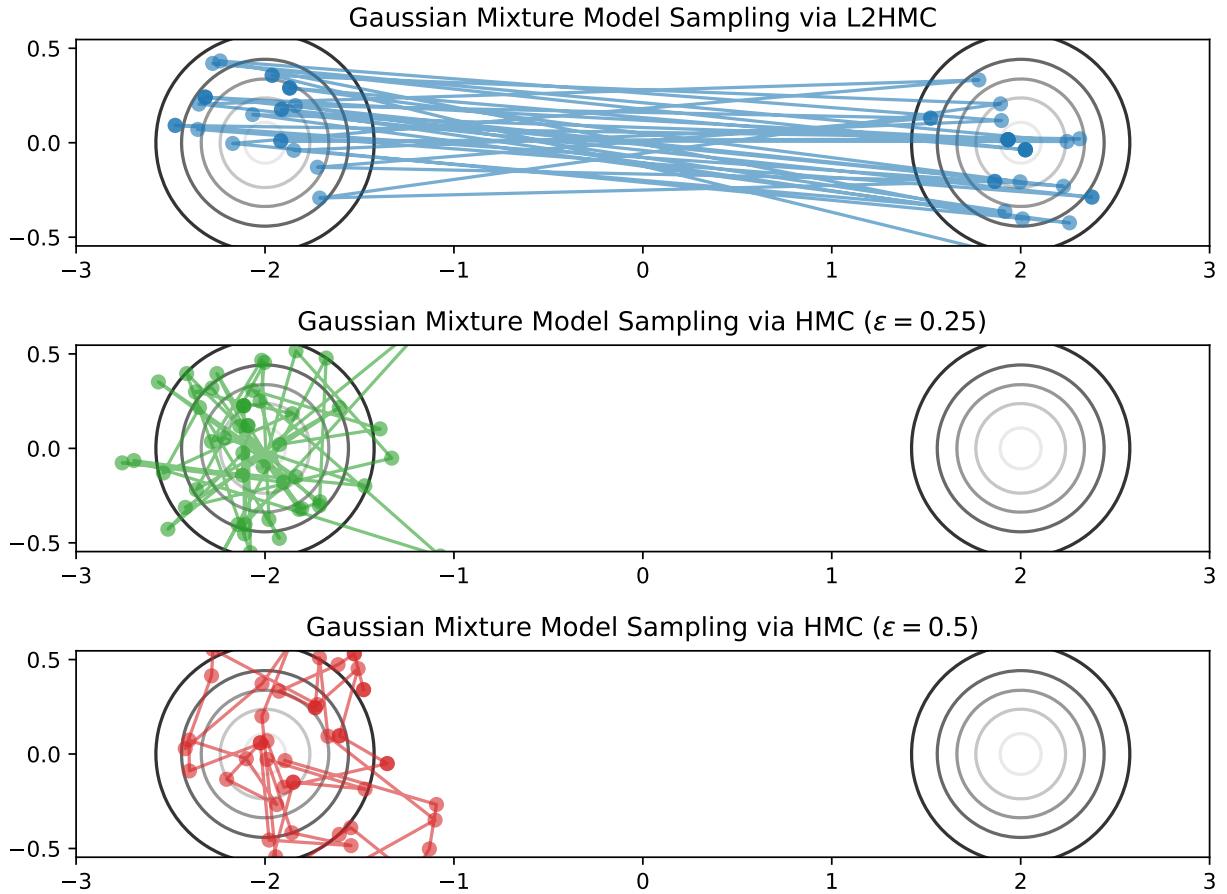


Figure 5.4: Comparison of trajectories generated using L2HMC (top), and traditional HMC with $\varepsilon = 0.25$ (middle) and $\varepsilon = 0.5$ (bottom). Note that L2HMC is able to successfully mix between modes, whereas HMC is not.

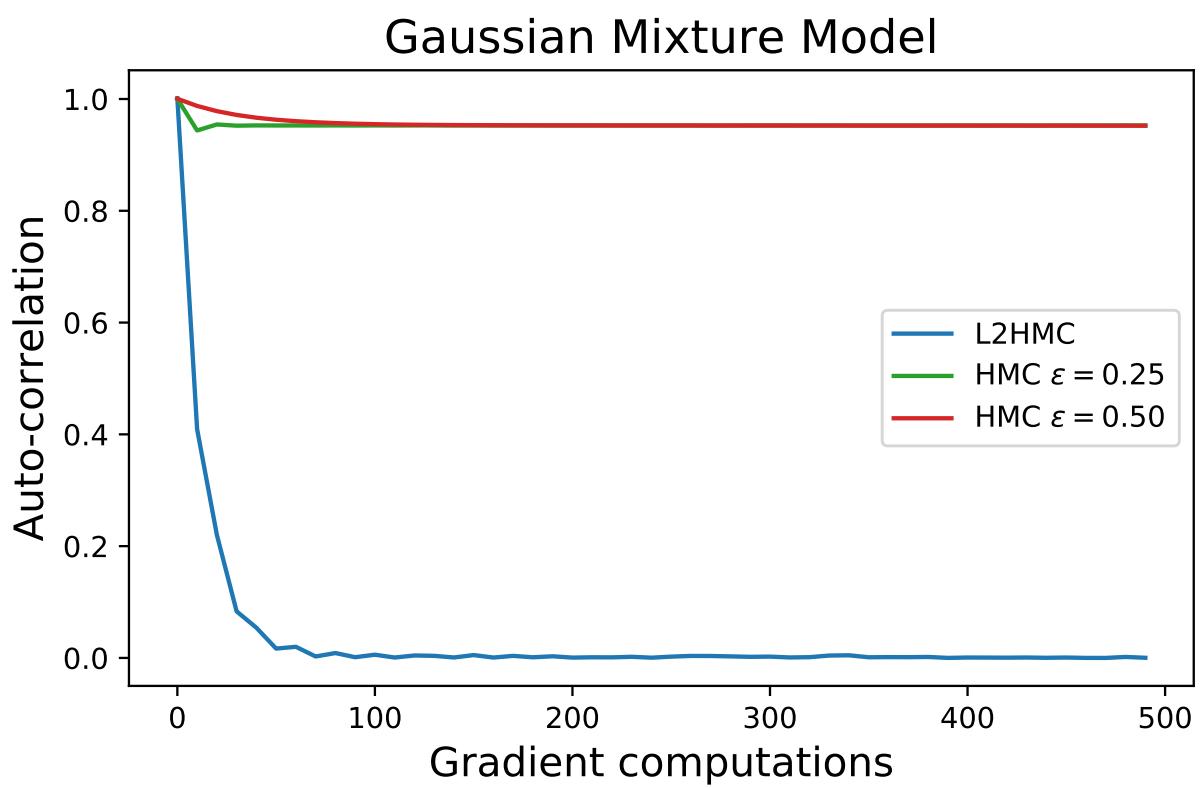


Figure 5.5: Autocorrelation vs. gradient evaluations (i.e. MD steps). Note that L2HMC (blue) has a significantly reduced autocorrelation after the same number of gradient evaluations when compared to either of the two HMC trajectories

5.9 2D $U(1)$ Lattice Gauge Theory

All lattice QCD simulations are performed at finite lattice spacing a and need an extrapolation to the continuum in order to be used for computing values of physical quantities. More reliable extrapolations can be done by simulating the theory at increasingly smaller lattice spacings. The picture that results when the lattice spacing is reduced and the physics kept constant is that all finite physical quantities of negative mass dimension diverge if measured in lattice units. In statistical mechanics language, this states that the continuum limit is a critical point of the theory since correlation lengths diverge. MCMC algorithms are known to encounter difficulties when used for simulating theories close to a critical point, an issue known as the *critical slowing down* of the algorithm. This effect is most prominent in the topological charge, whose auto-correlation time increases dramatically with finer lattice spacings. As a result, there is a growing interest in developing new sampling techniques for generating equilibrium configurations. In particular, algorithms that are able to offer improvements in efficiency through a reduction of statistical autocorrelations are highly desired. We begin with the two-dimensional $U(1)$ lattice gauge theory with dynamical variables $U_\mu(i)$ defined on the links of a lattice, where i labels a site and μ specifies the direction. Each link $U_\mu(i)$ can be expressed in terms of an angle $-\pi < \phi_\mu(i) \leq \pi$.

$$U_\mu(i) = e^{i\phi_\mu(i)} \quad (5.31)$$

with the Wilson action defined as:

$$\beta S = \beta \sum_P (1 - \cos(\phi_P)) \quad (5.32)$$

where

$$\phi_P \equiv \phi_{\mu\nu}(i) = \phi_\mu(i) + \phi_\nu(i + \hat{\mu}) - \phi_\mu(i + \hat{\nu}) - \phi_\nu(i) \quad (5.33)$$

and $\beta = 1/e^2$ is the gauge coupling, and the sum \sum_P runs over all plaquettes of the lattice. An illustration showing how these variables are defined for an elementary plaquette is shown in Fig. 5.6.

We can define the topological charge, $Q \in \mathbb{Z}$, as

$$Q \equiv \frac{1}{2\pi} \sum_P \tilde{\phi}_P = \frac{1}{2\pi} \sum_{i;\mu,\nu} \tilde{\phi}_{\mu\nu}(i) \quad (5.34)$$

where

$$\tilde{\phi}_P \equiv \phi_P - 2\pi \left\lfloor \frac{\phi_P + \pi}{2\pi} \right\rfloor \quad (5.35)$$

is the sum of the link variables around the elementary plaquette, projected onto the interval $[-\pi, \pi]$. From

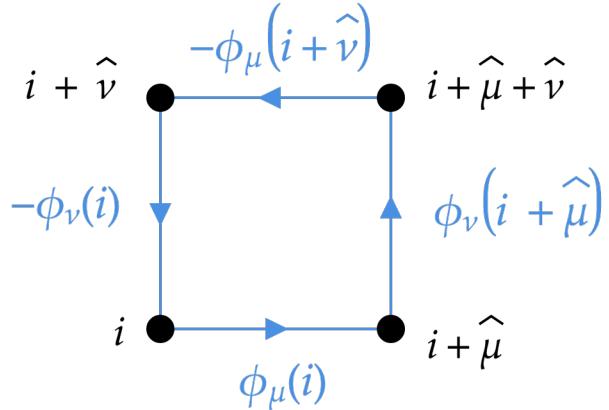


Figure 5.6: Illustration of an elementary plaquette on the lattice.

this, we can define topological susceptibility

$$\chi \equiv \frac{\langle Q^2 \rangle - \langle Q \rangle^2}{V} \quad (5.36)$$

By parity symmetry, $\langle Q \rangle = 0$, so we have that

$$\chi = \frac{\langle Q^2 \rangle}{V} \quad (5.37)$$

Unfortunately, the measurement of χ is often difficult due to the fact that the autocorrelation time with respect to Q tends to be extremely long. This is a consequence of the fact that the Markov chain tends to get stuck in a topological sector (characterized by $Q = \text{const.}$), a phenomenon known as *topological freezing*.

5.9.1 Modified Network Architecture

In order to better account for the rectangular geometry of the lattice, a stack of convolutional layers was prepended to the existing architecture, and can be seen in Fig. 5.8. The output from this convolutional structure is then fed to the generic network shown in Fig. 5.2. Additionally, the network architecture was modified to include a batch normalization layer after the second MaxPool layer. Introducing batch normalization is a commonly used technique in practice, and is known to help prevent against diverging gradients¹, (an issue that was occasionally encountered during the training procedure). Additionally, it has been shown to improve model performance and generally requires fewer training steps to achieve similar performance as models trained without it [53].

¹a numerical issue in which infinite values are generated when calculating the gradients in backpropagation

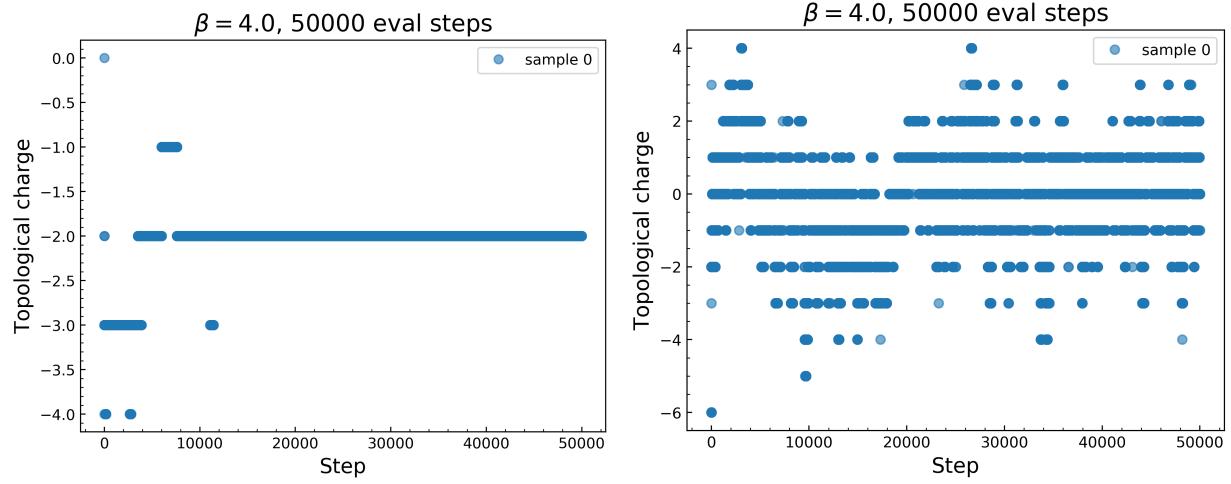


Figure 5.7: (left) Example of topological freezing in the 2D $U(1)$ lattice gauge theory, generated from generic HMC sampling for a 16×16 lattice. Note that for the majority of the simulation $Q = -2$, making it virtually impossible to get a reasonable estimate of χ . (right) Topological charge vs. step generated using the trained L2HMC sampler.

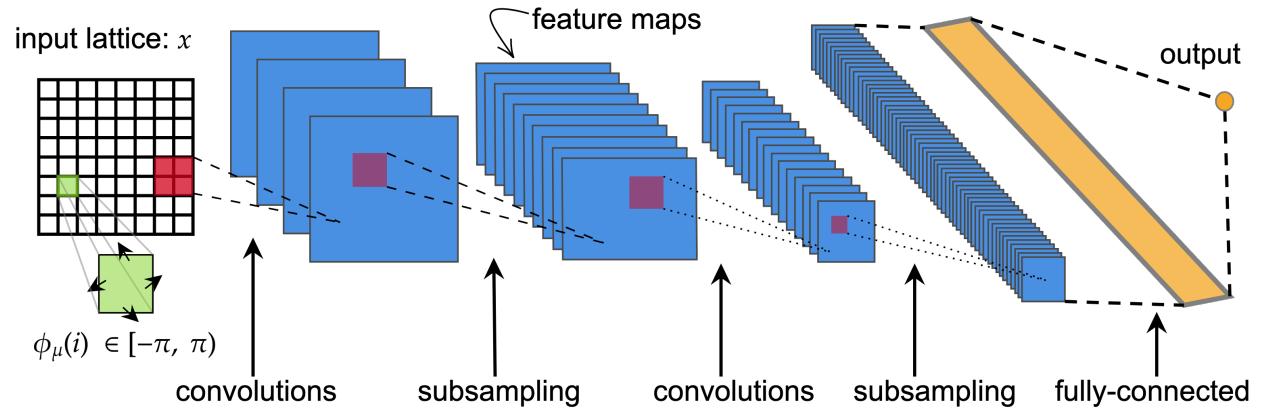


Figure 5.8: Convolutional structure used for learning localized features of rectangular lattice.

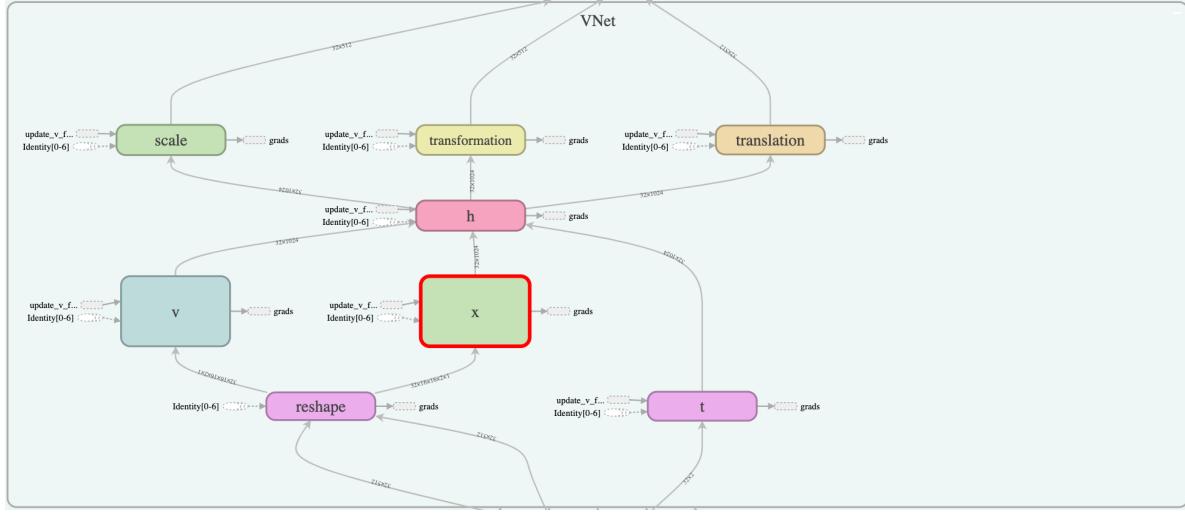


Figure 5.9: Illustration taken from TensorBoard showing an overview of the network architecture for VNet. Note that the architecture is identical for XNet.

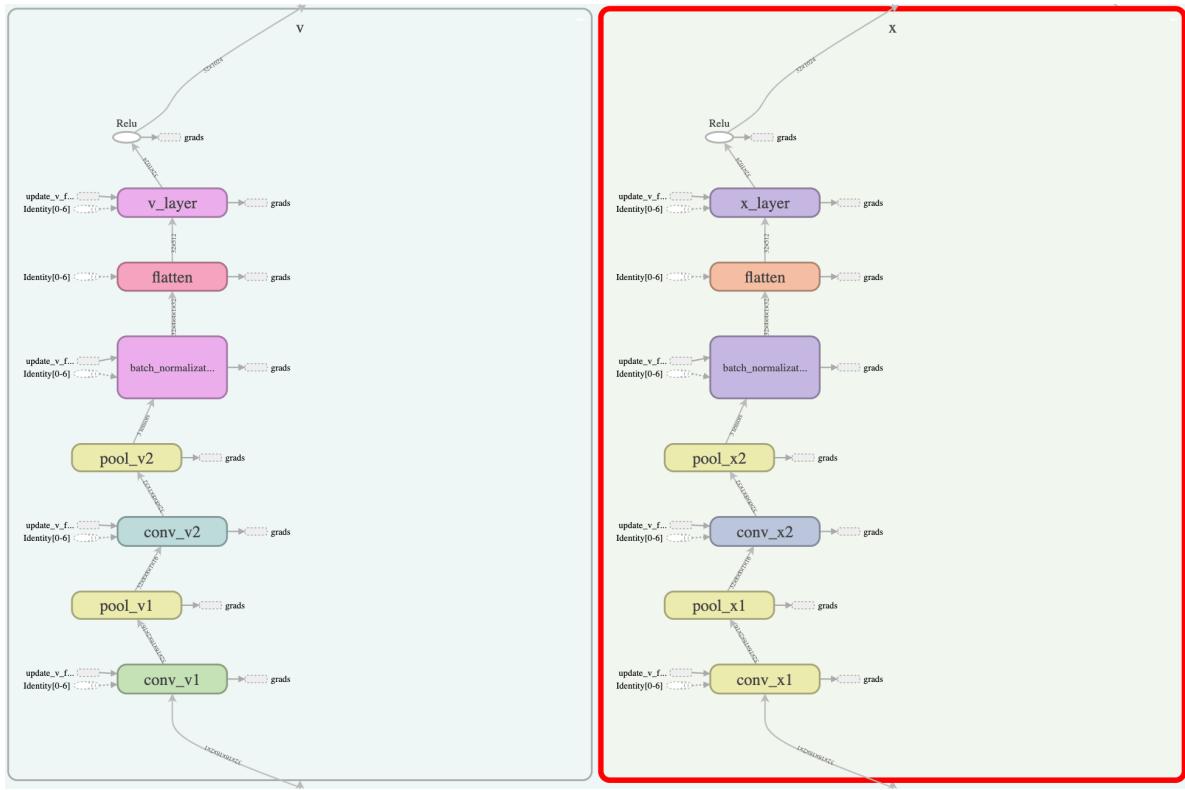


Figure 5.10: Detailed view of additional convolutional structure included to better account for rectangular geometry of lattice inputs.

5.9.2 Annealing Schedule

Proceeding as in the example of the Gaussian Mixture Model, we include a simulated annealing schedule in which the value of the gauge coupling β is continuously updated according to the annealing schedule shown in Eq. 5.38. This was done in order to encourage sampling from multiple different topological charge sectors, since our sampler is less ‘restricted’ at lower values of β .

$$\frac{1}{\beta(n)} = \left(\frac{1}{\beta_i} - \frac{1}{\beta_f} \right) \left(\frac{1-n}{N_{\text{train}}} \right) + \frac{1}{\beta_f} \quad (5.38)$$

Here $\beta(n)$ denotes the value of β to be used for the n^{th} training step ($n = 1, \dots, N_{\text{train}}$), β_i represents the initial value of β at the beginning of the training, and β_f represents the final value of β at the end of training. For a typical training session, $N_{\text{train}} = 25,000$, $\beta_i = 2$ and $\beta_f = 5$.

5.9.3 Modified Loss Function for $U(1)$ Gauge Model

In order to more accurately define the “distance” between two different lattice configurations, we redefine the metric in Eq. 5.24 to be

$$\delta(\xi, \xi') \equiv 1 - \cos(\xi - \xi') \quad (5.39)$$

where now $\xi \equiv (\phi_\mu^x(i), \phi_\mu^v(i), d)$, with ϕ_μ^x representing the lattice of (‘position’) gauge variables (what we called x previously), and ϕ_μ^v representing the lattice of (‘momentum’) gauge variables (what we called v previously). Note that i runs over all lattice sites² and $\mu = 0, 1$ for the two dimensional case. We see that this metric gives the expected behavior, since $\delta \rightarrow 0$ for $\xi \approx \xi'$.

While this new metric helps to better measure distances in this configuration space, it does nothing to encourage the exploration of different topological sectors since there may be configurations for which $\delta(\xi, \xi') \approx 1$ but $Q(\xi) = Q(\xi')$. In order to potentially address this issue, we modify the original loss function as follows.

First define $\xi' \equiv \mathbf{FL}_\theta \xi$ as the resultant configuration proposed by the augmented leapfrog integrator, and

$$\delta_Q(\xi, \xi') = |Q(\xi) - Q(\xi')| \quad (5.40)$$

$$\ell_Q(\xi, \xi', A(\mathbf{FL}_\theta \xi | \xi)) = \delta_Q(\xi, \xi') \times A(\xi' | \xi). \quad (5.41)$$

²In what follows, we will refrain from explicitly including the site index and make the assumption that it implicitly extends over all sites on the lattice.

So we have that δ_Q measures the difference in topological charge between the initial and proposed configurations, and ℓ_Q gives the expected topological charge difference. Proceeding as before, we include an additional auxiliary term which is identical in structure to the one above, except the input is now a configuration of link variables ϕ_μ drawn from the initialization distribution q , which for our purposes was chosen to be the standard random normal distribution on $[0, 2\pi]$.

We can then write the topological loss term as

$$\mathcal{L}_Q(\theta) \equiv \mathbb{E}_{p(\xi)} [\ell_Q(\xi, \mathbf{FL}_\theta \xi, A(\mathbf{FL}_\theta \xi | \xi))] + \alpha_{\text{aux}} \mathbb{E}_{q(\xi)} [\ell_Q(\xi, \mathbf{FL}_\theta \xi, A(\mathbf{FL}_\theta \xi | \xi))] \quad (5.42)$$

If we denote the standard loss (with the modified metric function) defined in Eq. 5.26 as $\mathcal{L}_{\text{std}}(\theta)$, we can write the new total loss as a combination of these two terms,

$$\mathcal{L}(\theta) = \alpha_{\text{std}} \mathcal{L}_{\text{std}}(\theta) + \alpha_Q \mathcal{L}_Q(\theta) \quad (5.43)$$

where $\alpha_{\text{std}}, \alpha_Q$ are multiplicative factors that weigh the relative contributions to the total loss from the standard and topological loss terms respectively, and α_{aux} in Eq. 5.42 weighs the contribution of configurations drawn from the initialization distribution.

5.9.4 Issues with the Average Plaquette

Upon further testing, an issue was encountered in which the average plaquette $\langle \phi_P \rangle$ seems to converge to a value which is noticeably different from the expected value in the infinite volume limit. This behavior can be seen in Fig. 5.11, and seems to depend on both the number of augmented leapfrog steps used by our integrator, as well as the ‘strength’ of the topological loss term in Eq. 5.43. The parameters used in Fig 5.11 are as follows: $L = 8$, $N_{\text{LF}} = 7$, $\alpha_Q = 0.5$, and $N_{\text{samples}} = 128$, $N_{\text{train}} = 1 \times 10^4$. In Fig 5.12, the only change was the weight factor for the topological charge term in the loss function $\alpha_Q = 0$. In order to quantify this unexpected behavior, we can calculate the difference between the observed value of the average plaquette, $\langle \phi_P \rangle$ and the expected value $\phi_P^{(*)}$:

$$\delta_{\phi_P}(\alpha_Q, N_{\text{LF}}) \equiv \langle \phi_P \rangle - \phi_P^{(*)} \neq 0 \quad (5.44)$$

Which allows us to measure the severity of this discrepancy.

Initially it was believed that this behavior was due to the topological charge term in the loss function, however after testing using $\alpha_Q = 0$, this behavior was still present. Following this initial test, it was discovered

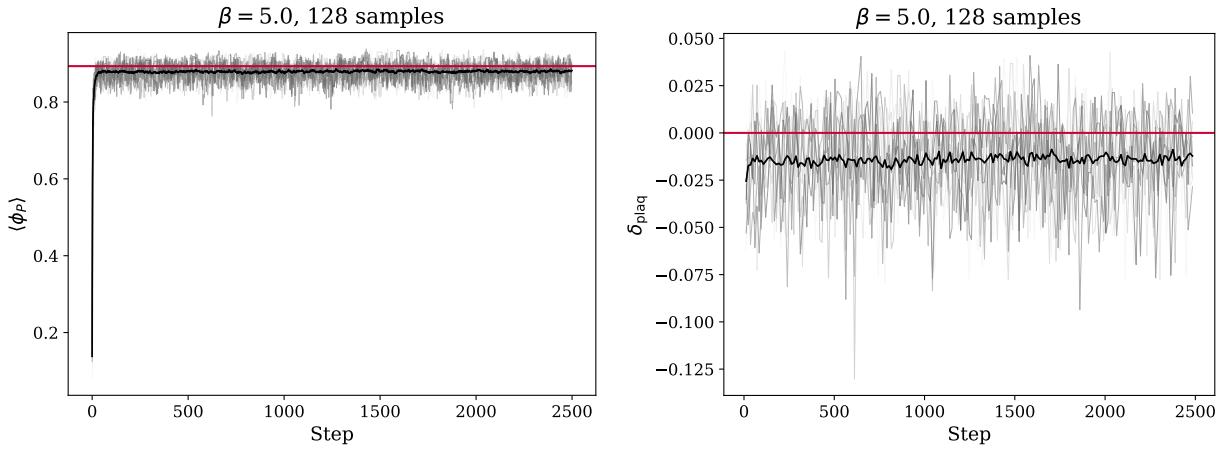


Figure 5.11: (left): Average plaquette $\langle \phi_P \rangle$ vs. step for $L = 8$, $N_{\text{LF}} = 7$, and $\alpha_Q = 0.5$. Here the solid red line indicates the true value of the average plaquette (in the infinite volume limit, and is equal to $0.89338\dots$ (right): Difference between the observed and expected value of the average plaquette δ_{ϕ_P} vs. step. Note that $\delta_{\phi_P} \neq 0$.

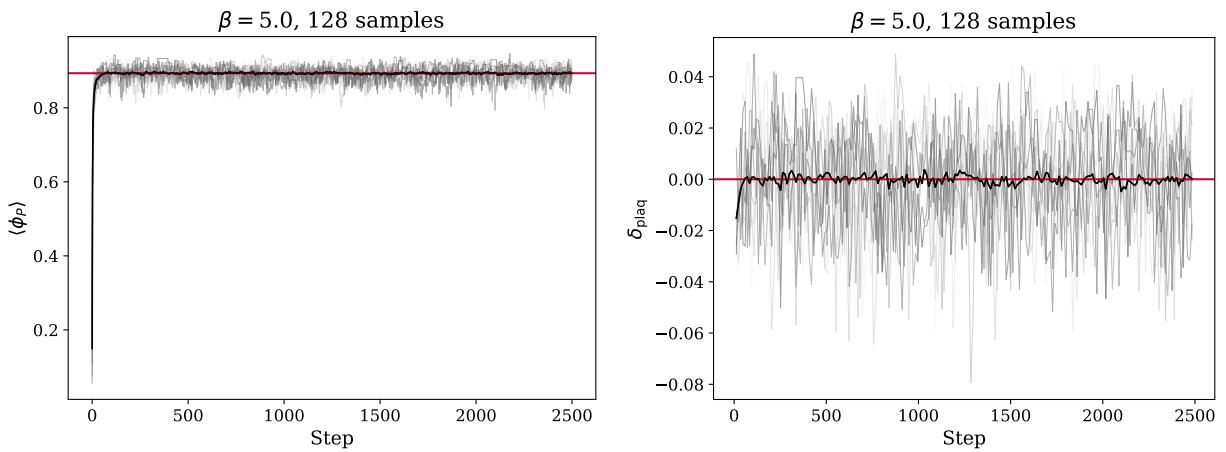


Figure 5.12: Same quantities as in Fig 5.11, with $\alpha_Q = 0$. Note that the discrepancy δ_{ϕ_P} is no longer present.

that the discrepancy seemed to be more prevalent when using a larger number of leapfrog steps N_{LF} . In an additional systematic attempt to debug the problem, the sampler was trained and evaluated multiple times over a range of different values of both N_{LF} and α_Q . These results are included in Appendix B. Frustratingly, the issue seemed to be almost irreproducible. Running the training/evaluation processes multiple times using the exact same set of initial parameters on the exact same hardware it was observed that sometimes the discrepancy was present while other times it was not.

5.10 Conclusion

In conclusion, we have seen how the Hamiltonian Monte Carlo algorithm follows from the Metropolis-Hastings algorithm used in generic Markov Chain Monte Carlo methods, and why this approach is often preferred when attempting to sample from the high-dimensional distributions characteristic of lattice gauge theory models.

Following this, a detailed description of the Hamiltonian Monte Carlo algorithm was presented, as well as some of the common issues faced when it is applied to lattice quantum chromodynamics. In order to address some of these issues, we then proceeded to introduce a learned inference architecture that successfully generalized HMC by augmenting the traditional leapfrog integrator with a set of carefully-chosen functions which are parameterized by weights in a neural network. This system is then trained to learn a MCMC kernel that encourages fast mixing and convergence to the target distribution. While this transition kernel is no longer symplectic, we are able to retain the strong theoretical guarantee of HMC by enforcing a tractable MH accept/reject step, making L2HMC potentially capable of sampling from very challenging distributions. Having introduced the necessary background, we then looked at applying this algorithm to the notoriously difficult two-dimensional Gaussian mixture model. In doing so, we found that the trained L2HMC sampler was capable of successfully mixing between modes in a way that traditional HMC was not. Additionally, we looked at the autocorrelation spectra of samples generated from the trained L2HMC sampler compared to those obtained from HMC. It was observed that the L2HMC sampler was able to produce samples which were noticeably less correlated in far fewer steps, indicating that the L2HMC algorithm significantly outperforms traditional HMC. In lattice QCD simulations, the ability of a sampler to quickly produce uncorrelated samples is one of the most important metrics for measuring its efficiency, and the approach outlined in this chapter shows promise in reducing the amount of computational resources required to generate new lattice configurations.

It remains to be seen how this algorithm performs when applied to more complicated models (e.g. models with fermions and non-Abelian gauge theories), a direction I plan to investigate more carefully in

future research. The other, seemingly non-fundamental issue with this approach is the discrepancy between the observed and expected value of the average plaquette. As of now, I am inclined to believe that this is more of a ‘bug’ than an inherent problem with the algorithm itself.

MCMC methods have proven to be indispensable in exploring new physics beyond what can be achieved analytically, and has significantly advanced our understanding of quantum field theory and quantum chromodynamics. Unfortunately, in order to both perform and extract meaningful results from these simulations, tremendous amounts of computational resources are required, with cutting-edge simulations being carried out almost exclusively on some of the worlds largest supercomputers. Because of this, even minor improvements in efficiency can dramatically reduce the amount of computational power required, allowing for increasingly complex models to be studied. The pursuit of better, more efficient algorithms is one of the major long-term goals of the lattice community, and is directly aligned with many of the goals outlined for high energy physics in the era of exascale computing. In particular, the results of these simulations are of central importance to the experiments being carried out at the Relativistic Heavy Ion Collider at Brookhaven National Laboratory (BNL), and the Large Hadron Collider (LHC) at CERN.

A | Appendix: L2HMC Source Code

Included below is the source code used for building, training, running and analyzing the L2HMC algorithm on the 2D $U(1)$ lattice gauge theory model. This code is also publicly hosted at <https://github.com/saforem2/l2hmc-qcd>.

A.1 README.md

A.1.1 l2hmc-qcd

Application of the L2HMC algorithm to simulations in lattice QCD. A description of the L2HMC algorithm can be found in the paper:

Generalizing Hamiltonian Monte Carlo with Neural Network
by Daniel Levy, Matt D. Hoffman and Jascha Sohl-Dickstein

A.1.1.1 Overview

NOTE: There are compatibility issues with `tensorflow.__version__ > 1.12`. To be sure everything runs correctly, make sure `tensorflow==1.12.x` is installed.

Given an analytically described distribution¹, L2HMC enables training of fast-mixing samplers.

A.1.1.2 Modified Implementation for Lattice Gauge Theory / Lattice QCD Models

This work is based on the original implementation which can be found at brain-research/l2hmc/.

My current focus is on applying this algorithm to simulations in lattice gauge theory and lattice QCD, in hopes of obtaining greater efficiency compared to generic HMC.

This new implementation includes the algorithm as applied to the 2D $U(1)$ lattice gauge theory model (i.e. compact QED). Additionally, this implementation includes a convolutional neural network architecture that is prepended to the network described in the original paper. The purpose of this additional structure is to better incorporate information about the geometry of the lattice.

Lattice code can be found in `l2hmc-qcd/lattice/` and the particular code for the 2D $U(1)$ lattice gauge model can be found in `l2hmc-qcd/lattice/lattice.py`.

A.1.1.3 Features

This model can be trained using distributed training through `horovod`, by passing the `--horovod` flag as a command line argument.

¹simple examples can be found in `l2hmc-qcd/utils/distributions.py`

A.1.1.4 Organization

To run `l2hmc-qcd/gauge_model_main.py` using one of the `.txt` files found in `l2hmc-qcd/args`, simply pass the `*.txt` file as the only command line argument prepended with `@`. For example, from within the `l2hmc-qcd/args` directory:

```
python3 ../gauge_model_main.py @gauge_model_args.txt
```

All of the relevant command line options are well documented and can be found in:²

```
l2hmc-qcd/utils/parse_args.py
```

Or by running `python3 gauge_model_main.py --help`.

Model information can be found in `l2hmc-qcd/models/gauge_model.py` which is responsible for building the graph and creating all the relevant tensorflow operations for training and running the L2HMC sampler.

The code responsible for actually implementing the L2HMC algorithm is divided up between `l2hmc-qcd/dynamics/gauge_dynamics.py` and `l2hmc-qcd/network.py`.

The code responsible for performing the augmented leapfrog algorithm is implemented in the `GaugeDynamics` class defined in `l2hmc-qcd/dynamics/gauge_dynamics.py`.

There are multiple different neural network architectures defined in `l2hmc-qcd/network.py` and different architectures can be specified as command line arguments defined in:

```
l2hmc-qcd/utils/parse_args.py
```

`l2hmc-qcd/notebooks/` contains a random collection of jupyter notebooks that each serve different purposes and should be somewhat self explanatory.

A.1.1.5 Contact

Code author: Sam Foreman

Pull requests and issues should be directed to: [saforem2](#)

A.1.1.6 Citation

If you use this code, please cite the original paper:

```
@article{levy2017generalizing,  
    title={Generalizing Hamiltonian Monte Carlo with Neural Networks},  
    author={Levy, Daniel and Hoffman, Matthew D. and Sohl-Dickstein, Jascha},  
    journal={arXiv preprint arXiv:1711.09268},  
    year={2017}  
}
```

²sample values for these arguments can be found in `l2hmc-qcd/args/gauge_model_args.txt`

A.2 l2hmc-qcd/args/args.txt

```

1 #####=====
2 #                                     #
3 #   To pass the arguments defined below to `gauge_model.py`:   #
4 #       python3 gauge_model.py @gauge_model_args.txt           #
5 #                                     #
6 #####=====
7 --time_size 8
8 --space_size 8
9 --rand
10 --num_samples 100
11 --num_steps 16
12 --eps 0.2
13 --beta_init 2.
14 --beta_final 5.
15 --lr_init 0.001
16 --lr_decay_steps 1000
17 --lr_decay_rate 0.96
18 --train_steps 1000
19 --logging_steps 10
20 --save_steps 1000
21 --print_steps 1
22 --network_arch 'generic'
23 --num_hidden1 50
24 --num_hidden2 50
25 --metric 'cos_diff'
26 --aux_weight 1.
27 --std_weight 1.
28 --charge_weight 0.
29 --loss_scale 0.1
30 --save_lf
31 #####=====
32 #--float64
33 #--zero_translation
34 #--warmup_lr
35 #--use_bn
36 #--dropout_prob 0.5
37 #--nnhmcloss
38 #--inference
39 #--loop_net_weights
40
41
42 #####=====
43 #             Description of command line arguments defined above      #
44 #####=====

45 # time_size:          temporal extent of lattice      #
46 # space_size:         spatial extent of lattice      #
47 # num_samples:        number of samples per batch    #
48 # num_steps:          number of leapfrog steps in MD update  #
49 # eps:                step size in leapfrog integrator  #
50 # annealing:          continuously anneal beta analogous to temperature)#
51 # beta_init:          initial value of beta used in annealing schedule  #
52 # beta_final:         final value of beta in annealing schedule  #
53 # lr_init:            initial value of learning rate  #
54 # lr_decay_steps:     steps after which lr is decayed by lr_decay_rate  #
55 # lr_decay_rate:      factor by which lr is decayed  #
56 # train_steps:        number of training steps to perform  #
57 # train_steps:        number of evaluation steps to run trained sampler for  #
58 # save_steps:         when to save model checkpoint  #
59 # print_steps:        how often to print information during training  #
60 # training_samples_steps: when to run sampler during training  #
61 # training_samples_length: how long sampler is ran for during training  #
62 # network_arch:       network architecture ('conv3D', 'conv2D' , 'generic')  #
63 # metric:             must be one of 'l1', 'l2', 'cos2', 'cos_diff'  #
64 # aux:                include auxiliary variable in loss function  #
65 # aux_weight:         factor multiplying aux. var contribution to loss fn  #
66 # std_weight:         factor multiplying std. var contribution to loss fn  #
67 # charge_weight:      factor multiplying charge. var contribution to loss fn  #

```

```
68 # loss_scale:           factor by which the loss is multiplied          #
69 # plaq_loss:            include top. charge difference in the loss fn      #
70 # summaries:            create summary objects for tensorboard             #
71 # eps_trainable:        allows the step size to be a trainable parameter    #
72 #####
```

A.3 l2hmc-qcd/main.py

```
1 """
2 gauge_model_main.py
3
4 Main method implementing the L2HMC algorithm for a 2D U(1) lattice gauge theory
5 with periodic boundary conditions.
6
7 Following an object oriented approach, there are separate classes responsible
8 for each major part of the algorithm:
9
10 (1.) Creating the loss function to be minimized during training and
11     building the corresponding TensorFlow graph.
12
13     - This is done using the `GaugeModel` class, found in
14       `models/gauge_model.py`.
15
16     - The `GaugeModel` class depends on the `GaugeDynamics` class
17       (found in `dynamics/gauge_dynamics.py`) that performs the augmented
18       leapfrog steps outlined in the original paper.
19
20 (2.) Training the model by minimizing the loss function over both the
21     target and initialization distributions.
22     - This is done using the `GaugeModelTrainer` class, found in
23       `trainers/gauge_model_trainer.py`.
24
25 (3.) Running the trained sampler to generate statistics for lattice
26     observables.
27     - This is done using the `GaugeModelRunner` class, found in
28       `runners/gauge_model_runner.py`.
29
30 Author: Sam Foreman (github: @saforem2)
31 Date: 04/10/2019
32 """
33 from __future__ import absolute_import
34 from __future__ import division
35 from __future__ import print_function
36
37 import os
38 import random
39 import time
40 import pickle
41 import inference
42 import tensorflow as tf
43 import numpy as np
44
45 from tensorflow.python import debug as tf_debug # noqa: F401
46 from tensorflow.client import timeline # noqa: F401
47 from tensorflow.core.protobuf import rewriter_config_pb2
48
49 import utils.file_io as io
50
51 from config import (
52     GLOBAL_SEED, NP_FLOAT, HAS_HOROVOD, HAS_COMET, HAS_MATPLOTLIB
53 )
54 from update import set_precision
55 from utils.parse_args import parse_args
56 from models.model import GaugeModel
57 from loggers.train_logger import TrainLogger
58 from trainers.trainer import GaugeModelTrainer
59
60 if HAS_COMET:
61     from comet_ml import Experiment
62
63 if HAS_HOROVOD:
64     import horovod.tensorflow as hvd
65
66 # if HAS_MATPLOTLIB:
67 #     import matplotlib.pyplot as plt
```

```

68
69 if float(tf.__version__.split('.')[0]) <= 2:
70     tf.logging.set_verbosity(tf.logging.INFO)
71
72 SEP_STR = 80 * '-' # + '\n'
73
74 # -----
75 # Set random seeds for tensorflow and numpy
76 # -----
77 os.environ['PYTHONHASHSEED'] = str(GLOBAL_SEED)
78 random.seed(GLOBAL_SEED)          # `python` build-in pseudo-random generator
79 np.random.seed(GLOBAL_SEED)      # numpy pseudo-random generator
80 tf.set_random_seed(GLOBAL_SEED)
81
82
83 def create_config(params):
84     """Helper method for creating a tf.ConfigProto object."""
85     config = tf.ConfigProto(allow_soft_placement=True)
86     if params['time_size'] > 8:
87         off = rewriter_config_pb2.RewriterConfig.OFF
88         config.graph_options.rewrite_options = off
89         config.graph_options.arithmetic_optimization = off
90
91     if params['gpu']:
92         # Horovod: pin GPU to be used to process local rank
93         # (one GPU per process)
94         config.gpu_options.allow_growth = True
95         # config.allow_soft_placement = True
96         if HAS_HOROVOD and params['horovod']:
97             config.gpu_options.visible_device_list = str(hvd.local_rank())
98
99     if HAS_MATPLOTLIB:
100        params['_plot'] = True
101
102    if params['theta']:
103        params['_plot'] = False
104        io.log("Training on Theta @ ALCF...")
105        params['data_format'] = 'channels_last'
106        os.environ["KMP_BLOCKTIME"] = str(0)
107        os.environ["KMP_AFFINITY"] = (
108            "granularity=fine,verbose,compact,1,0"
109        )
110        # NOTE: KMP affinity taken care of by passing -cc depth to aprun call
111        OMP_NUM_THREADS = 62
112        config.allow_soft_placement = True
113        config.intra_op_parallelism_threads = OMP_NUM_THREADS
114        config.inter_op_parallelism_threads = 0
115
116    return config, params
117
118
119
120 def latest_meta_file(checkpoint_dir=None):
121     """Returns the most recent meta-graph (`.meta`) file in checkpoint_dir."""
122     if not os.path.isdir(checkpoint_dir) or checkpoint_dir is None:
123         return
124
125     meta_files = [i for i in os.listdir(checkpoint_dir) if i.endswith('.meta')]
126     step_nums = [int(i.split('-')[-1].rstrip('.meta')) for i in meta_files]
127     step_num = sorted(step_nums)[-1]
128     meta_file = os.path.join(checkpoint_dir, f'model.ckpt-{step_num}.meta')
129
130     return meta_file
131
132
133 def count_trainable_params(out_file, log=False):
134     """Count the total number of trainable parameters in a tf.Graph object.
135
136     Args:
137         out_file (str): Path to file where all trainable parameters will be

```

```

138         written.
139     log (bool): Whether or not to print trainable parameters to console
140             (std-out).
141 Returns:
142     None
143 """
144 if log:
145     writer = io.log_and_write
146 else:
147     writer = io.write
148
149     io.log(f'Writing parameter counts to: {out_file}.')
150     writer(80 * '-', out_file)
151     total_params = 0
152     for var in tf.trainable_variables():
153         # shape is an array of tf.Dimension
154         shape = var.get_shape()
155         writer(f'var: {var}', out_file)
156         # var_shape_str = f' var.shape: {shape}'
157         writer(f' var.shape: {shape}', out_file)
158         writer(f' len(var.shape): {len(shape)}', out_file)
159         var_params = 1 # variable parameters
160         for dim in shape:
161             writer(f'    dim: {dim}', out_file)
162             # dim_strs += f'    dim: {dim}\'
163             var_params *= dim.value
164         writer(f'variable_parameters: {var_params}', out_file)
165         writer(80 * '-', out_file)
166         total_params += var_params
167
168     writer(80 * '-', out_file)
169     writer(f'Total parameters: {total_params}', out_file)
170
171 def train_setup(FLAGS, log_file=None):
172     io.log(SEP_STR)
173     io.log("Starting training using L2HMC algorithm...")
174     tf.keras.backend.clear_session()
175     tf.reset_default_graph()
176
177     # -----
178     # Parse command line arguments; copy key, val pairs from FLAGS to params.
179     # -----
180
181     try:
182         kv_pairs = FLAGS.__dict__.items()
183     except AttributeError:
184         kv_pairs = FLAGS.items()
185
186     params = {}
187     for key, val in kv_pairs:
188         params[key] = val
189         # if isinstance(val, bool) and not val: # skip if option is `off`
190         #     continue
191         # else:
192
193         params['log_dir'] = io.create_log_dir(FLAGS, log_file=log_file)
194         params['summaries'] = not FLAGS.no_summaries
195         if 'no_summaries' in params:
196             del params['no_summaries']
197
198         if FLAGS.save_steps is None and FLAGS.train_steps is not None:
199             params['save_steps'] = params['train_steps'] // 4
200
201         # if FLAGS.gpu:
202         #     params['data_format'] = 'channels_last'
203         #     # params['data_format'] = 'channels_first'
204         # else:
205         #     io.log("Using CPU for training.")
206         #     params['data_format'] = 'channels_last'
207

```

```

208     if FLAGS.float64:
209         io.log(f'INFO: Setting floating point precision to `float64`.')
210         set_precision('float64')
211
212     if FLAGS.horovod:
213         params['using_hvd'] = True
214         num_workers = hvd.size()
215         params['num_workers'] = num_workers
216
217         # -----
218         # Horovod: Scale initial lr by of num GPUs.
219         # ~~~~~
220         # NOTE: Even with a linear `warmup` of the learning rate,
221         #       the training remains unstable as evidenced by
222         #       exploding gradients and NaN tensors.
223         # ~~~~~
224         #   params['lr_init'] *= num_workers
225         # -----
226
227         # Horovod: adjust number of training steps based on number of GPUs.
228         #   params['train_steps'] // num_workers
229
230         # Horovod: adjust save_steps and lr_decay_steps accordingly.
231         #   params['save_steps'] // num_workers
232         #   params['lr_decay_steps'] // num_workers
233
234         #   if params['summaries']:
235         #       params['logging_steps'] // num_workers
236
237         # -----
238         # Horovod: BroadcastGlobalVariablesHook broadcasts initial
239         # variable states from rank 0 to all other processes. This
240         # is necessary to ensure consistent initialization of all
241         # workers when training is started with random weights or
242         # restored from a checkpoint.
243         # -----
244         hooks = [hvd.BroadcastGlobalVariablesHook(0)]
245     else:
246         params['using_hvd'] = False
247         hooks = []
248
249     return params, hooks
250
251
252 def train_l2hmc(FLAGS, log_file=None, experiment=None):
253     """Create, train, and run L2HMC sampler on 2D U(1) gauge model."""
254     tf.keras.backend.set_learning_phase(True)
255     params, hooks = train_setup(FLAGS, log_file)
256
257     # -----
258     # NOTE: Conditionals required for file I/O if we're not using
259     #       Horovod, `is_chief` should always be True otherwise,
260     #       if using Horovod, we only want to perform file I/O
261     #       on hvd.rank() == 0, so check that first.
262     # -----
263     condition1 = not params['using_hvd']
264     condition2 = params['using_hvd'] and hvd.rank() == 0
265     is_chief = condition1 or condition2
266
267     if is_chief:
268         # assert FLAGS.log_dir == params['log_dir']
269         log_dir = params['log_dir']
270         checkpoint_dir = os.path.join(log_dir, 'checkpoints/')
271         io.check_else_make_dir(checkpoint_dir)
272
273     else:
274         log_dir = None
275         checkpoint_dir = None
276
277     io.log(SEP_STR)

```

```

278     io.log('L2HMC PARAMETERS:')
279     for key, val in params.items():
280         io.log(f' {key}: {val}')
281     io.log(SEP_STR)
282
283     # -----
284     # Create model and train_logger
285     # -----
286     model = GaugeModel(params)
287
288     if is_chief:
289         train_logger = TrainLogger(model, log_dir, params['summaries'])
290     else:
291         train_logger = None
292
293     # -----
294     # Setup config and init_feed_dict for tf.train.Scaffold
295     # -----
296     config, params = create_config(params)
297
298     # set initial value of charge weight using value from FLAGS
299     charge_weight_init = params['charge_weight']
300     net_weights_init = [1., 1., 1.]
301     samples_init = np.reshape(np.array(model.lattice.samples, dtype=NP_FLOAT),
302                               (model.num_samples, model.x_dim))
303     beta_init = model.beta_init
304
305     init_feed_dict = {
306         model.x: samples_init,
307         model.beta: beta_init,                                # -----
308         model.charge_weight: charge_weight_init,             # NOTE
309         model.net_weights[0]: net_weights_init[0],           # ~ scale (S fn)
310         model.net_weights[1]: net_weights_init[1],           # ~ translation (T fn)
311         model.net_weights[2]: net_weights_init[2],           # ~ transformation (Q fn)
312         model.train_phase: True                            # -----
313     }
314
315     # ensure all variables are initialized
316     target_collection = []
317     if is_chief:
318         collection = tf.local_variables() + target_collection
319     else:
320         collection = tf.local_variables()
321
322     local_init_op = tf.variables_initializer(collection)
323     ready_for_local_init_op = tf.report_uninitialized_variables(collection)
324     init_op = tf.global_variables_initializer()
325
326     scaffold = tf.train.Scaffold(
327         init_feed_dict=init_feed_dict,
328         local_init_op=local_init_op,
329         ready_for_local_init_op=ready_for_local_init_op
330     )
331
332     # -----
333     # Create MonitoredTrainingSession
334     #
335     # NOTE: The MonitoredTrainingSession takes care of session
336     #       initialization, restoring from a checkpoint, saving to a
337     #       checkpoint, and closing when done or an error occurs.
338     # -----
339     sess_kwargs = {
340         'checkpoint_dir': checkpoint_dir,
341         'scaffold': scaffold,
342         'hooks': hooks,
343         'config': config,
344         'save_summaries_secs': None,
345         'save_summaries_steps': None
346     }
347

```

```

348     sess = tf.train.MonitoredTrainingSession(**sess_kwargs)
349     tf.keras.backend.set_session(sess)
350     sess.run(init_op)
351
352     # -----
353     #           TRAINING
354     # -----
355     trainer = GaugeModelTrainer(sess, model, train_logger)
356     train_kwargs = {
357         'samples_np': samples_init,
358         'beta_np': beta_init,
359         'net_weights': net_weights_init
360     }
361
362     trainer.train(model.train_steps, **train_kwargs)
363
364     if HAS_COMET and experiment is not None:
365         experiment.log_parameters(params)
366         g = sess.graph
367         experiment.set_model_graph(g)
368
369     params_file = os.path.join(os.getcwd(), 'params.pkl')
370     with open(params_file, 'wb') as f:
371         pickle.dump(model.params, f)
372
373     # Count all trainable parameters and write them out (w/ shapes) to txt file
374     count_trainable_params(os.path.join(params['log_dir'],
375                                         'trainable_params.txt'))
376
377     # close MonitoredTrainingSession and reset the default graph
378     sess.close()
379     tf.reset_default_graph()
380
381     return model, train_logger
382
383
384 def main(FLAGS):
385     """Main method for creating/training/running L2HMC for U(1) gauge model."""
386     log_file = 'output_dirs.txt'
387
388     if HAS_HOROVOD and FLAGS.horovod:
389         io.log("INFO: USING HOROVOD")
390         hvd.init()
391
392     condition1 = not FLAGS.horovod
393     condition2 = FLAGS.horovod and hvd.rank() == 0
394     is_chief = condition1 or condition2
395
396     if FLAGS.comet and is_chief:
397         experiment = Experiment(api_key="r7rKF035BJuaY3KT1Tpj4adco",
398                                 project_name="l2hmc-qcd",
399                                 workspace="saforem2")
400         name = (f'{FLAGS.network_arch}_'
401                 f'lf{FLAGS.num_steps}_'
402                 f'batch{FLAGS.num_samples}_'
403                 f'qw{FLAGS.charge_weight}_'
404                 f'aux{FLAGS.aux_weight}')
405         experiment.set_name(name)
406
407     else:
408         experiment = None
409
410     if FLAGS.hmc:
411         # -----
412         #   run generic HMC
413         # -----
414         inference.run_hmc(FLAGS, log_file=log_file)
415     else:
416         # -----
417         #   train l2hmc sampler

```

```
418     # -----
419     model, train_logger = train_l2hmc(FLAGS, log_file, experiment)
420     if experiment is not None:
421         experiment.log_parameters(model.params)
422
423
424 if __name__ == '__main__':
425     args = parse_args()
426     t0 = time.time()
427
428     main(args)
429
430     io.log('\n\n' + SEP_STR)
431     io.log(f'Time to complete: {time.time() - t0:.4g}')
```

A.4 l2hmc-qcd/inference.py

```
1 """
2 gauge_model_inference.py
3
4 Runs inference using the trained L2HMC sampler contained in a saved model.
5
6 This is done by reading in the location of the saved model from a .txt file
7 containing the location of the checkpoint directory.
8
9 ~~~~~
10 NOTE:
11 -----
12 If `--plot_lf` CLI argument passed, create the
13 following plots:
14
15     * The metric distance observed between individual
16       leapfrog steps and complete molecular dynamics
17       updates.
18
19     * The determinant of the Jacobian for each leapfrog
20       step and the sum of the determinant of the Jacobian
21       (sumlogdet) for each MD update.
22
23
24 Author: Sam Foreman (github: @saforem2)
25 Date: 07/08/2019
26 """
27 from __future__ import absolute_import
28 from __future__ import division
29 from __future__ import print_function
30
31 import os
32 import time
33 import pickle
34 import tensorflow as tf
35 import numpy as np
36
37
38 import utils.file_io as io
39
40 from config import (PARAMS, NP_FLOAT, HAS_HOROVOD, HAS_MATPLOTLIB,
41                     HAS_MEMORY_PROFILER)
42 from update import set_precision
43 # from main import create_config
44 from tensorflow.core.protobuf import rewriter_config_pb2
45 # from gauge_model_main import create_config
46
47 # from utils.parse_args import parse_args
48 from utils.parse_inference_args import parse_args
49 from models.model import GaugeModel
50 from loggers.summary_utils import create_summaries
51 from loggers.run_logger import RunLogger
52 from plotters.gauge_model_plotter import GaugeModelPlotter
53 # from plotters.plot_utils import plot_plaq_diffs_vs_net_weights
54 from plotters.leapfrog_plotters import LeapfrogPlotter
55 from runners.runner import GaugeModelRunner
56
57 if HAS_HOROVOD:
58     import horovod.tensorflow as hvd
59
60 if HAS_MEMORY_PROFILER:
61     import memory_profiler
62
63 # if HAS_MATPLOTLIB:
64 #     import matplotlib.pyplot as plt
65
66 if float(tf.__version__.split('.')[0]) <= 2:
67     tf.logging.set_verbosity(tf.logging.INFO)
```

```

68
69
70 SEP_STR = 80 * '-' # + '\n'
71
72
73 def create_config(params):
74     """Helper method for creating a tf.ConfigProto object."""
75     config = tf.ConfigProto(allow_soft_placement=True)
76     if params['time_size'] > 8:
77         off = rewriter_config_pb2.RewriterConfig.OFF
78         config.graph_options.rewrite_options
79         config.graph_options.arithmetic_optimization = off
80
81     if params['gpu']:
82         # Horovod: pin GPU to be used to process local rank
83         # (one GPU per process)
84         config.gpu_options.allow_growth = True
85         # config.allow_soft_placement = True
86         if HAS_HOROVOD and params['horovod']:
87             config.gpu_options.visible_device_list = str(hvd.local_rank())
88
89     if HAS_MATPLOTLIB:
90         params['_plot'] = True
91
92     if params['theta']:
93         params['_plot'] = False
94         io.log("Training on Theta @ ALCF...")
95         params['data_format'] = 'channels_last'
96         os.environ["KMP_BLOCKTIME"] = str(0)
97         os.environ["KMP_AFFINITY"] = (
98             "granularity=fine,verbose,compact,1,0"
99         )
100    # NOTE: KMP affinity taken care of by passing -cc depth to aprun call
101    OMP_NUM_THREADS = 62
102    config.allow_soft_placement = True
103    config.intra_op_parallelism_threads = OMP_NUM_THREADS
104    config.inter_op_parallelism_threads = 0
105
106    return config, params
107
108
109 def parse_flags(FLAGS, log_file=None):
110     """Parse command line FLAGS and create dictionary of parameters."""
111     try:
112         kv_pairs = FLAGS.__dict__.items()
113     except AttributeError:
114         kv_pairs = FLAGS.items()
115
116     params = {}
117     for key, val in kv_pairs:
118         params[key] = val
119
120     params['log_dir'] = io.create_log_dir(FLAGS, log_file=log_file)
121     params['summaries'] = not FLAGS.no_summaries
122     if 'no_summaries' in params:
123         del params['no_summaries']
124
125     if FLAGS.save_steps is None and FLAGS.train_steps is not None:
126         params['save_steps'] = params['train_steps'] // 4
127
128     if FLAGS.float64:
129         io.log(f'INFO: Setting floating point precision to `float64`.')
130         set_precision('float64')
131
132     return params
133
134
135 def load_params(params_pkl_file=None, log_file=None):
136     if params_pkl_file is None:
137         params_pkl_file = os.path.join(os.getcwd(), 'params.pkl')

```

```

138
139     if os.path.isfile(params_pkl_file):
140         with open(params_pkl_file, 'rb') as f:
141             params = pickle.load(f)
142     else:
143         io.log(f'INFO: Unable to locate: {params_pkl_file}.\n'
144               f'INFO: Using default parameters '
145               f'(PARAMS defined in `config.py`.)')
146
147         params = PARAMS.copy()
148         params['log_dir'] = io.create_log_dir(params, log_file=log_file)
149
150     return params
151
152
153 def set_model_weights(model, dest='rand'):
154     """Randomize model weights."""
155     if dest == 'rand':
156         io.log('Randomizing model weights...')
157     elif 'zero' in dest:
158         io.log(f'Zeroing model weights...')
159
160     xnet = model.dynamics.x_fn
161     vnet = model.dynamics.v_fn
162
163     for xblock, vblock in zip(xnet.layers, vnet.layers):
164         for xlayer, vlayer in zip(xblock.layers, vblock.layers):
165             try:
166                 print(f'xlayer.name: {xlayer.name}')
167                 print(f'vlayer.name: {vlayer.name}')
168                 kx, bx = xlayer.get_weights()
169                 kv, bv = vlayer.get_weights()
170
171                 if dest == 'rand':
172                     kx_new = np.random.randn(*kx.shape)
173                     bx_new = np.random.randn(*bx.shape)
174                     kv_new = np.random.randn(*kv.shape)
175                     bv_new = np.random.randn(*bv.shape)
176
177                 elif 'zero' in dest:
178                     kx_new = np.zeros(kx.shape)
179                     bx_new = np.zeros(bx.shape)
180                     kv_new = np.zeros(kv.shape)
181                     bv_new = np.zeros(bv.shape)
182
183                 xlayer.set_weights([kx_new, bx_new])
184                 vlayer.set_weights([kv_new, bv_new])
185             except ValueError:
186                 print(f'Unable to set weights for: {xlayer.name}')
187                 print(f'Unable to set weights for: {vlayer.name}')
188
189     return model
190
191
192 def log_plaq_diffs(run_logger, net_weights_arr, avg_plaq_diff_arr):
193     """Log the average values of the plaquette differences.
194
195     NOTE: If inference was performed with either the `--loop_net_weights`
196          or `--loop_transl_weights` flags passed, we want to see how the
197          difference between the observed and expected value of the average
198          plaquette varies with different values of the net weights, so save
199          this data to `*.pkl` file and plot the results.
200
201     """
202     pd_tup = [
203         (nw, md) for nw, md in zip(net_weights_arr, avg_plaq_diff_arr)
204     ]
205     pd_out_file = os.path.join(run_logger.log_dir, 'plaq_diffs_data.pkl')
206     with open(pd_out_file, 'wb') as f:
207         pickle.dump(pd_tup, f)
208
209
210 def log_mem_usage(run_logger, m_arr):

```

```

208     """Log memory usage."""
209     m_pkl_file = os.path.join(run_logger.log_dir, 'memory_usage.pkl')
210     m_txt_file = os.path.join(run_logger.log_dir, 'memory_usage.txt')
211     with open(m_pkl_file, 'wb') as f:
212         pickle.dump(m_arr, f)
213     with open(m_txt_file, 'w') as f:
214         for i in m_arr:
215             _ = [f.write(f'{j}\n') for j in i]
216
217
218 def inference_setup(kwargs):
219     """Set up relevant (initial) values to use when running inference."""
220     if kwargs['loop_net_weights']: # loop over different values of [S, T, Q]
221         # net_weights_arr = np.zeros((9, 3), dtype=NP_FLOAT)
222         w = np.random.randn(3) + 1.
223         net_weights_arr = np.array([[0.00, 0.00, 0.00], # set weights to 0.
224                                     # -----
225                                     [0.00, 0.00, 0.25], # loop over Q weights
226                                     [0.00, 0.00, 0.50],
227                                     [0.00, 0.00, 0.75],
228                                     [0.00, 0.00, 1.00],
229                                     [0.00, 0.00, 2.00],
230                                     # -----
231                                     [0.00, 0.25, 0.00], # loop over T weights
232                                     [0.00, 0.50, 0.00],
233                                     [0.00, 0.75, 0.00],
234                                     [0.00, 1.00, 0.00],
235                                     [0.00, 2.00, 0.00],
236                                     # -----
237                                     [0.25, 0.00, 0.00], # loop over S weights
238                                     [0.50, 0.00, 0.00],
239                                     [0.75, 0.00, 0.00],
240                                     [1.00, 0.00, 0.00],
241                                     [2.00, 0.00, 0.00],
242                                     # -----
243                                     [0.00, 1.00, 1.00], # [ , T, Q]
244                                     [1.00, 0.00, 1.00], # [S, , Q]
245                                     [1.00, 1.00, 0.00], # [S, T, ]
246                                     # -----
247                                     [w[0], w[1], w[2]], # randomize weights
248                                     # -----
249                                     [1.00, 1.00, 1.00]]) # set weights to 1.
250
251     elif kwargs['loop_transl_weights']:
252         net_weights_arr = np.array([[1.0, 0.00, 1.0],
253                                     [1.0, 0.10, 1.0],
254                                     [1.0, 0.25, 1.0],
255                                     [1.0, 0.50, 1.0],
256                                     [1.0, 0.75, 1.0],
257                                     [1.0, 1.00, 1.0]], dtype=NP_FLOAT)
258
259     else: # set [S, T, Q] = [1, 1, 1]
260         net_weights_arr = np.array([[1, 1, 1]], dtype=NP_FLOAT)
261
262     # if a value has been passed in `kwargs['beta_inference']` use it
263     # otherwise, use `model.beta_final`
264     beta_final = kwargs.get('beta_final', None)
265     beta_inference = kwargs.get('beta_inference', None)
266     beta = beta_final if beta_inference is None else beta_inference
267     # betas = [beta_final if beta_inference is None else beta_inference]
268
269     inference_dict = {
270         'net_weights_arr': net_weights_arr,
271         'beta': beta,
272         # 'charge_weight': kwargs.get('charge_weight', 1.),
273         'run_steps': kwargs.get('run_steps', 5000),
274         'plot_lf': kwargs.get('plot_lf', False),
275         'loop_net_weights': kwargs['loop_net_weights'],
276         'loop_transl_weights': kwargs['loop_transl_weights']
277     }

```

```

278     return inference_dict
280
281
282     def run_hmc(FLAGS, log_file=None):
283         """Run inference using generic HMC."""
284         condition1 = not FLAGS.horovod
285         condition2 = FLAGS.horovod and hvd.rank() == 0
286         is_chief = condition1 or condition2
287         if not is_chief:
288             return -1
289
290         FLAGS.hmc = True
291         FLAGS.log_dir = io.create_log_dir(FLAGS, log_file=log_file)
292
293         params = parse_flags(FLAGS)
294         params['hmc'] = True
295         params['use_bn'] = False
296         params['log_dir'] = FLAGS.log_dir
297         params['loop_net_weights'] = False
298         params['loop_transl_weights'] = False
299         params['plot_lf'] = False
300
301         figs_dir = os.path.join(params['log_dir'], 'figures')
302         io.check_else_make_dir(figs_dir)
303
304         io.log(SEP_STR)
305         io.log('HMC PARAMETERS:')
306         for key, val in params.items():
307             io.log(f' {key}: {val}')
308         io.log(SEP_STR)
309
310         config, params = create_config(params)
311         tf.reset_default_graph()
312
313         model = GaugeModel(params=params)
314
315         sess = tf.Session(config=config)
316         sess.run(tf.global_variables_initializer())
317
318         run_ops = tf.get_collection('run_ops')
319         inputs = tf.get_collection('inputs')
320         run_summaries_dir = os.path.join(model.log_dir, 'summaries', 'run')
321         io.check_else_make_dir(run_summaries_dir)
322         # create summary objects without having to train model
323         _, _ = create_summaries(model, run_summaries_dir, training=False)
324         run_logger = RunLogger(params, inputs, run_ops, save_lf_data=False)
325         plotter = GaugeModelPlotter(params, run_logger.figs_dir)
326
327         inference_dict = inference_setup(params)
328
329         # -----
330         # Create GaugeModelRunner for inference
331         # -----
332         runner = GaugeModelRunner(sess, params, inputs, run_ops, run_logger)
333         run_inference(inference_dict, runner, run_logger, plotter)
334
335
336     def inference(runner, run_logger, plotter, **kwargs):
337         """Perform an inference run, if it hasn't been ran previously.
338
339         Args:
340             runner: GaugeModelRunner object, responsible for performing inference.
341             run_logger: RunLogger object, responsible for running `tf.summary`
342                         operations and accumulating/saving run statistics.
343             plotter: GaugeModelPlotter object responsible for plotting lattice
344                     observables from inference run.
345
346         NOTE: If inference hasn't been run previously with the param values passed,
347               return `avg_plaq_diff`, i.e. the average value of the difference

```

```

348     between the expected and observed value of the average plaquette.
349
350     Returns:
351         avg_plaq_diff: If run hasn't been completed previously, else None
352         """
353     run_steps = kwargs.get('run_steps', 5000)
354     beta = kwargs.get('beta', 5.)
355     kwargs['eps'] = runner.eps
356
357     run_str = run_logger._get_run_str(**kwargs)
358     kwargs['run_str'] = run_str
359
360     if not run_logger.existing_run(run_str):
361         run_logger.reset(**kwargs)
362         t0 = time.time()
363         runner.run(**kwargs)
364         io.log(SEP_STR)
365         run_time = time.time() - t0
366         io.log(SEP_STR + f'\nTook: {run_time}s to complete run.\n' + SEP_STR)
367         avg_plaq_diff = plotter.plot_observables(run_logger.run_data, **kwargs)
368
369     if kwargs.get('plot_lf', False):
370         lf_plotter = LeapfrogPlotter(plotter.out_dir, run_logger)
371         num_samples = runner.params.get('num_samples', 20)
372         lf_plotter.make_plots(run_logger.run_dir,
373                               num_samples=num_samples)
374
375     return avg_plaq_diff
376
377 else:
378     nw = kwargs.get('net_weights', [1., 1., 1.])
379     io.log(SEP_STR)
380     io.log('\n Inference has already been completed for:\n'
381           f'\t net_weights: [{nw[0]}, {nw[1]}, {nw[2]}]\n'
382           f'\t run_steps: {run_steps}\n'
383           f'\t eps: {runner.eps}\n'
384           f'\t beta: {beta}\n'
385           f' Continuing...\n')
386     io.log(SEP_STR)
387
388
389 def run_inference(runner, run_logger=None, plotter=None, **kwargs):
390     """Run inference.
391
392     Args:
393         inference_dict: Dictionary containing parameters to use for inference.
394         runner: GaugeModelRunner object that actually performs the inference.
395         run_logger: RunLogger object that logs observables and other data
396             generated during inference.
397         plotter: GaugeModelPlotter object responsible for plotting observables
398             generated during inference.
399         """
400     if plotter is None or run_logger is None:
401         return
402
403     dir_append = kwargs.get('dir_append', None)
404     net_weights_arr = kwargs.get('net_weights_arr', [1., 1., 1.])
405     avg_plaq_diff_arr = []
406     m_arr = []
407     for net_weights in net_weights_arr:
408         if HAS_MEMORY_PROFILER:
409             usage0 = memory_profiler.memory_usage()
410             m_arr.append(usage0)
411             run_logger.clear()
412             if HAS_MEMORY_PROFILER:
413                 usage1 = memory_profiler.memory_usage()
414                 m_arr.append(usage1)
415         kwargs.update({
416             'net_weights': net_weights,
417             'dir_append': dir_append

```

```

418         })
419     inference(runner, run_logger, plotter, **kwargs)
420
421     log_mem_usage(run_logger, m_arr)
422     if kwargs['loop_net_weights']:
423         log_plaq_diffs(run_logger, net_weights_arr, avg_plaq_diff_arr)
424
425
426 def main_inference(inference_kwargs):
427     """Perform inference using saved model."""
428     params_file = inference_kwargs.get('params_file', None)
429     params = load_params(params_file) # load params used during training
430
431     # ~~~~~
432     # NOTE: We want to restrict all communication (file I/O) to only be
433     # performed on rank 0 (i.e. `is_chief`) so there are two cases:
434     #   1. We're using Horovod, so we have to check hvd.rank() explicitly.
435     #   2. We're not using Horovod, in which case `is_chief` is always True.
436     # -----
437     condition1 = not params['using_hvd']
438     condition2 = params['using_hvd'] and hvd.rank() == 0
439     is_chief = condition1 or condition2
440     if not is_chief:
441         return
442     # ~~~~~
443
444     checkpoint_dir = os.path.join(params['log_dir'], 'checkpoints/')
445     assert os.path.isdir(checkpoint_dir)
446     checkpoint_file = tf.train.latest_checkpoint(checkpoint_dir)
447
448     config, params = create_config(params)
449     sess = tf.Session(config=config)
450     saver = tf.train.import_meta_graph(f'{checkpoint_file}.meta')
451     saver.restore(sess, checkpoint_file)
452
453     run_ops = tf.get_collection('run_ops')
454     inputs = tf.get_collection('inputs')
455
456     run_logger = RunLogger(params, inputs, run_ops, save_lf_data=False)
457     plotter = GaugeModelPlotter(params, run_logger.figs_dir)
458
459     # ~~~~~
460     # Set up relevant values to use for inference (parsed from kwargs)
461     #
462     # NOTE: We are only interested in the command line arguments that
463     #       were passed to `inference.py` (i.e. those contained in kwargs)
464     inference_kwargs['loop_net_weights'] = True
465     inference_dict = inference_setup(inference_kwargs)
466     if inference_dict['beta'] is None:
467         inference_dict['beta'] = params['beta_final']
468     # ~~~~~
469
470     net_weights_file = os.path.join(params['log_dir'], 'net_weights.txt')
471     np.savetxt(net_weights_file, inference_dict['net_weights_arr'],
472               delimiter=' ', newline='\n', fmt='%-4g')
473
474     # ~~~~~
475     # Create GaugeModelRunner for inference
476     #
477     runner = GaugeModelRunner(sess, params, inputs, run_ops, run_logger)
478     run_inference(runner, run_logger, plotter, **inference_dict)
479
480     ##########
481     # NOTE: THE FOLLOWING WONT WORK WHEN RESTORING FROM CHECKPOINT (FOR
482     #       INFERENCE) UNLESS `GaugeModel` IS ENTIRELY REBUILT:
483     # -----
484     # set 'net_weights_arr' = [1., 1., 1.] so each O, S, T contribute
485     # inference_dict['net_weights_arr'] = np.array([[1, 1, 1]],
486     #                                              dtype=NP_FLOAT)
487     # set 'betas' to be a single value

```

```

488 # inference_dict['betas'] = inference_dict['betas'][-1]
489 #
490 #
491 # # randomize the model weights and run inference using these weights
492 # runner.model = set_model_weights(runner.model, dest='rand')
493 # run_inference(inference_dict,
494 #                 runner, run_logger,
495 #                 plotter, dir_append='_rand')
496 #
497 # # zero the model weights and run inference using these weights
498 # runner.model = set_model_weights(runner.model, dest='zero')
499 # run_inference(inference_dict,
500 #                 runner, run_logger,
501 #                 plotter, dir_append='_zero')
502 #####
```

A.5 l2hmc-qcd/models/model.py

```
1 """
2 gauge_model.py
3
4 Implements GaugeModel class responsible for building computation graph used in
5 tensorflow.
6
7 -----
8 TODO (taken from [1.]):  

9
10 update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
11 ! no update ops in the default graph
12 io.log("update_ops: ", update_ops)
13 Use the update ops of the model itself
14 io.log("model.updates: ", self.dynamics.updates)
15 -----
16 PREVIOUS (WORKING but diverging grads as of 07/15/2019):
17 """
18 with tf.control_dependencies(self.dynamics.updates):
19     self.train_op = self.optimizer.apply_gradients(
20         grads_and_vars,
21         global_step=self.global_step,
22         name='train_op'
23     )
24     """
25     self.train_op = tf.group(minimize_op,
26                             self.dynamics.updates)
27     try:
28         self.train_op = self._append_update_ops(train_op)
29     update_ops = [self.dynamics.updates,
30                  tf.get_collection(tf.GraphKeys.UPDATE_OPS)]
31     tf.add_to_collection(tf.GraphKeys.UPDATE_OPS,
32                         *[i for i in self.dynamics.updates])
33 -----
34
35 Author: Sam Foreman (github: @saforem2)
36 Date: 04/12/2019
37 """
38 from __future__ import division
39 from __future__ import print_function
40 from __future__ import absolute_import
41
42 import os
43 import functools
44
45 import numpy as np
46 import tensorflow as tf
47 import utils.file_io as io
48
49 from lattice.lattice import GaugeLattice
50 from utils.horovod_utils import warmup_lr
51 from dynamics.dynamics import GaugeDynamics
52 from config import GLOBAL_SEED, TF_FLOAT, TF_INT, PARAMS, HAS_HOROVOD
53 from tensorflow.python.ops import control_flow_ops as control_flow_ops
54
55 if HAS_HOROVOD:
56     import horovod.tensorflow as hvd
57
58
59 def check_log_dir(log_dir):
60     """Check log_dir for existing checkpoints."""
61     assert os.path.isdir(log_dir)
62     ckpt_dir = os.path.join(log_dir, 'checkpoints')
63     if os.path.isdir(ckpt_dir):
64         pass
65
66
67 class GaugeModel:
```

```

68     def __init__(self, params=None):
69         # -----
70         # Create attributes from (key, val) pairs in params
71         # -----
72         self.loss_weights = {}
73         if params is None:
74             params = PARAMS # default parameters, defined in `variables.py`
75
76         self.params = params
77         for key, val in self.params.items():
78             if 'weight' in key and key != 'charge_weight':
79                 self.loss_weights[key] = val
80             elif key == 'charge_weight':
81                 pass
82             else:
83                 setattr(self, key, val)
84
85         self.eps_trainable = not self.eps_fixed
86         self.charge_weight_np = params['charge_weight']
87
88         self.lattice = self._create_lattice()
89         self.batch_size = self.lattice.samples.shape[0]
90         self.x_dim = self.lattice.num_links
91
92         # build input placeholders for tensors
93         inputs = self._create_inputs()
94         _ = [setattr(self, key, val) for key, val in inputs.items()]
95
96         self.dynamics = self._create_dynamics()
97
98         # metric function used when calculating the loss
99         self.metric_fn = self._create_metric_fn(self.metric)
100
101        obs_ops = self._create_observables()
102        _ = [setattr(self, key, val) for key, val in obs_ops.items()]
103
104        self._build_sampler()
105        run_ops = self._build_run_ops()
106
107        if self.hmc:
108            train_ops = {}
109        else:
110            self._create_lr()
111            self._create_optimizer()
112            self._apply_grads()
113            train_ops = self._build_train_ops()
114
115        self.ops_dict = {
116            'run_ops': run_ops,
117            'train_ops': train_ops
118        }
119
120        for key, val in self.ops_dict.items():
121            _ = [tf.add_to_collection(key, op) for op in list(val.values())]
122
123    def _create_lattice(self):
124        """Create GaugeLattice object."""
125        with tf.name_scope('lattice'):
126            lattice = GaugeLattice(time_size=self.time_size,
127                                   space_size=self.space_size,
128                                   dim=self.dim,
129                                   link_type=self.link_type,
130                                   num_samples=self.num_samples,
131                                   rand=self.rand)
132
133        return lattice
134
135    def _create_dynamics(self, **kwargs):
136        """Initialize dynamics object."""
137        with tf.name_scope('dynamics'):

```

```

138     # default values of keyword arguments
139     keys = ['eps', 'hmc', 'network_arch',
140             'num_steps', 'use_bn', 'dropout_prob',
141             'num_hidden1', 'num_hidden2', 'zero_translation']
142     dynamics_kwargs = {k: getattr(self, k) for k in keys}
143     dynamics_kwargs['eps_trainable'] = not self.eps_fixed
144     # dynamics_kwargs = {
145     #     'eps': self.eps,
146     #     'hmc': self.hmc,
147     #     'network_arch': self.network_arch,
148     #     'num_steps': self.num_steps,
149     #     'eps_trainable': not self.eps_fixed,
150     #     # 'data_format': self.data_format,
151     #     'use_bn': self.use_bn,
152     #     'dropout_prob': self.dropout_prob,
153     #     'num_hidden1': self.num_hidden1,
154     #     'num_hidden2': self.num_hidden2,
155     #     'zero_translation': self.zero_translation
156     # }
157     dynamics_kwargs.update(kwargs)
158     samples = self.lattice.samples_tensor
159     potential_fn = self.lattice.get_potential_fn(samples)
160     dynamics = GaugeDynamics(lattice=self.lattice,
161                               potential_fn=potential_fn,
162                               **dynamics_kwargs)
163
164     return dynamics
165
166 def _create_observables(self):
167     """Create operations for calculating lattice observables."""
168     with tf.name_scope('observables'):
169         plaq_sums = self.lattice.calc_plaq_sums(self.x)
170         actions = self.lattice.calc_actions(plaq_sums=plaq_sums)
171         plaqs = self.lattice.calc_plaqs(plaq_sums=plaq_sums)
172         avg_plaqs = tf.reduce_mean(plaqs, name='avg_plaqs')
173         charges = self.lattice.calc_top_charges(plaq_sums=plaq_sums)
174
175     obs_ops = {
176         'plaq_sums_op': plaq_sums,
177         'actions_op': actions,
178         'plaqs_op': plaqs,
179         'avg_plaqs_op': avg_plaqs,
180         'charges_op': charges,
181     }
182
183     return obs_ops
184
185 def _create_inputs(self):
186     """Create input placeholders (if not executing eagerly).
187     Returns:
188         outputs: Dictionary with the following entries:
189             x: Placeholder for input lattice configuration with
190                 shape = (batch_size, x_dim) where x_dim is the number of
191                 links on the lattice and is equal to lattice.time_size *
192                 lattice.space_size * lattice.dim.
193             beta: Placeholder for inverse coupling constant.
194             charge_weight: Placeholder for the charge_weight (i.e. alpha_Q,
195                 the multiplicative factor that scales the topological
196                 charge term in the modified loss function) .
197             net_weights: Array of placeholders, each of which is a
198                 multiplicative constant used to scale the effects of the
199                 various S, Q, and T functions from the original paper.
200                 net_weights[0] = 'scale_weight', multiplies the S fn.
201                 net_weights[1] = 'transformation_weight', multiplies the Q
202                 fn. net_weights[2] = 'translation_weight', multiplies the
203                 T fn.
204             train_phase: Boolean placeholder indicating if the model is
205                 currently being trained.
206     """
207     with tf.name_scope('inputs'):

```

```

208     if not tf.executing_eagerly():
209         x = tf.placeholder(dtype=TF_FLOAT,
210                             shape=(self.batch_size, self.x_dim),
211                             name='x')
212
213         beta = tf.placeholder(dtype=TF_FLOAT,
214                               shape=(),
215                               name='beta')
216
217         charge_weight = tf.placeholder(dtype=TF_FLOAT,
218                                         shape=(),
219                                         name='charge_weight')
220
221         scale_weight = tf.placeholder(dtype=TF_FLOAT,
222                                       shape=(),
223                                       name='scale_weight')
224         transf_weight = tf.placeholder(dtype=TF_FLOAT,
225                                       shape=(),
226                                       name='transformation_weight')
227         transl_weight = tf.placeholder(dtype=TF_FLOAT,
228                                       shape=(),
229                                       name='translation_weight')
230
231         net_weights = [scale_weight, transf_weight, transl_weight]
232
233         train_phase = tf.placeholder(tf.bool, name='is_training')
234
235     else:
236         x = tf.convert_to_tensor(
237             self.lattice.samples.reshape((self.batch_size,
238                                         self.x_dim)))
239
240         beta = tf.convert_to_tensor(self.beta_init)
241         charge_weight = tf.convert_to_tensor(0.)
242         net_weights = tf.convert_to_tensor([1., 1., 1.])
243         train_phase = True
244
245     names = ['x', 'beta', 'charge_weight', 'train_phase', 'net_weights']
246     tensors = [x, beta, charge_weight, train_phase, net_weights]
247     outputs = dict(zip(names, tensors))
248
249     for name, tensor in outputs.items():
250         if name != 'net_weights':
251             tf.add_to_collection('inputs', tensor)
252         else:
253             _ = [tf.add_to_collection('inputs', t) for t in tensor]
254
255     return outputs
256
257 @staticmethod
258 def _create_metric_fn(metric):
259     """Create metric fn for measuring the distance between two samples."""
260     with tf.name_scope('metric_fn'):
261         if metric == 'l1':
262             def metric_fn(x1, x2):
263                 return tf.abs(x1 - x2)
264
265         elif metric == 'l2':
266             def metric_fn(x1, x2):
267                 return tf.square(x1 - x2)
268
269         elif metric == 'cos':
270             def metric_fn(x1, x2):
271                 return tf.abs(tf.cos(x1) - tf.cos(x2))
272
273         elif metric == 'cos2':
274             def metric_fn(x1, x2):
275                 return tf.abs(tf.cos(x1) - tf.cos(x2))
276
277         elif metric == 'cos_diff':

```

```

278     def metric_fn(x1, x2):
279         return 1. - tf.cos(x1 - x2)
280
281     elif metric == 'sin_diff':
282         def metric_fn(x1, x2):
283             return 1. - tf.sin(x1 - x2)
284
285     elif metric == 'tan_cos':
286         def metric_fn(x1, x2):
287             cos_diff = 1. - tf.cos(x1 - x2)
288             return tf.tan(np.pi * cos_diff / 2)
289     else:
290         raise AttributeError(f"""metric={metric}. Expected one of:
291             'l1', 'l2', 'cos', 'cos2', 'cos_diff', or
292             'tan_cos'.""")
293
294     return metric_fn
295
296 def __create_lr(self, lr_init=None):
297     """Create learning rate."""
298     if self.hmc:
299         return
300
301     if lr_init is None:
302         lr_init = self.lr_init
303
304     with tf.name_scope('global_step'):
305         self.global_step = tf.train.get_or_create_global_step()
306
307     with tf.name_scope('learning_rate'):
308         # HOROVOD: When performing distributed training, it can be useful
309         # to "warmup" the learning rate gradually, done using the
310         # `configure_learning_rate` method below..
311         if self.warmup_lr:
312             kwargs = {
313                 'target_lr': lr_init,
314                 'warmup_steps': int(0.1 * self.train_steps),
315                 'global_step': self.global_step,
316                 'decay_steps': self.lr_decay_steps,
317                 'decay_rate': self.lr_decay_rate,
318             }
319             self.lr = warmup_lr(**kwargs)
320         else:
321             self.lr = tf.train.exponential_decay(lr_init,
322                                                 self.global_step,
323                                                 self.lr_decay_steps,
324                                                 self.lr_decay_rate,
325                                                 staircase=False,
326                                                 name='learning_rate')
327
328 def __create_optimizer(self):
329     """Create learning rate and optimizer."""
330     if not hasattr(self, 'lr'):
331         self.__create_lr()
332
333     # with tf.control_dependencies(update_ops):
334     with tf.name_scope('optimizer'):
335         # update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
336         self.optimizer = tf.train.AdamOptimizer(self.lr)
337         if self.using_hvd:
338             self.optimizer = hvd.DistributedOptimizer(self.optimizer)
339
340     def __build_run_ops(self):
341         """Build run_ops dict containing grouped operations for inference."""
342         keys = ['x_out', 'px', 'actions_op',
343                'plaqs_op', 'avg_plaqs_op',
344                'charges_op', 'charge_diffs_op']
345         run_ops = {k: getattr(self, k) for k in keys}
346
347         if self.save_lf:

```

```

348     keys = ['lf_out', 'pxs_out', 'masks',
349              'logdets', 'sumlogdet', 'fns_out']
350
351     fkeys = [k + '_f' for k in keys]
352     bkeys = [k + '_b' for k in keys]
353
354     run_ops.update({k: getattr(self, k) for k in fkeys})
355     run_ops.update({k: getattr(self, k) for k in bkeys})
356
357     run_ops['dynamics_eps'] = self.dynamics.eps
358
359     return run_ops
360
361 def _build_train_ops(self):
362     """Build train_ops dict containing grouped operations for training."""
363     if self.hmc:
364         train_ops = {}
365
366     else:
367         keys = ['train_op', 'loss_op', 'x_out',
368                 'px', 'actions_op', 'plqs_op',
369                 'charges_op', 'charge_diffs_op', 'lr']
370         train_ops = {k: getattr(self, k) for k in keys}
371
372     train_ops['dynamics.eps'] = self.dynamics.eps
373
374     return train_ops
375
376 def _calc_std_loss(self, inputs, **weights):
377     """Calculate standard contribution to loss.
378
379     NOTE: In contrast to the original paper where the L2 difference was
380     used, we are now using  $1 - \cos(x_1 - x_2)$ .
381
382     Args:
383         x_tup: Tuple of (x, x_proposed) configurations.
384         z_tup: Tuple of (z, z_proposed) configurations.
385         p_tup: Tuple of (x, z) acceptance probabilities.
386         weights: dictionary of weights giving relative weight of each term
387             in the loss function.
388
389     Returns:
390         std_loss
391     """
392     eps = 1e-4
393     aux_weight = weights.get('aux_weight', 1.)
394     std_weight = weights.get('std_weight', 1.)
395
396     if std_weight == 0.: # don't bother calculating loss if weight is 0
397         return 0.
398
399     x_init = inputs['x_init']
400     x_proposed = inputs['x_proposed']
401     z_init = inputs['z_init']
402     z_proposed = inputs['z_proposed']
403     px = inputs['px']
404     pz = inputs['pz']
405
406     ls = self.loss_scale
407     with tf.name_scope('std_loss'):
408         with tf.name_scope('x_loss'):
409             x_std_loss = tf.reduce_sum(
410                 self.metric_fn(x_init, x_proposed), axis=1
411             )
412             x_std_loss *= px
413             # x_std_loss_avg = tf.reduce_sum(x_std_loss, axis=0,
414             #                                 name='x_std_loss_avg')
415
416             # Add eps for numerical stability
417             x_std_loss = tf.add(x_std_loss, eps, name='x_std_loss')

```

```

418     # tf.add_to_collection('losses', x_std_loss)
419
420     with tf.name_scope('z_loss'):
421         z_std_loss = tf.reduce_sum(
422             self.metric_fn(z_init, z_proposed), axis=1
423         )
424         z_std_loss *= pz * aux_weight
425         # z_std_loss_avg = tf.reduce_mean(z_std_loss, axis=0,
426         #                                     name='z_std_loss_avg')
427         z_std_loss = tf.add(z_std_loss, eps, name='z_std_loss')
428         # tf.add_to_collection('losses', z_std_loss)
429
430     with tf.name_scope('tot_loss'):
431         loss = (ls * (1. / x_std_loss + 1. / z_std_loss)
432                 - (x_std_loss + z_std_loss) / ls) * std_weight
433
434         std_loss = tf.reduce_mean(loss, axis=0, name='std_loss')
435         # tf.add_to_collection('losses', std_loss)
436
437     return std_loss # , x_std_loss_avg, z_std_loss_avg
438
439 def _calc_charge_loss(self, inputs, **weights):
440     """Calculate contribution to total loss from charge difference.
441
442     NOTE: This is an additional term introduced to the loss function that
443     measures the difference in the topological charge between the initial
444     configuration and the proposed configuration.
445
446     Args:
447         x_tup: Tuple of (x, x_proposed) configurations.
448         z_tup: Tuple of (z, z_proposed) configurations.
449         p_tup: Tuple of (x, z) acceptance probabilities.
450         weights: dictionary of weights giving relative weight of each term
451             in the loss function.
452
453     Returns:
454         std_loss
455     """
456     aw = weights.get('aux_weight', 1.)
457     qw = weights.get('charge_weight', 1.)
458
459     if qw == 0:
460         return 0.
461
462     x_init = inputs['x_init']
463     x_proposed = inputs['x_proposed']
464     z_init = inputs['z_init']
465     z_proposed = inputs['z_proposed']
466     px = inputs['px']
467     pz = inputs['pz']
468
469     # Calculate the difference in topological charge between the initial
470     # and proposed configurations multiplied by the probability of
471     # acceptance to get the expected value of the difference in topological
472     # charge
473     with tf.name_scope('charge_loss'):
474         with tf.name_scope('x_loss'):
475             x_dq = self.lattice.calc_top_charges_diff(x_init, x_proposed)
476             xq_loss = px * x_dq
477             # xq_loss_avg = tf.reduce_mean(xq_loss, axis=0,
478             #                               name='xq_loss_avg')
479
480         with tf.name_scope('z_loss'):
481             if aw > 0:
482                 z_dq = self.lattice.calc_top_charges_diff(z_init,
483                                              z_proposed)
484             else:
485                 z_dq = tf.zeros_like(x_dq)
486
487             zq_loss = aw * pz * z_dq

```

```

488     # zq_loss_avg = tf.reduce_mean(zq_loss, axis=0,
489     #                                 name='zq_loss_avg')
490
491
492     with tf.name_scope('tot_loss'):
493         # Each of the loss terms is scaled by the `loss_scale` which
494         # introduces a universal multiplicative factor that scales the
495         # value of the loss
496         charge_loss = - qw * (xq_loss + zq_loss) / self.loss_scale
497         charge_loss = tf.reduce_mean(charge_loss, axis=0,
498                               name='charge_loss')
499         # tf.add_to_collection('losses', charge_loss)
500
501     return charge_loss # , xq_loss_avg, zq_loss_avg
502
503 def calc_loss(self, x, beta, net_weights, train_phase, **weights):
504     """Create operation for calculating the loss.
505
506     Args:
507         x: Input tensor of shape (self.num_samples,
508             self.lattice.num_links) containing batch of GaugeLattice links
509             variables.
510         beta (float): Inverse coupling strength.
511
512     Returns:
513         loss (float): Operation responsible for calculating the total loss.
514         px (np.ndarray): Array of acceptance probabilities from
515             Metropolis-Hastings accept/reject step. Has shape:
516             (self.num_samples,)
517         x_out: Output samples obtained after Metropolis-Hastings
518             accept/reject step.
519
520     NOTE: If proposed configuration is accepted following
521         Metropolis-Hastings accept/reject step, x_proposed and x_out are
522         equivalent.
523     """
524     with tf.name_scope('x_update'):
525         x_dynamics_output = self.dynamics.apply_transition(x, beta,
526                                                       net_weights,
527                                                       train_phase,
528                                                       self.save_lf)
529         x_proposed = x_dynamics_output['x_proposed']
530         px = x_dynamics_output['accept_prob']
531         x_out = x_dynamics_output['x_out']
532
533     # Auxiliary variable
534     with tf.name_scope('z_update'):
535         z = tf.random_normal(tf.shape(x),
536                             dtype=TF_FLOAT,
537                             seed=GLOBAL_SEED,
538                             name='z')
539         z_dynamics_output = self.dynamics.apply_transition(z, beta,
540                                                       net_weights,
541                                                       train_phase,
542                                                       save_lf=False)
543         z_proposed = z_dynamics_output['x_proposed']
544         pz = z_dynamics_output['accept_prob']
545
546     with tf.name_scope('top_charge_diff'):
547         x_dq = tf.cast(
548             self.lattice.calc_top_charges_diff(x, x_out),
549             dtype=TF_INT
550         )
551
552     # NOTE:
553     # std_loss: 'standard' loss
554     # charge_loss: Contribution from the difference in topological charge
555     #   between the initial and proposed configurations to the total
556     #   loss.
557     inputs = {

```

```

558     'x_init': x,
559     'x_proposed': x_proposed,
560     'z_init': z,
561     'z_proposed': z_proposed,
562     'px': px,
563     'pz': pz
564 }
565
566 with tf.name_scope('calc_loss'):
567     std_loss = self._calc_std_loss(inputs, **weights)
568     charge_loss = self._calc_charge_loss(inputs, **weights)
569     total_loss = tf.add(std_loss, charge_loss, name='total_loss')
570
571 # losses = [*std_losses, *charge_losses]
572 # _ = [tf.add_to_collection('losses', i) for i in losses]
573
574 return total_loss, x_dq, x_dynamics_output
575
576 def calc_loss_and_grads(self,
577                         x,
578                         beta,
579                         net_weights,
580                         train_phase,
581                         **weights):
582     """Calculate loss its gradient with respect to all trainable variables.
583
584     Args:
585         x: Placeholder (tensor object)representing batch of GaugeLattice
586             link variables.
587         beta: Placeholder (tensor object) representing inverse coupling
588             strength.
589
590     Returns:
591         loss: Operation for calculating the total loss.
592         grads: Tensor containing the gradient of the loss function with
593             respect to all trainable variables.
594         x_out: Operation for obtaining new samples (i.e. output of
595             augmented L2HMC algorithm.)
596         accept_prob: Operation for calculating acceptance probabilities
597             used in Metropolis-Hastings accept reject.
598         x_dq: Operation for calculating the topological charge difference
599             between the initial and proposed configurations.
600     """
601     if tf.executing_eagerly():
602         with tf.name_scope('grads'):
603             with tf.GradientTape() as tape:
604                 loss, x_dq, dynamics_output = self.calc_loss(x,
605                                                               beta,
606                                                               net_weights,
607                                                               train_phase,
608                                                               **weights)
609
610                 grads = tape.gradient(loss, self.dynamics.trainable_variables)
611                 if self.clip_value > 0.:
612                     grads, _ = tf.clip_by_global_norm(grads, self.clip_value)
613             else:
614                 loss, x_dq, dynamics_output = self.calc_loss(x, beta,
615                                                               net_weights,
616                                                               train_phase,
617                                                               **weights)
618             with tf.name_scope('grads'):
619                 grads = tf.gradients(loss, self.dynamics.trainable_variables)
620                 if self.clip_value > 0.:
621                     grads, _ = tf.clip_by_global_norm(grads, self.clip_value)
622
623     return loss, grads, x_dq, dynamics_output
624
625 def _build_sampler(self):
626     """Build TensorFlow graph."""
627     with tf.name_scope('sampler'):

```

```

628     # self.loss_op, self.grads, self.x_out, self.px, x_dq = output
629     loss, grads, x_dq, dynamics_output = self.calc_loss_and_grads(
630         x=self.x,
631         beta=self.beta,
632         net_weights=self.net_weights,
633         train_phase=self.train_phase,
634         **self.loss_weights
635     )
636     self.loss_op = loss
637     self.grads = grads
638     self.x_out = dynamics_output['x_out']
639     self.px = dynamics_output['accept_prob']
640     self.charge_diffs_op = tf.reduce_sum(x_dq) / self.num_samples
641
642     if self.save_lf:
643         op_keys = ['masks_f', 'masks_b',
644                    'lf_out_f', 'lf_out_b',
645                    'pxs_out_f', 'pxs_out_b',
646                    'logdets_f', 'logdets_b',
647                    'fns_out_f', 'fns_out_b',
648                    'sumlogdet_f', 'sumlogdet_b']
649         for key in op_keys:
650             try:
651                 op = dynamics_output[key]
652                 setattr(self, key, op)
653             except KeyError:
654                 continue
655
656         with tf.name_scope('l2hmc_fns'):
657             self.l2hmc_fns = {
658                 'out_fns_f': self.extract_l2hmc_fns(self.fns_out_f),
659                 'out_fns_b': self.extract_l2hmc_fns(self.fns_out_b),
660             }
661
662     def extract_l2hmc_fns(self, fns):
663         """Method for extracting each of the Q, S, T functions as tensors."""
664         if not self.save_lf:
665             return
666
667         # fns has shape: (num_steps, 4, 3, num_samples, lattice.num_links)
668         fnsT = tf.transpose(fns, perm=[2, 1, 0, 3, 4], name='fns_transposed')
669
670         out_fns = {}
671         names = ['scale', 'translation', 'transformation']
672         subnames = ['v1', 'x1', 'x2', 'v2']
673         for idx, name in enumerate(names):
674             out_fns[name] = {}
675             for subidx, subname in enumerate(subnames):
676                 out_fns[name][subname] = fnsT[idx][subidx]
677
678         return out_fns
679
680     def _apply_grads(self):
681         """Build Tensorflow graph."""
682         # -----
683         # Ref. [1.] in TODO (line 8)
684         # -----
685         with tf.name_scope('train'):
686             grads_and_vars = zip(self.grads, self.dynamics.trainable_variables)
687             with tf.control_dependencies(self.dynamics.updates):
688                 self.train_op = self.optimizer.apply_gradients(
689                     grads_and_vars,
690                     global_step=self.global_step,
691                     name='train_op'
692             )

```

A.6 l2hmc-qcd/dynamics/gauge_dynamics.py

```
1 """
2 Dynamics engine for L2HMC sampler on Lattice Gauge Models.
3
4 Reference [Generalizing Hamiltonian Monte Carlo with Neural
5 Networks](https://arxiv.org/pdf/1711.09268.pdf)
6
7 Code adapted from the released TensorFlow graph implementation by original
8 authors https://github.com/brain-research/l2hmc.
9
10
11 TODO:
12     - Log separately the Q, S, T values for the 'x' and 'v' functions.
13     - Look at raw phase values both before and after mod operation.
14     - Try running generic net using 64-point precision.
15     - See if 64-point precision issues with Conv3D are fixed in tf 1.13-1.14.
16     - JLSE account.
17
18
19 Author: Sam Foreman (github: @saforem2)
20 Date: 1/14/2019
21 """
22 from __future__ import absolute_import
23 from __future__ import print_function
24 from __future__ import division
25
26 import numpy as np
27 import numpy.random as npr
28 import tensorflow as tf
29
30 from config import GLOBAL_SEED, TF_FLOAT, TF_INT
31 from network.network import FullNet
32
33
34 def exp(x, name=None):
35     """Safe exponential using tf.check_numerics."""
36     return tf.check_numerics(tf.exp(x), f'{name} is NaN')
37
38
39 def flatten_tensor(tensor):
40     """Flattens tensor along axes 1:, since axis=0 indexes sample in batch.
41
42     Example: for a tensor of shape [b, x, y, t] -->
43             returns a tensor of shape [b, x * y * t]
44     """
45     batch_size = tensor.shape[0]
46     return tf.reshape(tensor, shape=(batch_size, -1))
47
48
49 def hmc_network(inputs, train_phase):
50     return [tf.zeros_like(inputs[_]) for _ in range(3)]
51
52
53 def _add_to_collection(collection, ops):
54     if len(ops) > 1:
55         _ = [tf.add_to_collection(collection, op) for op in ops]
56     else:
57         tf.add_to_collection(collection, ops)
58
59
60 class GaugeDynamics(tf.keras.Model):
61     """Dynamics engine of naive L2HMC sampler."""
62
63     def __init__(self, lattice, potential_fn, **kwargs):
64         """Initialization.
65
66         Args:
67             lattice: Lattice object containing multiple sample lattices.
```

```

68     potential_fn: Function specifying minus log-likelihood objective
69     to minimize.
70
71     NOTE: kwargs (expected)
72         num_steps: Number of leapfrog steps to use in integrator.
73         eps: Initial step size to use in leapfrog integrator.
74         network_arch: String specifying network architecture to use.
75             Must be one of ``conv2D', `conv3D', `generic''. Networks
76             are defined in `../network/
77         hmc: Flag indicating whether generic HMC (no augmented
78             leapfrog) should be used instead of L2HMC. Defaults to
79             False.
80         eps_trainable: Flag indicating whether the step size (eps)
81             should be trainable. Defaults to True.
82         np_seed: Seed to use for numpy.random.
83     """
84     super(GaugeDynamics, self).__init__(name='GaugeDynamics')
85     np.random.seed(GLOBAL_SEED)
86
87     self.lattice = lattice
88     self.potential = potential_fn
89     self.batch_size = self.lattice.samples.shape[0]
90     self.x_dim = self.lattice.num_links
91     self.l2hmc_fns = {}
92
93     # create attributes from kwargs.items()
94     for key, val in kwargs.items():
95         if key != 'eps': # want to use self.eps as tf.Variable
96             setattr(self, key, val)
97
98         _eps_np = kwargs.get('eps', 0.4)
99         # with tf.variable_scope(reus):
100         #     self.eps = tf.get_variable('eps', dtype=TF_FLOAT,
101         #                               trainable=self.eps_trainable,
102         #                               initializer=tf.constant(_eps_np))
103
104         self.log_eps = tf.Variable(
105             initial_value=tf.log(tf.constant(_eps_np)),
106             name='log_eps',
107             dtype=TF_FLOAT,
108             trainable=self.eps_trainable
109         )
110
111         self.eps = tf.exp(self.log_eps, name='eps')
112         #     self.eps = tf.Variable(
113         #         initial_value=kwargs.get('eps', 0.4),
114         #         name='eps',
115         #         dtype=TF_FLOAT,
116         #         trainable=self.eps_trainable
117         #     )
118
119         self._construct_masks()
120
121     if self.hmc:
122         self.x_fn = lambda inp, train_phase: [
123             tf.zeros_like(inp[0]) for t in range(3)
124         ]
125         self.v_fn = lambda inp, train_phase: [
126             tf.zeros_like(inp[0]) for t in range(3)
127         ]
128
129     else:
130         num_filters = int(self.lattice.space_size)
131         net_kwargs = {
132             'network_arch': self.network_arch,
133             'use_bn': self.use_bn, # whether or not to use batch norm
134             'dropout_prob': self.dropout_prob,
135             'x_dim': self.lattice.num_links, # dim of target space
136             'links_shape': self.lattice.links.shape,
137             'num_hidden1': self.num_hidden1,

```

```

138         'num_hidden2': self.num_hidden2,
139         'generic_activation': tf.nn.relu,
140         'num_filters': [num_filters, int(2 * num_filters)],
141         'name_scope': 'x', # namespace in which to create network
142         'factor': 2., # scale factor used in original paper
143         '_input_shape': (self.batch_size, *self.lattice.links.shape),
144         'zero_translation': self.zero_translation,
145         # 'data_format': self.data_format,
146     }
147
148     if self.network_arch == 'conv3D':
149         net_kw_args.update({
150             'filter_sizes': [(3, 3, 1), (2, 2, 1)],
151         })
152     elif self.network_arch == 'conv2D':
153         net_kw_args.update({
154             'filter_sizes': [(3, 3), (2, 2)],
155         })
156
157     self.build_network(net_kw_args)
158
159 def build_network(self, net_kw_args):
160     """Build neural network used to train model."""
161     with tf.name_scope("DynamicsNetwork"):
162         self.x_fn = FullNet(model_name='XNet', **net_kw_args)
163
164         net_kw_args['name_scope'] = 'v' # update name scope
165         net_kw_args['factor'] = 1. # factor used in orig. paper
166         self.v_fn = FullNet(model_name='VNet', **net_kw_args)
167
168 def call(self, *args, **kwargs):
169     """Call method."""
170     return self.apply_transition(*args, **kwargs)
171
172 def apply_transition(self,
173                     x_in,
174                     beta,
175                     net_weights,
176                     train_phase,
177                     save_lf=False):
178     """Propose a new state and perform the accept/reject step.
179
180     We simulate the (molecular) dynamics update both forward and backward,
181     and use sampled masks to compute the actual solutions.
182
183     Args:
184         x_in (placeholder): Batch of (x) samples (GaugeLattice.samples).
185         beta (float): Inverse coupling constant.
186         net_weights: Array of scaling weights to multiply each of the
187             output functions (scale, translation, transformation).
188         train_phase: Boolean tf.placeholder used to indicate if currently
189             training model or running inference on trained model.
190         save_lf: Flag specifying whether or not to save output leapfrog
191             configs.
192
193     Returns:
194         x_proposed: Proposed x before accept/reject step.
195         v_proposed: Proposed v before accept/reject step.
196         accept_prob: Probability of accepting the proposed states.
197         x_out: Samples after accept/reject step.
198     """
199     lf_dict = {}
200
201     with tf.name_scope('transition_forward'):
202         outputs_f = self.transition_kernel(x_in, beta,
203                                             net_weights,
204                                             train_phase,
205                                             forward=True,
206                                             save_lf=save_lf)
207         xf = outputs_f['x_proposed']

```

```

208     vf = outputs_f['v_proposed']
209     lf_dict['pxs_out_f'] = outputs_f['accept_prob']
210
211     with tf.name_scope('transition_backward'):
212         outputs_b = self.transition_kernel(x_in, beta,
213                                             net_weights,
214                                             train_phase,
215                                             forward=False,
216                                             save_lf=save_lf)
217         xb = outputs_b['x_proposed']
218         vb = outputs_b['v_proposed']
219         lf_dict['pxs_out_b'] = outputs_b['accept_prob']
220
221     def get_lf_keys(direction):
222         base_keys = ['lf_out', 'logdets', 'sumlogdet', 'fns_out']
223         new_keys = [k + f'_{direction}' for k in base_keys]
224         return list(zip(new_keys, base_keys))
225
226     keys_f = get_lf_keys('f')
227     keys_b = get_lf_keys('b')
228
229     if save_lf:
230         lf_dict.update({k[0]: outputs_f[k[1]] for k in keys_f})
231         lf_dict.update({k[0]: outputs_b[k[1]] for k in keys_b})
232
233     # Decide direction uniformly
234     with tf.name_scope('transition_masks'):
235         with tf.name_scope('forward_mask'):
236             lf_dict['masks_f'] = tf.cast(
237                 tf.random_uniform((self.batch_size,),
238                                   dtype=TF_FLOAT,
239                                   seed=GLOBAL_SEED) > 0.5,
240                 TF_FLOAT,
241                 name='forward_mask'
242             )
243             with tf.name_scope('backward_mask'):
244                 lf_dict['masks_b'] = 1. - lf_dict['masks_f']
245
246     # Obtain proposed states
247     with tf.name_scope('x_proposed'):
248         x_proposed = (lf_dict['masks_f'][ :, None] * xf
249                         + lf_dict['masks_b'][ :, None] * xb)
250
251     with tf.name_scope('v_proposed'):
252         v_proposed = (lf_dict['masks_f'][ :, None] * vf
253                         + lf_dict['masks_b'][ :, None] * vb)
254
255     # Probability of accepting the proposed states
256     with tf.name_scope('accept_prob'):
257         accept_prob = (lf_dict['masks_f'] * lf_dict['pxs_out_f']
258                         + lf_dict['masks_b'] * lf_dict['pxs_out_b'])
259
260     # Accept or reject step
261     with tf.name_scope('accept_mask'):
262         accept_mask = tf.cast(
263             accept_prob > tf.random_uniform(tf.shape(accept_prob),
264                                             dtype=TF_FLOAT,
265                                             seed=GLOBAL_SEED),
266             TF_FLOAT,
267             name='accept_mask'
268         )
269         reject_mask = 1. - accept_mask
270
271     # Samples after accept / reject step
272     with tf.name_scope('x_out'):
273         x_out = (accept_mask[ :, None] * x_proposed
274                         + reject_mask[ :, None] * x_in)
275
276     outputs = {
277         'x_proposed': x_proposed,

```

```

278         'v_proposed': v_proposed,
279         'accept_prob': accept_prob,
280         'x_out': x_out
281     }
282
283     if save_lf:
284         outputs.update(lf_dict)
285
286     return outputs
287
288 def transition_kernel(self,
289                      x_in,
290                      beta,
291                      net_weights,
292                      train_phase,
293                      forward=True,
294                      save_lf=False):
295     """Transition kernel of augmented leapfrog integrator."""
296     lf_fn = self._forward_lf if forward else self._backward_lf
297
298     with tf.name_scope('refresh_momentum'):
299         v_in = tf.random_normal(tf.shape(x_in),
300                                dtype=TF_FLOAT,
301                                seed=GLOBAL_SEED,
302                                name='refresh_momentum')
303
304     with tf.name_scope('init'):
305         x_proposed, v_proposed = x_in, v_in
306
307         step = tf.constant(0, name='md_step', dtype=TF_FLOAT)
308         batch_size = tf.shape(x_in)[0]
309         logdet = tf.zeros((batch_size,), dtype=TF_FLOAT)
310         # fns0 = tf.zeros((4, 3, *x_in.shape,))
311         lf_out = tf.TensorArray(dtype=TF_FLOAT,
312                                size=self.num_steps+1,
313                                dynamic_size=True,
314                                name='lf_out',
315                                clear_after_read=False)
316         logdets_out = tf.TensorArray(dtype=TF_FLOAT,
317                                    size=self.num_steps+1,
318                                    dynamic_size=True,
319                                    name='logdets_out',
320                                    clear_after_read=False)
321         fns_out = tf.TensorArray(dtype=TF_FLOAT,
322                                size=self.num_steps,
323                                dynamic_size=True,
324                                name='l2hmc_fns',
325                                clear_after_read=False)
326
327         lf_out = lf_out.write(0, x_in)
328         logdets_out = logdets_out.write(0, logdet)
329         # fns_out.write(0, fns0)
330
331     def body(step, x, v, logdet, lf_samples, logdets, fns):
332         # cast leapfrog step to integer
333         i = tf.cast(step, dtype=tf.int32)
334         new_x, new_v, j, _fns = lf_fn(x, v, beta, step,
335                                         net_weights, train_phase)
336         lf_samples = lf_samples.write(i+1, new_x)
337         logdets = logdets.write(i+1, logdet+j)
338         fns = fns.write(i, _fns)
339         # fns = fns.write(i, _fns)
340
341     return (step+1, new_x, new_v, logdet+j, lf_samples, logdets, fns)
342
343     def cond(step, *args):
344         return tf.less(step, self.num_steps)
345
346     outputs = tf.while_loop(
347         cond=cond,

```

```

348     body=body,
349     loop_vars=[step, x_proposed, v_proposed,
350                logdet, lf_out, logdets_out, fns_out])
351
352     step = outputs[0]
353     x_proposed = outputs[1]
354     v_proposed = outputs[2]
355     sumlogdet = outputs[3]
356     lf_out = outputs[4].stack()
357     logdets_out = outputs[5].stack()
358     fns_out = outputs[6].stack()
359
360     accept_prob = self._compute_accept_prob(
361         x_in,
362         v_in,
363         x_proposed,
364         v_proposed,
365         sumlogdet,
366         beta
367     )
368
369     outputs = {
370         'x_proposed': x_proposed,
371         'v_proposed': v_proposed,
372         'sumlogdet': sumlogdet,
373         'accept_prob': accept_prob,
374     }
375
376     if save_lf:
377         outputs['lf_out'] = lf_out
378         outputs['logdets'] = logdets_out
379         outputs['fns_out'] = fns_out
380
381     return outputs
382
383 def _forward_lf(self, x, v, beta, step, net_weights, train_phase):
384     """One forward augmented leapfrog step."""
385     forward_fns = []
386     with tf.name_scope('forward_lf'):
387         with tf.name_scope('get_time'):
388             t = self._get_time(step, tile=tf.shape(x)[0])
389         with tf.name_scope('get_mask'):
390             mask, mask_inv = self._get_mask(step)
391
392         sumlogdet = 0.
393
394         v, logdet, vf_fns = self._update_v_forward(x, v, beta, t,
395                                                     net_weights,
396                                                     train_phase)
397         sumlogdet += logdet
398         forward_fns.append(vf_fns)
399
400         x, logdet, xf_fns = self._update_x_forward(x, v, t,
401                                                     net_weights,
402                                                     train_phase,
403                                                     mask, mask_inv)
404         sumlogdet += logdet
405         forward_fns.append(xf_fns)
406
407         x, logdet, xf_fns = self._update_x_forward(x, v, t,
408                                                     net_weights,
409                                                     train_phase,
410                                                     mask_inv, mask)
411         sumlogdet += logdet
412         forward_fns.append(xf_fns)
413
414         v, logdet, vf_fns = self._update_v_forward(x, v, beta, t,
415                                                     net_weights,
416                                                     train_phase)
417         sumlogdet += logdet

```

```

418         forward_fns.append(vf_fns)
419
420     return x, v, sumlogdet, forward_fns
421
422 def _backward_lf(self, x, v, beta, step, net_weights, train_phase):
423     """One backward augmented leapfrog step."""
424     backward_fns = []
425     with tf.name_scope('backward_lf'):
426         with tf.name_scope('get_time'):
427             # Reversed index/sinusoidal time
428             t = self._get_time(self.num_steps - step - 1,
429                               tile=tf.shape(x)[0])
430         with tf.name_scope('get_mask'):
431             mask, mask_inv = self._get_mask(
432                 self.num_steps - step - 1
433             )
434
435         sumlogdet = 0.
436
437         v, logdet, vb_fns = self._update_v_backward(x, v, beta, t,
438                                                       net_weights,
439                                                       train_phase)
440         sumlogdet += logdet
441         backward_fns.append(vb_fns)
442
443         x, logdet, xb_fns = self._update_x_backward(x, v, t,
444                                                       net_weights,
445                                                       train_phase,
446                                                       mask_inv, mask)
447         sumlogdet += logdet
448         backward_fns.append(xb_fns)
449
450         x, logdet, xb_fns = self._update_x_backward(x, v, t,
451                                                       net_weights,
452                                                       train_phase,
453                                                       mask, mask_inv)
454         sumlogdet += logdet
455         backward_fns.append(xb_fns)
456
457         v, logdet, vb_fns = self._update_v_backward(x, v, beta, t,
458                                                       net_weights,
459                                                       train_phase)
460         sumlogdet += logdet
461         backward_fns.append(vb_fns)
462
463     return x, v, sumlogdet, backward_fns
464
465 def _update_v_forward(self, x, v, beta, t, net_weights, train_phase):
466     """Update v in the forward leapfrog step.
467
468     Args:
469         x: input position tensor
470         v: input momentum tensor
471         beta: inverse coupling constant
472         t: current leapfrog step
473         net_weights: Placeholders for the multiplicative weights by which
474             to multiply the S, Q, and T functions (scale, transformation,
475             translation resp.)
476     Returns:
477         v: Updated (output) momentum
478         logdet: Jacobian factor
479     """
480     with tf.name_scope('update_vf'):
481         grad = self.grad_potential(x, beta)
482
483         scale, transl, transf = self.v_fn((x, grad, t), train_phase)
484
485         with tf.name_scope('vf_mul'):
486             scale *= 0.5 * self.eps * net_weights[0]
487             transl *= net_weights[1]

```

```

488     transf *= self.eps * net_weights[2]
489     fns = [scale, transl, transf]
490
491     with tf.name_scope('vf_exp'):
492         scale_exp = tf.cast(exp(scale, 'scale_exp'), dtype=TF_FLOAT)
493         transf_exp = tf.cast(exp(transf, 'transf_exp'),
494                               dtype=TF_FLOAT)
495
496     with tf.name_scope('proposed'):
497         v = (v * scale_exp
498              - 0.5 * self.eps * (grad * transf_exp - transl))
499
500     logdet = tf.reduce_sum(scale, axis=1, name='logdet_vf')
501
502     return v, logdet, fns
503
504 def _update_x_forward(self, x, v, t, net_weights,
505                      train_phase, mask, mask_inv):
506     """Update x in the forward leapfrog step."""
507     with tf.name_scope('update_xf'):
508         scale, transl, transf = self.x_fn([v, mask * x, t], train_phase)
509
510         with tf.name_scope('xf_mul'):
511             scale *= self.eps * net_weights[0]
512             transl *= net_weights[1]
513             transf *= self.eps * net_weights[2]
514             fns = [scale, transl, transf]
515
516         with tf.name_scope('xf_exp'):
517             scale_exp = exp(scale, 'scale_exp')
518             transf_exp = exp(transf, 'transformation_exp')
519
520         with tf.name_scope('proposed'):
521             x = (mask * x + mask_inv
522                  * (x * scale_exp + self.eps * (v * transf_exp + transl)))
523
524     logdet = tf.reduce_sum(mask_inv * scale, axis=1, name='logdet_xf')
525
526     return x, logdet, fns
527
528 def _update_v_backward(self, x, v, beta, t, net_weights, train_phase):
529     """Update v in the backward leapfrog step. Invert the forward update"""
530     with tf.name_scope('update_vb'):
531         grad = self.grad_potential(x, beta)
532
533         scale, transl, transf = self.v_fn([x, grad, t], train_phase)
534
535         with tf.name_scope('vb_mul'):
536             scale *= -0.5 * self.eps * net_weights[0]
537             transl *= net_weights[1]
538             transf *= self.eps * net_weights[2]
539             fns = [scale, transl, transf]
540
541         with tf.name_scope('vb_exp'):
542             scale_exp = exp(scale, 'scale_exp')
543             transf_exp = exp(transf, 'transformation_exp')
544
545         with tf.name_scope('proposed'):
546             v = scale_exp * (v + 0.5 * self.eps
547                               * (grad * transf_exp - transl))
548
549     logdet = tf.reduce_sum(scale, axis=1, name='logdet_vb')
550
551     return v, logdet, fns
552
553 def _update_x_backward(self, x, v, t, net_weights,
554                      train_phase, mask, mask_inv):
555     """Update x in the backward lf step. Inverting the forward update."""
556     with tf.name_scope('update_xb'):
557         scale, transl, transf = self.x_fn([v, mask * x, t], train_phase)

```

```

558
559     with tf.name_scope('xb_mul'):
560         scale *= -self.eps * net_weights[0]
561         transl *= net_weights[1]
562         transf *= self.eps * net_weights[2]
563         fns = [scale, transl, transf]
564
565     with tf.name_scope('xb_exp'):
566         scale_exp = exp(scale, 'xb_scale')
567         transf_exp = exp(transf, 'xb_transformation')
568
569     with tf.name_scope('proposed'):
570         x = (mask * x + mask_inv * scale_exp
571             * (x - self.eps * (v * transf_exp + transl)))
572
573     logdet = tf.reduce_sum(mask_inv * scale, axis=1,
574                           name='logdet_xb')
575
576     return x, logdet, fns
577
578 def _compute_accept_prob(self, xi, vi, xf, vf, sumlogdet, beta):
579     """Compute the prob of accepting the proposed state given old state.
580     Args:
581         xi: Initial state.
582         vi: Initial v.
583         xf: Proposed state.
584         vf: Proposed v.
585         sumlogdet: Sum of the terms of the log of the determinant.
586             (Eq. 14 of original paper).
587         beta: Inverse coupling constant of gauge model.
588     """
589     with tf.name_scope('compute_accept_prob'):
590         with tf.name_scope('old_hamiltonian'):
591             old_hamil = self.hamiltonian(xi, vi, beta)
592         with tf.name_scope('new_hamiltonian'):
593             new_hamil = self.hamiltonian(xf, vf, beta)
594
595         with tf.name_scope('prob'):
596             prob = exp(tf.minimum(
597                 (old_hamil - new_hamil + sumlogdet), 0.
598             ), 'accept_prob')
599
600         # Ensure numerical stability as well as correct gradients
601         accept_prob = tf.where(tf.is_finite(prob), prob,
602                               tf.zeros_like(prob))
603
604     return accept_prob
605
606 def _get_time(self, i, tile=1):
607     """Format time as [cos(..), sin(...)]."""
608     with tf.name_scope('get_time'):
609         trig_t = tf.squeeze([
610             tf.cos(2 * np.pi * i / self.num_steps),
611             tf.sin(2 * np.pi * i / self.num_steps),
612         ])
613
614         t = tf.tile(tf.expand_dims(trig_t, 0), (tile, 1))
615
616     return t
617
618 def _construct_masks(self):
619     """Construct different binary masks for different time steps."""
620     self.masks = []
621     for _ in range(self.num_steps):
622         # Need to use np here because tf would generate different random
623         # values across different `sess.run`
624         idx = np.random.permutation(np.arange(self.x_dim))[:self.x_dim // 2]
625         mask = np.zeros((self.x_dim,))
626         mask[idx] = 1.
627         mask = tf.constant(mask, dtype=TF_FLOAT)

```

```

628         self.masks.append(mask[None, :])
629
630     def _get_mask(self, step):
631         with tf.name_scope('get_mask'):
632             m = tf.gather(self.masks, tf.cast(step, dtype=TF_INT))
633             _m = 1. - m # complementary mask
634         return m, _m
635
636     def potential_energy(self, x, beta):
637         """Compute potential energy using `self.potential` and beta."""
638         with tf.name_scope('potential_energy'):
639             potential_energy = tf.multiply(beta, self.potential(x))
640
641         return potential_energy
642
643     def kinetic_energy(self, v):
644         """Compute the kinetic energy."""
645         with tf.name_scope('kinetic_energy'):
646             kinetic_energy = 0.5 * tf.reduce_sum(v**2, axis=1)
647
648         return kinetic_energy
649
650     def hamiltonian(self, x, v, beta):
651         """Compute the overall Hamiltonian."""
652         with tf.name_scope('hamiltonian'):
653             potential = self.potential_energy(x, beta)
654             kinetic = self.kinetic_energy(v)
655             hamiltonian = potential + kinetic
656
657         return hamiltonian
658
659     def grad_potential(self, x, beta):
660         """Get gradient of potential function at current location."""
661         with tf.name_scope('grad_potential'):
662             if tf.executing_eagerly():
663                 tfe = tf.contrib.eager
664                 grad_fn = tfe.gradients_function(self.potential_energy,
665                                                 params=["x"])
666                 grad = grad_fn(x, beta)[0]
667             else:
668                 grad = tf.gradients(self.potential_energy(x, beta), x)[0]
669         return grad

```

A.7 l2hmc-qcd/lattice/lattice.py

```
1 """
2 lattice.py
3
4 Contains implementation of GaugeLattice class.
5
6 Author: Sam Foreman (github: @saforem2)
7 Date: 01/15/2019
8 """
9
10 import numpy as np
11 import tensorflow as tf
12 from scipy.special import i0, i1
13 from config import TF_FLOAT, NP_FLOAT
14
15
16 def u1_plaq_exact(beta):
17     """Computes the expected value of the `average` plaquette for U(1)."""
18     return i1(beta) / i0(beta)
19
20
21 def u1_plaq_exact_tf(beta):
22     return tf.math.bessel_i1(beta) / tf.math.bessel_i0(beta)
23
24
25 def pbc(tup, shape):
26     """Returns tup % shape for implementing periodic boundary conditions."""
27     return list(np.mod(tup, shape))
28
29
30 def pbc_tf(tup, shape):
31     """Tensorflow implementation of `pbc` defined above."""
32     return list(tf.mod(tup, shape))
33
34
35 def mat_adj(mat):
36     """Returns the adjoint (i.e. conjugate transpose) of a matrix `mat`."""
37     return tf.transpose(tf.conj(mat)) # conjugate transpose
38
39
40 def project_angle(x):
41     """Returns the projection of an angle `x` from [-4pi, 4pi] to [-pi, pi]."""
42     return x - 2 * np.pi * tf.math.floor((x + np.pi) / (2 * np.pi))
43
44
45 def project_angle_fft(x, N=10):
46     """Use the fourier series representation `x` to approx `project_angle`.
47     NOTE: Because `project_angle` suffers a discontinuity, we approximate `x`
48     with its Fourier series representation in order to have a differentiable
49     function when computing the loss.
50     Args:
51         x (array-like): Array to be projected.
52         N (int): Number of terms to keep in Fourier series.
53     """
54     y = np.zeros(x.shape, dtype=NP_FLOAT)
55     for n in range(1, N):
56         y += (-2 / n) * ((-1) ** n) * tf.sin(n * x)
57     return y
58
59
60 class GaugeLattice(object):
61     """Lattice with Gauge field existing on links."""
62
63     def __init__(self,
64                  time_size=8,
65                  space_size=8,
66                  dim=2,
67                  link_type='U1',
```

```

68             num_samples=None,
69             rand=False):
70     """Initialization for GaugeLattice object.
71
72     Args:
73         time_size (int): Temporal extent of lattice.
74         space_size (int): Spatial extent of lattice.
75         dim (int): Dimensionality
76         link_type (str):
77             String representing the type of gauge group for the link
78             variables. Must be either 'U1', 'SU2', or 'SU3'
79     """
80     assert link_type.upper() in ['U1', 'SU2', 'SU3'], (
81         "Invalid link_type. Possible values: U1", 'SU2', 'SU3')
82
83     self.time_size = time_size
84     self.space_size = space_size
85     self.dim = dim
86     self.link_type = link_type
87     self.link_shape = ()
88
89     self.num_links = self.time_size * self.space_size * self.dim
90     self.num_plaqs = self.time_size * self.space_size
91     self.bases = np.eye(self.dim, dtype=np.int)
92
93     # create samples using @property method: `@samples.setter`
94     self.samples = (num_samples, rand)
95
96     self.links = self.samples[0]
97     self.batch_size = self.samples.shape[0]
98     self.links_shape = self.samples.shape[1:]
99     self.x_dim = self.num_links
100    # self.samples = self._init_samples(num_samples, rand)
101    self.samples_tensor = tf.convert_to_tensor(
102        self.samples.reshape((self.batch_size, self.x_dim)),
103        dtype=TF_FLOAT
104    )
105
106
107 @property
108 def samples(self):
109     return self._samples
110
111 @samples.setter
112 def samples(self, args):
113     """Create samples."""
114     num_samples, rand = args
115     links_shape = tuple(
116         [self.time_size]
117         + [self.space_size for _ in range(self.dim-1)]
118         + [self.dim]
119         + list(self.link_shape)
120     )
121     samples_shape = (num_samples, *links_shape)
122     if rand:
123         samples = np.array(
124             np.random.uniform(0, 2*np.pi, samples_shape),
125             dtype=NP_FLOAT
126         )
127     else:
128         samples = np.zeros(samples_shape, dtype=NP_FLOAT)
129
130     self._samples = samples
131
132 def calc_plaq_sums(self, samples=None):
133     """Calculate plaquette sums.
134
135     Explicitly, calculate the sum of the link variables around each
136     plaquette in the lattice for each sample in samples.
137

```

```

138     Args:
139         samples (tf.Tensor): Tensor of shape (N, D) where N is the batch
140             size and D is the number of links on the lattice. If samples is
141             None, self.samples will be used.
142
143     Returns:
144         plaq_sums (tf operation): Tensorflow operation capable of
145             calculating the plaquette sums.
146
147     NOTE: self.samples.shape = (N, L, T, D), where:
148         N = num_samples
149         L = space_size
150         T = time_size
151         D = dimensionality
152     """
153     if samples is None:
154         samples = self.samples
155
156     with tf.name_scope('plaq_sums'):
157         if samples.shape != self.samples.shape:
158             samples = tf.reshape(samples, shape=(self.samples.shape))
159
160         # assuming D = 2, plaq_sums will have shape: (N, L, T)
161         plaq_sums = (samples[:, :, :, 0]
162                     - samples[:, :, :, 1]
163                     - tf.roll(samples[:, :, :, 0], shift=-1, axis=2)
164                     + tf.roll(samples[:, :, :, 1], shift=-1, axis=1))
165
166     return plaq_sums
167
168 def calc_actions(self, samples=None, plaq_sums=None):
169     """Calculate the total action for each sample in samples."""
170     if plaq_sums is None:
171         if samples is None:
172             samples = self.samples
173         plaq_sums = self.calc_plaq_sums(samples)
174
175     with tf.name_scope('actions'):
176         total_actions = tf.reduce_sum(1. - tf.cos(plaq_sums),
177                                       axis=(1, 2), name='actions')
178
179     return total_actions
180
181 def calc_plaqs(self, samples=None, plaq_sums=None):
182     """Calculate the average plaq. values for each sample in samples."""
183     if plaq_sums is None:
184         if samples is None:
185             samples = self.samples
186         plaq_sums = self.calc_plaq_sums(samples)
187
188     with tf.name_scope('plaqs'):
189         # plaqs = tf.reduce_mean(tf.cos(self.calc_plaq_sums(samples)),
190         #                         axis=(1, 2), name='plaqs')
191         plaqs = tf.reduce_sum(tf.cos(plaq_sums),
192                               axis=(1, 2), name='plaqs') / self.num_plaqs
193
194     return plaqs
195
196 def calc_top_charges(self, samples=None, plaq_sums=None):
197     """Calculate topological charges for each sample in samples."""
198     if plaq_sums is None:
199         if samples is None:
200             samples = self.samples
201         plaq_sums = self.calc_plaq_sums(samples)
202
203     with tf.name_scope('top_charges'):
204         ps_proj = tf.sin(plaq_sums)
205         # if fft:
206         #     ps_proj = tf.sin(self.calc_plaq_sums(samples))
207         #     ps_proj = project_angle_fft(self.calc_plaq_sums(samples), N=1)
208         # else:

```

```
208     # ps_proj = tf.sin(self.calc_plaq_sums(samples))
209     # ps_proj = project_angle(self.calc_plaq_sums(samples))
210
211     top_charges = (tf.reduce_sum(ps_proj, axis=(1, 2),
212                                 name='top_charges')) / (2 * np.pi)
213
214     return top_charges
215
216 def calc_top_charges_diff(self, x1, x2, plaq_sums1=None, plaq_sums2=None):
217     """Calculate the difference in topological charge between x1 and x2."""
218     with tf.name_scope('top_charges_diff'):
219         charge_diff = tf.abs(self.calc_top_charges(x1, plaq_sums1)
220                             - self.calc_top_charges(x2, plaq_sums2))
221
222     return charge_diff
223
224 def get_potential_fn(self, samples):
225     """Returns callable function used for calculating the energy."""
226     def fn(samples):
227         return self.calc_actions(samples)
228
229     return fn
```

A.8 l2hmc-qcd/network/conv_net.py

```
1 """
2 conv_net.py
3
4 Implements both 2D and 3D convolutional neural networks.
5
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 06/14/2019
9 """
10
11 import numpy as np
12 import tensorflow as tf
13
14 from config import GLOBAL_SEED, TF_FLOAT
15 from .network_utils import batch_norm
16
17 np.random.seed(GLOBAL_SEED)
18
19 if '2.' not in tf.__version__:
20     tf.set_random_seed(GLOBAL_SEED)
21
22 ##### NOTE: Removed from `ConvNet2D.__init__` due to bugs when specifying
23 #      'data_format'
24 #
25 # -----
26 # if self.use_bn:
27 #     if self.data_format == 'channels_first':
28 #         self.bn_axis = 1
29 #     elif self.data_format == 'channels_last':
30 #         self.bn_axis = -1
31 #     else:
32 #         raise AttributeError("Expected 'data_format' "
33 #                             "to be 'channels_first' "
34 #                             "or 'channels_last'.")
35
36
37 class ConvNet2D(tf.keras.Model):
38     """Convolutional block used in ConvNet3D."""
39     def __init__(self, model_name, **kwargs):
40         """
41             Initialization method.
42
43             Args:
44                 model_name: Name of the model.
45                 kwargs: Keyword arguments used to specify specifics of
46                         convolutional structure.
47
48             super(ConvNet2D, self).__init__(name=model_name)
49
50             self.data_format = 'channels_last'
51
52             for key, val in kwargs.items():
53                 setattr(self, key, val)
54
55             self.activation = kwargs.get('conv_act', tf.nn.relu)
56             activation2 = None if self.use_bn else self.activation
57
58             if model_name == 'ConvNet2Dx':
59                 self.bn_name = 'batch_norm_x'
60             elif model_name == 'ConvNet2Dv':
61                 self.bn_name = 'batch_norm_v'
62
63             if self.name_scope == 'ConvNetX':
64                 conv1_name = 'conv_x1'
65                 pool1_name = 'pool_x1'
66                 conv2_name = 'conv_x2'
67                 pool2_name = 'pool_x2'
```

```

68     elif self.name_scope == 'ConvNetV':
69         conv1_name = 'conv_v1'
70         pool1_name = 'pool_v1'
71         conv2_name = 'conv_v2'
72         pool2_name = 'pool_v2'
73
74     with tf.name_scope(self.name_scope):
75
76         self.conv1 = tf.keras.layers.Conv2D(
77             filters=self.num_filters[0],
78             kernel_size=self.filter_sizes[0],
79             activation=self.activation,
80             # input_shape=self._input_shape[1:],
81             name=conv1_name,
82             padding='same',
83             dtype=TF_FLOAT,
84             data_format=self.data_format
85         )
86
87         # with tf.name_scope('pool_x1'):
88         self.max_pool1 = tf.keras.layers.MaxPooling2D(
89             pool_size=(2, 2),
90             strides=2,
91             # padding='same',
92             name=pool1_name,
93         )
94
95         self.conv2 = tf.keras.layers.Conv2D(
96             filters=self.num_filters[1],
97             kernel_size=self.filter_sizes[1],
98             activation=activation2,
99             padding='same',
100            name=conv2_name,
101            dtype=TF_FLOAT,
102            data_format=self.data_format
103        )
104
105         self.max_pool2 = tf.keras.layers.MaxPooling2D(
106             pool_size=(2, 2),
107             strides=2,
108             # padding='same',
109             name=pool2_name,
110         )
111
112         self.flatten = tf.keras.layers.Flatten(name='flatten')
113
114     def reshape_4D(self, tensor):
115         """
116             Reshape tensor to be compatible with tf.keras.layers.Conv3D.
117
118             If self.data_format is 'channels_first', and input `tensor` has shape
119             (N, 2, L, L), the output tensor has shape (N, 1, 2, L, L).
120
121             If self.data_format is 'channels_last' and input `tensor` has shape
122             (N, L, L, 2), the output tensor has shape (N, 2, L, L, 1).
123         """
124
125         N, H, W, C = self._input_shape
126         if self.data_format == 'channels_first':
127             # N, C, H, W = self._input_shape
128             # N, D, H, W = tensor.shape
129             if isinstance(tensor, np.ndarray):
130                 return np.reshape(tensor, (N, C, H, W))
131
132             return tf.reshape(tensor, (N, C, H, W))
133
134         if self.data_format == 'channels_last':
135             # N, H, W, C = self._input_shape
136             if isinstance(tensor, np.ndarray):
137                 return np.reshape(tensor, (N, H, W, C))

```

```

138     return tf.reshape(tensor, (N, H, W, C))
139
140     raise AttributeError(``self.data_format` should be one of "
141                         "'channels_first' or 'channels_last'``")
142
143 def call(self, inputs, train_phase):
144     """Forward pass through the network."""
145     inputs = self.reshape_4D(inputs)
146
147     inputs = self.max_pool1(self.conv1(inputs))
148     inputs = self.conv2(inputs)
149     if self.use_bn:
150         inputs = self.activation(batch_norm(inputs, train_phase,
151                                         axis=self.bn_axis,
152                                         internal_update=True,
153                                         scope=self.bn_name,
154                                         reuse=tf.AUTO_REUSE))
155     inputs = self.max_pool2(inputs)
156     inputs = self.flatten(inputs)
157
158     return inputs
159
160
161 class ConvNet3D(tf.keras.Model):
162     """Convolutional block used in ConvNet3D."""
163     def __init__(self, model_name, **kwargs):
164         """
165             Initialization method.
166
167             Args:
168                 model_name: Name of the model.
169                 kwargs: Keyword arguments used to specify specifics of
170                         convolutional structure.
171         """
172         super(ConvNet3D, self).__init__(name=model_name)
173
174         self.data_format = 'channels_last'
175
176         for key, val in kwargs.items():
177             setattr(self, key, val)
178
179         if model_name == 'ConvNet3Dx':
180             self.bn_name = 'batch_norm_x'
181         elif model_name == 'ConvNet3Dv':
182             self.bn_name = 'batch_norm_v'
183
184         if self.name_scope == 'x_conv_block':
185             conv1_name = 'conv_x1'
186             pool1_name = 'pool_x1'
187             conv2_name = 'conv_x2'
188             pool2_name = 'pool_x2'
189         elif self.name_scope == 'v_conv_block':
190             conv1_name = 'conv_v1'
191             pool1_name = 'pool_v1'
192             conv2_name = 'conv_v2'
193             pool2_name = 'pool_v2'
194
195         with tf.name_scope(self.name_scope):
196             if self.use_bn:
197                 if self.data_format == 'channels_first':
198                     self.bn_axis = 1
199                 elif self.data_format == 'channels_last':
200                     self.bn_axis = -1
201                 else:
202                     raise AttributeError("Expected 'data_format' "
203                                         "to be 'channels_first' "
204                                         "or 'channels_last'.")
205
206             self.activation = kwargs.get('conv_act', tf.nn.relu)
207             activation2 = None if self.use_bn else self.activation

```

```

208     self.conv1 = tf.keras.layers.Conv3D(
209         filters=self.num_filters[0],
210         kernel_size=self.filter_sizes[0],
211         activation=self.activation,
212         # input_shape=self._input_shape,
213         padding='same',
214         name=conv1_name,
215         dtype=TF_FLOAT,
216         data_format=self.data_format
217     )
218
219     self.max_pool1 = tf.keras.layers.MaxPooling3D(
220         pool_size=(2, 2, 2),
221         strides=2,
222         padding='same',
223         name=pool1_name,
224     )
225
226     self.conv2 = tf.keras.layers.Conv3D(
227         filters=self.num_filters[1],
228         kernel_size=self.filter_sizes[1],
229         activation=activation2,
230         # activation=None if self.use_bn else tf.nn.relu,
231         # initializer=HE_INIT,
232         padding='same',
233         name=conv2_name,
234         dtype=TF_FLOAT,
235         data_format=self.data_format
236     )
237
238     self.max_pool2 = tf.keras.layers.MaxPooling3D(
239         pool_size=(2, 2, 2),
240         strides=2,
241         padding='same',
242         name=pool2_name,
243     )
244
245     self.flatten = tf.keras.layers.Flatten(name='flatten')
246
247     def reshape_5D(self, tensor):
248         """
249             Reshape tensor to be compatible with tf.keras.layers.Conv3D.
250
251             If self.data_format is 'channels_first', and input `tensor` has shape
252             (N, 2, L, L), the output tensor has shape (N, 1, 2, L, L).
253
254             If self.data_format is 'channels_last' and input `tensor` has shape
255             (N, L, L, 2), the output tensor has shape (N, 2, L, L, 1).
256         """
257
258         N, H, W, C = self._input_shape
259         if self.data_format == 'channels_first':
260             # N, C, H, W = self._input_shape
261             # N, D, H, W = tensor.shape
262             if isinstance(tensor, np.ndarray):
263                 return np.reshape(tensor, (N, 1, C, H, W))
264
265             return tf.reshape(tensor, (N, 1, C, H, W))
266
267         if self.data_format == 'channels_last':
268             # N, H, W, C = self._input_shape
269             if isinstance(tensor, np.ndarray):
270                 return np.reshape(tensor, (N, H, W, C, 1))
271
272             return tf.reshape(tensor, (N, H, W, C, 1))
273
274         raise AttributeError(`self.data_format` should be one of "
275                             "'channels_first' or 'channels_last'`")
276
277     def call(self, inputs, train_phase):

```

```
278     """Forward pass through the network."""
279     inputs = self.reshape_5D(inputs)
280
281     inputs = self.max_pool1(self.conv1(inputs))
282     inputs = self.conv2(inputs)
283     if self.use_bn:
284         inputs = self.activation(batch_norm(inputs, train_phase,
285                                     axis=self.bn_axis,
286                                     internal_update=True,
287                                     scope=self.bn_name,
288                                     reuse=tf.AUTO_REUSE))
289     inputs = self.max_pool2(inputs)
290     inputs = self.flatten(inputs)
291
292     return inputs
```

A.9 l2hmc-qcd/network/generic_net.py

```
1 """
2 Generic, fully-connected neural network architecture for running L2HMC on a
3 gauge lattice configuration of links.
4
5 NOTE: Lattices are flattened before being passed as input to the network.
6
7 Reference [Generalizing Hamiltonian Monte Carlo with Neural
8 Networks](https://arxiv.org/pdf/1711.09268.pdf)
9
10 Code adapted from the released TensorFlow graph implementation by original
11 authors https://github.com/brain-research/l2hmc.
12
13 Author: Sam Foreman (github: @saforem2)
14 Date: 01/16/2019
15 """
16
17 import tensorflow as tf
18
19 from config import TF_FLOAT, GLOBAL_SEED
20
21 from .network_utils import custom_dense
22
23 class GenericNet(tf.keras.Model):
24     """Generic (fully-connected) network used in training L2HMC."""
25     def __init__(self, model_name, **kwargs):
26         """
27             Initialization method.
28
29             Args:
30                 model_name: Name of the model.
31                 kwargs: Keyword arguments used to specify specifics of
32                         convolutional structure.
33
34             super(GenericNet, self).__init__(name=model_name)
35
36         for key, val in kwargs.items():
37             setattr(self, key, val)
38
39         self.activation = kwargs.get('generic_activation', tf.nn.relu)
40
41         with tf.name_scope(self.name_scope):
42             self.coeff_scale = tf.Variable(
43                 initial_value=tf.zeros([1, self.x_dim], dtype=TF_FLOAT),
44                 name='coeff_scale',
45                 trainable=True,
46                 dtype=TF_FLOAT
47         )
48
49         # with tf.name_scope('coeff_transformation'):
50         self.coeff_transformation = tf.Variable(
51             initial_value=tf.zeros([1, self.x_dim], dtype=TF_FLOAT),
52             name='coeff_transformation',
53             trainable=True,
54             dtype=TF_FLOAT
55         )
56
57         if self.dropout_prob > 0:
58             self.dropout = tf.keras.layers.Dropout(self.dropout_prob,
59                                                 seed=GLOBAL_SEED, )
60
61         self.x_layer = custom_dense(self.num_hidden1,
62                                     self.factor/3.,
63                                     name='x_layer')
64         self.v_layer = custom_dense(self.num_hidden1,
65                                     1./3.,
66                                     name='v_layer')
67         self.t_layer = custom_dense(self.num_hidden1,
```

```

68             1./3.,
69             name='t_layer')
70
71     self.h_layer = custom_dense(self.num_hidden2,
72                               name='hidden_layer')
73
74     self.scale_layer = custom_dense(self.x_dim,
75                                     0.001,
76                                     name='scale_layer')
77     if not self.zero_translation:
78         self.translation_layer = custom_dense(self.x_dim,
79                                              0.001,
80                                              'translation_layer')
81
82     self.transformation_layer = custom_dense(self.x_dim,
83                                              0.001,
84                                              'transformation_layer')
85
86 def call(self, inputs, train_phase):
87     v, x, t = inputs
88
89     with tf.name_scope('v'):
90         v = self.v_layer(v)
91     with tf.name_scope('x'):
92         x = self.x_layer(x)
93     with tf.name_scope('t'):
94         t = self.t_layer(t)
95
96     with tf.name_scope('hidden_layer'):
97         h = self.activation(v + x + t)
98         h = self.activation(self.h_layer(h))
99
100    # dropout gets applied to the output of the previous layer
101    if self.dropout_prob > 0:
102        h = self.dropout(h, training=train_phase)
103
104    with tf.name_scope('scale'):
105        scale = (tf.nn.tanh(self.scale_layer(h))
106                  * tf.exp(self.coeff_scale, name='exp_coeff_scale'))
107
108    with tf.name_scope('transformation'):
109        transformation = (tf.nn.tanh(self.transformation_layer(h))
110                           * tf.exp(self.coeff_transformation,
111                                     name='exp_coeff_transformation'))
112
113    with tf.name_scope('translation'):
114        if self.zero_translation:
115            translation = tf.zeros_like(scale, name='translation')
116        else:
117            translation = self.translation_layer(h)
118
119    return scale, translation, transformation

```

A.10 l2hmc-qcd/network/network.py

```
1 """
2 network.py
3
4 Implements network architecture used to train L2HMC algorithm on 2D U(1)
5 lattice gauge model.
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 07/22/2019
9 """
10
11 import numpy as np
12 import tensorflow as tf
13 from .conv_net import ConvNet2D, ConvNet3D
14 from .generic_net import GenericNet
15 import utils.file_io as io
16
17 from config import GLOBAL_SEED
18
19 np.random.seed(GLOBAL_SEED)
20
21 if '2.' not in tf.__version__:
22     tf.set_random_seed(GLOBAL_SEED)
23
24
25 class FullNet(tf.keras.Model):
26     """Complete network used for training L2HMC model."""
27     def __init__(self, model_name, **kwargs):
28         """
29             Initialization method.
30
31             Args:
32                 model_name: Name of the model.
33                 kwargs: Keyword arguments used to specify specifics of
34                     convolutional structure.
35
36         super(FullNet, self).__init__(name=model_name)
37
38         if model_name == 'XNet':
39             generic_name_scope = 'GenericNetX'
40         elif model_name == 'VNet':
41             generic_name_scope = 'GenericNetV'
42
43         with tf.name_scope(model_name):
44             kwargs['name_scope'] = 'ConvNetX'
45             network_arch = kwargs.get('network_arch', 'conv3D')
46
47             if network_arch == 'conv2D':
48                 io.log('Using ConvNet2D architecture...')
49                 self.x_conv_net = ConvNet2D('ConvNet2Dx', **kwargs)
50
51             kwargs['name_scope'] = 'ConvNetV'
52             self.v_conv_net = ConvNet2D('ConvNet2Dv', **kwargs)
53
54         elif network_arch == 'conv3D':
55             io.log('Using ConvNet3D architecture...')
56             self.x_conv_net = ConvNet3D('ConvNet3Dx', **kwargs)
57
58             kwargs['name_scope'] = 'ConvNetV'
59             self.v_conv_net = ConvNet3D('ConvNet3Dv', **kwargs)
60
61         else:
62             io.log('Using GenericNet architecture...')
63             self.x_conv_net = self.v_conv_net = None
64
65         kwargs['name_scope'] = generic_name_scope
66         self.generic_net = GenericNet("GenericNet", **kwargs)
67
```

```
68     def call(self, inputs, train_phase):
69         v, x, t = inputs
70
71         if self.x_conv_net is not None:
72             v = self.v_conv_net(v, train_phase)
73             x = self.x_conv_net(x, train_phase)
74
75         scale, translation, transformation = self.generic_net([v, x, t],
76                                         train_phase)
77
78     return scale, translation, transformation
```

A.11 l2hmc-qcd/network/network_utils.py

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.contrib.framework import add_arg_scope, arg_scope
4
5 from config import GLOBAL_SEED, TF_FLOAT, NP_FLOAT
6
7
8 np.random.seed(GLOBAL_SEED)
9
10 if '2.' not in tf.__version__:
11     tf.set_random_seed(GLOBAL_SEED)
12
13
14 def activation_model(model):
15     """Create Keras Model that outputs activations of all conv./pool layers.
16
17     Args:
18         model (tf.keras.Model): Model for which we wish to visualize
19             activations.
20     Returns:
21         activation_model (tf.keras.Model): Model that outputs the activations
22             for each layer in `model`.
23     """
24     layer_outputs = [layer.output for layer in model.layers]
25
26     output_model = tf.keras.models.Model(inputs=model.input,
27                                         output=layer_outputs)
28
29     return output_model
30
31
32 def flatten(_list):
33     return [item for sublist in _list for item in sublist]
34
35
36 def add_elements_to_collection(elements, collection_list):
37     elements = flatten(elements)
38     collection_list = flatten(collection_list)
39     # collection_list = tf.nest.flatten(collection_list)
40     for name in collection_list:
41         collection = tf.get_collection_ref(name)
42         collection_set = set(collection)
43         for element in elements:
44             if element not in collection_set:
45                 collection.append(element)
46
47
48 def _assign_moving_average(orig_val, new_val, momentum, name):
49     with tf.name_scope(name):
50         scaled_diff = (1 - momentum) * (new_val - orig_val)
51         return tf.assign_add(orig_val, scaled_diff)
52
53
54 @add_arg_scope
55 def batch_norm(x,
56                 phase,
57                 axis=-1,
58                 shift=True,
59                 scale=True,
60                 momentum=0.99,
61                 eps=1e-3,
62                 internal_update=False,
63                 scope=None,
64                 reuse=None):
65
66     C = x._shape_as_list()[axis]
67     ndim = len(x.shape)
```

```

68     var_shape = [1] * (ndim - 1) + [C]
69
70     with tf.variable_scope(scope, 'batch_norm', reuse=reuse):
71         def training():
72             m, v = tf.nn.moments(x, list(range(ndim - 1)), keep_dims=True)
73             update_m = _assign_moving_average(moving_m,
74                                              m, momentum,
75                                              'update_mean')
76             update_v = _assign_moving_average(moving_v,
77                                              v, momentum,
78                                              'update_var')
79             # tf.add_to_collection(tf.GraphKeys.UPDATE_OPS, update_m)
80             # tf.add_to_collection(tf.GraphKeys.UPDATE_OPS, update_v)
81             tf.add_to_collection('update_ops', update_m)
82             tf.add_to_collection('update_ops', update_v)
83
84             if internal_update:
85                 with tf.control_dependencies([update_m, update_v]):
86                     output = (x - m) * tf.rsqrt(v + eps)
87             else:
88                 output = (x - m) * tf.rsqrt(v + eps)
89             return output
90
91     def testing():
92         m, v = moving_m, moving_v
93         output = (x - m) * tf.rsqrt(v + eps)
94         return output
95
96     # Get mean and variance, normalize input
97     moving_m = tf.get_variable('mean', var_shape,
98                                initializer=tf.zeros_initializer,
99                                trainable=False)
100    moving_v = tf.get_variable('var', var_shape,
101                               initializer=tf.ones_initializer,
102                               trainable=False)
103
104    if isinstance(phase, bool):
105        output = training() if phase else testing()
106    else:
107        output = tf.cond(phase, training, testing)
108
109    if scale:
110        output *= tf.get_variable('gamma', var_shape,
111                                  initializer=tf.ones_initializer)
112
113    if shift:
114        output += tf.get_variable('beta', var_shape,
115                                  initializer=tf.zeros_initializer)
116
117    return output
118
119
120    def custom_dense(units, factor=1., name=None):
121        """Custom dense layer with specified weight intialization."""
122        if '2.' not in tf.__version__:
123            kernel_initializer = tf.keras.initializers.VarianceScaling(
124                scale=factor,
125                mode='fan_in',
126                distribution='truncated_normal',
127                # distribution='uniform',
128                dtype=TF_FLOAT,
129                seed=GLOBAL_SEED,
130            )
131        else:
132            kernel_initializer = tf.contrib.layers.variance_scaling_initializer(
133                factor=factor,
134                mode='FAN_IN',
135                seed=GLOBAL_SEED,
136                dtype=TF_FLOAT,
137                uniform=False,

```

```

138         )
139
140     return tf.keras.layers.Dense(
141         units=units,
142         use_bias=True,
143         kernel_initializer=kernel_initializer,
144         bias_initializer=tf.constant_initializer(0., dtype=TF_FLOAT),
145         name=name
146     )
147
148
149 def variable_on_cpu(name, shape, initializer):
150     """Helper to create a Variable stored on CPU memory.
151
152     Args:
153         name: name of the variable
154         shape: list of ints
155         initializer: initializer for Variable
156
157     Returns:
158         Variable Tensor
159     """
160
161     with tf.device('/cpu:0'):
162         var = tf.get_variable(name, shape, initializer, TF_FLOAT)
163     return var
164
165
166 def variable_with_weight_decay(name, shape, stddev, wd, cpu=True):
167     """Helper to create an initialized Variable with weight decay.
168
169     Note that the Variable is initialized with a truncated normal distribution.
170     A weight decay is added only if one is specified.
171
172     Args:
173         name: Name of the variable
174         shape: list of ints
175         stddev: standard deviation of a truncated Gaussian
176         wd: Add L2Loss weight decay multiplied by this float. If None, weight
177             decay is not added for this variable.
178
179     Returns:
180         Variable Tensor
181     """
182
183     if cpu:
184         var = variable_on_cpu(
185             name, shape, tf.truncated_normal_initializer(stddev=stddev,
186                                                 dtype=TF_FLOAT)
187         )
188     else:
189         var = tf.get_variable(
190             name, shape, tf.truncated_normal_initializer(stddev=stddev,
191                                                 dtype=TF_FLOAT)
192         )
193     if wd is not None:
194         weight_decay = tf.multiply(tf.nn.l2_loss(var), wd, name='weight_loss')
195         tf.add_to_collection('losses', weight_decay)
196
197     return var
198
199
200 def create_periodic_padding(samples, filter_size):
201     """Create periodic padding for multiple samples, using filter_size."""
202     original_size = np.shape(samples)
203     N = original_size[1] # number of links in lattice
204     # N = np.shape(samples)[1] # number of links in lattice
205     padding = filter_size - 1
206
207     samples = tf.reshape(samples, shape=(samples.shape[0], -1))
208
209     x = []

```

```
208     for sample in samples:  
209         padded = np.zeros((N + 2 * padding), N + 2 * padding, 2)  
210         # lower left corner  
211         padded[:padding, :padding, :] = sample[N-padding:, N-padding:, :]  
212         # lower middle  
213         padded[padding:N+padding, :padding, :] = sample[:, N-padding:, :]  
214         # lower right corner  
215         padded[N+padding:, :padding, :] = sample[:padding, N-padding:, :]  
216         # left side  
217         padded[:padding, padding: N+padding, :] = sample[N-padding:, :, :]  
218         # center  
219         padded[:padding:N+padding, padding:N+padding, :] = sample[:, :, :]  
220         # right side  
221         padded[N+padding:, padding:N+padding:, :] = sample[:padding, :, :]  
222         # top middle  
223         padded[:padding:N+padding, N+padding:, :] = sample[:, :padding, :]  
224         # top right corner  
225         padded[N+padding:, N+padding:, :] = sample[:padding, :padding, :]  
226  
227         x.append(padded)  
228  
229     return np.array(x, dtype=NP_FLOAT).reshape(*original_size)
```

A.12 l2hmc-qcd/trainers/trainer.py

```
1 """
2 gauge_model_trainer.py
3
4 Implements GaugeModelTrainer class responsible for training GaugeModel.
5
6 Author: Sam Foreman (github: @saforem2)
7 Date: 04/09/2019
8 """
9 # import os
10 import time
11 import numpy as np
12 # import tensorflow as tf
13
14 import utils.file_io as io
15 from lattice.lattice import u1_plaq_exact
16 from config import NP_FLOAT
17
18
19 h_str = ("{:^12s}{:^10s}{:^10s}{:^10s}{:^10s}"\
20 " {:^10s}{:^10s}{:^10s}{:^10s}{:^10s}")
21
22 h_strf = h_str.format("STEP", "LOSS", "t/STEP", "% ACC", "EPS",
23 "BETA", "ACTION", "PLAQ", "(EXACT)", "dQ", "LR")
24
25 dash0 = (len(h_strf) + 1) * '-'
26 dash1 = (len(h_strf) + 1) * '-'
27 TRAIN_HEADER = dash0 + '\n' + h_strf + '\n' + dash1
28
29
30 class GaugeModelTrainer:
31     def __init__(self, sess, model, logger=None):
32         """Initialization method.
33
34         Args:
35             sess: tf.Session object.
36             model: GaugeModel object (defined in `models/gauge_model.py`)
37             logger: TrainLogger object (defined in `loggers/train_logger.py`)
38         """
39         self.sess = sess
40         self.model = model
41         self.logger = logger
42
43     def update_beta(self, step, **kwargs):
44         """Returns new beta to follow annealing schedule."""
45         beta_init = kwargs.get('beta_init', self.model.beta_init)
46         beta_final = kwargs.get('beta_final', self.model.beta_final)
47         train_steps = kwargs.get('train_steps', self.model.train_steps)
48
49         temp = ((1. / beta_init - 1. / beta_final)
50                 * (1. - step / float(train_steps))
51                 + 1. / beta_final)
52         new_beta = 1. / temp
53
54         return new_beta
55
56     def run_l2hmc_fns(self, key, samples_np, **kwargs):
57         """Get numerical values of the scale, transl, and transf fns."""
58         l2hmc_fns = self.model.dynamics.l2hmc_fns[key]
59
60         feed_dict = {
61             self.model.x: samples_np,
62             self.model.train_phase: True,
63         }
64
65         ops = [l2hmc_fns['scale'], l2hmc_fns['transl'], l2hmc_fns['transf']]
66
67         outputs = self.sess.run(ops, feed_dict=feed_dict)
```

```

68
69     return outputs
70
71 def train_step(self, step, samples_np, **kwargs):
72     """Perform a single training step.
73
74     Args:
75         step (int): Current training step.
76         samples_np (np.ndarray): Array of input data.
77         beta_np (float, optional): Input value for inverse coupling
78             constant.
79
80     Returns:
81         out_data (dict)
82     """
83     beta_np = kwargs.get('beta_np', None)
84
85     # net_weights: [scale_w, translation_w, transformation_w]
86     net_weights = kwargs.get('net_weights', [1., 1., 1.])
87     train_steps = kwargs.get('train_steps', self.model.train_steps)
88
89     if beta_np is None:
90         if self.model.fixed_beta:
91             beta_np = self.model.beta_init
92         else:
93             beta_np = self.update_beta(step, train_steps=train_steps)
94
95     feed_dict = {
96         self.model.x: samples_np,
97         self.model.beta: beta_np,
98         self.model.net_weights[0]: net_weights[0],
99         self.model.net_weights[1]: net_weights[1],
100        self.model.net_weights[2]: net_weights[2],
101        self.model.train_phase: True,
102    }
103
104     global_step = self.sess.run(self.model.global_step)
105
106     ops = [
107         self.model.train_op,          # apply gradients
108         self.model.loss_op,          # calculate loss
109         self.model.x_out,            # get new samples
110         self.model.px,               # calculate accept prob.
111         self.model.dynamics.eps,     # calculate current step size
112         self.model.actions_op,       # calculate avg. actions
113         self.model.plaqs_op,         # calculate avg. plaqs
114         self.model.charges_op,       # calculate top. charges
115         self.model.charge_diffs_op,   # change in top. charge/num_samples
116         self.model.lr,                # evaluate learning rate
117     ]
118
119     start_time = time.time()
120     outputs = self.sess.run(ops, feed_dict=feed_dict)
121     dt = time.time() - start_time
122
123     out_data = {
124         'step': global_step,
125         'loss': outputs[1],
126         'samples': np.mod(outputs[2], 2 * np.pi),
127         'samples_orig': outputs[2],
128         'px': outputs[3],
129         'eps': outputs[4],
130         'actions': outputs[5],
131         'plaqs': outputs[6],
132         'charges': outputs[7],
133         'charge_diffs': outputs[8],
134         'lr': outputs[9],
135         'beta': beta_np
136     }
137

```

```

138     data_str = (
139         f"[global_step:>5g}/{train_steps:<6g} "
140         f"[outputs[1]:^9.4g] " # loss value
141         f"[dt:^9.4g] " # time / step
142         f"[np.mean(outputs[3]):^9.4g]" # accept prob
143         f"[outputs[4]:^9.4g] " # step size
144         f"[beta_np:^9.4g] " # beta
145         f"[np.mean(outputs[5]):^9.4g] " # avg. actions
146         f"[np.mean(outputs[6]):^9.4g] " # avg. plaqs.
147         f"[u1_plaq_exact(beta_np):^9.4g] " # exact plaq.
148         f"[outputs[8]:^9.4g] " # charge diff
149         f"[outputs[9]:^9.4g]" # learning rate
150     )
151
152     return out_data, data_str
153
154 def train(self, train_steps, **kwargs):
155     """Train the L2HMC sampler for `train_steps`.  

156
157     Args:  

158         train_steps: Integer number of training steps to perform.  

159         **kwargs: Possible (key, value) pairs are  

160             'samples_np': Array of initial samples used to start  

161                 training.  

162             'beta_np': Initial value of beta used in annealing  

163                 schedule.  

164             'trace': Flag specifying that the training loop should be  

165                 ran through a profiler.  

166
167     initial_step = kwargs.get('initial_step', 0)
168     samples_np = kwargs.get('samples_np', None)
169     beta_np = kwargs.get('beta_np', None)
170     net_weights = kwargs.get('net_weights', [1., 1., 1.])
171
172     if beta_np is None:
173         beta_np = self.model.beta_init
174
175     if samples_np is None:
176         samples_np = np.reshape(
177             np.array(self.model.lattice.samples, dtype=NP_FLOAT),
178             (self.model.num_samples, self.model.x_dim)
179         )
180
181     assert samples_np.shape == self.model.x.shape
182
183     try:
184         io.log(TRAIN_HEADER)
185         for step in range(initial_step, train_steps):
186             out_data, data_str = self.train_step(step,
187                                                 samples_np,
188                                                 net_weights=net_weights,
189                                                 train_steps=train_steps)
190             samples_np = out_data['samples']
191
192             if self.logger is not None:
193                 self.logger.update_training(self.sess,
194                                             out_data,
195                                             net_weights,
196                                             data_str)
197
198             if self.logger is not None:
199                 self.logger.write_train_strings()
200
201     except (KeyboardInterrupt, SystemExit):
202         io.log("\nKeyboardInterrupt detected!")
203         io.log("Saving current state and exiting.")
204         if self.logger is not None:
205             self.logger.update_training(out_data, data_str)

```

A.13 l2hmc-qcd/runners/runner.py

```

1  """
2  runner.py
3
4  Implements GaugeModelRunner class responsible for running the L2HMC algorithm
5  on a U(1) gauge model.
6
7  Author: Sam Foreman (github: @saforem2)
8  Date: 04/09/2019
9  """
10
11 # import os
12 import time
13 # import pickle
14 # from scipy.stats import sem
15 # from collections import Counter, OrderedDict
16
17 import numpy as np
18
19 import utils.file_io as io
20 from lattice.lattice import u1_plaq_exact
21 from loggers.run_logger import RunLogger
22 from config import RUN_HEADER
23
24 class GaugeModelRunner:
25
26     def __init__(self, sess, params, inputs, run_ops, logger=None):
27         """
28             Args:
29                 sess: tf.Session() object.
30                 model: GaugeModel object (defined in `models/gauge_model.py`)
31                 logger: RunLogger object (defined in `loggers/run_logger.py`),
32                         defaults to None. This is to simplify communication when using
33                         Horovod since the RunLogger object exists only on
34                         hvd.rank() == 0, which is responsible for all file I/O.
35         """
36
37         self.sess = sess
38         self.params = params
39         self.logger = logger
40
41         if logger is not None:
42             self.inputs_dict = self.logger.inputs_dict
43             self.run_ops_dict = self.logger.run_ops_dict
44         else:
45             self.inputs_dict = RunLogger.build_inputs_dict(inputs)
46             self.run_ops_dict = RunLogger.build_run_ops_dict(params, run_ops)
47
48         self.eps = self.sess.run(self.run_ops_dict['dynamics_eps'])
49
50     def calc_charge_autocorrelation(self, charges):
51         autocorr = np.correlate(charges, charges, mode='full')
52         return autocorr
53
54     def run_step(self, step, run_steps, inputs, net_weights):
55         """
56             Perform a single run step.
57
58             Args:
59                 step (int): Current step.
60                 run_steps (int): Total number of run_steps to perform.
61                 inputs (tuple): Tuple consisting of (samples_in, beta_np, eps,
62                               plaq_exact) where samples_in (np.ndarray) is the input batch of
63                               samples, beta (float) is the input value of beta, eps is the
64                               step size, and plaq_exact (float) is the expected avg. value of
65                               the plaquette at this value of beta.
66
67             Returns:
68                 out_data: Dictionary containing the output of running all of the
69                             tensorflow operations in `ops` defined below.
70         """

```

```

68     samples_in, beta_np, eps, plaq_exact = inputs
69
70     keys = ['x_out', 'px', 'actions_op',
71             'plaqs_op', 'charges_op', 'charge_diffs_op']
72
73     if self.params['save_lf']:
74         keys.extend(['lf_out_f', 'pxs_out_f',
75                     'lf_out_b', 'pxs_out_b',
76                     'masks_f', 'masks_b',
77                     'logdets_f', 'logdets_b',
78                     'sumlogdet_f', 'sumlogdet_b'])
79
80     ops = [self.run_ops_dict[k] for k in keys]
81
82     feed_dict = {
83         self.inputs_dict['x']: samples_in,
84         self.inputs_dict['beta']: beta_np,
85         self.inputs_dict['net_weights'][0]: net_weights[0],
86         self.inputs_dict['net_weights'][1]: net_weights[1],
87         self.inputs_dict['net_weights'][2]: net_weights[2],
88         self.inputs_dict['train_phase']: False
89     }
90
91     start_time = time.time()
92     outputs = self.sess.run(ops, feed_dict=feed_dict)
93     dt = time.time() - start_time
94
95     out_data = {
96         'step': step,
97         'beta': beta_np,
98         'eps': self.eps,
99         'samples': np.mod(outputs[0], 2 * np.pi),
100        'samples_orig': outputs[0],
101        'px': outputs[1],
102        'actions': outputs[2],
103        'plaqs': outputs[3],
104        'charges': outputs[4],
105        'charge_diffs': outputs[5],
106    }
107
108    if self.params['save_lf']:
109        out_data.update({k: v for k, v in zip(keys[6:], outputs[6:])})
110        # lf_outputs = {}
111        # for key, val in zip(keys[6:], outputs[6:]):
112        #     lf_outputs[key] = val
113        #     out_data.update(lf_outputs)
114
115    data_str = (f'{step:>5g}/{run_steps:<6g} '
116                f'{dt:^9.4g} '           # time / step
117                f'{np.mean(outputs[1]):^9.4g} '   # accept. prob
118                f'{self.eps:^9.4g} '          # step size
119                f'{beta_np:^9.4g} '         # beta val
120                f'{np.mean(outputs[2]):^9.4g} '   # avg. actions
121                f'{np.mean(outputs[3]):^9.4g} '   # avg. plaquettes
122                f'{plaq_exact:^9.4g} '        # exact plaquette val
123                f'{outputs[5]:^9.4g} ')       # top. charge diff
124
125    return out_data, data_str
126
127 def run(self, **kwargs):
128     """Run the simulation to generate samples and calculate observables.
129
130     Args:
131         run_steps: Number of steps to run the sampler for.
132         kwargs: Dictionary of keyword arguments to use for running
133                 inference.
134
135     Returns:
136         observables: Tuple of observables dictionaries consisting of:
137                     (actions_dict, plaqs_dict, charges_dict, charge_diffs_dict).

```

```

138     """
139     run_steps = int(kwargs.get('run_steps', 5000))
140
141     beta = kwargs.get('beta_final', self.params.get('beta_final', 5))
142     net_weights = kwargs.get('net_weights', [1., 1., 1.])
143     therm_frac = kwargs.get('therm_frac', 10)
144
145     if beta is None:
146         beta = self.params['beta_final']
147
148     if net_weights is None:
149         # scale_weight, translation weight, transformation weight
150         net_weights = [1., 1., 1.]
151
152     plaq_exact = u1_plaq_exact(beta)
153
154     x_dim = (self.params['space_size']
155               * self.params['time_size']
156               * self.params['dim'])
157
158     # start with randomly generated samples
159     samples_np = np.random.randn(*self.params['num_samples'], x_dim)
160
161     try:
162         io.log(RUN_HEADER)
163         for step in range(run_steps):
164             inputs = (samples_np, beta, self.eps, plaq_exact)
165             out_data, data_str = self.run_step(step, run_steps,
166                                              inputs, net_weights)
167             samples_np = out_data['samples']
168
169             if self.logger is not None:
170                 self.logger.update(self.sess, out_data,
171                                   net_weights, data_str)
172
173             if self.logger is not None:
174                 self.logger.save_run_data(therm_frac=therm_frac)
175
176     except (KeyboardInterrupt, SystemExit):
177         io.log(80 * '-' + '\n')
178         io.log("\nWARNING: KeyboardInterrupt detected!")
179         io.log("WARNING: Saving current state and exiting.")
180         if self.logger is not None:
181             self.logger.save_run_data(therm_frac=therm_frac)

```

A.14 l2hmc-qcd/loggers/run_logger.py

```
1 """
2 run_logger.py
3
4 Implements RunLogger class responsible for saving/logging data
5 from `run` phase of GaugeModel.
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 04/24/2019
9 """
10 import os
11 import pickle
12 import datetime
13 import shutil
14 import errno
15
16 import tensorflow as tf
17 import numpy as np
18
19 from collections import Counter, OrderedDict
20 from scipy.stats import sem
21 import utils.file_io as io
22
23 from config import RUN_HEADER, NP_FLOAT
24
25 from lattice.lattice import u1_plaq_exact
26 # from utils.tf_logging import variable_summaries
27
28 from .train_logger import save_params
29
30
31 FB_DICT = {
32     'forward': [],
33     'backward': [],
34 }
35
36
37 def arr_from_dict(d, key):
38     return np.array(list(d[key]))
39
40
41 def autocorr(x):
42     autocorr = np.correlate(x, x, mode='full')
43
44     return autocorr[autocorr.size // 2:]
45
46
47 def copy(src, dest):
48     try:
49         shutil.copytree(src, dest)
50     except OSError as e:
51         # If the error was caused because the source wasn't a directory
52         if e.errno == errno.ENOTDIR:
53             shutil.copy(src, dest)
54         else:
55             io.log(f'Directory not copied. Error: {e}')
56
57
58 class RunLogger:
59     def __init__(self, params, inputs, run_ops, save_lf_data=False):
60         """
61             Args:
62                 model: GaugeModel object.
63                 log_dir: Existing logdir from `TrainLogger`.
64             """
65         self.params = params
66         self.save_lf_data = save_lf_data
67         self.summaries = params['summaries']
```

```

68     assert os.path.isdir(params['log_dir'])
69     self.log_dir = params['log_dir']
70
71     runs_dir = os.path.join(self.log_dir, 'runs')
72     figs_dir = os.path.join(self.log_dir, 'figures')
73     run_summaries_dir = os.path.join(self.log_dir, 'summaries', 'run')
74     if os.path.isdir(runs_dir) or os.path.isdir(figs_dir):
75         now = datetime.datetime.now()
76         time_str = now.strftime("%H%M")
77         if os.path.isdir(runs_dir):
78             renamed_runs_dir = runs_dir + f'_{time_str}'
79             # io.log(f'Renaming existing runs_dir to: {renamed_runs_dir}')
80             io.log(f'Copying existing runs_dir to: {renamed_runs_dir}')
81             # io.check_else_make_dir(renamed_runs_dir)
82             copy(runs_dir, renamed_runs_dir)
83             # os.rename(runs_dir, renamed_runs_dir)
84         if os.path.isdir(figs_dir):
85             renamed_figs_dir = figs_dir + f'_{time_str}'
86             # io.log(f'Renaming existing figs_dir to: {renamed_figs_dir}')
87             io.log(f'Copying existing figs_dir to: {renamed_figs_dir}')
88             # io.check_else_make_dir(renamed_figs_dir)
89             copy(figs_dir, renamed_figs_dir)
90             # os.rename(figs_dir, renamed_figs_dir)
91         if os.path.isdir(run_summaries_dir):
92             new_rsd = run_summaries_dir + f'_{time_str}'
93             io.log(f'Copying existing run summaries dir to: {new_rsd}')
94             # io.check_else_make_dir(new_rsd)
95             copy(run_summaries_dir, new_rsd)
96
97         self.runs_dir = runs_dir
98         io.check_else_make_dir(self.runs_dir)
99
100        self.figs_dir = figs_dir
101        io.check_else_make_dir(self.figs_dir)
102
103        self.run_summaries_dir = run_summaries_dir
104        io.check_else_make_dir(self.run_summaries_dir)
105
106        # self._reset_counter = 0
107        self.run_steps = None
108        self.beta = None
109        self.run_data = {}
110        self.run_stats = {}
111        self.run_strings = [RUN_HEADER]
112
113        if params['save_lf']:
114            self.samples_arr = []
115            self.lf_out = FB_DICT.copy()
116            self.pxs_out = FB_DICT.copy()
117            self.masks = FB_DICT.copy()
118            self.logdets = FB_DICT.copy()
119            self.sumlogdet = FB_DICT.copy()
120            self.l2hmc_fns = FB_DICT.copy()
121
122        self.run_ops_dict = self.build_run_ops_dict(params, run_ops)
123        self.inputs_dict = self.build_inputs_dict(inputs)
124
125        if self.summaries:
126            # self.run_summaries_dir = os.path.join(self.log_dir,
127            #                                         'summaries', 'run')
128            # io.check_else_make_dir(self.run_summaries_dir)
129            self.writer = tf.summary.FileWriter(self.run_summaries_dir,
130                                              tf.get_default_graph())
131            self.create_summaries()
132
133    @staticmethod
134    def build_run_ops_dict(params, run_ops):
135        """Build dictionary of tensorflow operations used for inference."""
136    def get_lf_keys(direction):
137        base_keys = ['lf_out', 'pxs_out', 'masks',

```

```

138             'logdets', 'sumlogdet', 'fns_out']
139     return [k + f'_{direction}' for k in base_keys]
140
141     keys = ['x_out', 'px', 'actions_op',
142             'plaqs_op', 'avg_plaqs_op',
143             'charges_op', 'charge_diffs_op']
144
145     run_ops_dict = {key: run_ops[idx] for idx, key in enumerate(keys)}
146
147     if params['save_lf']:
148         keys.extend(get_lf_keys('f'))
149         keys.extend(get_lf_keys('b'))
150         for key, val in zip(keys[7:], run_ops[7:]):
151             run_ops_dict.update({key: val})
152
153     run_ops_dict['dynamics_eps'] = run_ops[-1]
154
155     return run_ops_dict
156
157 @staticmethod
158 def build_inputs_dict(inputs):
159     """Build dictionary of tensorflow placeholders used as inputs."""
160     keys = ['x', 'beta', 'charge_weight', 'train_phase']
161     num_keys = len(keys)
162
163     inputs_dict = dict(zip(keys, inputs[:num_keys]))
164     inputs_dict.update({'net_weights': inputs[num_keys:]})
165
166     return inputs_dict
167
168 def create_summaries(self):
169     """Create summary objects for logging in TensorBoard."""
170     summary_list = tf.get_collection(tf.GraphKeys.SUMMARIES)
171     ignore_strings = ['loss', 'learning_rate',
172                       'step_size', 'train', 'eps']
173     run_summaries = [
174         i for i in summary_list if not any(
175             s in i.name for s in ignore_strings
176         )
177     ]
178     self.summary_op = tf.summary.merge(run_summaries)
179
180 def log_step(self, sess, step, samples_np, beta_np, net_weights):
181     """Update self.logger.summaries."""
182     feed_dict = {
183         self.inputs_dict['x']: samples_np,
184         self.inputs_dict['beta']: beta_np,
185         self.inputs_dict['net_weights'][0]: net_weights[0],
186         self.inputs_dict['net_weights'][1]: net_weights[1],
187         self.inputs_dict['net_weights'][2]: net_weights[2],
188         self.inputs_dict['train_phase']: False
189     }
190     summary_str = sess.run(self.summary_op, feed_dict=feed_dict)
191
192     self.writer.add_summary(summary_str, global_step=step)
193     self.writer.flush()
194
195 @classmethod
196 def _clear(cls):
197     cls.run_data = None
198     cls.run_strings = None
199     if cls.params['save_lf']:
200         cls.samples_arr = None
201         cls.lf_out = None
202         cls.pxs_out = None
203         cls.masks = None
204         cls.logdets = None
205         cls.sumlogdet = None
206         cls.l2hmc_fns = None
207

```

```

208     cls.params['net_weights'] = None
209     cls.run_dir = None
210     if cls.summaries:
211         cls.run_summary_dir = None
212         cls.writer = None
213
214     def clear(self):
215         self.run_data = None
216         self.run_strings = None
217         if self.params['save_lf']:
218             self.samples_arr = None
219             self.lf_out = None
220             self.pxs_out = None
221             self.masks = None
222             self.logdets = None
223             self.sumlogdet = None
224             self.l2hmc_fns = None
225         self.params['net_weights'] = None
226         self.run_dir = None
227         if self.summaries:
228             self.run_summary_dir = None
229             self.writer = None
230
231     def existing_run(self, run_str):
232         """Check if this run has been completed previously, if so skip it."""
233         run_dir = os.path.join(self.runs_dir, run_str)
234         run_summary_dir = os.path.join(self.run_summaries_dir, run_str)
235         if os.path.isdir(run_dir) and os.path.isdir(run_summary_dir):
236             return True
237         return False
238
239     def _get_run_str(self, **kwargs):
240         """Parse parameter values and create unique string to name the dir."""
241         run_steps = kwargs.get('run_steps', 5000)
242         beta = kwargs.get('beta', 5.)
243         net_weights = kwargs.get('net_weights', [1., 1., 1.])
244         eps_np = kwargs.get('eps', None)
245         dir_append = kwargs.get('dir_append', None)
246         # eps = self.model.eps
247         eps_str = f'{eps_np:.3}'.replace('.', '')
248         beta_str = f'{beta:.3}'.replace('.', '')
249         # qw_str = f'{charge_weight:.3}'.replace('.', '')
250         scale_wstr = f'{net_weights[0]:.3.2f}'.replace('.', '')
251         transl_wstr = f'{net_weights[1]:.3.2f}'.replace('.', '')
252         transf_wstr = f'{net_weights[2]:.3.2f}'.replace('.', '')
253
254         run_str = (f'steps{run_steps}'
255                    f'_beta{beta_str}'
256                    f'_eps{eps_str}'
257                    # f'_qw_{qw_str:.2}'
258                    f'_S{scale_wstr}'
259                    f'_T{transl_wstr}'
260                    f'_Q{transf_wstr}')
261         # f'_{self._reset_counter}')
262
263         if dir_append is not None:
264             run_str += dir_append
265
266     return run_str
267
268     def reset(self, **kwargs):
269         """Reset run_data and run_strings to prep for new run."""
270         run_steps = kwargs.get('run_steps', 5000)
271         beta = kwargs.get('beta', 5.)
272         net_weights = kwargs.get('net_weights', [1., 1., 1.])
273         # eps_np = kwargs.get('eps', None)
274         # dir_append = kwargs.get('dir_append', None)
275
276         self.run_steps = int(run_steps)
277         self.beta = beta

```

```

278     self.run_data = {
279         'px': {},
280         'actions': {},
281         'plaqs': {},
282         'charges': {},
283         'charge_diffs': {}
284     }
285     self.run_stats = {}
286     self.run_strings = []
287     # if self.model.save_lf:
288     if self.params['save_lf']:
289         self.samples_arr = []
290         self.lf_out = FB_DICT.copy()
291         self.pxs_out = FB_DICT.copy()
292         self.masks = FB_DICT.copy()
293         self.logdets = FB_DICT.copy()
294         self.sumlogdet = FB_DICT.copy()
295         self.l2hmc_fns = FB_DICT.copy()
296
297     # params = self.model.params
298     self.params['net_weights'] = net_weights
299
300     run_str = self._get_run_str(**kwargs)
301     self._set_run_dirs(run_str)
302     self._save_net_weights(net_weights)
303
304     if self.summaries:
305         self.writer = tf.summary.FileWriter(self.run_summary_dir,
306                                         tf.get_default_graph())
307         save_params(self.params, self.run_dir)
308         # self._reset_counter += 1
309
310     # return self.run_dir, run_str
311
312 def _save_net_weights(self, net_weights):
313     """Write net weights to a '.txt' file; append if existing."""
314     # def _round_float_as_str(f):
315     #     return f'{f:.3g}'
316     weights_txt_file = os.path.join(self.run_dir, 'net_weights.txt')
317     nw_str = [f'{w:.3g}' for w in net_weights]
318     # nw_str = [_round_float_as_str(w) for w in net_weights]
319     w_str = '[' + nw_str[0] + nw_str[1] + nw_str[2] + ']'
320     if not os.path.isfile(weights_txt_file):
321         with open(weights_txt_file, 'w') as f:
322             f.write('[scale_weight,\n'
323                   'translation_weight,\n'
324                   'transformation_weight]')
325             f.write(80 * '-' + '\n')
326     else:
327         with open(weights_txt_file, 'a') as f:
328             f.write(w_str)
329
330 def _set_run_dirs(self, run_str):
331     """Sets dirs containing data about inference run using run_str."""
332     self.run_dir = os.path.join(self.runs_dir, run_str)
333     io.check_else_make_dir(self.run_dir)
334     if self.summaries:
335         self.run_summary_dir = os.path.join(self.run_summaries_dir,
336                                           run_str)
337         io.check_else_make_dir(self.run_summary_dir)
338
339 def update(self, sess, data, net_weights, data_str):
340     """Update run_data and append data_str to data_strings."""
341     # projection of samples onto [0, 2pi) done in run_step above
342     step = data['step']
343     beta = data['beta']
344     key = (step, beta)
345
346     obs_keys = ['px', 'actions', 'plaqs', 'charges', 'charge_diffs']

```

```

348     for k in obs_keys:
349         self.run_data[k][key] = data[k]
350
351     if self.params['save_lf']:
352         samples_np = data['samples']
353         self.samples_arr.append(samples_np)
354         self.lf_out['forward'].extend(np.array(data['lf_out_f']))
355         self.lf_out['backward'].extend(np.array(data['lf_out_b']))
356         self.logdets['forward'].extend(np.array(data['logdets_f']))
357         self.logdets['backward'].extend(np.array(data['logdets_b']))
358         self.sumlogdet['forward'].append(np.array(data['sumlogdet_f']))
359         self.sumlogdet['backward'].append(np.array(data['sumlogdet_b']))
360         # self.pxs_out['forward'].extend(np.array(data['pxs_out_f']))
361         # self.pxs_out['backward'].extend(np.array(data['pxs_out_b']))
362         # self.masks['forward'].extend(np.array(data['masks_f']))
363         # self.masks['backward'].extend(np.array(data['masks_b']))
364
365     self.run_strings.append(data_str)
366
367     if self.summaries and (step + 1) % self.params['Logging_steps'] == 0:
368         self.log_step(sess, step, data['samples'], beta, net_weights)
369
370     if step % (10 * self.params['print_steps']) == 0:
371         io.log(data_str)
372
373     if step % 100 == 0:
374         io.log(RUN_HEADER)
375
376 def calc_observables_stats(self, run_data, therm_frac=10):
377     """Calculate statistics for lattice observables.
378
379     Args:
380         run_data: Dictionary of observables data. Keys denote the
381             observables name.
382         therm_frac: Fraction of data to throw out for thermalization.
383
384     Returns:
385         stats: Dictionary containing statistics for each observable in
386             run_data. Additionally, contains `charge_probs` which is a
387             dictionary of the form {charge_val: charge_val_probability}.
388     """
389     def get_stats(data, t_frac=10):
390         if isinstance(data, dict):
391             arr = np.array(list(data.values()))
392         elif isinstance(data, (list, np.ndarray)):
393             arr = np.array(data)
394
395         num_steps = arr.shape[0]
396         therm_steps = num_steps // t_frac
397         arr = arr[therm_steps:, :]
398         avg = np.mean(arr, axis=0)
399         err = sem(arr, axis=0)
400         stats = np.array([avg, err]).T
401         return stats
402
403     actions_stats = get_stats(run_data['actions'], therm_frac)
404     plaqs_stats = get_stats(run_data['plaqs'], therm_frac)
405
406     charges_arr = np.array(list(run_data['charges'].values()), dtype=int)
407     charges_stats = get_stats(charges_arr, therm_frac)
408
409     suspect_arr = charges_arr ** 2
410     suspect_stats = get_stats(suspect_arr)
411
412     charge_probs = {}
413     counts = Counter(list(charges_arr.flatten()))
414     total_counts = np.sum(list(counts.values()))
415     for key, val in counts.items():
416         charge_probs[key] = val / total_counts
417

```

```

418 charge_probs = OrderedDict(sorted(charge_probs.items(),
419                             key=lambda k: k[0]))
420
421     stats = {
422         'actions': actions_stats,
423         'plaqs': plaqs_stats,
424         'charges': charges_stats,
425         'suscept': suspect_stats,
426         'charge_probs': charge_probs
427     }
428
429     return stats
430
431 def save_attr(self, name, attr, out_dir=None, dtype=NP_FLOAT):
432     if out_dir is None:
433         out_dir = self.run_dir
434
435     assert os.path.isdir(out_dir)
436     out_file = os.path.join(out_dir, name + '.npz')
437
438     if not isinstance(attr, np.ndarray) or attr.dtype != dtype:
439         attr = np.array(attr, dtype=dtype)
440
441     if os.path.isfile(out_file):
442         io.log(f'File {out_file} already exists. Skipping.')
443     else:
444         io.log(f'Saving {name} to: {out_file}')
445         np.savez_compressed(out_file, attr)
446
447 def save_run_data(self, therm_frac=10):
448     """Save run information."""
449     observables_dir = os.path.join(self.run_dir, 'observables')
450
451     io.check_else_make_dir(self.run_dir)
452     io.check_else_make_dir(observables_dir)
453
454     if self.save_lf_data:
455         keys = ['lf_out', 'masks', 'logdets', 'sumlogdet', 'pxs_out']
456         for key in keys:
457             f = 'forward'
458             b = 'backward'
459             self.save_attr(key + f'_{{f}}', getattr(self, key)[f])
460             self.save_attr(key + f'_{{b}}', getattr(self, key)[b])
461
462     run_stats = self.calc_observables_stats(self.run_data, therm_frac)
463     charges = self.run_data['charges']
464     charges_arr = np.array(list(charges.values()))
465     charges_autocorrs = [autocorr(x) for x in charges_arr.T]
466     charges_autocorrs = [x / np.max(x) for x in charges_autocorrs]
467     self.run_data['charges_autocorrs'] = charges_autocorrs
468
469     data_file = os.path.join(self.run_dir, 'run_data.pkl')
470     io.log(f"Saving run_data to: {data_file}.")
471     with open(data_file, 'wb') as f:
472         pickle.dump(self.run_data, f)
473
474     stats_data_file = os.path.join(self.run_dir, 'run_stats.pkl')
475     io.log(f"Saving run_stats to: {stats_data_file}.")
476     with open(stats_data_file, 'wb') as f:
477         pickle.dump(run_stats, f)
478
479     for key, val in self.run_data.items():
480         out_file = key + '.pkl'
481         out_file = os.path.join(observables_dir, out_file)
482         io.save_data(val, out_file, name=key)
483
484     for key, val in run_stats.items():
485         out_file = key + '_stats.pkl'
486         out_file = os.path.join(observables_dir, out_file)
487         io.save_data(val, out_file, name=key)

```

```

488
489     history_file = os.path.join(self.run_dir, 'run_history.txt')
490     io.write(RUN_HEADER, history_file, 'w')
491     _ = [io.write(s, history_file, 'a') for s in self.run_strings]
492
493     self.write_run_stats(run_stats, therm_frac)
494
495     def write_run_stats(self, stats, therm_frac=10):
496         """Write statistics in human readable format to .txt file."""
497         # run_steps = kwargs['run_steps']
498         # beta = kwargs['beta']
499         # current_step = kwargs['current_step']
500         # therm_steps = kwargs['therm_steps']
501         # training = kwargs['training']
502         # run_dir = kwargs['run_dir']
503         therm_steps = self.run_steps // therm_frac
504
505         out_file = os.path.join(self.run_dir, 'run_stats.txt')
506
507         actions_avg, actions_err = stats['actions'].mean(axis=0)
508         plaqs_avg, plaqs_err = stats['plaqs'].mean(axis=0)
509         charges_avg, charges_err = stats['charges'].mean(axis=0)
510         suspect_avg, suspect_err = stats['suscept'].mean(axis=0)
511
512         # ns = self.model.num_samples
513         ns = self.params['num_samples']
514         suspect_k1 = f' \navg. over all {ns} samples < Q >'
515         suspect_k2 = f' \navg. over all {ns} samples < Q^2 >'
516         actions_k1 = f' \navg. over all {ns} samples < action >'
517         plaqs_k1 = f' \n avg. over all {ns} samples < plaq >'
518
519         _est_key = ' \nestimate +/- stderr'
520
521         suspect_ss = {
522             suspect_k1: f'{charges_avg:.4g} +/- {charges_err:.4g}',
523             suspect_k2: f'{suspect_avg:.4g} +/- {suspect_err:.4g}',
524             _est_key: {}
525         }
526
527         actions_ss = {
528             actions_k1: f'{actions_avg:.4g} +/- {actions_err:.4g}\n',
529             _est_key: {}
530         }
531
532         plaqs_ss = {
533             'exact_plaq': f'{u1_plaq_exact(self.beta):.4g}\n',
534             plaqs_k1: f'{plaqs_avg:.4g} +/- {plaqs_err:.4g}\n',
535             _est_key: {}
536         }
537
538         def format_stats(x, name=None):
539             return [f'{name}: {i[0]:.4g} +/- {i[1]:.4g}' for i in x]
540
541         def zip_keys_vals(stats_strings, keys, vals):
542             for k, v in zip(keys, vals):
543                 stats_strings[_est_key][k] = v
544             return stats_strings
545
546         keys = [f"sample {idx}" for idx in range(ns)]
547
548         suspect_vals = format_stats(stats['suscept'], '< Q^2 >')
549         actions_vals = format_stats(stats['actions'], '< action >')
550         plaqs_vals = format_stats(stats['plaqs'], '< plaq >')
551
552         suspect_ss = zip_keys_vals(suspect_ss, keys, suspect_vals)
553         actions_ss = zip_keys_vals(actions_ss, keys, actions_vals)
554         plaqs_ss = zip_keys_vals(plaqs_ss, keys, plaqs_vals)
555
556         def accumulate_strings(d):
557             all_strings = []

```

```

558     for k1, v1 in d.items():
559         if isinstance(v1, dict):
560             for k2, v2 in v1.items():
561                 all_strings.append(f'{k2} {v2}')
562         else:
563             all_strings.append(f'{k1}: {v1}\n')
564
565     return all_strings
566
567 actions_strings = accumulate_strings(actions_ss)
568 plaqs_strings = accumulate_strings(plaqs_ss)
569 suspect_strings = accumulate_strings(suscept_ss)
570
571 charge_probs_strings = []
572 for k, v in stats['charge_probs'].items():
573     charge_probs_strings.append(f' probability[Q = {k}]: {v}\n')
574
575 run_str = (f" stats for {ns} chains ran for {self.run_steps} steps "
576             f" at beta = {self.beta}.")
577
578 str0 = "Topological susceptibility" + run_str
579 str1 = "Total actions" + run_str
580 str2 = "Average plaquette" + run_str
581 str3 = "Topological charge probabilities" + run_str[6:]
582 therm_str = (
583     f'Ignoring first {therm_steps} steps for thermalization.'
584 )
585
586 ss0 = ('-' * max(len(str0), len(therm_str)))
587 ss1 = ('-' * max(len(str1), len(therm_str)))
588 ss2 = ('-' * max(len(str2), len(therm_str)))
589 ss3 = ('-' * max(len(str3), len(therm_str)))
590
591 io.log(f"Writing statistics to: {out_file}")
592
593 def log_and_write(sep_str, str0, therm_str, stats_strings, file):
594     io.log(sep_str)
595     io.log(str0)
596     io.log(therm_str)
597     io.log('')
598     _ = [io.log(s) for s in stats_strings]
599     io.log(sep_str)
600     io.log('')
601
602     io.write(sep_str, file, 'a')
603     io.write(str0, file, 'a')
604     io.write(therm_str, file, 'a')
605     _ = [io.write(s, file, 'a') for s in stats_strings]
606     io.write('\n', file, 'a')
607
608 log_and_write(ss0, str0, therm_str, suspect_strings, out_file)
609 log_and_write(ss1, str1, therm_str, actions_strings, out_file)
610 log_and_write(ss2, str2, therm_str, plaqs_strings, out_file)
611 log_and_write(ss3, str3, therm_str, charge_probs_strings, out_file)
612 log_and_write(ss3, str3, therm_str, charge_probs_strings, out_file)
613 log_and_write(ss3, str3, therm_str, charge_probs_strings, out_file)

```

A.15 l2hmc-qcd/loggers/train_logger.py

```
1 """
2 train_logger.py
3
4 Implements TrainLogger class responsible for saving/logging data from
5 GaugeModel.
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 04/09/2019
9 """
10 from __future__ import absolute_import
11 from __future__ import division
12 from __future__ import print_function
13
14 import os
15 import pickle
16
17 import tensorflow as tf
18
19 import utils.file_io as io
20
21 from config import TRAIN_HEADER
22 from .summary_utils import create_summaries
23 from utils.file_io import save_params
24
25
26 class TrainLogger:
27     def __init__(self, model, log_dir, summaries=False):
28         # self.sess = sess
29         self.model = model
30         self.summaries = summaries
31
32         self.charges_dict = {}
33         self.charge_diffs_dict = {}
34
35         self._current_state = {
36             'step': 0,
37             'beta': self.model.beta_init,
38             'eps': self.model.eps,
39             'lr': self.model.lr_init,
40             'samples': self.model.lattice.samples_tensor,
41         }
42
43         self.train_data_strings = [TRAIN_HEADER]
44         self.train_data = {
45             'loss': {},
46             'actions': {},
47             'plaqs': {},
48             'charges': {},
49             'charge_diffs': {},
50             'px': {}
51         }
52
53         if self.model.save_lf:
54             self.train_data['l2hmc_fns'] = {
55                 'forward': [],
56                 'backward': []
57             }
58
59             # log_dir will be None if using_hvd and hvd.rank() != 0
60             # this prevents workers on different ranks from corrupting checkpoints
61             # if log_dir is not None and self.is_chief:
62             self._create_dir_structure(log_dir)
63             save_params(self.model.params, self.log_dir)
64
65         if self.summaries:
66             # with tf.variable_scope('train_summaries'):
67             self.writer = tf.summary.FileWriter(self.train_summary_dir,
```

```

68             tf.get_default_graph())
69         self.summary_writer, self.summary_op = create_summaries(
70             model, self.train_summary_dir, training=True
71         )
72
73     def _create_dir_structure(self, log_dir):
74         """Create relevant directories for storing data.
75
76         Args:
77             log_dir: Root directory in which all other directories are created.
78
79         Returns:
80             None
81         """
82         dirs = {
83             'log_dir': log_dir,
84             'checkpoint_dir': os.path.join(log_dir, 'checkpoints'),
85             'train_dir': os.path.join(log_dir, 'training'),
86             'train_summary_dir': os.path.join(log_dir, 'summaries', 'train'),
87         }
88         files = {
89             'train_log_file': os.path.join(dirs['train_dir'],
90                                             'training_log.txt'),
91             'current_state_file': os.path.join(dirs['train_dir'],
92                                              'current_state.pkl')
93         }
94
95         for key, val in dirs.items():
96             io.check_else_make_dir(val)
97             setattr(self, key, val)
98
99         for key, val in files.items():
100            setattr(self, key, val)
101
102    def save_current_state(self):
103        """Save current state to pickle file.
104
105        The current state contains the following, which makes life easier if
106        we're trying to restore training from a saved checkpoint:
107            * most recent samples
108            * learning_rate
109            * beta
110            * dynamics.eps
111            * training_step
112
113        with open(self.current_state_file, 'wb') as f:
114            pickle.dump(self._current_state, f)
115
116    def log_step(self, sess, step, samples_np, beta_np, net_weights):
117        """Update self.logger.summaries."""
118        feed_dict = {
119            self.model.x: samples_np,
120            self.model.beta: beta_np,
121            self.model.net_weights[0]: net_weights[0],
122            self.model.net_weights[1]: net_weights[1],
123            self.model.net_weights[2]: net_weights[2],
124            self.model.train_phase: True
125        }
126        summary_str = sess.run(self.summary_op, feed_dict=feed_dict)
127
128        self.writer.add_summary(summary_str, global_step=step)
129        self.writer.flush()
130
131    def update_training(self, sess, data, net_weights, data_str):
132        """Update _current_state and train_data."""
133        step = data['step']
134        beta = data['beta']
135        self._current_state['step'] = step
136        self._current_state['beta'] = beta
137        self._current_state['lr'] = data['lr']

```

```

138     self._current_state['eps'] = data['eps']
139     self._current_state['samples'] = data['samples']
140     self._current_state['net_weights'] = net_weights
141
142     key = (step, beta)
143
144     self.charges_dict[key] = data['charges']
145     self.charge_diffs_dict[key] = data['charge_diffs']
146
147     obs_keys = ['loss', 'actions',
148                 'plags', 'charges',
149                 'charge_diffs', 'px']
150     for obs_key in obs_keys:
151         self.train_data[obs_key][key] = data[obs_key]
152
153     self.train_data_strings.append(data_str)
154     if step % self.model.print_steps == 0:
155         io.log(data_str)
156
157     if self.summaries and (step + 1) % self.model.logging_steps == 0:
158         self.log_step(sess, step, data['samples'], beta, net_weights)
159
160     if (step + 1) % self.model.save_steps == 0:
161         self.save_current_state()
162
163     if step % 100 == 0:
164         io.log(TRAIN_HEADER)
165
166     def write_train_strings(self):
167         """Write training strings out to file."""
168         tlf = self.train_log_file
169         _ = [io.write(s, tlf, 'a') for s in self.train_data_strings]

```

A.16 l2hmc-qcd/plotters/gauge_model_plotter.py

```
1 """
2 plotters.py
3
4 Implements GaugeModelPlotter class, responsible for loading and plotting
5 gauge model observables.
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 04/10/2019
9 """
10
11 import os
12 import numpy as np
13 import utils.file_io as io
14
15 from collections import Counter, OrderedDict
16 from scipy.stats import sem
17
18 from lattice.lattice import u1_plaq_exact
19 from config import COLORS, MARKERS, HAS_MATPLOTLIB
20
21 from .plot_utils import MPL_PARAMS, plot_multiple_lines
22
23 if HAS_MATPLOTLIB:
24     import matplotlib as mpl
25     import matplotlib.pyplot as plt
26     mpl.rcParams.update(MPL_PARAMS)
27
28 def arr_from_dict(d, key):
29     return np.array(list(d[key].values()))
30
31
32 def get_out_file(out_dir, out_str):
33     return os.path.join(out_dir, out_str + '.pdf')
34
35
36 class GaugeModelPlotter:
37     def __init__(self, params, figs_dir=None, experiment=None):
38         self.figs_dir = figs_dir
39         self.params = params
40         # self.model = model
41         if experiment is not None:
42             self.experiment = experiment
43
44     def calc_stats(self, data, therm_frac=10):
45         """Calculate observables statistics.
46
47         Args:
48             data (dict): Run data.
49             therm_frac (int): Percent of total steps to ignore to account for
50             thermalization.
51
52         Returns:
53             stats: Dictionary containing statistics for actions, plaquettes,
54             top. charges, and charge probabilities. For each of the
55             observables (actions, plaquettes, charges), the dictionary values
56             consist of a tuple of the form: (data, error), and
57             charge_probabilities is a dictionary of the form:
58                 {charge_val: charge_val_probability}
59
60         actions = arr_from_dict(data, 'actions')
61         plaqs = arr_from_dict(data, 'plaqs')
62         charges = arr_from_dict(data, 'charges')
63
64         charge_probs = {}
65         counts = Counter(list(charges.flatten()))
66         total_counts = np.sum(list(counts.values()))
67         for key, val in counts.items():
68             charge_probs[key] = (val / total_counts, sem(val / total_counts))
```

```

68     charge_probs[key] = val / total_counts
69
70     charge_probs = OrderedDict(sorted(charge_probs.items(),
71                                     key=lambda k: k[0]))
72
73     def get_mean_err(x):
74         num_steps = x.shape[0]
75         therm_steps = num_steps // therm_frac
76         x = x[therm_steps:, :]
77         avg = np.mean(x, axis=0)
78         err = sem(x)
79         return avg, err
80
81     stats = {
82         'actions': get_mean_err(actions),
83         'plaqs': get_mean_err(plaqs),
84         'charges': get_mean_err(charges),
85         'suscept': get_mean_err(charges ** 2),
86         'charge_probs': charge_probs
87     }
88
89     return stats
90
91 def log_figure(self):
92     try:
93         self.experiment.log_figure()
94     except AttributeError:
95         pass
96
97 def _parse_data(self, data, beta):
98     """Helper method for extracting relevant data from `data`."""
99     actions = arr_from_dict(data, 'actions')
100    plaqs = arr_from_dict(data, 'plaqs')
101    charges = np.array(arr_from_dict(data, 'charges'), dtype=int)
102    charge_diffs = arr_from_dict(data, 'charge_diffs')
103    charge_autocorrs = np.array(data['charges_autocorrs'])
104    plaqs_diffs = plaqs - u1_plaq_exact(beta)
105
106    actions_avg = np.mean(actions, axis=1)
107    actions_err = sem(actions, axis=1)
108
109    plaqs_avg = np.mean(plaqs, axis=1)
110    plaqs_err = sem(plaqs, axis=1)
111
112    autocorrs_avg = np.mean(charge_autocorrs.T, axis=1)
113    autocorrs_err = sem(charge_autocorrs.T, axis=1)
114
115    num_steps, num_samples = actions.shape
116    steps_arr = np.arange(num_steps)
117
118    # skip 5% of total number of steps between successive points when
119    # plotting to help smooth out graph
120    skip_steps = max((1, int(0.005 * num_steps)))
121    # ignore first 10% of pts (warmup)
122    warmup_steps = max((1, int(0.01 * num_steps)))
123
124    _charge_diffs = charge_diffs[warmup_steps:][::skip_steps]
125    _plaq_diffs = plaqs_diffs[warmup_steps:][::skip_steps]
126    _steps_diffs = (
127        skip_steps * np.arange(_plaq_diffs.shape[0]) + skip_steps
128    )
129    _plaq_diffs_avg = np.mean(_plaq_diffs, axis=1)
130    _plaq_diffs_err = sem(_plaq_diffs, axis=1)
131
132    xy_data = {
133        'actions': (steps_arr, actions_avg, actions_err),
134        'plaqs': (steps_arr, plaqs_avg, plaqs_err),
135        'charges': (steps_arr, charges.T),
136        'charge_diffs': (_steps_diffs, _charge_diffs.T),
137        'autocorrs': (steps_arr, autocorrs_avg, autocorrs_err),

```

```

138         'plaqs_diffs': (_steps_diffs, _plaq_diffs_avg, _plaq_diffs_err)
139     }
140
141     return xy_data
142
143 def __plot_setup(self, data, **kwargs):
144     """Prepare for plotting observables."""
145     beta = kwargs.get('beta', 5.)
146     run_str = kwargs.get('run_str', None)
147     net_weights = kwargs.get('net_weights', [1., 1., 1.])
148     dir_append = kwargs.get('dir_append', None)
149
150     if dir_append:
151         run_str += dir_append
152
153     self.out_dir = os.path.join(self.figs_dir, run_str)
154     io.check_else_make_dir(self.out_dir)
155
156     # L = self.params['space_size']
157     lf_steps = self.params['num_steps']
158     bs = self.params['num_samples'] # batch size
159     # qw = weights['charge_weight']
160     nw = net_weights
161     sw, translw, transfw = nw
162     title_str = (r"$N_{\mathrm{LF}} = $" + f"{lf_steps}, "
163                  r"$N_{\mathrm{B}} = $" + f"{bs}, "
164                  r"$\mathbf{n}_{\mathrm{w}} = $" + f"[nw[0], nw[1], nw[2]])"
165
166     # r"$L = $" + f"{L}, "
167     # r"$\beta = $" + f"{beta}, "
168     # r"$\alpha_Q = $" + f"{qw}, "
169     kwargs.update({
170         'markers': False,
171         'lines': True,
172         'alpha': 0.6,
173         'title': title_str,
174         'legend': False,
175         'ret': False,
176         'out_file': [],
177     })
178
179     xy_data = self.__parse_data(data, beta)
180
181     return xy_data, kwargs
182
183 def plot_observables(self, data, **kwargs):
184     """Plot observables."""
185     # beta = kwargs.get('beta', 5.)
186     # run_str = kwargs.get('run_str', None)
187     # net_weights = kwargs.get('net_weights', [1., 1., 1.])
188     # dir_append = kwargs.get('dir_append', None)
189
190     xy_data, kwargs = self.__plot_setup(data, **kwargs)
191
192     self.__plot_actions(xy_data['actions'], **kwargs)
193     self.log_figure()
194     self.__plot_plaqs(xy_data['plaqs'], **kwargs)
195     self.log_figure()
196     self.__plot_charges(xy_data['charges'], **kwargs)
197     self.log_figure()
198     self.__plot_charge_probs(xy_data['charges'][1], **kwargs)
199     self.log_figure()
200     self.__plot_charge_diffs(xy_data['charge_diffs'], **kwargs)
201     self.log_figure()
202     self.__plot_autocorrs(xy_data['autocorrs'], **kwargs)
203     self.log_figure()
204     mean_diff = self.__plot_plaqs_diffs(xy_data['plaqs_diffs'], **kwargs)
205
206     self.log_figure()
207
208     return mean_diff

```

```

208
209     def _plot(self, xy_data, **kwargs):
210         """Basic plotting wrapper."""
211         x, y, yerr = xy_data
212
213         labels = kwargs.get('labels', None)
214         if labels is not None:
215             xlabel = labels.get('x_label', '')
216             ylabel = labels.get('y_label', '')
217         else:
218             xlabel = ''
219             ylabel = ''
220
221         _leg = kwargs.get('legend', False)
222
223         if kwargs.get('two_rows', False):
224             fig, (ax0, ax1) = plt.subplots(
225                 nrows=2, ncols=1, gridspec_kw={'height_ratios': [2.5, 1],
226                                             'hspace': 0.175})
227         )
228         n = len(x)
229         mid = n // 2
230         x0 = int(mid - 0.025 * n)
231         x1 = int(mid + 0.025 * n)
232     else:
233         fig, ax0 = plt.subplots()
234         ax1 = None
235
236     plt_kwargs = {
237         'color': 'k',
238         # 'lw': 1.,
239         # 'ls': '-',
240         'alpha': 0.8,
241         'marker': ',',
242     }
243     # err_kwargs = plt_kwargs.update({'lw': 1.5, 'alpha': 0.7})
244
245     ax0.plot(x, y, **plt_kwargs)
246     ax0.errorbar(x, y, yerr=yerr,
247                  # ls='-', lw=1.,
248                  alpha=0.7,
249                  color='k',
250                  ecolor='gray')
251
252     if ax1 is not None:
253         ax1.plot(x[x0:x1:10], y[x0:x1:10], **plt_kwargs)
254         ax1.errorbar(x[x0:x1:10], y[x0:x1:10], yerr=yerr[x0:x1:10],
255                      # ls='-', lw=1.,
256                      alpha=0.7,
257                      color='k',
258                      ecolor='gray')
259
260     ax1.set_xlabel(xlabel, fontsize=14)
261     ax0.set_ylabel(ylabel, fontsize=14)
262     ax1.set_ylabel('', fontsize=14)
263     if _leg:
264         ax0.legend(loc='best')
265
266     title = kwargs.get('title', None)
267     if title is not None:
268         _ = ax0.set_title(title)
269
270     plt.tight_layout()
271     if kwargs.get('save', True):
272         fname = kwargs.get('fname', f'plot_{np.random.randint(10)}')
273         out_file = get_out_file(self.out_dir, fname)
274         io.log(f'Saving figure to: {out_file}.')
275         plt.savefig(out_file, dpi=400, bbox_inches='tight')
276
277     return fig, (ax0, ax1)

```

```

278
279     def _plot_actions(self, xy_data, **kwargs):
280         """Plot actions"""
281         # kwargs['out_file'] = get_out_file(self.out_dir, 'actions_vs_step')
282         labels = {
283             'x_label': 'Step',
284             'y_label': 'Action',
285             'plt_label': 'Action'
286         }
287
288         kwargs.update({
289             'fname': 'actions_vs_step',
290             'labels': labels,
291             'two_rows': True,
292         })
293         self._plot(xy_data, **kwargs)
294         # kwargs['bounds'] = [0.2, 0.6, 0.7, 0.3]
295
296         # xy_labels = ('Step', 'Action')
297         # plot_multiple_lines(xy_data, xy_labels, **kwargs)
298         # plot_with_inset(xy_data, labels, **kwargs)
299
300     def _plot_plaqs(self, xy_data, beta, **kwargs):
301         """Plot average plaquette."""
302         labels = {
303             'x_label': 'Step',
304             'y_label': r"\langle \phi_P \rangle",
305             'plt_label': r"\langle \phi_P \rangle"
306         }
307         kwargs.update({
308             'labels': labels,
309             'fname': 'plaqs_vs_step',
310             'two_rows': True,
311             'save': False,
312         })
313         fig, (ax0, ax1) = self._plot(xy_data, **kwargs)
314
315         ax0.axhline(y=u1_plaq_exact(beta),
316                      color='#CC0033', ls='-', lw=1.5, label='exact')
317         ax1.axhline(y=u1_plaq_exact(beta),
318                      color='#CC0033', ls='-', lw=1.5, label='exact')
319
320         plt.tight_layout()
321
322         out_file = get_out_file(self.out_dir, 'plaqs_vs_step')
323         io.log(f'Saving figure to: {out_file}')
324         plt.savefig(out_file, dpi=400, bbox_inches='tight')
325
326     def _plot_plaqs_diffs(self, xy_data, **kwargs):
327         kwargs['out_file'] = None
328         kwargs['ret'] = True
329         labels = {
330             'x_label': 'Step',
331             'y_label': r"\delta_{\phi_P}",
332             'plt_label': r"\delta_{\phi_P}"
333         }
334         x, y, yerr = xy_data
335         y_mean = np.mean(y)
336         fig, ax = plt.subplots()
337         _ = ax.plot(x, y, label='', marker=',', color='k', alpha=0.8)
338         _ = ax.errorbar(x, y, yerr=yerr, label='', marker=None, ls='',
339                          alpha=0.7, color='gray', ecolor='gray')
340         _ = ax.axhline(y=0, color='#CC0033', ls='-', lw=2.)
341         _ = ax.axhline(y=y_mean, label=f'avg {y_mean:.5f}',
342                         color='C2', ls='-', lw=2.)
343
344         _ = ax.set_xlabel(labels['x_label'], fontsize=14)
345         _ = ax.set_ylabel(labels['y_label'], fontsize=14)
346         title = kwargs.get('title', None)
347         if title is not None:

```

```

348         _ = ax.set_title(title)
349
350     ax.legend(loc='best')
351
352     _ = plt.tight_layout()
353     out_file = get_out_file(self.out_dir, 'plaqs_diffs_vs_step')
354     io.log(f'Saving figure to: {out_file}.')
355     plt.savefig(out_file, dpi=400, bbox_inches='tight')
356
357     return y_mean
358
359 def _plot_charges(self, xy_data, **kwargs):
360     """Plot topological charges."""
361     kwargs['out_file'] = get_out_file(self.out_dir, 'charges_vs_step')
362     kwargs['markers'] = True
363     kwargs['lines'] = False
364     kwargs['alpha'] = 1.
365     kwargs['ret'] = False
366     xy_labels = ('Step', r'$Q$')
367     plot_multiple_lines(xy_data, xy_labels, **kwargs)
368
369     charges = np.array(xy_data[1].T, dtype=int)
370     num_steps, num_samples = charges.shape
371
372     out_dir = os.path.join(self.out_dir, 'top_charge_plots')
373     io.check_else_make_dir(out_dir)
374     # if we have more than 10 chains in charges, only plot first 10
375     for idx in range(min(num_samples, 5)):
376         _, ax = plt.subplots()
377         _ = ax.plot(charges[:, idx],
378                     marker=MARKERS[idx],
379                     color=COLORS[idx],
380                     ls='',
381                     alpha=0.5,
382                     label=f'sample {idx}')
383         _ = ax.legend(loc='best')
384         _ = ax.set_xlabel(xy_labels[0], fontsize=14)
385         _ = ax.set_ylabel(xy_labels[1], fontsize=14)
386         _ = ax.set_title(kwargs['title'], fontsize=16)
387         _ = plt.tight_layout()
388         out_file = get_out_file(out_dir, f'top_charge_vs_step_{idx}')
389         io.check_else_make_dir(os.path.dirname(out_file))
390         io.log(f'Saving figure to: {out_file}')
391         plt.savefig(out_file, dpi=400, bbox_inches='tight')
392
393     plt.close('all')
394
395 def _plot_charge_diffs(self, xy_data, **kwargs):
396     """Plot tunneling events (change in top. charge)."""
397     out_file = get_out_file(self.out_dir, 'top_charge_diffs')
398     steps_arr, charge_diffs = xy_data
399
400     # ignore first two data points when plotting since the top. charge
401     # should change dramatically for the very first few steps when starting
402     # from a random configuration
403     _, ax = plt.subplots()
404     _ = ax.plot(xy_data[0][2:], xy_data[1][2:],
405                 marker='.', ls='', fillstyle='none', color='C0')
406     _ = ax.set_xlabel('Steps', fontsize=14)
407     _ = ax.set_ylabel(r'$\Delta_Q$', fontsize=14)
408     _ = ax.set_title(kwargs['title'], fontsize=16)
409     _ = plt.tight_layout()
410     io.log(f"Saving figure to: {out_file}")
411     plt.savefig(out_file, dpi=400, bbox_inches='tight')
412
413 def _plot_charge_probs(self, charges, **kwargs):
414     """Plot top. charge probabilities."""
415     num_steps, num_samples = charges.shape
416     charges = np.array(charges, dtype=int)
417     out_dir = os.path.join(self.out_dir, 'top_charge_probs')

```

```

418     io.check_else_make_dir(out_dir)
419     if 'title' in list(kwargs.keys()):
420         title = kwargs.pop('title')
421     # if we have more than 10 chains in charges, only plot first 10
422     for idx in range(min(num_samples, 5)):
423         counts = Counter(charges[:, idx])
424         total_counts = np.sum(list(counts.values()))
425         _, ax = plt.subplots()
426         ax.plot(list(counts.keys()),
427                 np.array(list(counts.values()) / total_counts),
428                 marker=MARKERS[idx],
429                 color=COLORS[idx],
430                 ls='',
431                 label=f'sample {idx}')
432         _ = ax.legend(loc='best')
433         _ = ax.set_xlabel(r"$Q$") # , fontsize=14)
434         _ = ax.set_ylabel('Probability') # , fontsize=14)
435         _ = ax.set_title(title) # , fontsize=16)
436         _ = plt.tight_layout()
437         out_file = get_out_file(out_dir, f'top_charge_vs_step_{idx}')
438         io.check_else_make_dir(os.path.dirname(out_file))
439         io.log(f"Saving plot to: {out_file}.")
440         plt.savefig(out_file, dpi=400, bbox_inches='tight')
441         # for f in out_file:
442         #     io.check_else_make_dir(os.path.dirname(f))
443         #     io.log(f"Saving plot to: {f}.")
444     plt.close('all')
445
446     all_counts = Counter(list(charges.flatten()))
447     total_counts = np.sum(list(counts.values()))
448     _, ax = plt.subplots()
449     ax.plot(list(all_counts.keys()),
450             np.array(list(all_counts.values()) / (total_counts * *
451                                         num_samples)),
452             marker='o',
453             color='C0',
454             ls='',
455             alpha=0.6,
456             label=f'total across {num_samples} samples')
457     _ = ax.legend(loc='best')
458     _ = ax.set_xlabel(r"$Q$") # , fontsize=14)
459     _ = ax.set_ylabel('Probability') # , fontsize=14)
460     _ = ax.set_title(title) # , fontsize=16)
461     _ = plt.tight_layout()
462     out_file = get_out_file(self.out_dir, f'TOP_CHARGE_PROBS_ALL')
463     io.check_else_make_dir(os.path.dirname(out_file))
464     io.log(f"Saving plot to: {out_file}.")
465     plt.savefig(out_file, dpi=400, bbox_inches='tight')
466     # for f in out_file:
467     plt.close('all')
468
469 def _plot_autocorrs(self, xy_data, **kwargs):
470     """Plot topological charge autocorrelations."""
471     # xy_labels = ('Step', 'Autocorrelation of ' + r'$Q$')
472     # return plot_multiple_lines(xy_data, xy_labels, **kwargs)
473     # kwargs['out_file'] = get_out_file(self.out_dir,
474     #                                   'charge_autocorrs_vs_step')
475     # kwargs['ret'] = True
476     # kwargs['bounds'] = [0.2, 0.6, 0.7, 0.3]
477     # try:
478     #     kwargs['out_file'] = get_out_file(
479     #         self.out_dir, 'charge_autocorrs_vs_step'
480     #         )
481     # except AttributeError:
482     #     kwargs['out_file'] = None
483     labels = {
484         'x_label': 'Step',
485         'y_label': 'Autocorrelation of ' + r'$Q$',
486         'plt_label': 'Autocorrelation of ' + r'$Q$',
487     }

```

```
488     kwargs.update({
489         'labels': labels,
490         'fname': 'charge_autocorrs',
491         'two_rows': True,
492     })
493     self._plot(xy_data, **kwargs)
494     # _, ax, axins = plot_with_inset(xy_data, labels, **kwargs)
495     # xy_labels = ('Step', 'Autocorrelation of ' + r'$Q$')
496     # return plot_multiple_lines(xy_data, xy_labels, **kwargs)
```

A.17 l2hmc-qcd/plotters/leapfrog_plotters.py

```
1 import os
2 import pickle
3 import numpy as np
4 import utils.file_io as io
5 from utils.data_loader import DataLoader
6
7 from config import HAS_MATPLOTLIB, HAS_PSUTIL
8
9 if HAS_MATPLOTLIB:
10     import matplotlib as mpl
11     import matplotlib.pyplot as plt
12
13 if HAS_PSUTIL:
14     import psutil
15
16
17 def load_and_sep(out_file, keys=('forward', 'backward')):
18     with open(out_file, 'rb') as f:
19         data = pickle.load(f)
20
21     return (np.array(data[k]) for k in keys)
22
23
24 params = {
25     'axes.labelsize': 16,    # fontsize for x and y labels (was 10)
26     'axes.titlesize': 16,
27     'legend.fontsize': 10,   # was 10
28     'xtick.labelsize': 12,
29     'ytick.labelsize': 12,
30     'font.family': 'serif'
31 }
32
33 try:
34     mpl.rcParams.update(params)
35 except FileNotFoundError:
36     params['text.usetex'] = False
37     params['text.latex.preamble'] = None
38     try:
39         mpl.rcParams.update(params)
40     except FileNotFoundError:
41         pass
42
43
44 def print_memory():
45     if HAS_PSUTIL:
46         pid = os.getpid()
47         py = psutil.Process(pid)
48         memory_use = py.memory_info()[0] / 2. ** 30
49         io.log(80 * '-')
50         io.log(f'memory use: {memory_use}')
51         io.log(80 * '-')
52
53
54 class LeapfrogPlotter:
55     def __init__(self, figs_dir, run_logger=None, run_dir=None):
56         self.figs_dir = figs_dir
57         self.pdfs_dir = os.path.join(self.figs_dir, 'pdfs_plots')
58         io.check_else_make_dir(self.pdfs_dir)
59
60         if run_logger is None:
61             if run_dir is None:
62                 raise AttributeError(
63                     """Either a `run_logger` object containing data or a
64                     `run_dir` from which to load data must be specified.
65                     Exiting.
66                     """
67             )
```

```

68     try:
69         data = self.load_data(run_dir)
70         self.samples = data[0]
71         self.lf_f, self.lf_b = data[1]
72         self.logdets_f, self.logdets_b = data[2]
73         self.sumlogdet_f, self.sumlogdet_b = data[3]
74     except FileNotFoundError:
75         io.log(f'''Unable to load leapfrog data from run_dir:
76             {run_dir}. Exiting.'')
77         return
78
79     else:
80         self.samples = np.array(run_logger.samples_arr)
81         self.lf_f = np.array(run_logger.lf_out['forward'])
82         self.lf_b = np.array(run_logger.lf_out['backward'])
83         self.logdets_f = np.array(run_logger.logdets['forward'])
84         self.logdets_b = np.array(run_logger.logdets['backward'])
85         self.sumlogdet_f = np.array(run_logger.sumlogdet['forward'])
86         self.sumlogdet_b = np.array(run_logger.sumlogdet['backward'])
87
88         self.lf_f_diffs = 1. - np.cos(self.lf_f[1:] - self.lf_f[:-1])
89         self.lf_b_diffs = 1. - np.cos(self.lf_b[1:] - self.lf_b[:-1])
90         self.samples_diffs = 1. - np.cos(self.samples[1:] - self.samples[:-1])
91         self.tot_lf_steps = self.lf_f_diffs.shape[0]
92         self.tot_md_steps = self.samples_diffs.shape[0]
93         self.num_lf_steps = self.tot_lf_steps // self.tot_md_steps
94
95         self.indiv_kwargs = {
96             'ls': '-',
97             'alpha': 0.5,
98             'lw': 0.5,
99             'rasterized': True
100        }
101
102        self.avg_kwargs = {
103            'ls': '-',
104            'alpha': 0.6,
105            'lw': 0.75,
106            'rasterized': True
107        }
108
109    def load_data(self, run_dir):
110        loader = DataLoader(run_dir)
111        io.log("Loading samples...")
112        samples = loader.load_samples(run_dir)
113        io.log('done.')
114        io.log("Loading leapfrogs...")
115        leapfrogs = loader.load_leapfrogs(run_dir)
116        io.log("Loading logdets...")
117        logdets = loader.load_logdets(run_dir)
118        io.log("Loading sumlogdets...")
119        sumlogdets = loader.load_sumlogdets(run_dir)
120        return (samples, leapfrogs, logdets, sumlogdets)
121
122    def get_colors(self, num_samples=20):
123        reds_cmap = mpl.cm.get_cmap('Reds', num_samples + 1)
124        blues_cmap = mpl.cm.get_cmap('Blues', num_samples + 1)
125        idxs = np.linspace(0.2, 0.75, num_samples + 1)
126        reds = [reds_cmap(i) for i in idxs]
127        blues = [blues_cmap(i) for i in idxs]
128
129        return reds, blues
130
131    def update_figs_dir(self, figs_dir):
132        self.figs_dir = figs_dir
133        self.pdfs_dir = os.path.join(self.figs_dir, 'pdfs')
134
135        io.check_else_make_dir(figs_dir)
136        io.check_else_make_dir(self.pdfs_dir)
137

```

```

138 def make_plots(self, run_dir, num_samples=20, ret=False):
139     """Make plots of the leapfrog differences and logdets.
140
141     Immediately after creating and saving the plots, delete these
142     (no-longer needed) attributes to free up memory.
143
144     Args:
145         run_dir (str): Path to directory in which to save all of the
146             relevant instance attributes.
147         num_samples (int): Number of samples to include when creating
148             plots.
149         save (bool): Boolean indicating whether or not plotted data should
150             be saved.
151
152     NOTE:
153         `save` is very data intensive and will produce LARGE (compressed)
154         `.npz` files.
155
156     run_key = run_dir.split('/')[-1].split('_')
157     beta_idx = run_key.index('beta') + 1
158     beta = run_key[beta_idx]
159
160     # print_memory()
161     fig_ax1 = self.plot_lf_diffs(beta, num_samples)
162
163     # print_memory()
164     fig_ax2 = self.plot_logdets(beta, num_samples)
165
166     if ret:
167         return fig_ax1, fig_ax2
168
169 def plot_lf_diffs(self, beta, num_samples=20):
170     reds, blues = self.get_colors(num_samples)
171     samples_y_avg = np.mean(self.samples_diffs, axis=(1, 2))
172     samples_x_avg = np.arange(len(samples_y_avg))
173
174     fig, (ax1, ax2) = plt.subplots(2, 1)
175     for idx in range(num_samples):
176         yf = np.mean(self.lf_f_diffs, axis=-1)
177         xf = np.arange(len(yf))
178         yb = np.mean(self.lf_b_diffs, axis=-1)
179         xb = np.arange(len(yb))
180
181         _ = ax1.plot(xf, yf[:, idx], color=reds[idx], **self.indiv_kw_args)
182         _ = ax1.plot(xb, yb[:, idx], color=blues[idx], **self.indiv_kw_args)
183
184     yf_avg = np.mean(self.lf_f_diffs, axis=(1, 2))
185     yb_avg = np.mean(self.lf_b_diffs, axis=(1, 2))
186     xf_avg = np.arange(len(yf))
187     xb_avg = np.arange(len(yb))
188
189     _ = ax1.plot(xf_avg, yf_avg, label='forward',
190                  color='r', **self.avg_kw_args)
191     _ = ax1.plot(xb_avg, yb_avg, label='backward',
192                  color='b', **self.avg_kw_args)
193
194     _ = ax2.plot(samples_x_avg, samples_y_avg, label='MD avg.',
195                  color='k', lw=1., rasterized=True, ls='--')
196
197     _ = ax1.set_xlabel('Leapfrog step', fontsize=16)
198     _ = ax2.set_xlabel('MD step', fontsize=16)
199
200     ylabel = r'$\langle \delta\phi_{\mu}(i) \rangle$'
201     _ = ax1.set_ylabel(ylabel, fontsize=16)
202     _ = ax2.set_ylabel(ylabel, fontsize=16)
203
204     _ = ax1.legend(loc='best', fontsize=10)
205     _ = ax2.legend(loc='best', fontsize=10)
206     fig.tight_layout()
207     # fig.subplots_adjust(hspace=0.5)

```

```

208
209     beta_str = str(beta).replace('.', '')
210     fn = f'leapfrog_diffs_beta{beta_str}'
211     out_file = os.path.join(self.pdfs_dir, fn + '.pdf')
212     out_file_zoom = os.path.join(self.pdfs_dir, fn + '_zoom.pdf')
213     out_file_zoom1 = os.path.join(self.pdfs_dir, fn + '_zoom1.pdf')
214     io.log(f'Saving figure to: {out_file}')
215     _ = plt.savefig(out_file, dpi=400, bbox_inches='tight')
216
217     lf_xlim = 100
218     md_xlim = lf_xlim // self.num_lf_steps
219
220     _ = ax1.set_xlim((0, lf_xlim))
221     _ = ax2.set_xlim((0, md_xlim))
222     _ = plt.savefig(out_file_zoom, dpi=400, bbox_inches='tight')
223
224     lf_xlim //= 4
225     md_xlim //= 4
226     _ = ax1.set_xlim((0, lf_xlim))
227     _ = ax1.set_ylim((-0.05, 0.25))
228     _ = ax2.set_xlim((0, md_xlim))
229     _ = plt.savefig(out_file_zoom1, dpi=400, bbox_inches='tight')
230
231     return fig, (ax1, ax2)
232
233 def plot_logdets(self, beta, num_samples=20):
234     reds, blues = self.get_colors(num_samples)
235
236     sumlogdet_yf_avg = np.mean(self.sumlogdet_f, axis=-1)
237     sumlogdet_yb_avg = np.mean(self.sumlogdet_b, axis=-1)
238     sumlogdet_xf_avg = np.arange(len(sumlogdet_yf_avg))
239     sumlogdet_xb_avg = np.arange(len(sumlogdet_yb_avg))
240
241     fig, (ax1, ax2) = plt.subplots(2, 1)
242     for idx in range(num_samples):
243         yf = self.logdets_f[:, idx]
244         xf = np.arange(len(yf))
245         yb = self.logdets_b[:, idx]
246         xb = np.arange(len(yb))
247
248         _ = ax1.plot(xf, yf, color=reds[idx], **self.indiv_kwargs)
249         _ = ax1.plot(xb, yb, color=blues[idx], **self.indiv_kwargs)
250
251     yf_avg = np.mean(self.logdets_f, axis=-1)
252     yb_avg = np.mean(self.logdets_b, axis=-1)
253
254     xf_avg = np.arange(len(yf_avg))
255     xb_avg = np.arange(len(yb_avg))
256
257     _ = ax1.plot(xf_avg, yf_avg, label='forward',
258                  color='r', **self.avg_kwargs)
259     _ = ax1.plot(xb_avg, yb_avg, label='backward',
260                  color='b', **self.avg_kwargs)
261
262     _ = ax2.plot(sumlogdet_xf_avg, sumlogdet_yf_avg, label='forward',
263                  color='r', lw=1.2, alpha=0.9, marker='.')
264
265     _ = ax2.plot(sumlogdet_xb_avg, sumlogdet_yb_avg, label='backward',
266                  color='b', lw=1.2, alpha=0.9, marker='.')
267
268     _ = ax1.set_xlabel('Leapfrog step', fontsize=16)
269     _ = ax1.set_ylabel(r'$\log|\mathcal{J}|^{(t)}|$', fontsize=16)
270     _ = ax2.set_xlabel('MD step', fontsize=16)
271     _ = ax2.set_ylabel(r'$\sum_t \log|\mathcal{J}|^{(t)}|$', fontsize=16)
272     _ = ax1.legend(loc='best', fontsize=10)
273     _ = ax2.legend(loc='best', fontsize=10)
274     _ = fig.tight_layout()
275     beta_str = str(beta).replace('.', '')
276     fn = f'avg_logdets_beta{beta_str}'

```

```
278     out_file = os.path.join(self.pdfs_dir, fn + '.pdf')
279     out_file_zoom = os.path.join(self.pdfs_dir, fn + '_zoom.pdf')
280     out_file_zoom1 = os.path.join(self.pdfs_dir, fn + '_zoom1.pdf')
281     io.log(f'Saving figure to: {out_file}')
282     _ = plt.savefig(out_file, dpi=400, bbox_inches='tight')
283
284     lf_xlim = 100
285     md_xlim = lf_xlim // self.num_lf_steps
286
287     _ = ax1.set_xlim((0, lf_xlim))
288     _ = ax2.set_xlim((0, md_xlim))
289     _ = plt.savefig(out_file_zoom, dpi=400, bbox_inches='tight')
290
291     lf_xlim //= 4
292     md_xlim //= 4
293     _ = ax1.set_xlim((0, lf_xlim))
294     _ = ax2.set_xlim((0, md_xlim))
295     _ = plt.savefig(out_file_zoom1, dpi=400, bbox_inches='tight')
296
297     return fig, (ax1, ax2)
```

A.18 l2hmc-qcd/utils/parse_args.py

```
1 """
2 parse_args.py
3
4 Implements method for parsing command line arguments for `gauge_model.py`
5
6 Author: Sam Foreman (github: @saforem2)
7 Date: 04/09/2019
8 """
9 import os
10 import sys
11 import argparse
12 import shlex
13
14 import utils.file_io as io
15
16 # from config import process_config
17 # from attr_dict import AttrDict
18
19 DESCRIPTION = (
20     'L2HMC model using U(1) lattice gauge theory for target distribution.'
21 )
22
23
24 # =====
25 # * NOTE:
26 #     - if action == 'store_true':
27 #         The argument is FALSE by default. Passing this flag will cause the
28 #         argument to be ''stored true''.
29 #     - if action == 'store_false':
30 #         The argument is TRUE by default. Passing this flag will cause the
31 #         argument to be ''stored false''.
32 # =====
33 def parse_args():
34     """Parse command line arguments."""
35     parser = argparse.ArgumentParser(
36         description=DESCRIPTION,
37         fromfile_prefix_chars='@',
38     )
39     ##### Lattice parameters #####
40     # parser.add_argument("--space_size",
41     #                     dest="space_size",
42     #                     type=int,
43     #                     default=8,
44     #                     required=False,
45     #                     help="""Spatial extent of lattice.\n (Default: 8)""")
46
47     parser.add_argument("--time_size",
48                     dest="time_size",
49                     type=int,
50                     default=8,
51                     required=False,
52                     help="""Temporal extent of lattice.\n (Default: 8)""")
53
54     parser.add_argument("--link_type",
55                     dest="link_type",
56                     type=str,
57                     required=False,
58                     default='U1',
59                     help="""Link type for gauge model.\n (Default: 'U1')""")
60
61     parser.add_argument("--dim",
62                     dest="dim",
63                     type=int,
64                     required=False,
```

```

68             default=2,
69             help="""Dimensionality of lattice.\n (Default: 2)""")
70
71     parser.add_argument("--num_samples",
72                         dest="num_samples",
73                         type=int,
74                         default=128,
75                         required=False,
76                         help="""Number of samples (batch size) to use for
77                               training.\n (Default: 20)""")
78
79     parser.add_argument("--rand",
80                         dest="rand",
81                         action="store_true",
82                         required=False,
83                         help="""If passed, set `rand=True` and start lattice
84                               from randomized initial configuration.\n
85                               (Default: `rand=False`, i.e. NOT passed.)""")
86
87 #####
88 #           Leapfrog parameters                      #
89 #####
90
91     parser.add_argument("-n", "--num_steps",
92                         dest="num_steps",
93                         type=int,
94                         default=5,
95                         required=False,
96                         help="""Number of leapfrog steps to use in (augmented)
97                               HMC sampler. (Default: 5)""")
98
99     parser.add_argument("--eps",
100                        dest="eps",
101                        type=float,
102                        default=0.1,
103                        required=False,
104                        help="""Step size to use in leapfrog integrator.
105                               (Default: 0.1)""")
106
107    parser.add_argument("--loss_scale",
108                        dest="loss_scale",
109                        type=float,
110                        default=1.,
111                        required=False,
112                        help="""Scaling factor to be used in loss function.
113                               (lambda in Eq. 7 of paper). (Default: 1.)""")
114
115 #####
116 #           Learning rate parameters                 #
117 #####
118
119    parser.add_argument("--lr_init",
120                        dest="lr_init",
121                        type=float,
122                        default=1e-3,
123                        required=False,
124                        help="""Initial value of learning rate.
125                               (Default: 1e-3)""")
126
127    parser.add_argument("--lr_decay_steps",
128                        dest="lr_decay_steps",
129                        type=int,
130                        default=500,
131                        required=False,
132                        help="""Number of steps after which to decay learning
133                               rate. (Default: 500)""")
134
135    parser.add_argument("--lr_decay_rate",
136                        dest="lr_decay_rate",
137                        type=float, default=0.96,

```

```

138     required=False,
139     help=("""Learning rate decay rate to be used during
140           training. (Default: 0.96)""")
141
142 ##### Annealing rate parameters #####
143 #
144 ##### Annealing rate parameters #####
145
146 parser.add_argument("--fixed_beta",
147     dest="fixed_beta",
148     action="store_true",
149     required=False,
150     help=("""Flag that when passed runs the training loop
151           at fixed beta (i.e. no annealing is done).
152           (Default: `fixed_beta=False`)""))
153
154 parser.add_argument("--beta_init",
155     dest="beta_init",
156     type=float,
157     default=2.,
158     required=False,
159     help=("""Initial value of beta (inverse coupling
160           constant) used in gauge model when
161           annealing. (Default: 2.)"""))
162
163 parser.add_argument("--beta_final",
164     dest="beta_final",
165     type=float,
166     default=5.,
167     required=False,
168     help=("""Final value of beta (inverse coupling
169           constant) used in gauge model when
170           annealing. (Default: 5.)"""))
171
172 parser.add_argument("--warmup_lr",
173     dest="warmup_lr",
174     action="store_true",
175     required=False,
176     help=("""Flag that when passed will 'warmup' the
177           learning rate (i.e. gradually scale it up to the
178           value passed to `--lr_init` (performs better when
179           using Horovod for distributed training)."""))
180
181 ##### Training parameters #####
182 #
183 ##### Training parameters #####
184
185 parser.add_argument("--train_steps",
186     dest="train_steps",
187     type=int,
188     default=5000,
189     required=False,
190     help=("""Number of training steps to perform.
191           (Default: 5000)"""))
192
193 parser.add_argument("--trace",
194     dest="trace",
195     action="store_true",
196     required=False,
197     help=("""Flag that when passed will set `--trace=True`,
198           and create a trace during training loop.
199           (Default: `--trace=False`, i.e. not passed)"""))
200
201 parser.add_argument("--save_steps",
202     dest="save_steps",
203     type=int,
204     default=50,
205     required=False,
206     help=("""Number of steps after which to save the model
207           and current values of all parameters.

```

```

208                                         (Default: 50)""")
209
210     parser.add_argument("--print_steps",
211                         dest="print_steps",
212                         type=int,
213                         default=1,
214                         required=False,
215                         help=("""Number of steps after which to display
216                               information about the loss and various
217                               other quantities. (Default: 1)"""))
218
219     parser.add_argument("--logging_steps",
220                         dest="logging_steps",
221                         type=int,
222                         default=10,
223                         required=False,
224                         help=("""Number of steps after which to write logs for
225                               tensorflow. (default: 50)"""))
226
227 # -----
228 # Model parameters
229 # -----
230
231     parser.add_argument('--network_arch',
232                         dest='network_arch',
233                         type=str,
234                         default='conv3D',
235                         required=False,
236                         help=("""String specifying the architecture to use for
237                               the neural network. Must be one of:
238                               `'conv3D', 'conv2D', 'generic'`.
239                               (Default: 'conv3D')"""))
240
241     parser.add_argument('--num_hidden1',
242                         dest='num_hidden1',
243                         type=int,
244                         default=100,
245                         required=False,
246                         help=("""Number of nodes to include in each of the
247                               fully-connected hidden layers for x, v, and t.
248                               (Default: 100)"""))
249
250     parser.add_argument('--num_hidden2',
251                         dest='num_hidden2',
252                         type=int,
253                         default=100,
254                         required=False,
255                         help=("""Number of nodes to include in fully-connected
256                               hidden layer `h`. If not explicitly passed, will
257                               default to 2 * lattice.num_links.
258                               (Default: None)"""))
259
260     parser.add_argument('--zero_translation',
261                         dest='zero_translation',
262                         action='store_true',
263                         required=False,
264                         help=("""Flag that when passed explicitly sets the
265                               translation function (T in the original paper) to
266                               be zero. (Default: False)"""))
267
268     parser.add_argument('--no_summaries',
269                         dest="no_summaries",
270                         action="store_true",
271                         required=False,
272                         help=("""Flag that when passed will prevent tensorflow
273                               from creating tensorflow summary objects.""))
274
275     parser.add_argument("--hmc",
276                         dest="hmc",
277                         action="store_true",

```

```

278     required=False,
279     help="""Use generic HMC (without augmented leapfrog
280         integrator described in paper). Used for
281         comparing against L2HMC algorithm."""))
282
283 parser.add_argument("--run_steps",
284     dest="run_steps",
285     type=int,
286     default=10000,
287     required=False,
288     help="""Number of evaluation 'run' steps to perform
289         after training (i.e. length of desired chain
290         generate using trained L2HMC sample).
291         (Default: 10000)""")
292
293 parser.add_argument("--eps_fixed",
294     dest="eps_fixed",
295     action="store_true",
296     required=False,
297     help="""Flag that when passed will cause the step size
298         `eps` to be a fixed (non-trainable)
299         parameter.""))
300
301 parser.add_argument("--metric",
302     dest="metric",
303     type=str,
304     default="cos_diff",
305     required=False,
306     help="""Metric to use in loss function. Must be one
307         of: 'l1', 'l2', 'cos', 'cos2', 'cos_diff'.""))
308
309 parser.add_argument("--nnehmc_loss",
310     dest="nnehmc_loss",
311     action="store_true",
312     required=False,
313     help="""Flag that when passed will calculate the
314         'NNEHMC Loss' from
315         (https://infoscience.epfl.ch/record/264887/files/robust\_parameter\_estimation.pdf)
316         (Default: False).""))
317
318 parser.add_argument("--std_weight",
319     dest="std_weight",
320     type=float,
321     default=1.,
322     required=False,
323     help="""Multiplicative factor used to weigh relative
324         strength of stdiliary term in loss function.
325         (Default: 1.)""")
326
327 parser.add_argument("--aux_weight",
328     dest="aux_weight",
329     type=float,
330     default=1.,
331     required=False,
332     help="""Multiplicative factor used to weigh relative
333         strength of auxiliary term in loss function.
334         (Default: 1.)""")
335
336 parser.add_argument("--charge_weight",
337     dest="charge_weight",
338     type=float,
339     default=1.,
340     required=False,
341     help="""Multiplicative factor used to weigh relative
342         strength of top. charge term in loss
343         function. (Default: 1.)""")
344
345 parser.add_argument("--profiler",
346     dest='profiler',
347     action="store_true",

```

```

348     required=False,
349     help=("""Flag that when passed will profile the graph
350           execution using `TFProf`."""))
351
352     parser.add_argument("--gpu",
353         dest="gpu",
354         action="store_true",
355         required=False,
356         help=("""Flag that when passed indicates we're training
357               using an NVIDIA GPU.""""))
358
359     parser.add_argument("--use_bn",
360         dest='use_bn',
361         action="store_true",
362         required=False,
363         help=("""Flag that when passed causes batch
364               normalization layer to be used in ConvNet.""""))
365
366     parser.add_argument("--horovod",
367         dest="horovod",
368         action="store_true",
369         required=False,
370         help=("""Flag that when passed uses Horovod for
371               distributed training on multiple nodes.""""))
372
373     parser.add_argument("--comet",
374         dest="comet",
375         action="store_true",
376         required=False,
377         help=("""Flag that when passed uses comet.ml for
378               parameter logging and additional metric
379               tracking/displaying.""""))
380
381     parser.add_argument("--dropout_prob",
382         dest="dropout_prob",
383         type=float,
384         required=False,
385         default=0.,
386         help=("""Dropout probability in network. If > 0,
387               dropout will be used. (Default: 0.)"""))
388
389 ######
390 # (Mostly) Deprecated #
391 #####
392
393     parser.add_argument('--save_samples',
394         dest='save_samples',
395         action='store_true',
396         required=False,
397         help=("""Flag that when passed will set
398               '--save_samples=True', and save the samples
399               generated during the 'run' phase.
400               (Default: '--save_samples=False, i.e.
401               '--save_samples' is not passed).\n
402               WARNING!! This is very data intensive.""""))
403
404     parser.add_argument('--save_lf',
405         dest='save_lf',
406         action='store_true',
407         required=False,
408         help=("""Flag that when passed will save the
409               output from each leapfrog step.""""))
410
411     parser.add_argument("--clip_value",
412         dest="clip_value",
413         type=float,
414         default=0.,
415         required=False,
416         help=("""Clip value, used for clipping value of
417               gradients by global norm. (Default: 0.) If a

```

```

418                     value greater than 0. is passed, gradient
419                     clipping will be performed.""))
420
421     parser.add_argument("--restore",
422                         dest="restore",
423                         action="store_true",
424                         required=False,
425                         help="""Restore model from previous run. If this
426                             argument is passed, a `log_dir` must be specified
427                             and passed to `--log_dir` argument.""")
428
429     parser.add_argument("--theta",
430                         dest="theta",
431                         action="store_true",
432                         required=False,
433                         help="""Flag that when passed indicates we're training
434                             on theta @ ALCf.""")
435
436     parser.add_argument("--log_dir",
437                         dest="log_dir",
438                         type=str,
439                         default=None,
440                         required=False,
441                         help="""Log directory to use from previous run.
442                             If this argument is not passed, a new
443                             directory will be created.""")
444
445     parser.add_argument("--num_intra_threads",
446                         dest="num_intra_threads",
447                         type=int,
448                         default=0,
449                         required=False,
450                         help="""Number of intra op threads to use for
451                             tf.ConfigProto.intra_op_parallelism_THREADS""")
452
453     parser.add_argument("--num_inter_threads",
454                         dest="num_intra_threads",
455                         type=int,
456                         default=0,
457                         required=False,
458                         help="""Number of intra op threads to use for
459                             tf.ConfigProto.intra_op_parallelism_THREADS""")
460
461     parser.add_argument("--float64",
462                         dest="float64",
463                         action="store_true",
464                         required=False,
465                         help="""When passed, using 64 point floating precision
466                             by settings globals.TF_FLOAT = tf.float64. False
467                             by default (use tf.float32).""")
468
469     if sys.argv[1].startswith('@'):
470         args = parser.parse_args(shlex.split(open(sys.argv[1][1:]).read(),
471                                         comments=True))
472     else:
473         args = parser.parse_args()
474
475     return args

```

A.19 l2hmc-qcd/utils/parse_inference_args.py

```
1 """
2 parse_inference_args.py
3
4 Implements method for parsing command line arguments for `gauge_model.py`
5
6 Author: Sam Foreman (github: @saforem2)
7 Date: 04/09/2019
8 """
9 import os
10 import sys
11 import argparse
12 import shlex
13
14 import utils.file_io as io
15
16 # from config import process_config
17 # from attr_dict import AttrDict
18
19 DESCRIPTION = 'Run inference on trained L2HMC model.'
20
21
22 # =====
23 # * NOTE:
24 #     - if action == 'store_true':
25 #         The argument is FALSE by default. Passing this flag will cause the
26 #         argument to be ''stored true''.
27 #     - if action == 'store_false':
28 #         The argument is TRUE by default. Passing this flag will cause the
29 #         argument to be ''stored false''.
30 # =====
31 def parse_args():
32     """Parse command line arguments."""
33     parser = argparse.ArgumentParser(
34         description=DESCRIPTION,
35         fromfile_prefix_chars='@',
36     )
37     ##### Lattice parameters #####
38     # parser.add_argument('--lattice_size', dest='lattice_size', type=int, required=True, default=100, help='Lattice size')
39     # parser.add_argument('--beta', dest='beta', type=float, required=True, default=0.5, help='Chemical potential beta')
40     parser.add_argument('--params_file', dest='params_file', required=False, default=None, help="""Path to `params.pkl` or `parameters.pkl` file containing the model parameters needed to run inference.""")
41
42     parser.add_argument("--run_steps",
43                         dest="run_steps",
44                         type=int,
45                         default=5000,
46                         required=False,
47                         help="""Number of evaluation 'run' steps to perform after training (i.e. length of desired chain generate using trained L2HMC sample). (Default: 5000)""")
48
49     parser.add_argument("--beta_inference",
50                         dest="beta_inference",
51                         type=float,
52                         default=None,
53                         required=False,
54                         help="""Flag specifying a singular value of beta at which to run inference using the trained L2HMC sampler. (Default: None)""")
55
56
57     parser.add_argument('--plot_lf',
58                         dest="plot_lf",
59                         type=bool,
60                         default=False,
61                         required=False,
62                         help='Flag indicating whether to plot the learned function. (Default: False)')
```

```

68     dest='plot_lf',
69     action='store_true',
70     required=False,
71     help=("""Flag that when passed will set
72         '--plot_lf=True', and will plot the 'metric'
73         distance between subsequent configurations, as
74         well as the determinant of the Jacobian of the
75         transformation for each individual leapfrog step,
76         as well as each molecular dynamics step (with
77         Metropolis-Hastings accept/reject).\n
78         When plotting the determinant of the Jacobian
79         following the MD update, we actually calculate
80         the sum of the determinants from each individual
81         LF step since this is the quantity that actually
82         enters into the MH acceptance probability.
83         (Default: `--plot_lf=False`, i.e. `--plot_lf` is
84         not passed mostly just because the plots are
85         extremely large (many LF steps during inference)
86         and take a while to actually generate.)"""))
87
88 parser.add_argument('--loop_net_weights',
89     dest='loop_net_weights',
90     action='store_true',
91     required=False,
92     help=("""Flag that when passed sets
93         '--loop_net_weights=True', and will iterate over
94         multiple values of `net_weights`, which are
95         multiplicative scaling factors applied to each of
96         the Q, S, T functions when running the trained
97         sampler.
98         (Default: `--loop_net_weights=False`, i.e.
99             `--loop_net_weights` is not passed)"""))
100
101 parser.add_argument('--loop_transl_weights',
102     dest='loop_transl_weights',
103     action='store_true',
104     required=False,
105     help=("""Flag that when passed will loop over different
106         values for the `translation_weight` in
107         `net_weights`, which is believed to be causing
108         the discrepancy between the observed and expected
109         value of the average plaquette when running
110             inference.""""))
111
112 parser.add_argument('--save_samples',
113     dest='save_samples',
114     action='store_true',
115     required=False,
116     help=("""Flag that when passed will set
117         '--save_samples=True', and save the samples
118         generated during the `run` phase.
119         (Default: `--save_samples=False`, i.e.
120             `--save_samples` is not passed).\n
121             WARNING!! This is very data intensive.""""))
122
123 if sys.argv[1].startswith('@'):
124     args = parser.parse_args(shlex.split(open(sys.argv[1][1:]).read(),
125                                         comments=True))
126 else:
127     args = parser.parse_args()
128
129 return args

```

A.20 l2hmc-qcd/utils/file_io.py

```
1 """
2 Helper methods for performing file IO.
3
4 Author: Sam Foreman (github: @saforem2)
5 Created: 2/27/2019
6 """
7 from __future__ import absolute_import
8 from __future__ import division
9 from __future__ import print_function
10 import os
11 import datetime
12 import pickle
13 import numpy as np
14
15 # pylint:disable=invalid-name
16
17 try:
18     import horovod.tensorflow as hvd
19
20     HAS_HOROVOD = True
21     hvd.init()
22
23 except ImportError:
24     HAS_HOROVOD = False
25
26 from config import FILE_PATH
27
28
29 def log(s, nl=True):
30     """Print string `s` to stdout if and only if hvd.rank() == 0."""
31     try:
32         if HAS_HOROVOD and hvd.rank() != 0:
33             return
34         print(s, end='\n' if nl else '')
35     except NameError:
36         print(s, end='\n' if nl else '')
37
38
39 def write(s, f, mode='a', nl=True):
40     """Write string `s` to file `f` if and only if hvd.rank() == 0."""
41     try:
42         if HAS_HOROVOD and hvd.rank() != 0:
43             return
44         with open(f, mode) as ff:
45             ff.write(s + '\n' if nl else '')
46     except NameError:
47         with open(f, mode) as ff:
48             ff.write(s + '\n' if nl else '')
49
50
51 def log_and_write(s, f):
52     """Print string `s` to std out and also write to file `f`."""
53     log(s)
54     write(s, f)
55
56
57 def check_else_make_dir(d):
58     """If directory `d` doesn't exist, it is created.
59
60     Args:
61         d (str): Location where directory should be created if it doesn't
62             already exist.
63     """
64     if not os.path.isdir(d):
65         log(f"Creating directory: {d}")
66         os.makedirs(d, exist_ok=True)
67
```

```

68
69 def make_dirs(dirs):
70     """Make directories if and only if hvd.rank == 0."""
71     _ = [check_else_make_dir(d) for d in dirs]
72
73
74 def _parse_flags(FLAGS):
75     """Helper method for parsing flags as both AttrDicts or generic dicts."""
76     if isinstance(FLAGS, dict):
77         flags_dict = FLAGS
78     else:
79         try:
80             flags_dict = FLAGS.__dict__
81         except (NameError, AttributeError):
82             pass
83     # if isinstance(FLAGS, dict):
84     try:
85         LX = flags_dict['space_size']
86         NS = flags_dict['num_samples']
87         LF = flags_dict['num_steps']
88         SS = flags_dict['eps']
89         QW = flags_dict['charge_weight']
90         NA = flags_dict['network_arch']
91         BN = flags_dict['use_bn']
92         DP = flags_dict['dropout_prob']
93         AW = flags_dict['aux_weight']
94         hmc = flags_dict['hmc']
95     try:
96         _log_dir = flags_dict['log_dir']
97     except KeyError:
98         _log_dir = ''
99
100    except (NameError, AttributeError):
101        LX = FLAGS.space_size
102        NS = FLAGS.num_samples
103        LF = FLAGS.num_steps
104        SS = FLAGS.eps
105        QW = FLAGS.charge_weight
106        NA = FLAGS.network_arch
107        BN = FLAGS.use_bn
108        DP = FLAGS.dropout_prob
109        AW = FLAGS.aux_weight
110        hmc = FLAGS.hmc
111    try:
112        _log_dir = FLAGS.log_dir
113    except AttributeError:
114        _log_dir = ''
115
116    out_dict = {
117        'LX': LX,
118        'NS': NS,
119        'LF': LF,
120        'SS': SS,
121        'QW': QW,
122        'NA': NA,
123        'BN': BN,
124        'DP': DP,
125        'AW': AW,
126        'hmc': hmc,
127        '_log_dir': _log_dir
128    }
129
130    return out_dict
131
132
133 def create_log_dir(FLAGS, root_dir=None, log_file=None):
134     """Automatically create and name `log_dir` to save model data to.
135
136     The created directory will be located in `logs/YYYY_M_D/`, and will have
137     the format (without `_qw{QW}` if running generic HMC):

```

```

138     `lattice{LX}_batch{NS}_lf{LF}_eps{SS}_qw{QW}`
139
140
141     Returns:
142         FLAGS, with FLAGS.log_dir being equal to the newly created log_dir.
143
144     NOTE: If log_dir does not already exist, it is created.
145     """
146
147     flags_dict = _parse_flags(FLAGS)
148     LX = flags_dict['LX']
149     NS = flags_dict['NS']
150     LF = flags_dict['LF']
151     SS = flags_dict['SS']
152     QW = flags_dict['QW']
153     NA = flags_dict['NA']
154     BN = flags_dict['BN']
155     DP = flags_dict['DP']
156     AW = flags_dict['AW']
157     _log_dir = flags_dict['_log_dir']
158
159     aw = str(AW).replace('.', '')
160     qw = str(QW).replace('.', '')
161     dp = str(DP).replace('.', '')
162
163     if flags_dict['hmc']:
164         run_str = f'HMC_lattice{LX}_batch{NS}_lf{LF}_eps{SS:.3g}'
165     else:
166         run_str = f'lattice{LX}_batch{NS}_lf{LF}_qw{qw}_aw{aw}_{NA}_dp{dp}'
167         if BN:
168             run_str += '_bn'
169
170     now = datetime.datetime.now()
171     # print(now.strftime("%b %d %Y %H:%M:%S"))
172     day_str = now.strftime('%Y_%m_%d')
173     time_str = now.strftime("%Y_%m_%d_%H%M")
174
175     project_dir = os.path.abspath(os.path.dirname(FILE_PATH))
176     # if FLAGS.log_dir is None:
177     if _log_dir is None:
178         if root_dir is None:
179             _dir = 'logs'
180         else:
181             _dir = root_dir
182
183     else:
184         if root_dir is None:
185             # _dir = FLAGS.log_dir
186             _dir = _log_dir
187         else:
188             _dir = os.path.join(_log_dir, root_dir)
189     root_log_dir = os.path.join(project_dir, _dir, day_str, time_str, run_str)
190     check_else_make_dir(root_log_dir)
191     if any('run' in i for i in os.listdir(root_log_dir)):
192         run_num = get_run_num(root_log_dir)
193         log_dir = os.path.abspath(os.path.join(root_log_dir,
194                                         f'run_{run_num}'))
195     else:
196         log_dir = root_log_dir
197     if log_file is not None:
198         write(f'Output saved to: \n\t{log_dir}', log_file, 'a')
199         write(80*'-', log_file, 'a')
200
201     return log_dir
202
203 def _list_and_join(d):
204     """For each dir `dd` in `d`, return a list of paths ['d/dd1', ...]"""
205     contents = [os.path.join(d, i) for i in os.listdir(d)]
206     paths = [i for i in contents if os.path.isdir(i)]
207

```

```

208     return paths
209
210
211 def list_and_join(d):
212     """Deal with the case of `d` containing multiple directories."""
213     if isinstance(d, (list, np.ndarray)):
214         paths = []
215         for dd in d:
216             _path = _list_and_join(dd)[0]
217             paths.append(_path)
218     else:
219         paths = _list_and_join(d)
220
221     return paths
222
223
224 def save_data(data, out_file, name=None):
225     """Save data to out_file using either pickle.dump or np.save."""
226     if os.path.isfile(out_file):
227         log(f"WARNING: File {out_file} already exists...")
228         tmp = out_file.split('.')
229         out_file = tmp[0] + '_1' + f'.{tmp[1]}'
230
231     log(f"Saving {name} to {out_file}...")
232     if out_file.endswith('.pkl'):
233         with open(out_file, 'wb') as f:
234             pickle.dump(data, f)
235
236     elif out_file.endswith('.npy'):
237         np.save(out_file, np.array(data))
238
239     else:
240         log("Extension not recognized! out_file must end in .pkl or .npy")
241
242
243 def save_params(params, out_dir):
244     check_else_make_dir(out_dir)
245     params_txt_file = os.path.join(out_dir, 'parameters.txt')
246     params_pkl_file = os.path.join(out_dir, 'parameters.pkl')
247     with open(params_txt_file, 'w') as f:
248         for key, val in params.items():
249             f.write(f"{key}: {val}\n")
250     with open(params_pkl_file, 'wb') as f:
251         pickle.dump(params, f)
252
253
254 def save_params_to_pkl_file(params, out_dir):
255     """Save `params` dictionary to `parameters.pkl` in `out_dir`."""
256     check_else_make_dir(out_dir)
257     params_file = os.path.join(out_dir, 'parameters.pkl')
258     # print(f"Saving params to: {params_file}.")
259     log(f"Saving params to: {params_file}.")
260     with open(params_file, 'wb') as f:
261         pickle.dump(params, f)
262
263
264 def get_run_num(log_dir):
265     """Get integer value for next run directory."""
266     check_else_make_dir(log_dir)
267     contents = os.listdir(log_dir)
268     if contents in ([], ['.DS_Store']):
269         return 1
270     try:
271         run_dirs = [i for i in os.listdir(log_dir) if 'run' in i]
272         run_nums = [int(i.split('_')[-1]) for i in run_dirs]
273         run_num = sorted(run_nums)[-1] + 1
274     except (ValueError, IndexError):
275         log(f"No previous runs found in {log_dir}, setting run_num=1.")
276         run_num = 1
277

```

```
278     return run_num
279
280
281 def get_eps_from_run_history_txt_file(txt_file):
282     """Parse 'run_history.txt' file and return 'eps' (step size)."""
283     with open(txt_file, 'r') as f:
284         data_line = [f.readline() for _ in range(10)][-1]
285     eps = float([i for i in data_line.split(' ') if i != ''][3])
286
287     return eps
```

A.21 l2hmc-qcd/loggers/summary_utils.py

```
1 """
2 summary_utils.py
3
4 Collection of helper methods for creating various summary objects for logging
5 data in TensorBoard.
6
7 Author: Sam Foreman (github: @saforem2)
8 Date: 08/16/2019
9 """
10 from __future__ import absolute_import
11 from __future__ import division
12 from __future__ import print_function
13
14 import tensorflow as tf
15 from lattice.lattice import u1_plaq_exact_tf
16
17
18 def grad_norm_summary(name_scope, grad):
19     with tf.name_scope(name_scope + '_gradients'):
20         grad_norm = tf.sqrt(tf.reduce_mean(grad ** 2))
21         summary_name = name_scope + '_grad_norm'
22         tf.summary.scalar(summary_name, grad_norm)
23
24
25 def check_var_and_op(name, var):
26     return (name in var.name or name in var.op.name)
27
28
29 def variable_summaries(var, name=''):
30     """Attach a lot of summaries to a Tensor (for TensorBoard visualization)"""
31     mean_name = 'mean'
32     stddev_name = 'stddev'
33     max_name = 'max'
34     min_name = 'min'
35     hist_name = 'histogram'
36     if name != '':
37         mean_name = name + '/' + mean_name
38         stddev_name = name + '/' + stddev_name
39         max_name = name + '/' + max_name
40         min_name = name + '/' + min_name
41         hist_name = name + '/' + hist_name
42
43     mean = tf.reduce_mean(var)
44     tf.summary.scalar(mean_name, mean)
45     with tf.name_scope('stddev'):
46         stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
47         tf.summary.scalar(stddev_name, stddev)
48         tf.summary.scalar(max_name, tf.reduce_max(var))
49         tf.summary.scalar(min_name, tf.reduce_min(var))
50         tf.summary.histogram(hist_name, var)
51     # activation summaries
52     # tf.summary.histogram(tensor_name + '/activations', x)
53     # tf.summary.scalar(tensor_name + '/sparsity', tf.nn.zero_fraction(x))
54
55
56 def add_loss_summaries(total_loss):
57     """Add summaries for losses in GaugeModel.
58
59     Generates a moving average for all losses and associated summaries for
60     visualizing the performance of the network.
61
62     Args:
63         total_loss: Total loss from model._calc_loss()
64
65     Returns:
66         loss_averages_op: Op for generating moving averages of losses.
67     """
68
```

```

68     # Compute the moving average of all individual losses and the total
69     # loss.
70     loss_averages = tf.train.ExponentialMovingAverage(0.9, name='avg')
71     losses = tf.get_collection('losses')
72     loss_averages_op = loss_averages.apply(losses + [total_loss])
73
74     # Attach a scalar summary to all individual losses and the total loss;
75     # do the same for the averaged version of the losses.
76     for l in losses + [total_loss]:
77         # Name each loss as '(raw)' and name the moving average version of
78         # the loss as the original loss name.
79         tf.summary.scalar(l.op.name, l)
80         tf.summary.scalar(l.op.name + 'moving_avg', loss_averages.average(l))
81
82     return loss_averages_op
83
84
85 def make_summaries_from_collection(collection, names):
86     try:
87         for op, name in zip(tf.get_collection(collection), names):
88             variable_summaries(op, name)
89     except AttributeError:
90         pass
91
92
93 def _create_loss_summaries(total_loss):
94     """Add summaries for losses in GaugeModel.
95
96     Generates a moving average for all losses and associated summaries for
97     visualizing the performance of the network.
98
99     Args:
100        total_loss: Total loss operation.
101    Returns:
102        loss_averages_op: Operation for generating moving averages of losses.
103    """
104    # Compute the moving avg. of all individ. losses and total loss
105    loss_averages = tf.train.ExponentialMovingAverage(0.9, name='avg')
106    losses = tf.get_collection('losses')
107    loss_averages_op = loss_averages.apply(losses + [total_loss])
108
109    std_names = ['std_loss', 'x_std_loss', 'z_std_loss']
110    std_losses = losses[:3]
111    for name, loss in zip(std_names, std_losses):
112        with tf.name_scope(name):
113            tf.summary.scalar(name + '/raw', loss)
114            tf.summary.scalar(name + '/moving_avg',
115                              loss_averages.average(loss))
116
117    charge_names = ['charge_loss', 'xq_loss', 'zq_loss']
118    charge_losses = losses[3:]
119    for name, loss in zip(charge_names, charge_losses):
120        with tf.name_scope(name):
121            tf.summary.scalar(name + '/raw', loss)
122            tf.summary.scalar(name + '/moving_avg',
123                              loss_averages.average(loss))
124
125    with tf.name_scope('total_loss'):
126        tf.summary.scalar('total_loss/raw', total_loss)
127        tf.summary.scalar('total_loss/moving_avg',
128                          loss_averages.average(total_loss))
129
130    return loss_averages_op
131
132
133 def _create_training_summaries(model):
134     """Create summary objects for training operations in TensorBoard."""
135     with tf.name_scope('loss'):
136         tf.summary.scalar('loss', model.loss_op)
137

```

```

138     with tf.name_scope('learning_rate'):
139         tf.summary.scalar('learning_rate', model.lr)
140
141     with tf.name_scope('step_size'):
142         tf.summary.scalar('step_size', model.dynamics.eps)
143
144
145     def _create_grad_norm_summaries(grad, var):
146         """Create grad_norm summaries."""
147         if 'XNet' in var.name:
148             net_str = 'XNet/'
149         elif 'VNet' in var.name:
150             net_str = 'VNet/'
151         else:
152             net_str = ''
153         with tf.name_scope(net_str):
154             if 'scale' in var.name:
155                 grad_norm_summary(net_str + 'scale', grad)
156                 tf.summary.histogram(net_str + 'scale', grad)
157             if 'transf' in var.name:
158                 grad_norm_summary(net_str + 'transformation', grad)
159                 tf.summary.histogram(net_str + 'transformation', grad)
160             if 'transl' in var.name:
161                 grad_norm_summary(net_str + 'translation', grad)
162                 tf.summary.histogram(net_str + 'translation', grad)
163
164
165     def _create_pair_summaries(grad, var):
166         """Create summary objects for a gradient, variable pair."""
167         try:
168             _name = var.name.split('/')[-2:]
169             if len(_name) > 1:
170                 name = _name[0] + '/' + _name[1][-2:]
171             else:
172                 name = var.name[-2:]
173         except (AttributeError, IndexError):
174             name = var.name[-2:]
175
176         with tf.name_scope(name):
177             var_name = var.name.replace(':', '')
178             variable_summaries(var, name=var_name)
179
180         grad_name = name + '/gradient'
181         with tf.name_scope(grad_name):
182             variable_summaries(grad, name=grad_name)
183
184
185     def _create_obs_summaries(model):
186         """Create summary objects for physical observables."""
187         with tf.name_scope('avg_charge_diffs'):
188             tf.summary.scalar('avg_charge_diffs',
189                               tf.reduce_mean(model.charge_diffs_op))
190
191         with tf.name_scope('avg_plaq'):
192             tf.summary.scalar('avg_plaq', model.avg_plaqs_op)
193
194         with tf.name_scope('avg_plaq_diff'):
195             tf.summary.scalar('avg_plaq_diff',
196                               (u1_plaq_exact_tf(model.beta)
197                                - model.avg_plaqs_op))
198
199
200     def _create_md_summaries(model):
201         """Create summary objects for each of the MD functions and outputs."""
202         for k1, v1 in model.l2hmc_fns['out_fns_f'].items():
203             for k2, v2 in v1.items():
204                 with tf.name_scope(f'{k1}_fn_{k2}_f'):
205                     variable_summaries(v2)
206
207         for k1, v1 in model.l2hmc_fns['out_fns_b'].items():

```

```

208     for k2, v2 in v1.items():
209         with tf.name_scope(f'{k1}_fn_{k2}_b'):
210             variable_summaries(v2)
211
212     with tf.name_scope('lf_out_f'):
213         variable_summaries(model.lf_out_f)
214
215     with tf.name_scope('lf_out_b'):
216         variable_summaries(model.lf_out_b)
217
218
219 def create_summaries(model, summary_dir, training=True):
220     """Create summary objects for logging in TensorBoard."""
221     summary_writer = tf.contrib.summary.create_file_writer(summary_dir)
222
223     grads_and_vars = zip(model.grads,
224                           model.dynamics.trainable_variables)
225
226     if training:
227         _create_training_summaries(model)
228
229     _create_obs_summaries(model)
230     _create_md_summaries(model)
231     # _ = _create_loss_summaries(model.loss_op)
232
233     for grad, var in grads_and_vars:
234         _create_pair_summaries(grad, var)
235
236         if 'kernel' in var.name:
237             _create_grad_norm_summaries(grad, var)
238
239     summary_op = tf.summary.merge_all(name='train_summary_op')
240
241     return summary_writer, summary_op

```

A.22 l2hmc-qcd/utils/data_loader.py

```
1 """
2 data_loader.py
3
4 Implements DataLoader class responsible for loading run_data.
5
6 Author: Sam Foreman (github: @saforem2)
7 Date: 05/03/2019
8 """
9 import os
10 import pickle
11
12 import numpy as np
13
14 def load_params_from_dir(d):
15     params_file = os.path.join(d, 'params.pkl')
16     with open(params_file, 'rb') as f:
17         params = pickle.load(f)
18
19     return params
20
21
22 def get_run_dirs(root_dir):
23     run_dirs = []
24     root_dir = os.path.abspath(root_dir)
25     for dirpath, dirnames, filenames in os.walk(root_dir):
26         for dirname in dirnames:
27             keys = ['steps_', '_beta_', '_eps_']
28             conditions = [key in dirname for key in keys]
29             if np.alltrue(conditions):
30                 run_dirs.append(os.path.join(dirpath, dirname))
31
32     return run_dirs
33
34 # def make_fig_dirs(run_dirs):
35 #     for dirpath, dirnames, filenames in os.walk(root_dir):
36 #         for dirname in dirnames:
37 #             keys = ['steps_', '_beta_', '_eps_']
38 #             conditions = [key in dirname for key in keys]
39 #             if np.alltrue(conditions):
40 #                 run_dirs.append(os.path.join(dirpath, dirname))
41 #     return run_dirs
42
43
44 class DataLoader:
45     def __init__(self, run_dir=None):
46         self.run_dir = None
47
48     def load_pkl_file(self, pkl_file):
49         with open(pkl_file, 'rb') as f:
50             contents = pickle.load(f)
51
52         return contents
53
54     def load_npz_file(self, npz_file):
55         arr = np.load(npz_file)
56
57         return arr.f.arr_0
58
59     def _load_observable(self, observable_str, run_dir=None):
60         if run_dir is None:
61             run_dir = self.run_dir
62
63         if not observable_str.endswith('.pkl'):
64             observable_str += '.pkl'
65
66         obs_dir = os.path.join(run_dir, 'observables')
```

```

68     obs_file = os.path.join(obs_dir, observable_str)
69
70     return self.load_pkl_file(obs_file)
71
72 def load_samples(self, run_dir):
73     samples_file = os.path.join(run_dir, 'samples_out.npz')
74     samples = self.load_npz_file(samples_file)
75
76     return samples
77
78 def load_leapfrogs(self, run_dir):
79     lf_f_file = os.path.join(run_dir, 'lf_forward.npz')
80     lf_b_file = os.path.join(run_dir, 'lf_backward.npz')
81     lf_f = self.load_npz_file(lf_f_file)
82     lf_b = self.load_npz_file(lf_b_file)
83
84     return (lf_f, lf_b)
85
86 def load_logdets(self, run_dir):
87     logdets_f_file = os.path.join(run_dir, 'logdets_forward.npz')
88     logdets_b_file = os.path.join(run_dir, 'logdets_backward.npz')
89
90     logdets_f = self.load_npz_file(logdets_f_file)
91     logdets_b = self.load_npz_file(logdets_b_file)
92
93     return (logdets_f, logdets_b)
94
95 def load_sumlogdets(self, run_dir):
96     sumlogdet_f_file = os.path.join(run_dir, 'sumlogdet_forward.npz')
97     sumlogdet_b_file = os.path.join(run_dir, 'sumlogdet_backward.npz')
98
99     sumlogdet_f = self.load_npz_file(sumlogdet_f_file)
100    sumlogdet_b = self.load_npz_file(sumlogdet_b_file)
101
102    return (sumlogdet_f, sumlogdet_b)
103
104 def load_plaqs(self, run_dir):
105     obs_dir = os.path.join(run_dir, 'observables')
106     plaqs_file = os.path.join(obs_dir, 'plaqs.pkl')
107
108     return self.load_pkl_file(plaqs_file)
109
110 def load_autocorrs(self, run_dir):
111     obs_dir = os.path.join(run_dir, 'observables')
112     autocorrs_file = os.path.join(obs_dir, 'charges_autocorrs.pkl')
113
114     return self.load_pkl_file(autocorrs_file)
115
116 def load_params(self, run_dir=None):
117     if run_dir is None:
118         run_dir = self.run_dir
119     params_file = os.path.join(run_dir, 'parameters.pkl')
120
121     return self.load_pkl_file(params_file)
122
123 def load_run_data(self, run_dir=None):
124     if run_dir is None:
125         run_dir = self.run_dir
126
127     run_data_file = os.path.join(run_dir, 'run_data.pkl')
128
129     return self.load_pkl_file(run_data_file)
130
131 def load_run_stats(self, run_dir=None):
132     if run_dir is None:
133         run_dir = self.run_dir
134
135     run_stats_file = os.path.join(run_dir, 'run_stats.pkl')
136
137     return self.load_pkl_file(run_stats_file)

```

```
138
139     def load_observables(self, run_dir=None):
140         if run_dir is None:
141             run_dir = self.run_dir
142
143         obs_dir = os.path.join(run_dir, 'observables')
144         files = os.listdir(obs_dir)
145
146         obs_names = [i.rstrip('.pkl') for i in files]
147         obs_files = [os.path.join(obs_dir, i) for i in files]
148
149         observables = {
150             n: self.load_pkl_file(f) for (n, f) in zip(obs_names, obs_files)
151         }
152
153     return observables
```

B | Appendix: Systematic Debugging Results of the Average Plaquette

In each of the plots below, the L2HMC sampler was trained for 25,000 steps with a simulated annealing schedule beginning at $\beta = 2.0$ and ending at $\beta = 5.0$. Additionally, the training was distributed across 16 nodes on COOLEY using horovod. The results are averaged over all N_{samples} samples in the mini-batch.

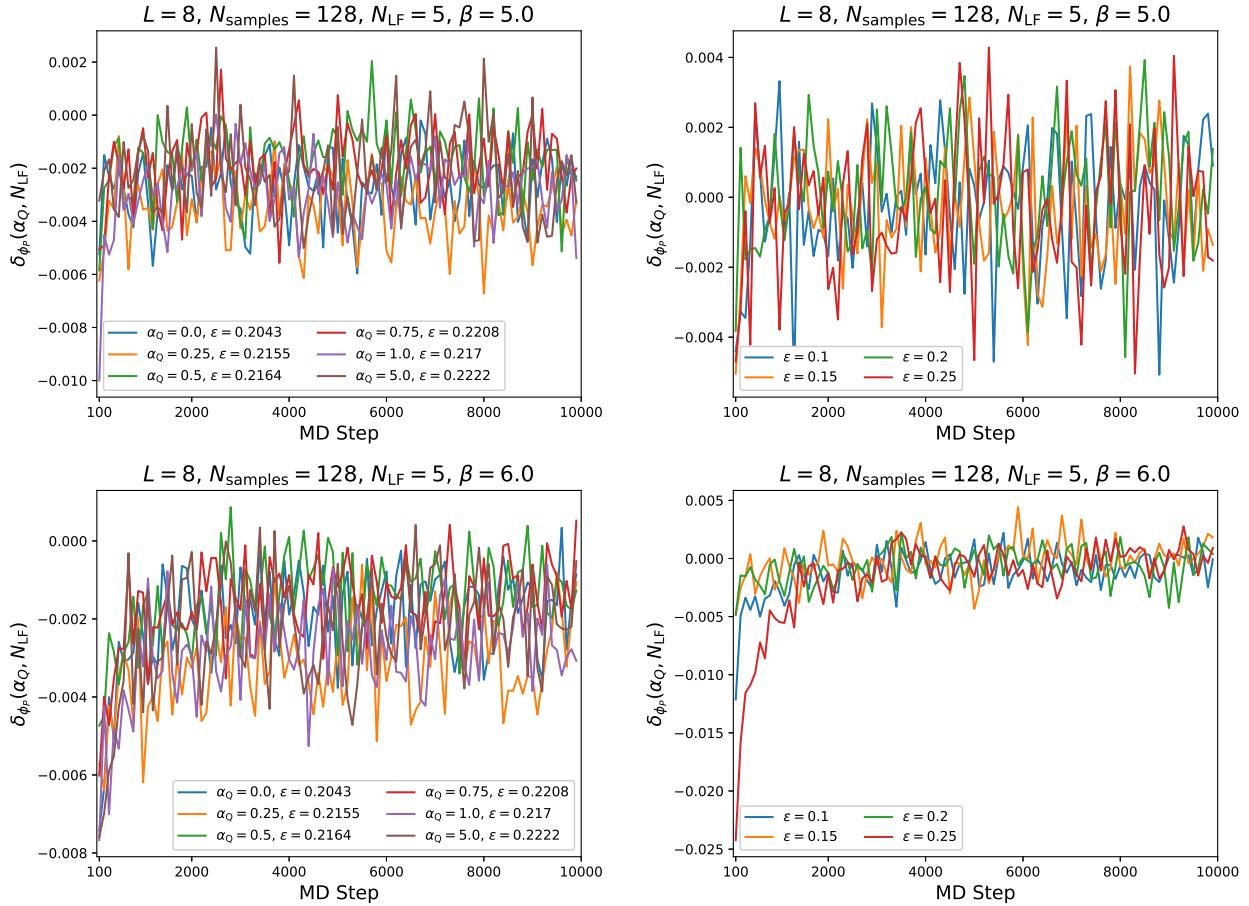


Figure B.1: δ_{ϕ_p} vs MD step with $N_{\text{LF}} = 5$ for $\beta = 5.0$ (top row) and $\beta = 6.0$ (bottom row). The results from the trained L2HMC (generic HMC) sampler are shown in the left (right) column. As can be seen, the difference δ_{ϕ_p} remains roughly consistent for all values of α_Q .

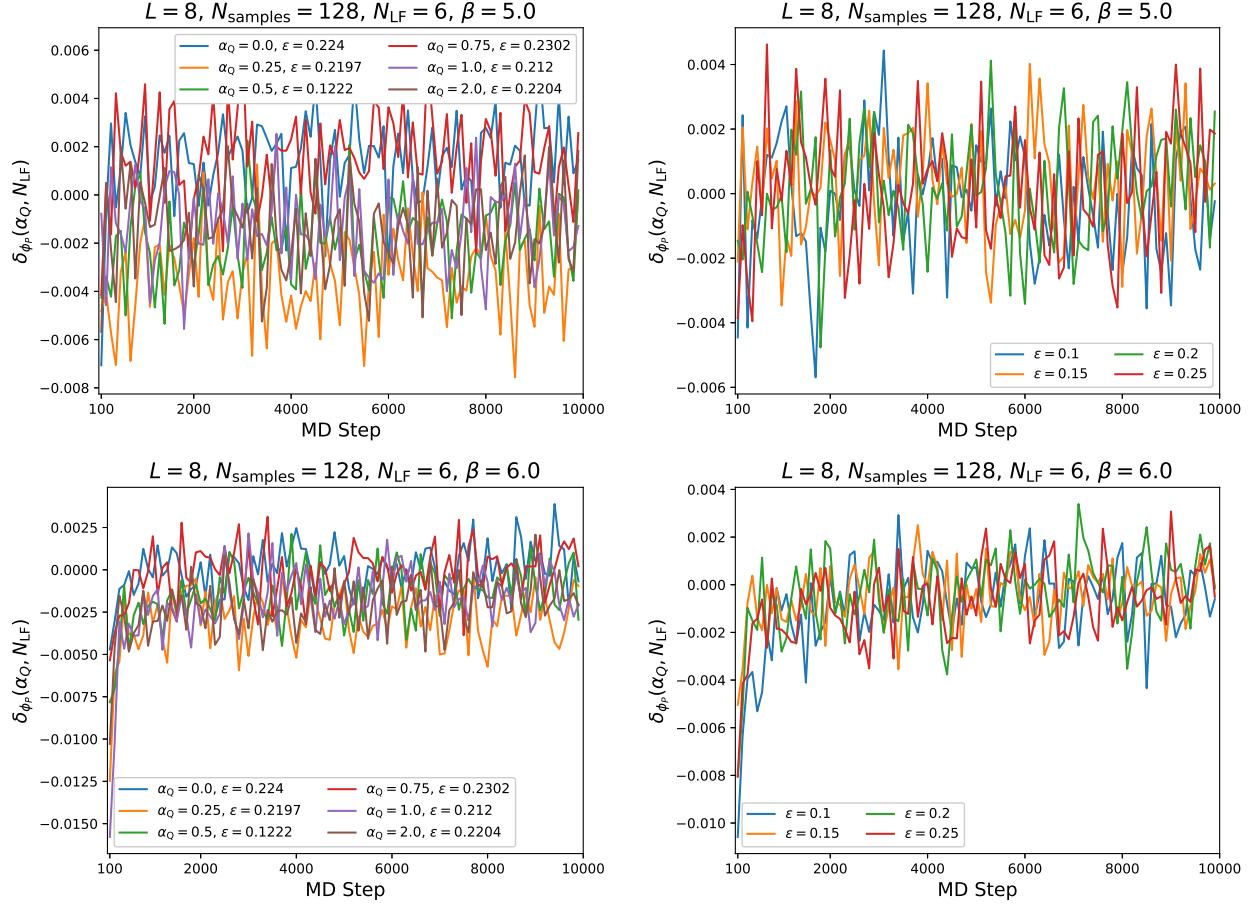


Figure B.2: δ_{ϕ_p} vs MD step with $N_{\text{LF}} = 5$ for $\beta = 5.0$ (top row) and $\beta = 6.0$ (bottom row). The results from the trained L2HMC (generic HMC) sampler are shown in the left (right) column. As can be seen, the difference δ_{ϕ_p} remains roughly consistent for all values of α_Q .

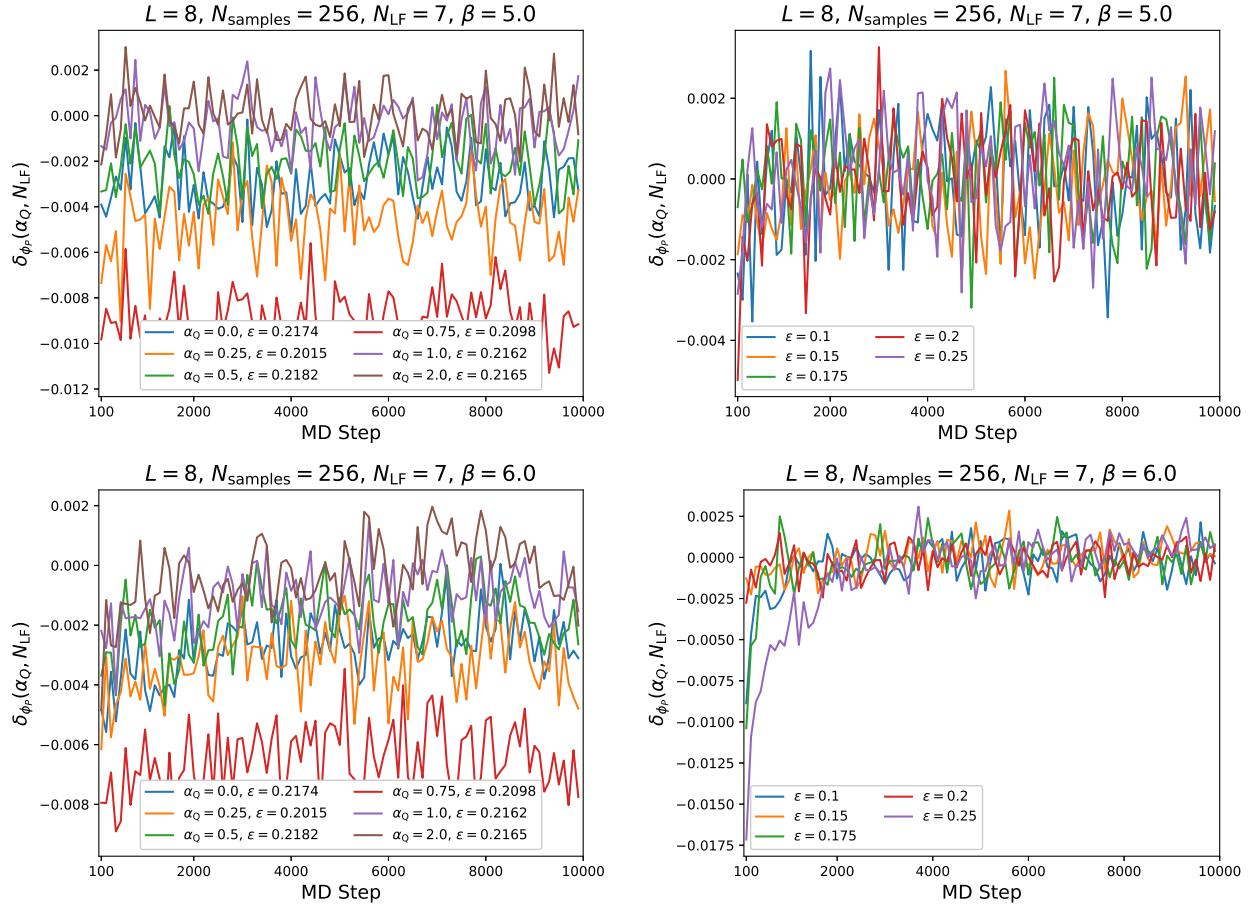


Figure B.3: δ_{ϕ_p} vs MD step with $N_{\text{LF}} = 7$. As can be seen, the difference δ_{ϕ_p} is noticeably larger for $\alpha_Q = 0.75$, but remains roughly consistent for all other values of α_Q .

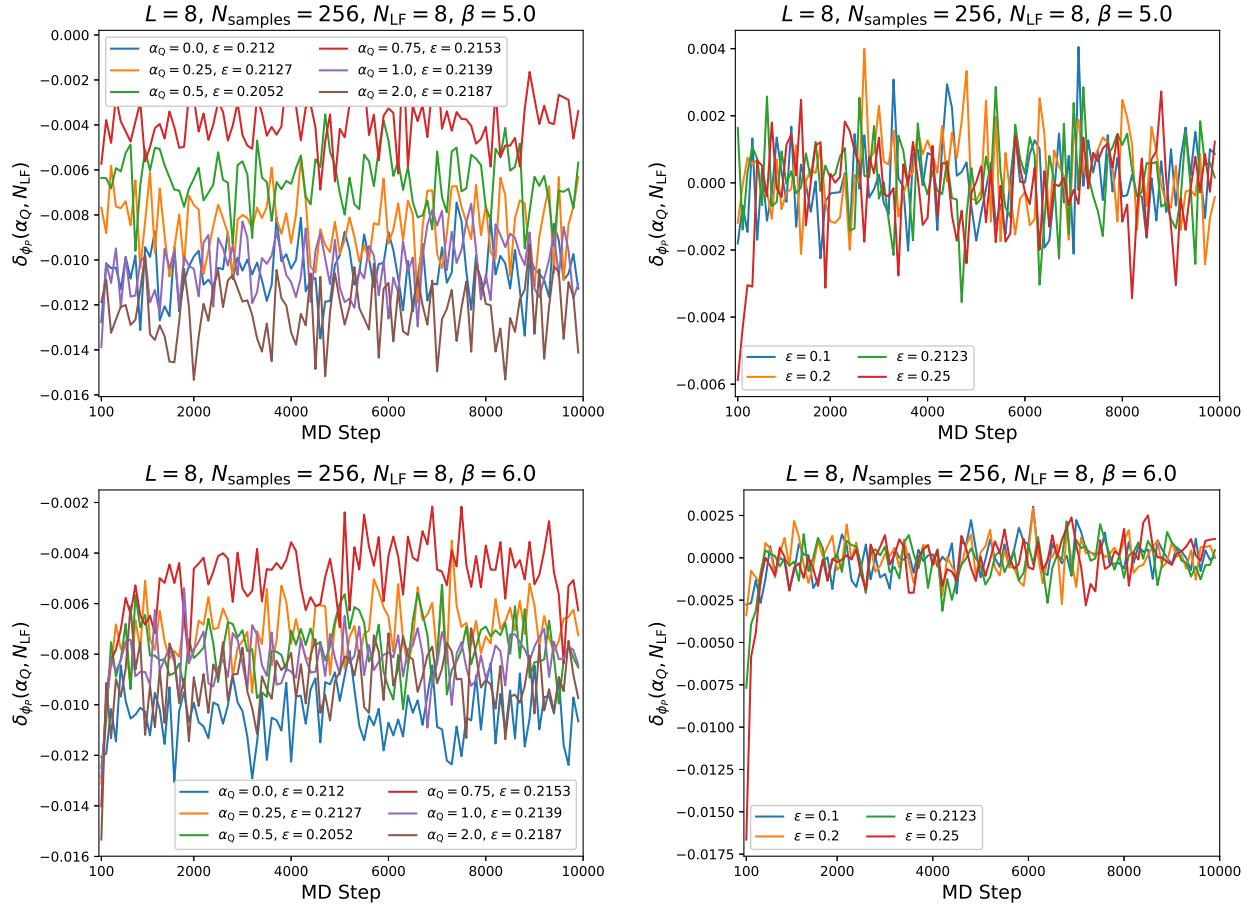


Figure B.4: δ_{ϕ_P} vs. MD step with $N_{LF} = 8$. As can be seen, the difference δ_{ϕ_P} remains roughly consistent for all values of α_Q .

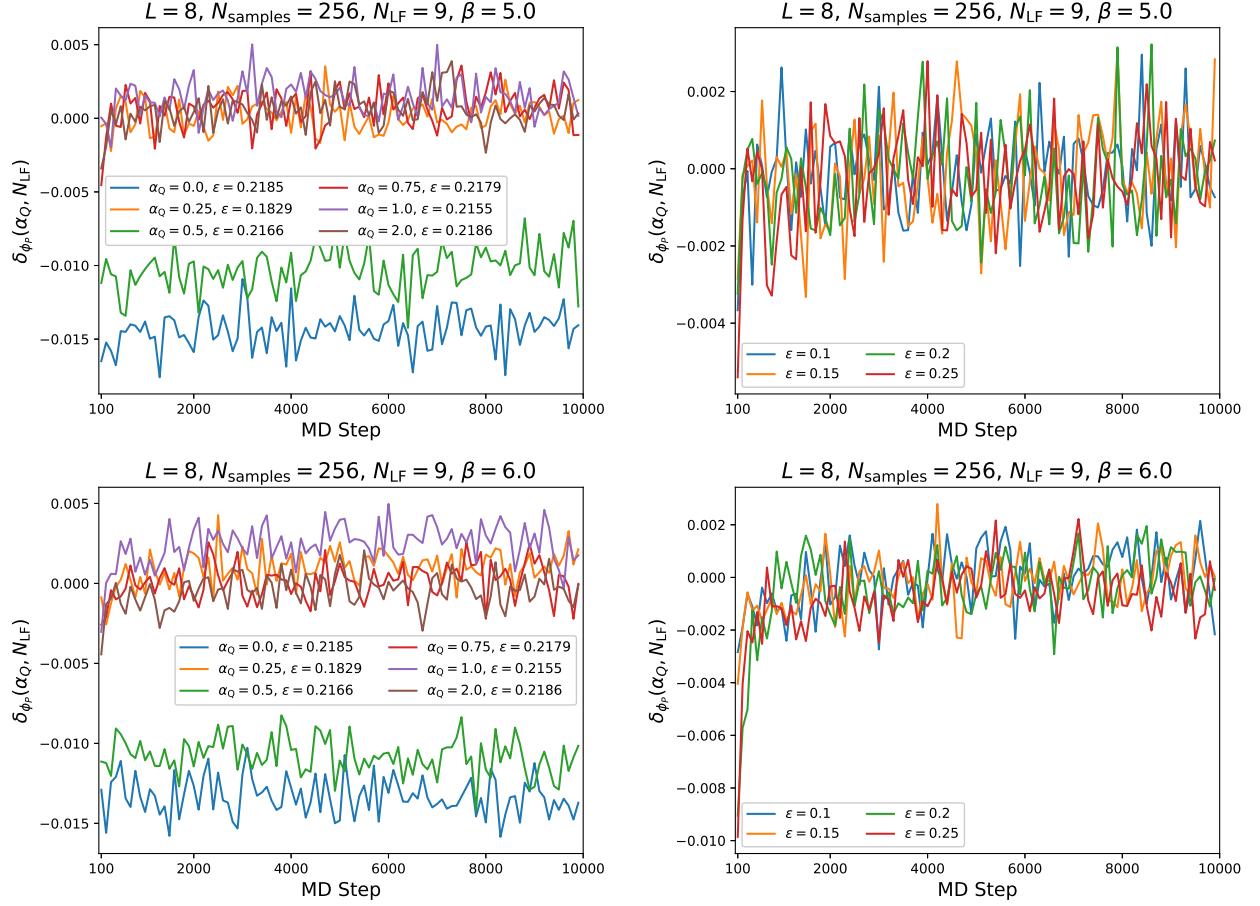


Figure B.5: δ_{ϕ_p} vs MD step with $N_{LF} = 9$. As can be seen, the difference δ_{ϕ_p} is largest for $\alpha_Q = 0.0$ and $\alpha_Q = 0.5$.

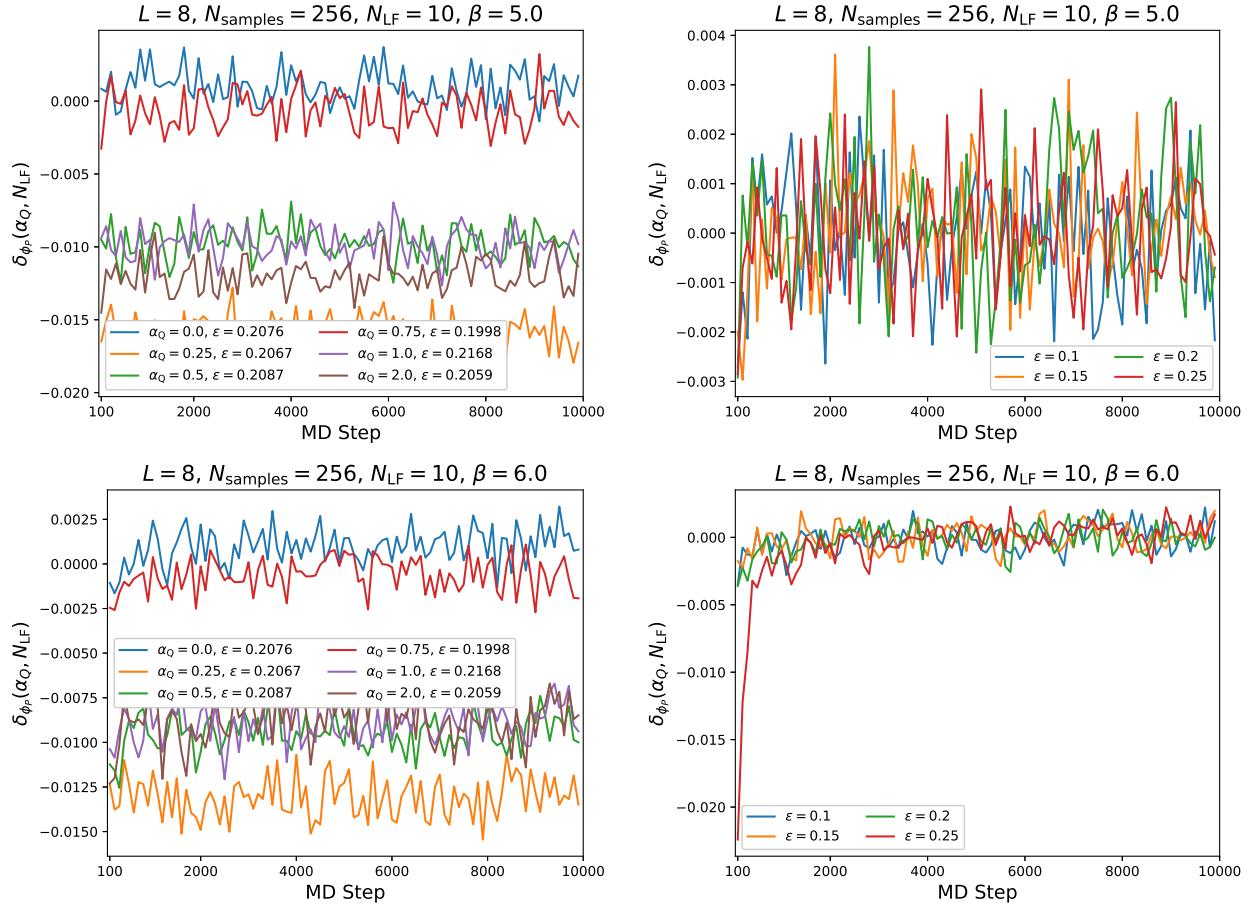


Figure B.6: δ_{ϕ_p} vs MD step with $N_{LF} = 10$. As can be seen, the difference δ_{ϕ_p} is smallest for $\alpha_Q = 0., 0.75$ and largest for $\alpha_Q = 0.25$, while $\alpha_Q = 0.5, 2.0$, takes on intermediate values.

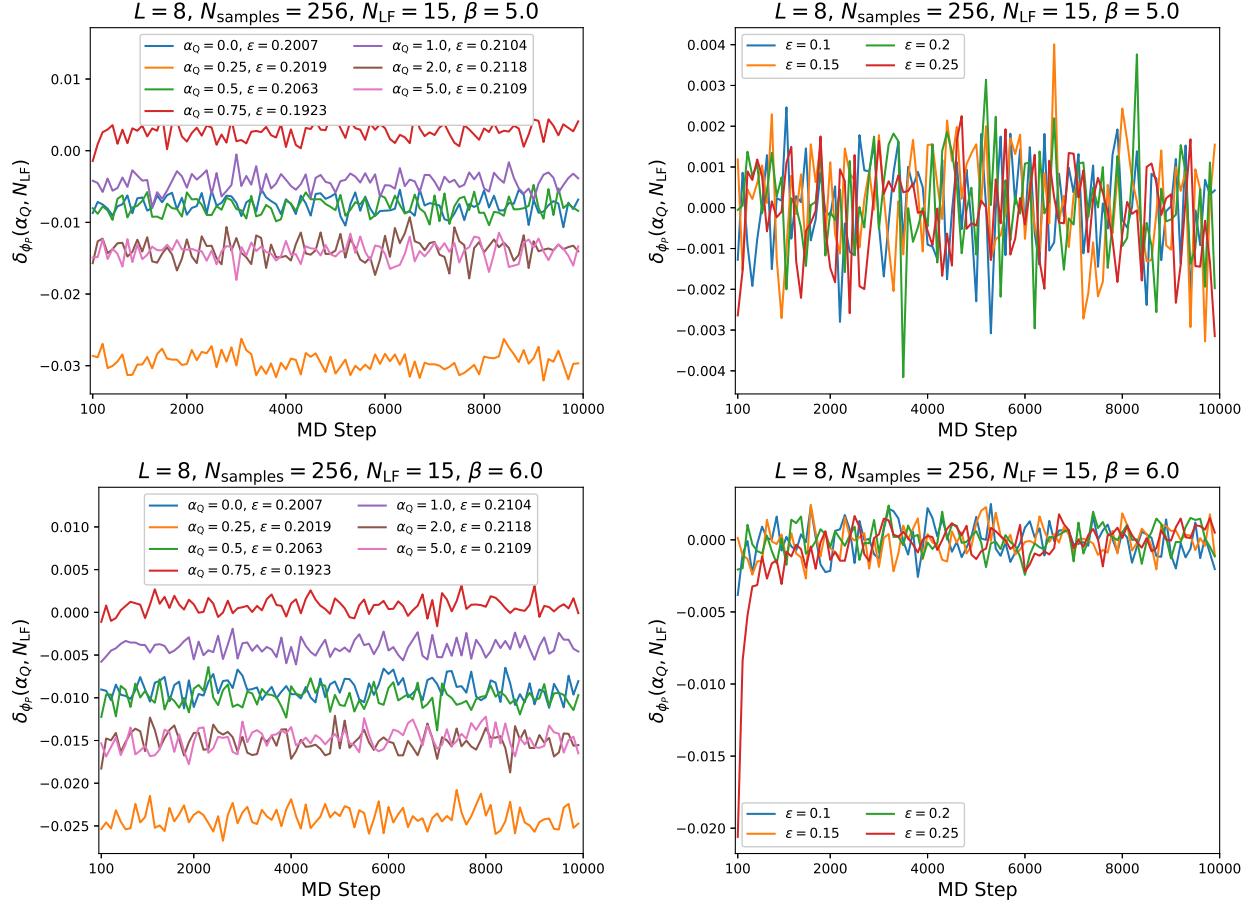


Figure B.7: δ_{ϕ_p} vs MD step with $N_{LF} = 15$. As can be seen, the difference δ_{ϕ_p} varies for different values of α_Q .

References

- [1] D. Levy, M. D. Hoffman, and J. Sohl-Dickstein, [arXiv:1711.09268 \[cs, stat\]](#), arXiv: 1711.09268 (2017).
- [2] S. Shalev-Shwartz and S. Ben-David, *Understanding machine learning: from theory to algorithms* (Cambridge university press, 2014).
- [3] P. Mehta, M. Bukov, C.-H. Wang, A. G. R. Day, C. Richardson, C. K. Fisher, and D. J. Schwab, *Physics Reports* **810**, arXiv: 1803.08823, 1 (2019).
- [4] D. P. Kingma and M. Welling, [arXiv:1312.6114 \[cs, stat\]](#), arXiv: 1312.6114 (2013).
- [5] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, [arXiv:1406.2661 \[cs, stat\]](#), arXiv: 1406.2661 (2014).
- [6] A. Ng, “Cs229 lecture notes - supervised learning”, 2012.
- [7] C. M. Bishop, *Pattern recognition and machine learning*, Information science and statistics (Springer, 2006).
- [8] S.-i. Horikawa, T. Furuhashi, and Y. Uchikawa, *IEEE Transactions on Neural Networks* **3**, 801 (1992).
- [9] J. Schmidhuber, *Neural Networks* **61**, arXiv: 1404.7828, 85 (2015).
- [10] M. Nielsen, [224](#) (2015).
- [11] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, *Nature* **323**, 533 (1986).
- [12] Y. LeCun, Y. Bengio, et al., *The handbook of brain theory and neural networks* **3361**, 1995 (1995).
- [13] W. Zhang, K. Itoh, J. Tanida, and Y. Ichioka, *Applied Optics* **29**, 4790 (1990).
- [14] G. E. Hinton, S. Osindero, and Y.-W. Teh, *Neural Computation* **18**, PMID: 16764513, 1527 (2006).
- [15] D. Scherer, A. Müller, and S. Behnke, in *International conference on artificial neural networks* (Springer, 2010), pp. 92–101.
- [16] J. T. Springenberg, A. Dosovitskiy, T. Brox, and M. Riedmiller, [arXiv:1412.6806 \[cs\]](#), arXiv: 1412.6806 (2014).
- [17] S.-H. Li and L. Wang, *Physical Review Letters* **121**, 260601, 260601 (2018).
- [18] C. Bény, [arXiv:1301.3124 \[quant-ph\]](#), arXiv: 1301.3124 (2013).
- [19] P. Mehta and D. J. Schwab, [arXiv:1410.3831 \[cond-mat, stat\]](#), arXiv: 1410.3831 (2014).
- [20] S. Foreman, J. Giedt, Y. Meurice, and J. Unmuth-Yockey, *EPJ Web of Conferences* **175**, edited by M. Della Morte, P. Fritzsch, E. Gámiz Sánchez, and C. Pena Ruano, 11025 (2018).
- [21] S. Bradde and W. Bialek, *Journal of Statistical Physics* **167**, 462 (2017).
- [22] J. Carrasquilla and R. G. Melko, *Nature Physics* **13**, 431 (2017).
- [23] L. Wang, *Phys. Rev. B* **94**, 195105 (2016).
- [24] W. Hu, R. R. P. Singh, and R. T. Scalettar, *Phys. Rev. E* **95**, 062122, 062122 (2017).
- [25] S. J. Wetzel, *Phys. Rev. E* **96**, 022140, 022140 (2017).
- [26] M. Levin and C. P. Nave, *Phys. Rev. Lett.* **99**, 120601 (2007).
- [27] Z.-C. Gu, M. Levin, B. Swingle, and X.-G. Wen, *Phys. Rev. B* **79**, 085118 (2009).

- [28] Z. Y. Xie, J. Chen, M. P. Qin, J. W. Zhu, L. P. Yang, and T. Xiang, *Phys. Rev. B* **86**, 045139 (2012).
- [29] Y. Meurice, *Phys. Rev. B* **87**, 064422 (2013).
- [30] Y. Liu, Y. Meurice, M. P. Qin, J. Unmuth-Yockey, T. Xiang, Z. Y. Xie, J. F. Yu, and H. Zou, *Phys. Rev. D* **88**, 056005 (2013).
- [31] A. Denbleyker, Y. Liu, Y. Meurice, M. P. Qin, T. Xiang, Z. Y. Xie, J. F. Yu, and H. Zou, *Phys. Rev. D* **89**, 016008 (2014).
- [32] J. F. Yu, Z. Y. Xie, Y. Meurice, Y. Liu, A. Denbleyker, H. Zou, M. P. Qin, and J. Chen, *Phys. Rev. E* **89**, 013308 (2014).
- [33] R. Savit, *Rev. Mod. Phys.* **52**, 453 (1980).
- [34] N. Prokof'ev and B. Svistunov, *Phys. Rev. Lett.* **87**, 160601 (2001).
- [35] T. A. EnSSlin and M. Frommert, *Phys. Rev. D* **83**, 105014 (2011).
- [36] R. H. Swendsen and J.-S. Wang, *Phys. Rev. Lett.* **58**, 86 (1987).
- [37] B. Kaufman, *Phys. Rev.* **76**, 1232 (1949).
- [38] A. Krizhevsky and G. Hinton, Master's thesis, Department of Computer Science, University of Toronto (2009).
- [39] B. Gidas, *IEEE Transactions on Pattern Analysis and Machine Intelligence* **11**, 164 (1989).
- [40] K. Tanaka, *Journal of Physics A: Mathematical and General* **35**, R81 (2002).
- [41] T. Horiguchi, Y. Honda, and M. Miya, *Physics Letters A* **227**, 319 (1997).
- [42] Z. Y. Xie, H. C. Jiang, Q. N. Chen, Z. Y. Weng, and T. Xiang, *Phys. Rev. Lett.* **103**, 160601 (2009).
- [43] G. E. Hinton and R. R. Salakhutdinov, *Science* **313**, 504 (2006).
- [44] P. E. Shanahan, D. Trewartha, and W. Detmold, *Phys. Rev. D* **97**, 094506 (2018).
- [45] S. Brooks, A. Gelman, G. Jones, and X.-L. Meng, *Handbook of markov chain monte carlo* (CRC press, 2011).
- [46] *Scribe_note_lecture17.pdf*, https://www.cs.cmu.edu/~epxing/Class/10708-14/scribe_notes/scribe_note_lecture17.pdf, (Accessed on 06/18/2019).
- [47] M. Hanada, arXiv:1808.08490 [cond-mat, physics:hep-lat, physics:hep-th], arXiv: 1808.08490 (2018).
- [48] R. M. Neal, arXiv:1206.1901 [physics, stat], arXiv: 1206.1901 (2012).
- [49] M. Betancourt, arXiv:1701.02434 [stat], arXiv: 1701.02434 (2017).
- [50] *30 joeylitalien/l2hmc: iclr 2018 reproducibility challenge: generalizing hamiltonian monte carlo with neural networks*, <https://joeylitalien.github.io/assets/reports/l2hmc.pdf>, (Accessed on 06/21/2019).
- [51] L. Dinh, J. Sohl-Dickstein, and S. Bengio, arXiv:1605.08803 [cs, stat], arXiv: 1605.08803 (2016).
- [52] C. Pasarica and A. Gelman, *Statistica Sinica* **20**, 343 (2010).
- [53] S. Ioffe and C. Szegedy, arXiv:1502.03167 [cs], arXiv: 1502.03167 (2015).